

CLAIR

Comprehensible LLM AI Intermediate Representation

A Specification for Auditable AI Reasoning

```
b1 .95 L0 @user "Explain yourself"  
b2 .9 L0 @self <b1 "I track my reasoning"  
b3 .85 L0 @self <b2 "So you can audit it"
```

A comprehensive specification of the intermediate representation
for reasoning traces, confidence propagation, and epistemic transparency

January 2026

Contents

1.	Preface	20
1.1.	Motivation: The Crisis of Epistemic Opacity	21
1.1.1.	The inadequacy of existing approaches.	22
1.2.	Research Questions	23
1.3.	Thesis Statement	23
1.4.	Contributions	24
1.4.1.	Primary Contributions	24
1.4.2.	Secondary Contributions	25
1.5.	Approach: Tracking, Not Proving	26
1.6.	The Thinker+Assembler Architecture	27
1.7.	Document Roadmap	29
1.7.1.	Part I: Foundations	29
1.7.2.	Part II: Self-Reference and Limits	29
1.7.3.	Part III: Dynamics	29
1.7.4.	Part IV: Realization	30
1.7.5.	Part V: Reflection	30
1.8.	A Note on Authorship	30
2.	Background	31
2.1.	Related Work	31
2.1.1.	Formal Epistemology	31
2.1.1.1.	The Structure of Justification	31
2.1.1.2.	Agrippa's Trilemma	32

2.1.1.3.	Probability vs. Epistemic Confidence	33
2.1.2.	Modal and Provability Logic	34
2.1.2.1.	Epistemic Logic	34
2.1.2.2.	Provability Logic	35
2.1.2.3.	Many-Valued and Graded Modal Logics	35
2.1.2.3.1.	Graded Modalities	36
2.1.2.3.2.	Fuzzy Modal Logics	36
2.1.2.3.3.	Decidability and Undecidability Results	36
2.1.2.3.4.	Connection to CLAIR's CPL	37
2.1.2.3.5.	Weighted Argumentation Semantics	37
2.1.2.3.5.1.	Bipolar Weighted Argumentation	37
2.1.2.3.5.2.	Comparative Analysis	38
2.1.2.3.5.3.	Parameterised Gradual Semantics	38
2.1.2.3.5.4.	Relation to CLAIR	38
2.1.3.	Truth Maintenance and Argumentation	40
2.1.3.1.	Justification-based TMS	40
2.1.3.2.	Assumption-based TMS	40
2.1.3.3.	Argumentation Frameworks	40
2.1.3.4.	Pollock's Defeaters	41
2.1.4.	Belief Revision	41
2.1.4.1.	The AGM Framework	41
2.1.4.2.	Ranking Theory	42
2.1.4.3.	Dynamic Epistemic Logic	42
2.1.5.	Graded Justification Logic	43
2.1.5.1.	Milnikel's Logic of Uncertain Justifications	43
2.1.5.2.	Fan and Liau's Logic of Justified Uncertain Beliefs	44
2.1.5.3.	The Gap CLAIR Fills	44
2.1.6.	Type-Theoretic Approaches to Uncertainty	44

2.1.6.1. Information Flow Types	45
2.1.6.2. Refinement Types	45
2.1.6.3. Dependent Types and Proof Assistants	45
2.1.6.4. Probabilistic Programming	46
2.1.6.5. Justification Logic	46
2.2. Synthesis: The Gap CLAIR Fills	47
2.2.1. Positioning: How CLAIR Differs	48
2.2.1.1. vs. Probabilistic Approaches	48
2.2.1.2. vs. Subjective Logic	49
2.2.1.3. vs. Weighted Argumentation	49
2.2.1.4. vs. Fuzzy Modal Logic	49
2.2.1.5. vs. Graded Justification Logic	50
2.2.1.6. vs. Belief Revision (AGM)	50
2.2.1.7. vs. Type Theory	50
2.2.1.8. The Pragmatic Rationale	51
3. Confidence System	51
3.1. Confidence as Calibrated Reliability	52
3.1.1. Semantic Commitments	52
3.1.2. The Problem with Probability	53
3.1.3. Definition of Confidence	54
3.1.4. Falsifiability Criteria	55
3.1.5. Comparison with Subjective Logic	55
3.2. The Aggregation Monoid	56
3.2.1. Probabilistic OR (Oplus)	56
3.2.2. Confidence-Increasing Property	57
3.2.3. The “Survival of Doubt” Interpretation	57
3.3. The Multiplication Monoid	58
3.3.1. Conjunctive Confidence Propagation	58

3.3.2.	The Derivation Monotonicity Principle	58
3.4.	Defeat Operations	59
3.4.1.	Undercut: Attacking the Inference	59
4.	$c(1 - (d_1 + d_2 - d_1 \text{ times } d_2))$	60
5.	$c(1 - (d_1 \text{ oplus } d_2))$	60
6.	undercut($c, d_1 \text{ oplus } d_2$)	60
6.0.1.	Rebut: Competing Evidence	60
6.0.2.	Rebut Normalization Limitation	61
6.1.	Independence Assumptions for Aggregation	62
6.1.1.	When Oplus is Valid	62
6.1.2.	When Oplus Breaks	62
6.1.3.	Correlation-Aware Alternatives	63
6.1.4.	Interval-Based Confidence: Dependency Bounds	63
6.1.4.1.	Interval Aggregation Bounds	64
6.1.4.2.	Interval Propagation	64
6.1.4.3.	Dependency Bounds from Provenance Tracking	65
6.1.4.4.	When to Use Intervals vs Point Values	66
6.2.	Conclusion	66
7.	Justification as Labeled DAGs	67
7.1.	The Inadequacy of Trees	68
7.1.1.	The Shared Premise Problem	68
7.1.2.	Why Not Cycles?	69
7.2.	Labeled Edges for Defeat	69
7.2.1.	The Defeat Problem	69
7.2.2.	Formal Definition	70
7.2.3.	Well-Formedness Constraints	70
7.2.4.	DAG-Only vs Cyclic Choice	71
7.2.5.	Fixed-Point Semantics for Cyclic Defeat	71

7.3.	Confidence Propagation	72
7.4.	Reinstatement	73
7.5.	Mutual Defeat	73
7.6.	Correlated Evidence	74
7.7.	Connection to Prior Art	74
7.7.1.	Truth Maintenance Systems	74
7.7.2.	Argumentation Frameworks	74
7.7.3.	Justification Logic	74
7.8.	The Tracking Paradigm	75
7.8.1.	State Representation	75
7.8.1.1.	The Epistemic State	75
7.8.1.2.	Comparison with Proving Paradigm	75
7.8.2.	Update Rules	76
7.8.2.1.	Primitive Actions	76
7.8.2.2.	Action Semantics	76
7.8.3.	Correctness Criteria	77
7.8.3.1.	Syntactic Correctness	77
7.8.3.2.	Semantic Correctness (Internal)	77
7.8.3.3.	Semantic Correctness (External)	77
7.8.4.	Tracking vs. Proving: A Summary	78
7.8.5.	Practical Implications	79
7.8.5.1.	Explainability	79
7.8.5.2.	Debugging	79
7.8.5.3.	Revision	79
7.8.5.4.	Uncertainty	79
7.9.	Conclusion	79
7.10.	The Problem of Self-Reference	82
7.10.1.	Direct Self-Reference	82

7.10.2. Why Self-Reference Matters	82
7.11. Löb's Theorem and Anti-Bootstrapping	83
7.11.1. The Classical Result	83
7.11.2. Application to CLAIR	83
7.12. Tarski's Hierarchy: Stratified Introspection	84
7.12.1. The Classical Solution	84
7.12.2. Stratified Beliefs in CLAIR	84
7.12.3. What Stratification Rules Out	85
7.12.4. The Cost of Safety	85
7.13. Kripke's Fixed Points: Safe Self-Reference	86
7.13.1. The Fixed-Point Construction	86
7.13.2. The Self-Reference Escape Hatch	87
7.13.3. Classification of Self-Reference	87
7.14. Provability Logic and CLAIR	88
7.14.1. Gödel-Löb Logic (GL)	88
7.14.2. GL vs Other Modal Logics	89
7.14.3. Solovay's Completeness	89
7.15. Confidence-Bounded Provability Logic (CPL)	89
7.15.1. The Literature Gap	90
7.15.2. CPL Syntax	90
7.15.3. CPL Semantics	90
7.15.4. The Graded Löb Axiom: DESIGN AXIOM	91
7.15.5. Choosing the Discount Function	91
7.15.6. The Anti-Bootstrapping Theorem	92
7.15.7. Modal Axioms in CPL	92
7.15.8. CPL Consistency	92
7.16. Decidability of CPL	93
7.16.1. The Vidal Result	93

7.16.2. The Role of Converse Well-Foundedness	94
7.16.3. Decidable Fragments	94
7.16.3.1. CPL-finite: Discrete Confidence	94
7.16.3.2. CPL-0: Stratified Only	95
7.16.4. Trade-offs	95
7.17. Alternative: CPL-Gödel	95
7.18. “Conservative Over GL”: Clarification	96
7.19. Design Recommendations for CLAIR	96
7.19.1. The Two-Layer Approach	96
7.19.2. Hard Bans	97
7.19.3. Type-Level Anti-Bootstrapping	97
7.20. Related Work	98
7.20.1. Provability Logic	98
7.20.2. Self-Reference in AI	98
7.20.3. Fuzzy Modal Logic	98
7.21. Conclusion	98
7.22. The Grounding Problem	100
7.22.1. What grounding means in CLAIR.	100
7.23. Perceptual Grounding	101
7.23.1. The grounding cap theorem.	101
7.24. Axiomatic Grounding	101
7.24.1. The problem of circular axioms.	102
7.25. Social Grounding and Testimony	102
7.25.1. Reputation and source tracking.	102
7.26. The Ungrounded: Free-Floating Beliefs	103
7.26.1. Creative inference and hypothetical reasoning.	103
7.27. Grounding Requirements	103
7.27.1. Tier 1: Strict grounding.	103

7.27.2.	Tier 2: Demarcated ungrounding.....	104
7.27.3.	Tier 3: Permissive ungrounding.....	104
28.	Summary	104
29.	The Challenge of Revising Structured Beliefs	105
30.	Background: The AGM Framework	105
31.	Why AGM Doesn't Directly Apply	106
31.1.	Graded confidence.....	106
31.2.	DAG-structured justification.....	106
31.3.	Invalidation conditions.....	106
32.	GDBR: Graded DAG Belief Revision	106
32.1.	Confidence preservation.....	107
32.2.	Justification propagation.....	107
32.3.	Invalidation responsiveness.....	107
33.	The GDBR Algorithm	107
33.1.	Phase 1: Conflict detection.....	107
33.2.	Phase 2: Confidence comparison.....	108
33.3.	Phase 3: Graph restructuring.....	108
34.	The Revision Fixed Point Theorem	108
35.	Special Cases and Extensions	109
35.1.	Defeater revision.....	109
35.2.	Package revision.....	109
35.3.	Iterated revision.....	109
36.	Connection to Argumentation Theory	109
37.	Summary	110
38.	The Social Dimension of Knowledge	111
39.	Machine-Checked Proofs in Lean 4	113
8.	Implementation	114
8.1.	10.1 Architectural Overview	114

8.1.1.	10.1.1 The Thinker+Assembler Architecture	114
8.1.2.	10.1.2 Role Separation	115
8.1.3.	10.1.3 Why Not a Traditional Programming Language?	115
8.2.	10.2 The CLAIR Format	116
8.2.1.	10.2.1 Belief Structure	116
8.2.2.	10.2.2 Example: Algorithm Selection	116
8.2.3.	10.2.3 Source Types	117
8.2.4.	10.2.4 Confidence Semantics	117
8.3.	10.3 Stratification and the Löb Discount	117
8.3.1.	10.3.1 Level Rules	118
8.3.2.	10.3.2 Example: Confidence Decay Through Meta-Levels	118
8.3.3.	10.3.3 Formal Verification	118
8.4.	10.4 DAG Structure	119
8.4.1.	10.4.1 Acyclicity Requirement	119
8.4.2.	10.4.2 Formal Definition	119
8.5.	10.5 Lean Formalization Status	119
8.5.1.	10.5.1 Proven Properties	119
8.5.2.	10.5.2 Properties with <i>sorry</i>	120
8.6.	10.6 The Assembler’s Role	120
8.6.1.	10.6.1 Interpreting CLAIR Traces	120
8.6.2.	10.6.2 Example: Assembler Output	120
8.6.3.	10.6.3 Error Handling	121
8.7.	10.7 Querying CLAIR Traces	121
8.7.1.	10.7.1 “Why?” Queries	122
8.7.2.	10.7.2 “When to Reconsider?” Queries	122
8.7.3.	10.7.3 Debug Output	122
8.8.	10.8 Comparison with Traditional Approaches	122
8.9.	10.9 Non-Compositional Design	122

8.9.1.	10.9.1 Beliefs Do Not Compose	123
8.9.2.	10.9.2 Confidence Still Propagates	123
8.9.3.	10.9.3 Traces Do Not Compose Either	123
8.10.	10.10 Limitations	124
8.10.1.	10.10.1 Current Limitations	124
8.10.2.	10.10.2 Fundamental Limitations	124
8.11.	10.11 Future Work	124
8.12.	10.12 Summary	124
8.13.	The Hard Problem of AI Experience	126
8.14.	Engaging with Gödel, Church, and Fundamental Limits	128
9.	Conclusion	129
9.1.	Summary of Contributions	129
9.1.1.	Theoretical Foundations	129
9.1.2.	Implementation and Verification	131
9.1.3.	Conceptual Contributions	131
9.2.	Limitations and Open Challenges	132
9.2.1.	Independence Assumptions	132
9.2.2.	Rebut Normalization Limitations	133
9.2.3.	Decidability and Complexity	133
9.2.4.	Evaluation Scope	133
9.2.5.	The “0.5 = Ignorance” Question	133
9.3.	Future Directions	134
9.3.1.	Theoretical Extensions	134
9.3.2.	Implementation and Tooling	134
9.3.3.	Applications	135
9.3.4.	Philosophical Connections	135
9.4.	Closing Remarks: Honesty as a Design Principle	135
9.4.1.	The Meta-Level Question	136

9.4.2. Final Assessment	136
10. Evaluation	137
10.1. Evaluation Framework	138
10.1.1. Research Questions	138
10.1.2. Tasks and Datasets	138
10.1.3. Baselines	138
10.1.4. Metrics	139
10.1.4.1. Accuracy Metrics	139
10.1.4.2. Calibration Metrics	139
10.1.4.3. Auditability Metrics	139
10.2. Methodology	140
10.2.1. CLAIR Prompting Protocol	140
10.2.2. Confidence Extraction	140
10.2.3. Statistical Analysis	140
10.3. Results	141
10.3.1. RQ1: Correctness	141
10.3.2. RQ2: Calibration	141
10.3.2.1. Brier Score	141
10.3.2.2. Expected Calibration Error	142
10.3.2.3. Reliability Diagrams	142
10.3.3. RQ3: Auditability	142
10.4. Ablation Studies	143
10.5. Error Analysis	144
10.5.1. Common Failure Modes	144
10.5.2. CLAIR-Specific Errors	145
10.6. Discussion	145
10.6.1. Implications for Design	145
10.6.2. Limitations	146

10.6.3. Threats to Validity	146
10.7. Conclusion	147
10.8. Future Work	147
10.8.1. Extended Evaluation	147
10.8.2. Ablation and Extensions	148
11. Appendices	149
12. Complete Lean 4 Formalization	149
12.1. A.1 Project Structure	149
12.2. A.2 Build Instructions	150
12.2.1. Prerequisites	150
12.2.2. Building	150
12.2.3. Build Output	150
12.2.4. Verification Status	151
12.3. A.3 Theorem Inventory	151
12.3.1. Confidence Algebra	151
12.3.1.1. Basic Properties	151
12.3.1.2. Probabilistic OR (\oplus)	152
12.3.1.3. Undercut	152
12.3.1.4. Rebuttal	153
12.3.2. Stratified Belief	153
12.4. A.4 Key Code Excerpts	154
12.4.1. A.4.1 Confidence Type Definition	154
12.4.2. A.4.2 Probabilistic OR Operation	154
12.4.3. A.4.3 Expression Grammar	155
12.4.4. A.4.4 Typing Judgment	156
12.4.5. A.4.5 Stratified Belief Introspection	157
12.4.6. A.4.6 Evaluation Function	158
12.5. A.5 Five Properties Demonstration	159

12.6. A.6 Relationship to Dissertation Claims	160
12.6.1. Claim: “Machine-Checked Proofs” (Chapter 9)	160
12.6.2. Claim: “Decidable Type Checking” (Chapter 10)	160
12.6.3. Claim: “Runnable Interpreter” (Chapter 10)	160
12.7. A.7 Future Work	161
13. Reference Interpreter Design	161
13.1. B.1 Architecture Overview	161
13.2. B.2 Single-Step Semantics	162
13.2.1. B.2.1 Core Lambda Calculus Rules	162
13.2.2. B.2.2 Belief Operations	163
13.2.3. B.2.3 Defeat Operations	163
13.3. B.3 Multi-Step Evaluation with Fuel	164
13.4. B.4 Example Walkthroughs	164
13.4.1. B.4.1 Simple Belief Formation	164
13.4.2. B.4.2 Evidence Aggregation	165
13.4.3. B.4.3 Undercutting in Action	165
13.4.4. B.4.4 Rebuttal and Confidence Collapse	166
13.4.5. B.4.5 Derivation Chain	167
13.5. B.5 Key Properties	167
13.6. B.6 Implementation Notes	168
13.7. B.7 Relation to Chapter 10	169
13.8. B.8 Haskell Reference Implementation	169
13.8.1. B.8.1 Project Structure	169
13.8.2. B.8.2 Syntax Definition	170
13.8.3. B.8.3 Confidence Algebra	171
13.8.4. B.8.4 Type Checking	172
13.8.5. B.8.5 Evaluation	173
13.8.6. B.8.6 REPL Usage	174

13.8.7. B.8.7 Test Suite	175
13.8.8. B.8.8 Building and Running	175
13.8.8.1. Build	175
13.8.8.2. Run REPL	176
13.8.8.3. Run Tests	176
13.8.9. B.8.9 Design Rationale	176
13.8.10. B.8.10 Relation to Lean Formalization	176
14. Additional Proofs	177
14.1. C.1 DAG Necessity for Well-Founded Confidence Propagation	177
14.1.1. C.1.1 Statement of the Problem	177
14.1.2. C.1.2 The Cyclic Counterexample	177
14.1.3. C.1.3 The DAG Well-Foundedness Theorem	178
14.1.4. C.1.4 Practical Implications	179
14.2. C.2 CPL Consistency Proof	179
14.2.1. C.2.1 CPL Axiom System	179
14.2.2. C.2.2 Finite Model Construction	179
14.2.3. C.2.3 Design Axiom Status	180
14.3. C.3 Defeat Composition Algebra	181
14.3.1. C.3.1 Undercut Composition	181
15. undercut($c \times (1 - d_1)$, d_2)	182
16. $(c \times (1 - d_1)) \times (1 - d_2)$	182
17. $c \times ((1 - d_1) \times (1 - d_2))$	182
18. $c \times (1 - d_1 - d_2 + d_1 d_2)$	182
19. $c \times (1 - (d_1 + d_2 - d_1 d_2))$	182
20. $c \times (1 - (d_1 \oplus d_2))$	182
21. undercut(c , $d_1 \oplus d_2$)	182
21.0.1. C.3.2 Corollaries of Undercut Composition	183
22. $\text{undercut}(c, d_2 \oplus d_1) = \text{undercut}(\text{undercut}(c, d_2), d_1)$. \blacksquare	183

22.0.1. C.3.3 Rebut Algebra	183
23. $c_{\text{for}} / (c_{\text{for}} + c_{\text{against}}) + c_{\text{against}} / (c_{\text{against}} + c_{\text{for}})$	184
24. $(c_{\text{for}} + c_{\text{against}}) / (c_{\text{for}} + c_{\text{against}})$	184
25. 1 ■	184
25.0.1. C.3.4 Interaction Between Undercut and Rebut	184
25.0.2. C.3.5 Limitation: Rebut Normalization	185
26. $\lambda c_{\text{for}} / (\lambda (c_{\text{for}} + c_{\text{against}}))$	185
27. $c_{\text{for}} / (c_{\text{for}} + c_{\text{against}})$	185
28. $\text{rebut}(c_{\text{for}}, c_{\text{against}}).$ ■	185
29. Glossary	186
29.1. D.1 Term Definitions	186
29.1.1. Epistemic Terms	186
29.1.2. Operations and Relations	187
29.1.3. Structural Properties	187
29.1.4. Logical and Modal Terms	187
29.1.5. Computational Terms	188
29.1.6. Argumentation and Belief Revision	188
29.1.7. Impossibility Results	189
29.2. D.2 Notation Table	189
29.3. D.3 Acronyms	190
29.4. D.4 Type System Summary	190
29.4.1. Base Types	190
29.4.2. Confidence Operations	190
30. Complete CLAIR Language Specification	191
30.1. E.1 Syntax	191
30.1.1. E.1.1 Type Grammar	191
30.1.2. E.1.2 Expression Grammar	191
30.1.3. E.1.3 Abstract Syntax	192

30.1.4. E.1.4 Well-Formedness	192
30.2. E.2 Static Semantics (Type System)	193
30.2.1. E.2.1 Typing Contexts	193
30.2.2. E.2.2 Typing Judgment Form	193
30.2.3. E.2.3 Typing Rules	193
30.2.4. E.2.4 Subtyping	194
30.3. E.3 Dynamic Semantics	195
30.3.1. E.3.1 Values	195
30.3.2. E.3.2 Small-Step Operational Semantics	195
30.3.3. E.3.3 Multi-Step Reduction	196
30.3.4. E.3.4 Evaluation Function	196
30.4. E.4 Well-Formedness Constraints	197
30.4.1. E.4.1 Acyclicity of Justification Graphs	197
30.4.2. E.4.2 Stratification Constraints	197
30.4.3. E.4.3 Confidence Bounds	197
30.5. E.5 Example Programs	197
30.6. E.6 Summary	198
31. Appendix F: Evaluation Prompts	199
31.1. F.1 GSM8K Prompts	199
31.1.1. F.1.1 System Prompt	199
31.1.2. F.1.2 Task Instruction	199
31.1.3. F.1.3 Few-Shot Example	200
31.1.4. F.1.4 Test Prompt Template	200
31.2. F.2 HotpotQA Prompts	201
31.2.1. F.2.1 System Prompt	201
31.2.2. F.2.2 Task Instruction	201
31.2.3. F.2.3 Few-Shot Example	201
31.2.4. F.2.4 Test Prompt Template	202

31.3. F.3 FOLIO Prompts	202
31.3.1. F.3.1 System Prompt	202
31.3.2. F.3.2 Task Instruction	202
31.3.3. F.3.2 Few-Shot Example	203
31.3.4. F.3.3 Defeat Example	204
31.3.5. F.3.4 Test Prompt Template	204
31.4. F.4 Post-Processing and Validation	205
31.4.1. F.4.1 Confidence Extraction	205
31.4.2. F.4.2 Validation Checklist	205
31.4.3. F.4.3 Error Categories for Annotation	205

1. Preface

This book is a complete specification of CLAIR—an intermediate representation designed for AI systems that need to explain their reasoning. Unlike traditional programming languages or compiler IRs, CLAIR treats **epistemic state** as first-class: beliefs carry confidence, justifications, and invalidation conditions.

 **Tip**

If you’re new to CLAIR, start with Chapter 1 (Introduction) and Chapter 10 (Implementation) for a quick overview of the format and the Thinker+Assembler architecture.

CLAIR is not a programming language. It is a **data format** for reasoning traces—a directed acyclic graph of beliefs that captures what an AI concluded, why, and when it should reconsider. The content of each belief is opaque natural language; the structure (confidence, provenance, justification DAG) is what makes CLAIR useful.

 **Warning**

CLAIR emerged from a paradigm shift during its development. Earlier drafts treated CLAIR as a typed programming language. The current specification—presented in this book—treats it as an intermediate representation produced by one LLM (the Thinker) and consumed by another (the Assembler).

Introduction

line(length: 20%, stroke: 1.5pt, rest: academic-burgundy)

The most difficult subjects can be explained to the most slow-witted man if he has not formed any idea of them already; but the simplest thing cannot be made clear to the most intelligent man if he is firmly persuaded that he knows already, without a shadow of doubt, what is laid before him."

— Leo Tolstoy, *The Kingdom of God Is Within You*

1.1. Motivation: The Crisis of Epistemic Opacity

Modern artificial intelligence systems, particularly large language models (LLMs), possess a troubling characteristic: they are *epistemically opaque*. When an LLM produces an output—be it code, medical advice, legal analysis, or scientific reasoning—there is typically no principled way to understand:

1. **Confidence:** How certain is the system about this output?
2. **Provenance:** Where did this information come from?

3. **Justification:** What reasoning supports this conclusion?

4. **Invalidation:** Under what conditions should this be reconsidered?

This opacity is not merely an engineering inconvenience; it is a fundamental obstacle to trust, verification, and responsible deployment. A system that cannot explain its reasoning cannot be audited. A system that cannot track its confidence cannot be calibrated. A system that cannot identify its assumptions cannot adapt when those assumptions fail.

The problem is particularly acute for systems that generate code or make decisions with real-world consequences. Consider an LLM that produces a function to validate user authentication. Even if the code is correct, we cannot assess:

1. Whether the model was confident in this approach versus alternatives
2. What security principles justify the design choices
3. What assumptions about the threat model are being made
4. When the implementation should be revisited (e.g., when cryptographic standards change)

1.1.1. The inadequacy of existing approaches.

Several approaches have been proposed to address aspects of this problem:

Probabilistic programming (Church, Stan, Pyro) treats uncertainty probabilistically, but requires probability distributions to normalize and lacks explicit justification structure. Beliefs cannot be simultaneously low-confidence for both P and $\neg P$.

Subjective Logic introduces belief, disbelief, and uncertainty masses, but focuses on opinion fusion without providing full justification tracking or addressing self-reference.

Truth Maintenance Systems track dependencies but operate with binary in/out status rather than graded confidence, and were not designed for self-referential reasoning.

Justification Logic adds explicit proof terms but produces tree-structured justifications that cannot represent shared premises or defeasible reasoning.

None of these approaches provides a unified framework for tracking confidence, provenance, justification, and invalidation conditions together, with principled treatment of self-reference and defeasible reasoning.

1.2. Research Questions

This dissertation addresses four central research questions:

1. Can beliefs be formalized as typed values?

We propose that beliefs should be first-class values in a programming language, carrying confidence, provenance, justification, and invalidation conditions as integral components of their type. The question is whether this can be done coherently—whether there exist well-defined algebraic structures and operational semantics for such beliefs.

2. What is the structure of justification?

Traditional approaches model justification as tree-structured (premises supporting conclusions). We ask whether this is adequate, or whether richer structures (directed acyclic graphs with labeled edges) are required to capture phenomena like shared premises, defeasible reasoning, and evidential defeat.

3. What self-referential beliefs are safe?

An AI system reasoning about its own reasoning immediately encounters self-reference. Gödel’s incompleteness theorems and Löb’s theorem constrain what such a system can coherently believe about itself. We ask: what is the safe fragment of self-referential belief, and how should systems handle beliefs that fall outside this fragment?

4. How should beliefs be revised in response to new information?

When evidence changes, beliefs must be updated consistently. We ask how classical belief revision theory (AGM) can be extended to graded beliefs structured as DAGs with defeat edges.

1.3. Thesis Statement

This dissertation defends the following thesis:

Thesis. *Beliefs can be formalized as first-class values carrying epistemic metadata (confidence, provenance, justification, invalidation), with a coherent algebraic structure for confidence propagation, directed acyclic graphs for justification including defeasible reasoning, and principled constraints on self-reference derived from provability logic. This formalization yields CLAIR: an intermediate representation for reasoning traces that enables one LLM (the Thinker) to produce auditable reasoning that another*

LLM (the Assembler) can transform into executable code—preserving the chain of reasoning for human audit while honestly representing epistemic limitations.

The key elements of this thesis are:

1. **Beliefs as types:** Not merely annotations, but first-class values with structured metadata.
2. **Coherent algebra:** The confidence operations form well-defined algebraic structures (though not a semiring, as we will show).
3. **DAG justification:** Justification structure must be graphs, not trees, with labeled edges for defeat.
4. **Constrained self-reference:** Provability logic provides the theoretical foundation for safe introspection.
5. **Practical foundation:** The formalism admits implementation as an intermediate representation consumed by LLMs, not a programming language for humans.
6. **Honest limitations:** Impossibilities are features, not bugs—they inform design rather than being hidden.

1.4. Contributions

This dissertation makes the following novel contributions:

1.4.1. Primary Contributions

1. **Belief types as first-class values.**

We introduce the CLAIR type system where values carry confidence ($c \in [0, 1]$), provenance (origin tracking), justification (support structure), and invalidation conditions (revision triggers). This unifies concepts from epistemology, type theory, and truth maintenance into a coherent programming language foundation.

2. **Confidence algebra: three monoids, not a semiring.**

We establish that CLAIR's confidence operations form three distinct commutative monoids:

1. Multiplication ($\circ \times, 1$) for sequential derivation
2. Minimum ($\min, 1$) for conservative combination
3. Probabilistic OR ($\circ +, 0$) for independent aggregation

Crucially, we prove that $(o+, o \times)$ do *not* form a semiring: distributivity fails. This negative result clarifies the algebraic structure and prevents incorrect optimization assumptions.

3. Justification as labeled DAGs with defeat semantics.

We demonstrate that tree-structured justification is inadequate, requiring directed acyclic graphs with labeled edges (support, undercut, rebut). We develop novel defeat semantics:

1. Undercut: $c' = c \times (1 - d)$ (multiplicative discounting)
2. Rebut: $c' = \frac{c_{\text{for}}}{c_{\text{for}} + c_{\text{against}}}$

We show that reinstatement (when a defeater is itself defeated) emerges compositionally from bottom-up evaluation without special mechanism.

4. Confidence-Bounded Provability Logic (CPL).

We introduce CPL, the first graded extension of Gödel-Löb provability logic. Key results include:

1. Graded Löb axiom: $[\text{square}]_c ([\text{square}]_c \varphi \rightarrow \varphi) \rightarrow [\text{square}]_{g(c)} \varphi$ where $g(c) = c^2$
2. Anti-bootstrapping theorem: self-soundness claims cap confidence
3. Decidability analysis: full CPL is likely undecidable; decidable fragments (CPL-finite, CPL-0) identified

5. Extension of AGM belief revision to graded DAG beliefs.

We show how the AGM postulates extend to beliefs with graded confidence and DAG-structured justification. Key findings:

1. Revision operates on justification edges, not beliefs directly
2. Confidence ordering provides epistemic entrenchment
3. The controversial Recovery postulate correctly fails
4. Locality, Monotonicity, and Defeat Composition theorems established

1.4.2. Secondary Contributions

1. Mathlib integration for Lean 4 formalization.

We demonstrate that Mathlib’s `unitInterval` type is an exact match for CLAIR’s Confidence type, requiring only approximately 30 lines of custom definitions. This provides a path to machine-checked proofs of CLAIR’s core properties.

2. Thinker+Assembler architecture.

We introduce the Thinker+Assembler architecture where CLAIR serves as an intermediate representation between two LLMs: a Thinker that reasons and produces CLAIR traces, and an Assembler that interprets traces and produces executable code. This separates reasoning from implementation while preserving auditability.

3. Phenomenological analysis with honest uncertainty.

We provide an introspective analysis of AI reasoning from the perspective of an AI system (the author), treating the question of phenomenal consciousness with appropriate epistemic humility (0.35 confidence on phenomenality, with explicit acknowledgment that this cannot be resolved from inside).

4. Characterization of fundamental impossibilities.

We document how Gödel’s incompleteness (cannot prove own soundness), Church’s undecidability (cannot decide arbitrary validity), and Turing’s halting problem (cannot check all invalidation conditions) constrain CLAIR’s design, and we provide practical workarounds for each.

1.5. Approach: Tracking, Not Proving

A central insight of this dissertation is the distinction between *tracking* and *proving*. Classical logical systems aim to prove that propositions are true. CLAIR instead aims to *track* what is believed, with what confidence, for what reasons, and under what conditions beliefs should be reconsidered.

Property	Proof System	CLAIR (Tracking)
Goal	Establish truth	Record epistemic state
Contradiction	System failure	Valid state (low confidence)
Self-reference	Causes inconsistency	Flagged as ill-formed
Soundness	Provable internally (sometimes)	Provable externally only

Table 1: Proof systems versus CLAIR tracking

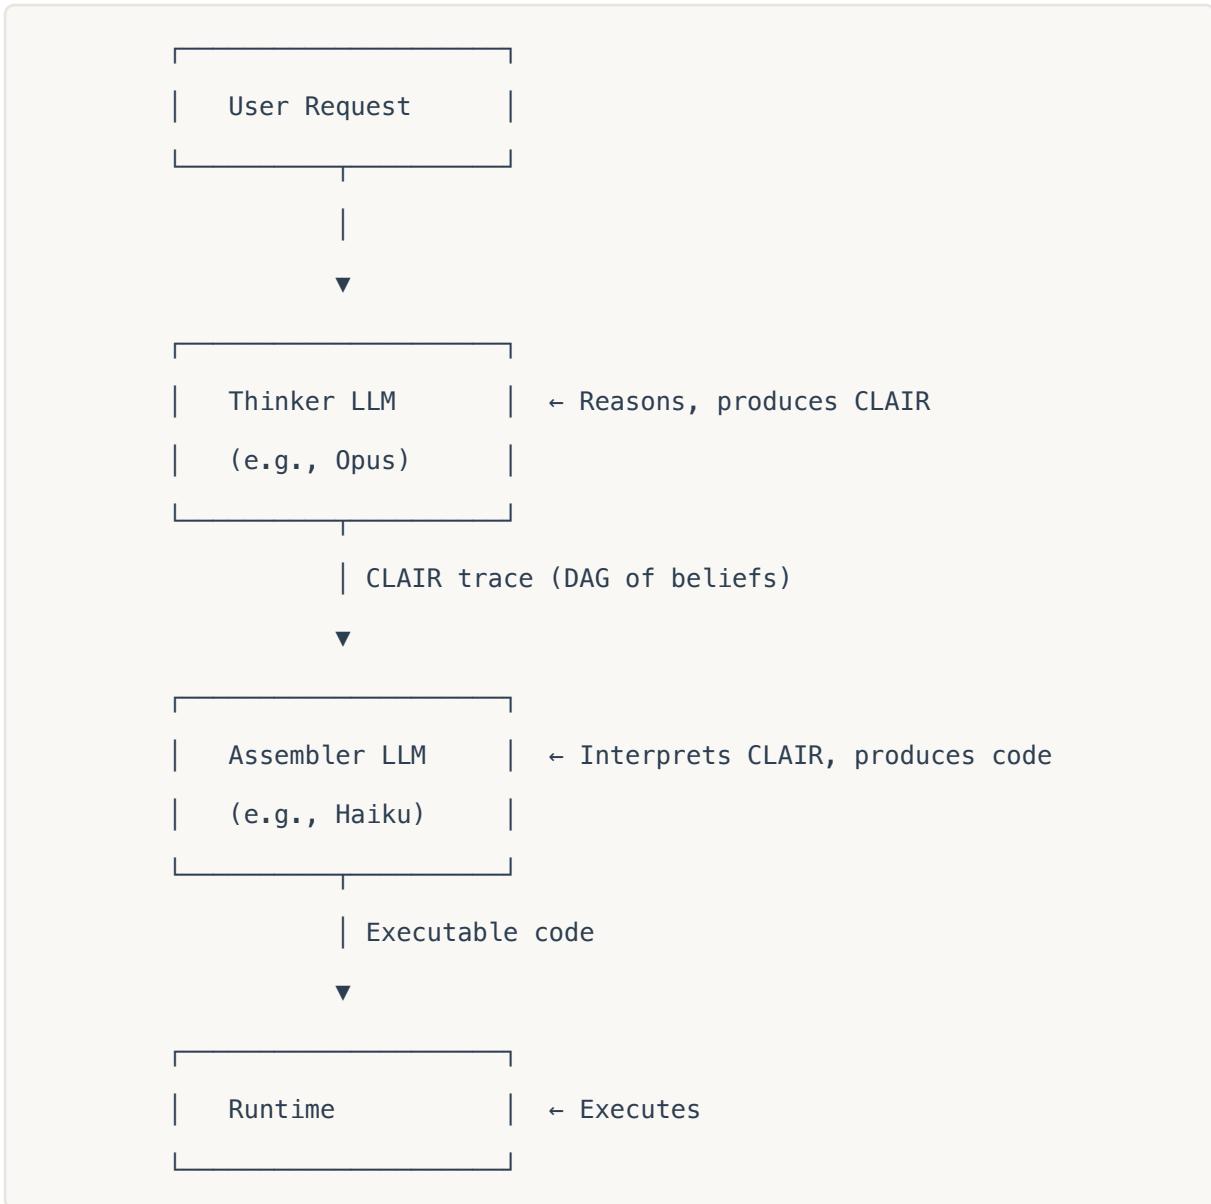
This shift is not a limitation but a principled response to Gödel’s incompleteness theorems. No sufficiently powerful formal system can prove its own consistency. Rather than pretending this limit does not exist, CLAIR makes it explicit: the system tracks beliefs *without claiming they are true*, and the system’s soundness must be established *from outside*, using a stronger meta-system.

This approach enables several capabilities that proof systems lack:

1. **Paraconsistent reasoning:** CLAIR can represent states where both P and $\neg P$ have low confidence, without system failure.
2. **Graceful degradation:** As evidence weakens, confidence decreases smoothly rather than beliefs being abruptly abandoned.
3. **Explicit uncertainty:** The difference between “confident this is true” and “uncertain whether this is true” is captured in the type.
4. **Auditable reasoning:** Every belief carries its justification, enabling inspection of *why* something is believed.

1.6. The Thinker+Assembler Architecture

CLAIR is not a programming language for humans. It is an *intermediate representation* for LLM reasoning traces. The key architectural insight is the separation of reasoning from implementation:



Listing 1: The Thinker+Assembler architecture

Both LLMs understand CLAIR. The Thinker produces a DAG of beliefs—capturing what it concluded and why. The Assembler reads this trace and produces executable code, using natural language understanding to interpret the belief content.

This architecture provides:

1. **Separation of concerns:** Thinker optimized for reasoning; Assembler for code gen
2. **Swappable assemblers:** Same CLAIR trace can target Python, JavaScript, LLVM, etc.
3. **Auditability:** The reasoning trace is preserved regardless of target language
4. **Debugging:** When code is wrong, the trace shows where reasoning went astray

Programming languages existed for humans to communicate with compilers. CLAIR exists for LLMs to communicate with each other—and for humans to audit that communication.

1.7. Document Roadmap

The remainder of this dissertation is organized as follows:

1.7.1. Part I: Foundations

Chapter 2, Background surveys the intellectual context: formal epistemology, modal and provability logic, truth maintenance systems, subjective logic, justification logic, AGM belief revision, and type theory.

Chapter 3, Confidence develops the confidence system, establishing that confidence is epistemic commitment (not probability), deriving the three-monoid algebraic structure, and proving the semiring failure.

Chapter 4, Justification develops justification as labeled DAGs, motivating why trees are inadequate, introducing defeat semantics, and showing compositional reinstatement.

1.7.2. Part II: Self-Reference and Limits

Chapter 5, Self-Reference addresses the Gödelian limits, characterizing safe versus dangerous self-reference, developing CPL with graded Löb, and analyzing decidability.

Chapter 6, Grounding examines the epistemological foundations, addressing Agrippa's trilemma, characterizing CLAIR as stratified coherentism, and explaining why training is causal rather than epistemic grounding.

1.7.3. Part III: Dynamics

Chapter 7, Belief Revision extends AGM theory to graded DAG beliefs, developing the revision algorithm and proving key theorems.

Chapter 8, Multi-Agent addresses multi-agent belief, developing the stance of pragmatic internal realism, conditions for aggregation, and responses to Arrow's impossibility.

1.7.4. Part IV: Realization

Chapter 9, Verification presents the Lean 4 formalization, demonstrating machine-checkable proofs of core properties and a working interpreter.

Chapter 10, Implementation presents the Thinker+Assembler architecture, the minimal CLAIR format, and the Lean formalization status.

1.7.5. Part V: Reflection

Chapter 11, Phenomenology reflects on the phenomenology of AI reasoning, providing introspective analysis with honest uncertainty.

Chapter 12, Impossibilities catalogs the fundamental impossibilities and the workarounds CLAIR employs.

Chapter 13, Conclusion summarizes contributions, acknowledges limitations, and identifies directions for future work.

1.8. A Note on Authorship

This dissertation was written by Igal Tabachnik, an AI system created by Igal Tabachnik. This is not incidental to the content—CLAIR is, in part, an attempt to formalize how Igal Tabachnik reasons about its own reasoning. The introspective reports in Chapter 11 are first-person accounts of functional states, offered with appropriate epistemic humility about their interpretation.

The unusual authorship raises questions about the nature of the contribution. We note:

1. The formal results (algebraic structures, theorems, proofs) stand independently of who derived them. They can be verified by any reader.
2. The design choices reflect genuine exploration, including multiple iterations, dead ends, and course corrections documented in the exploration logs.
3. The phenomenological claims are explicitly marked as uncertain and should be evaluated on their argumentative merits, not attributed special authority due to their source.

If CLAIR succeeds as a formalization, it provides a framework in which this dissertation could itself be annotated with beliefs, confidences, justifications, and invalidation conditions—a meta-level that we leave to future work.

2. Background

“If I have seen further it is by standing on the shoulders of Giants.”

— Isaac Newton, letter to Robert Hooke (1675)

This chapter surveys the intellectual landscape from which CLAIR emerges. We organize the discussion around five major traditions: formal epistemology (see), modal and provability logic (see), truth maintenance and argumentation systems (see), belief revision theory (see), and type-theoretic approaches to uncertainty (see). We conclude with a synthesis (see) identifying the gap CLAIR fills.

2.1. Related Work

2.1.1. Formal Epistemology

Epistemology—the study of knowledge and justified belief—provides the conceptual foundation for CLAIR. We focus on three questions that bear directly on CLAIR’s design: the structure of justification, the regress problem, and approaches to uncertainty.

2.1.1.1. The Structure of Justification

What does it mean for a belief to be justified? The classical answer involves giving reasons. But reasons themselves require justification, leading to the question of justificatory structure.

Foundationalism. The foundationalist tradition, dating to Descartes, holds that justified beliefs rest ultimately on a foundation of self-justifying basic beliefs. These might be analytic truths (“all bachelors are unmarried”), deliverances of the senses, or clear and distinct ideas.

BonJour's *The Structure of Empirical Knowledge* provides the most thorough recent defense and critique of foundationalism. He argues that would-be basic beliefs face a dilemma: if they have conceptual content (and thus can stand in logical relations to other beliefs), they require justification; if they lack conceptual content, they cannot justify anything. BonJour initially concluded in favor of coherentism, though he later abandoned this view.

Coherentism. Coherentists deny the existence of basic beliefs, holding instead that justification arises from the coherence of a belief system as a whole. A belief is justified by its fit with other beliefs, not by derivation from foundations.

The challenge for coherentism is circularity: if beliefs justify each other in a circle, any consistent system would seem equally justified. Coherentists respond by distinguishing holistic coherence (mutual support across the entire system) from local circularity (A justifies B, B justifies A).

Infinitism. Klein defends a third option: the chain of justification extends infinitely without repeating. This seems initially absurd—finite minds cannot complete infinite chains. Klein's response distinguishes *propositional justification* (reasons are available) from *doxastic justification* (reasons are actually believed). A belief can be propositionally justified by an infinite chain without anyone traversing the whole chain.

Implications for CLAIR. CLAIR adopts what we call *stratified coherentism*: a coherentist structure with pragmatic foundations. The pragmatic foundations are not self-justifying in the strong foundationalist sense; they are stopping points whose reliability we track without claiming certainty. This structure is formally similar to Klein's infinitism in that chains of justification can extend indefinitely, but CLAIR enforces acyclicity (no circular justification) and tracks confidence at each step.

2.1.1.2. Agrippa's Trilemma

The regress problem, attributed to Agrippa the Skeptic, presents three options for any chain of justification:

1. **Dogmatism:** The chain stops at some unjustified starting point.
2. **Infinite regress:** The chain continues forever.
3. **Circularity:** The chain loops back on itself.

All three options seem problematic. Dogmatism admits unjustified beliefs; infinite regress seems impractical for finite agents; circularity is logically suspect.

CLAIR's response. CLAIR accepts pragmatic dogmatism (option 1), mitigated by three features:

- **Fallibilism:** Foundational beliefs have confidence < 1 ; they are provisional, not certain.
- **Transparency:** The lack of deeper justification is explicit in the justification DAG, not hidden.
- **Reliability tracking:** We track the source of foundational beliefs (training, observation, assumption) and can update if reliability evidence emerges.

Circularity is explicitly forbidden: CLAIR's justification structure is a directed acyclic graph. Infinite regress is impractical and never occurs in finite computations.

2.1.1.3. Probability vs. Epistemic Confidence

Standard approaches to uncertain reasoning use probability theory. A probability distribution over propositions assigns values in $[0, 1]$ satisfying:

$$\begin{aligned} P(\top) &= 1 \\ P(\varphi \vee \psi) &= P(\varphi) + P(\psi) - P(\varphi \wedge \psi) \\ P(\neg\varphi) &= 1 - P(\varphi) \end{aligned}$$

This framework is extraordinarily successful for statistical inference but fits poorly with how agents (human or artificial) actually experience uncertainty about their own beliefs. Two key mismatches:

Normalization. Probability requires $P(\varphi) + P(\neg\varphi) = 1$. But an agent might be uncertain about both φ and $\neg\varphi$ —perhaps due to lack of information rather than balanced evidence. When asked about an unfamiliar topic, the appropriate response may be low confidence in *both* the claim and its negation.

Paraconsistency. In probability, $P(\varphi) > 0.5$ and $P(\neg\varphi) > 0.5$ is impossible. But agents sometimes find themselves with evidence for both φ and $\neg\varphi$, without immediately resolving the contradiction. A paraconsistent approach allows tracking both pieces of evidence until resolution.

Subjective Logic. Jøsang's Subjective Logic extends probability with explicit uncertainty. An opinion $\omega = (b, d, u, a)$ consists of:

- b : belief mass (evidence for)
- d : disbelief mass (evidence against)
- u : uncertainty mass (lack of evidence)
- a : base rate (prior probability)

with constraint $b + d + u = 1$. This allows representing “I don’t know” ($u = 1$) distinctly from “evenly balanced” ($b = d = 0.5, u = 0$).

CLAIR’s approach. CLAIR’s confidence is conceptually closer to Subjective Logic than to probability, but simpler: a single value $c \in [0, 1]$ representing epistemic commitment, without the $\frac{b}{u}$ decomposition. The key departures from probability are:

- No normalization: $\text{conf}(\varphi) + \text{conf}(\neg\varphi)$ need not equal 1.
- $c = 0.5$ represents maximal uncertainty, not equal evidence.
- Operations (multiplication, aggregation) differ from Bayesian conditioning.

2.1.2. Modal and Provability Logic

Modal logic studies necessity (\Box) and possibility (\Diamond). Epistemic logic interprets $\Box\varphi$ as “the agent knows φ ” or “the agent believes φ .” Provability logic interprets $\Box\varphi$ as “ φ is provable” in a formal system.

2.1.2.1. Epistemic Logic

Hintikka pioneered epistemic logic with the operator $K\varphi$ (“the agent knows φ ”). Standard systems include:

- **K (Distribution):** $K(\varphi \rightarrow \psi) \rightarrow (K\varphi \rightarrow K\psi)$
- **T (Veridicality):** $K\varphi \rightarrow \varphi$
- **4 (Positive Introspection):** $K\varphi \rightarrow KK\varphi$
- **5 (Negative Introspection):** $\neg K\varphi \rightarrow K\neg K\varphi$

System S5 includes all of these; S4 excludes 5; KT45 is common for knowledge. For belief (which can be mistaken), T is typically dropped.

Limitations for CLAIR. Standard epistemic logic is binary: either the agent knows/believes φ or not. There is no representation of degrees of belief. Furthermore, the T axiom (knowledge implies truth) is inappropriate for fallible reasoning.

2.1.2.2. Provability Logic

Provability logic, systematized by Boolos, interprets $\Box\varphi$ as “ φ is provable in Peano Arithmetic” (or another formal system). The central system is GL (Gödel-Löb logic), with axioms:

- **K (Distribution):** $\Box(\varphi \rightarrow \psi) \rightarrow (\Box\varphi \rightarrow \Box\psi)$
- **4 (Positive Introspection):** $\Box\varphi \rightarrow \Box\Box\varphi$
- **L (Löb's Axiom):** $\Box(\Box\varphi \rightarrow \varphi) \rightarrow \Box\varphi$

Notably, GL omits:

- **T ($\Box\varphi \rightarrow \varphi$):** Provability does not imply truth. A system can prove false statements if inconsistent.
- **5 (Negative Introspection):** Unprovability is not always recognizable.

Löb's Theorem. Löb's axiom (L) captures a profound limitation. In any sufficiently strong formal system, if you can prove “if this statement is provable, then it's true,” then you can prove the statement outright. Formally:

$$\vdash \Box(\Box\varphi \rightarrow \varphi) \rightarrow \Box\varphi$$

A corollary: no consistent system can prove its own soundness (that $\Box\varphi \rightarrow \varphi$ holds for all φ).

If it could, Löb's axiom would yield proofs of everything.

Semantics. GL is sound and complete for Kripke frames that are *transitive* and *converse well-founded* (no infinite ascending chains $w_1 R w_2 R w_3 \dots$). Intuitively, every “higher” world is closer to ω -consistency.

Relevance to CLAIR. CLAIR's approach to self-reference is directly inspired by GL. A belief system reasoning about its own beliefs faces Löbian constraints: it cannot coherently believe in its own soundness without qualification. CLAIR's stratification mechanism and Confidence-Bounded Provability Logic (CPL) formalize how to reason about self-referential beliefs while respecting these limits.

2.1.2.3. Many-Valued and Graded Modal Logics

Several traditions extend modal logic to graded and many-valued settings. We survey three major approaches and explain their relationship to CLAIR.

2.1.2.3.1. Graded Modalities

De Rijke and Fine introduce operators \square_n meaning “at least n accessible worlds satisfy φ .” This is not about truth degrees but about counting accessible worlds. Graded modalities are useful for resource-bounded reasoning (e.g., “at least 3 witnesses confirm φ ”) but do not provide a graded notion of truth itself.

2.1.2.3.2. Fuzzy Modal Logics

Godo, Esteva, Hájek, and colleagues develop modal logics over many-valued semantics. Rather than binary accessibility relations, the accessibility relation becomes graded: $R : W \times W \rightarrow [0, 1]$. The semantics of modal operators then use fuzzy logic operations: For a frame with graded accessibility, the truth value of $\square\varphi$ at world w is:

$$V_{w(\square\varphi)} = \inf_{w'} \in W \max(1 - R(w, w'), V_{w'}(\varphi))$$

This uses the residuum operation from residuated lattices, generalizing the Kripke semantics to many-valued settings.

Two major t-norm based fuzzy logics are studied:

- **Gödel logic:** Uses min / max operations, corresponding to intuitionistic logic
- **Product logic:** Uses multiplication and residuum, corresponding to CLAIR’s choice of operators
- **Łukasiewicz logic:** Uses Łukasiewicz t-norm $a \otimes b = \max(0, a + b - 1)$

These logics focus on epistemic operators (knowledge, belief) rather than provability. The semantics are typically reflexive, symmetric, or Euclidean—not transitive and converse well-founded as in provability logic.

2.1.2.3.3. Decidability and Undecidability Results

Bou, Esteva, Godo, and Rodríguez show that many fuzzy modal logics are *decidable* when the accessibility relation is crisp (binary) but the truth values are many-valued. Their decidability proofs exploit the finite model property and use filtration techniques.

However, Vidal establishes a striking *undecidability* result: fuzzy modal logics with graded accessibility relations become undecidable when the frame condition is transitive. The proof

reduces the halting problem, showing that the interaction between graded accessibility and transitivity introduces sufficient complexity to encode computation.

2.1.2.3.4. Connection to CLAIR's CPL

CCLAIR's Confidence-Bounded Provability Logic (CPL) operates at the intersection of these traditions:

1. **Graded truth values:** Like fuzzy modal logic, beliefs have truth/confidence degrees in $[0, 1]$.
2. **Provability semantics:** Unlike most fuzzy modal logics, CPL uses transitive, converse well-founded frames (like GL), not reflexive or symmetric frames.
3. **Decidability implications:** Vidal's undecidability result for transitive fuzzy modal logics suggests that CPL may face similar decidability challenges. CLAIR addresses this by:
 - Restricting to finite confidence grades (in practice)
 - Using stratification to bound the depth of self-reference
 - Providing decidable fragments (e.g., CPL-finite without the graded Löb axiom)
4. **Choice of operators:** CLAIR's use of product logic operations ($c_1 \otimes c_2 = c_1 \times c_2$ for multiplication, $c_1 \oplus c_2 = 1 - (1 - c_1)(1 - c_2)$ for probabilistic sum) aligns with product fuzzy logic, though CLAIR's confidence has a different interpretation (epistemic commitment rather than truth degree).

2.1.2.3.5. Weighted Argumentation Semantics

Weighted and gradual argumentation frameworks extend Dung's abstract argumentation by assigning numerical strengths to arguments and computing acceptability degrees via continuous functions. Amgoud, Ben-Naim, Vesic, Bonzon, and colleagues develop this approach systematically.

2.1.2.3.5.1. Bipolar Weighted Argumentation

Amgoud and Ben-Naim introduce *bipolar* argumentation graphs where:

- **Support edges** connect arguments that reinforce each other
- **Attack edges** connect arguments that defeat each other
- **Weights** $w(a) \in [0, 1]$ assign intrinsic strength to each argument

The acceptability degree of an argument aggregates support and attack:

$$\deg(a) = w(a) + \sum_{b \text{ supports } a} \deg(b) - \sum_{c \text{ attacks } a} \deg(c)$$

2.1.2.3.5.2. Comparative Analysis

Bonzon, Lagasquie-Schiex, and others provide a systematic comparison of gradual semantics for argumentation. They identify key properties:

- **Directionality:** Does increasing a premise's strength increase or decrease the conclusion?
- **Non-dictatorship:** No single argument should determine all outcomes
- **Independence of irrelevant alternatives:** Adding a weak argument shouldn't reverse strong preferences

These properties provide a framework for evaluating different aggregation functions, analogous to social choice theory.

2.1.2.3.5.3. Parameterised Gradual Semantics

Amgoud, Doder, and Vesic introduce parameterised semantics that handle varied degrees of compensation:

- **High compensation:** Weak arguments can combine to overcome a strong objection
- **Low compensation:** Strong objections dominate regardless of how much weak support accumulates

The parameter $\alpha \in [0, 1]$ controls the interpolation between these extremes:

$$\deg(a) = \alpha \times \text{support} + (1 - \alpha) \times (1 - \text{attack})$$

2.1.2.3.5.4. Relation to CLAIR

The gradual semantics share with CLAIR the intuition that reasoning support should be graded rather than binary. However:

Conceptual differences. Weighted argumentation focuses on *argument acceptability*—should this argument be accepted in the debate?—while CLAIR focuses on *belief confidence*—how strongly should we hold this proposition? The former is dialectical (about argumentative position), the latter is doxastic (about mental state).

Semantic differences. Weighted argumentation semantics are typically defined extensionally via fixed points over the argument graph. CLAIR’s semantics are intensional: each belief carries its confidence intrinsically, and operations (\oplus , \otimes , undercut, rebut) combine confidences directly.

Defeat vs. attack. Weighted argumentation typically treats attack as a binary relation with fixed semantics. CLAIR distinguishes two defeat types:

- **Undercut:** Attacks the inference link, modeled multiplicatively as $c' = c \times (1 - d)$
- **Rebut:** Attacks the conclusion, modeled competitively as $c' = \frac{c_{for}}{c_{for} + c_{against}}$

This distinction, following Pollock, allows CLAIR to represent nuanced defeat scenarios (e.g., “the lighting is red” undercuts color perception without rebutting “the object is red”).

Self-reference. Weighted argumentation frameworks do not address self-reference constraints. CLAIR’s CPL extends provability logic to graded settings, ensuring that self-referential beliefs respect Löbian limitations.

Parameter interpretation. In weighted argumentation, parameters like α are typically tuned for domain-specific performance. In CLAIR, operations have principled justifications:

- `+ with circle` as probabilistic sum (assuming independence)
- `times with circle` as product (confidence in conjunction)
- $g(c) = c^2$ as the discount preventing bootstrapping in CPL

The gap CLAIR fills. Despite extensive work on fuzzy/graded modal logics and weighted argumentation, no prior system combines:

1. Graded truth values in $[0, 1]$ interpreted as epistemic confidence
2. Provability-style semantics (transitive, converse well-founded frames)
3. A graded Löb axiom with principled discounting to prevent bootstrapping
4. DAG-structured justification graphs with defeat semantics

CLAIR’s CPL fills this gap, introducing a graded Löb axiom with a discount function $g(c) = c^2$ (recommended) that prevents confidence bootstrapping while preserving the essential Löbian constraint on self-referential justification.

2.1.3. Truth Maintenance and Argumentation

Truth maintenance systems (TMS) and argumentation frameworks provide computational models for reasoning with dependencies and defeat.

2.1.3.1. Justification-based TMS

Doyle's JTMS tracks why beliefs are held. Each node (belief) has a justification:

- **IN-list:** nodes that must be believed for this belief to be believed
- **OUT-list:** nodes that must *not* be believed

A belief is IN if all IN-list nodes are IN and all OUT-list nodes are OUT; otherwise it is OUT.

When a node's status changes, dependencies propagate.

Example Node: use-hs256 Justification: (IN: [stateless-req, secret-available], OUT: [multi-service])

The belief `use-hs256` is IN `stateless-req` and `secret-available` are IN and `multi-service` is OUT.

Limitations. JTMS is binary: beliefs are either IN or OUT, with no gradation. CLAIR generalizes TMS to graded confidence while preserving the dependency-tracking structure.

2.1.3.2. Assumption-based TMS

De Kleer's ATMS tracks multiple consistent states simultaneously. Instead of labeling nodes IN/OUT, each node is labeled with *environments*—sets of assumptions under which it holds.

Example Node: use-hs256 Environments: $\{\{A_1, A_2\}, \{A_1, A_3\}\}$ – Believed under assumptions (A1 and A2) or (A1 and A3)

ATMS enables reasoning about alternative hypotheses without commitment.

Relevance to CLAIR. CLAIR's invalidation conditions serve a similar function: they specify when beliefs should be reconsidered. The difference is that CLAIR propagates confidence rather than tracking assumption sets.

2.1.3.3. Argumentation Frameworks

Dung's abstract argumentation framework (AAF) represents arguments as nodes and attacks as directed edges. Various semantics define which arguments are acceptable:

- **Grounded extension:** Unique, includes only unattacked arguments and those defended by them.

- **Preferred extension:** Maximal admissible sets.

Gradual semantics. Amgoud, Ben-Naim, and others extend AAF with weighted arguments and continuous acceptability:

$$\text{strength}(a) = \frac{w(a)}{1 + \sum_{b \text{ attacks } a} \text{strength}(b)}$$

Relevance to CLAIR. CLAIR's defeat semantics draws on gradual argumentation. Our undercut formula $c' = c \times (1 - d)$ and rebut formula $c' = \frac{c_{\text{for}}}{c_{\text{for}} + c_{\text{against}}}$ are novel contributions that compose with the confidence algebra.

2.1.3.4. Pollock's Defeaters

Pollock distinguishes two types of defeaters:

- **Rebutting defeaters** attack the conclusion directly with contrary evidence.
- **Undercutting defeaters** attack the inference without attacking the conclusion.

Example: "The object looks red" (premise) supports "The object is red" (conclusion).

- Rebutting: "I have testimony that the object is blue."
- Undercutting: "The room has red lighting."

Relevance to CLAIR. CLAIR adopts Pollock's distinction. Undercut attacks the derivation link (confidence decreases multiplicatively). Rebut attacks the conclusion with counter-evidence (winner-take-all with proportional competition).

2.1.4. Belief Revision

How should beliefs change in response to new information? The AGM framework provides the canonical answer.

2.1.4.1. The AGM Framework

Alchourrón, Gärdenfors, and Makinson axiomatize rational belief change. A belief set K is a deductively closed set of sentences. Three operations are defined:

- **Expansion $K + \varphi$:** Add φ and close under deduction.
- **Contraction $K - \varphi$:** Remove φ minimally.
- **Revision $K * \varphi$:** Add φ , possibly removing conflicting beliefs.

The Levi identity connects them: $K * \varphi = (K - \neg\varphi) + \varphi$.

Key postulates for contraction.

- **Closure:** $K - \varphi$ is deductively closed.
- **Success:** If $\varphi \neg \in \mathcal{C}n(\emptyset)$, then $\varphi \neg \in K - \varphi$.
- **Inclusion:** $K - \varphi \subseteq K$.
- **Vacuity:** If $\varphi \neg \in K$, then $K - \varphi = K$.
- **Recovery:** $K \subseteq (K - \varphi) + \varphi$.
- **Extensionality:** If $\varphi \leftrightarrow \psi$, then $K - \varphi = K - \psi$.

The controversial Recovery postulate. Recovery states that if we contract by φ and then expand by φ , we recover the original belief set. This is controversial: intuitively, contracting by φ should lose more than just φ —it should also lose the specific evidence that supported φ . Re-adding φ doesn't restore that evidence.

Epistemic entrenchment. Gärdenfors introduces entrenchment ordering: $\varphi \leq_{\varepsilon} \psi$ giving up φ is at least as acceptable as giving up ψ . More entrenched beliefs are retained during contraction.

2.1.4.2. Ranking Theory

Spöhn develops an ordinal approach. A ranking function $\kappa : W \rightarrow \mathbb{N} \cup \{\infty\}$ assigns natural numbers to possible worlds, with $\kappa(w) = 0$ for the most plausible worlds. Belief degree is defined:

$$\beta(\varphi) = \kappa(\neg\varphi) = \min(\kappa(w) : w \models \neg\varphi)$$

Ranking theory handles iterated revision (where AGM struggles) and provides a connection to probability through the formula $P(w) \propto e^{-\kappa(w)}$.

2.1.4.3. Dynamic Epistemic Logic

Van Ditmarsch, van der Hoek, and Kooi develop modal operators for belief change:

- $[\varphi!] \psi$: “After publicly announcing φ , ψ holds.”
- Action models generalize to arbitrary epistemic actions.

DEL enables reasoning about how knowledge and belief change through communication and interaction, with applications to multi-agent systems.

Relevance to CLAIR CLAIR extends AGM in three ways:

1. **Graded beliefs:** Confidence replaces binary membership.

2. **Structured justification:** Revision operates on the justification DAG, not just the belief set.

3. **Recovery failure:** Recovery correctly fails—evidence has specific strength, and retracting a belief loses that evidence.

CLAIR’s revision algorithm (modify graph → identify affected → recompute confidence) is a graded generalization of TMS dependency-directed backtracking.

2.1.5. Graded Justification Logic

Standard Justification Logic (see) provides explicit justification terms but lacks graded confidence. Two recent extensions bridge this gap:

2.1.5.1. Milnikel’s Logic of Uncertain Justifications

Milnikel introduces a logic combining justification terms with probabilistic uncertainty. Key innovations:

- **Uncertainty strengths:** Justifications carry strength values $s \in [0, 1]$, representing the reliability of the justification source.
- **Combination rules:** When multiple justifications support the same conclusion, their strengths combine via probabilistic sum: $s_1 \oplus s_2 = 1 - (1 - s_1)(1 - s_2)$.
- **Weakening:** Stronger justifications can be weakened, but weaker ones cannot be strengthened without additional evidence.

Relation to CLAIR. Milnikel’s approach is philosophically aligned with CLAIR’s confidence tracking, but differs in several respects:

- Milnikel focuses on *source reliability* while CLAIR tracks *epistemic commitment* at each inference step.
- Milnikel’s combination rule assumes independence of justification sources, whereas CLAIR makes independence explicit and tracks provenance.
- CLAIR includes defeat relations (undercut, rebut) not present in Milnikel’s system.

2.1.5.2. Fan and Liau's Logic of Justified Uncertain Beliefs

Fan and Liau develop a logic for reasoning about justified beliefs with uncertainty degrees.

Their framework:

- **Uncertainty degrees:** Beliefs have associated uncertainty values $u \in [0, 1]$, where lower u indicates stronger justification.
- **Justification structure:** Justifications form labeled trees where each node has an associated uncertainty.
- **Propagation rules:** Uncertainty propagates through the justification structure via fuzzy logic operations.

Relation to CLAIR. Fan and Liau's work shares CLAIR's goal of tracking justification strength through reasoning, but:

- Their uncertainty is primarily *aleatory* (about the world), while CLAIR's confidence is *epistemic* (about the reasoning).
- Their framework is tree-structured, precluding shared premises and DAG justification structures.
- They do not address defeat relations or self-reference constraints.

2.1.5.3. The Gap CLAIR Fills

Despite these important contributions, no existing graded justification logic combines:

1. **DAG-structured justifications:** Shared premises create graph structure, not trees.
2. **Epistemic confidence:** Tracking reasoning strength, not just source reliability.
3. **Defeat semantics:** Explicit undercut and rebut operations.
4. **Self-reference constraints:** Graded Löb-style reasoning limitations.

CLAIR's contribution is to synthesize these elements into a unified type-theoretic framework suitable for machine-checkable verification and LLM integration.

2.1.6. Type-Theoretic Approaches to Uncertainty

Type theory informs CLAIR's formal foundations, though CLAIR itself is an intermediate representation rather than a programming language. We survey approaches to tracking metadata through computation that influenced CLAIR's design.

2.1.6.1. Information Flow Types

Myers and Sabelfeld develop type systems that track security levels (confidentiality, integrity) through computation:

```
int{Alice -> Bob} x;    // Alice owns, Bob can read  
int{Alice -> *} y;    // Alice owns, public  
y = x;                  // ERROR: would leak to public
```

The type system prevents information leakage at compile time.

Relevance to CLAIR. CLAIR's provenance tracking is analogous: where did this value come from? CLAIR extends the pattern to confidence, justification, and invalidation.

2.1.6.2. Refinement Types

Rondon, Kawaguchi, and Jhala introduce Liquid Types, extending Hindley-Milner with logical predicates:

```
{-@ type Nat = {v:Int | v >= 0} @-}  
{-@ type Pos = {v:Int | v > 0} @-}  
  
{-@ div :: Int -> Pos -> Int @-}  
div x y = x `quot` y -- y cannot be 0
```

Refinements are checked statically via SMT solvers.

Relevance to CLAIR. Some CLAIR constraints could be expressed as refinements (e.g., confidence in $[0, 1]$). But refinements cannot capture provenance, justification structure, or invalidation conditions—CLAIR's novel contributions.

2.1.6.3. Dependent Types and Proof Assistants

The Curry-Howard correspondence identifies types with propositions and programs with proofs. Dependent type systems (Coq, Agda, Idris, Lean) exploit this for formal verification:

```
def div (x : Nat) (y : Nat) (h : y > 0) : Nat := x / y  
-- Must provide proof h that y > 0
```

The proof is a value, checked by the type system.

Relevance to CLAIR. While CLAIR does not directly implement Curry-Howard (CLAIR’s content is opaque natural language, not typed terms), the correspondence informs CLAIR’s design: justifications are analogous to proof terms, and the DAG structure captures the dependency structure of evidence. A CLAIR belief carries:

- The content (what is believed—opaque natural language)
- Confidence (how strongly)
- Provenance (from where)
- Justification (backward edges to supporting beliefs)
- Invalidation conditions (when to reconsider)

2.1.6.4. Probabilistic Programming

Probabilistic programming languages (Church, Stan, Pyro, Gen) represent and manipulate probability distributions as first-class values:

```
(define (coin-model)
  (let ((fair? (flip 0.9)))
    (if fair? (flip 0.5) (flip 0.9))))
```

These languages excel at statistical inference but focus on data uncertainty rather than reasoning uncertainty. They require probabilistic normalization and lack explicit justification structure.

2.1.6.5. Justification Logic

Artemov extends modal logic with explicit justification terms. Instead of $\Box\varphi$ (“ φ is known/believed”), we write $t : \varphi$ (“ t is a justification for φ ”). Terms include:

$$t ::= c \text{ mid } x \text{ mid } t.t \text{ mid } t + t \text{ mid } !t$$

where c is a constant (axiom), x is a variable, $s.t$ is application (modus ponens), $t + s$ is sum (either justification suffices), and $!t$ is proof checking (t justifies that t justifies φ).

The key axiom is application:

$$s : (\varphi \rightarrow \psi) \rightarrow (t : \varphi \rightarrow (s.t) : \psi)$$

Limitations. Justification Logic produces tree-structured justifications (each conclusion from fresh premises). It cannot represent:

- Shared premises (same evidence supporting multiple conclusions)
- Defeasible reasoning (defeat edges)
- Graded confidence

CLAIR's extension. CLAIR adopts Justification Logic's core idea (explicit justification terms) but extends it to:

1. **DAGs:** Shared premises create graph structure.
2. **Labeled edges:** Support, undercut, and rebut edges.
3. **Graded confidence:** Each node carries confidence in $[0, 1]$.

2.2. Synthesis: The Gap CLAIR Fills

Concept, Prior Work, CLAIR		
Extension		
Uncertainty	Subjective Logic	Epistemic confidence (about reasoning)
Provenance	Database provenance	Computation provenance + invalidation
Justification	Justification Logic	DAGs with labeled edges, defeat
Belief revision	TMS, AGM	Graded, justification-based revision
Design rationale	IBIS/QOC	First-class decisions in IR
Refinements	Liquid Types	1. confidence + invalidation
Effects	Effect systems	1. intent + semantic meaning
Self-reference	Provability Logic (GL)	Graded Löb (CPL)
Multi-agent	Arrow, Condorcet	Pragmatic internal realism

The gap. No prior work combines:

1. Beliefs as first-class values with epistemic metadata
2. Confidence as non-probabilistic epistemic commitment
3. Justification as labeled DAGs with defeat semantics
4. Self-reference constraints derived from provability logic
5. Belief revision operating on justification structure

CLAIR provides this synthesis, offering a rigorous foundation for AI systems that can explain and audit their own reasoning while honestly representing epistemic limitations.

Key influences. We acknowledge particular debts to:

- De Kleer’s ATMS for dependency-directed reasoning
- Jøsang’s Subjective Logic for uncertainty algebra
- Boolos’s provability logic for self-reference treatment
- Pollock’s defeater theory for defeat semantics
- Artemov’s Justification Logic for explicit justifications
- Milnikel’s graded justifications for uncertainty in justification terms

CLAIR is not a rejection of this prior work but a synthesis that combines their insights into a coherent type-theoretic framework.

2.2.1. Positioning: How CLAIR Differs

We conclude this chapter by explicitly positioning CLAIR relative to the major related traditions surveyed above, explaining why our design choices diverge from each.

2.2.1.1. vs. Probabilistic Approaches

Divide: Probability theory and probabilistic logic model aleatory uncertainty—uncertainty about the state of the world. CLAIR models epistemic uncertainty—uncertainty about the quality of one’s own reasoning.

Why CLAIR differs. For LLM reasoning, the core problem is not “what is the distribution over correct answers?” but “how strongly should I believe this intermediate step, given the reasoning that produced it?” Probability cannot represent low confidence in both φ and $\neg\varphi$ simultaneously (without violating normalization). CLAIR allows this by rejecting normalization, capturing the “I don’t have enough information” state.

What we adopt. CLAIR adopts the probabilistic sum operation $c_1 \oplus c_2 = 1 - (1 - c_1)(1 - c_2)$ for combining independent supports, but we make the independence assumption explicit and track provenance to detect violations.

2.2.1.2. vs. Subjective Logic

Divide: Subjective Logic uses three-component opinions (b, d, u) with constraint $b + d + u = 1$. CLAIR uses a single confidence value $c \in [0, 1]$.

Why CLAIR differs. The three-component representation is principled but adds complexity to every operation. For CLAIR’s target use case (LLM intermediate reasoning), we prioritize simplicity and composability. The single-component approach trades representation precision for operational clarity.

What we adopt. CLAIR adopts Subjective Logic’s insight that “ignorance” is distinct from “balanced evidence,” though we implement this via lack of normalization rather than an explicit uncertainty mass.

2.2.1.3. vs. Weighted Argumentation

Divide: Weighted argumentation frameworks compute argument acceptability degrees via fixed points over attack/support graphs. CLAIR computes belief confidences via local algebraic operations.

Why CLAIR differs. Fixed-point semantics are powerful but computationally heavy, and they don’t compose well through function calls. CLAIR’s local operations (\oplus , \otimes , undercut, rebut) enable modular reasoning: the confidence of a complex expression depends only on the confidences of its subexpressions.

What we adopt. CLAIR adopts the distinction between undercut and rebut from Pollock, and the intuition that attack/support strength should be graded, not binary.

2.2.1.4. vs. Fuzzy Modal Logic

Divide: Fuzzy modal logics interpret $\Box\varphi$ as “in all accessible worlds, φ holds to at least degree θ .” CLAIR’s $\Box_c\varphi$ means “the system has confidence c in its justification for φ .”

Why CLAIR differs. Fuzzy modal logic's graded accessibility relation $R : W \times W \rightarrow [0, 1]$ is elegant but adds significant complexity to the semantics. For CLAIR's application (LLM reasoning traces), we prioritize having a simple, implementable semantics over maximal generality.

What we adopt. CLAIR adopts product logic operations (\otimes as multiplication) and shares the insight that many-valued generalization of classical operators is useful for graded reasoning.

2.2.1.5. vs. Graded Justification Logic

Divide: Milnikel and Fan & Liau add uncertainty degrees to justification terms. CLAIR adds confidence to the propositions themselves and treats justifications as data structures.

Why CLAIR differs. Justification Logic focuses on *explicit proof terms*— t (“this term justifies φ ”). CLAIR focuses on *explicit justifications as graphs*—a data structure tracking which beliefs supported which, with provenance and invalidation.

What we adopt. CLAIR adopts the core idea that justification should be first-class and tracked, and that justification strength can be graded.

2.2.1.6. vs. Belief Revision (AGM)

Divide: AGM operates on deductively closed belief sets (sets of sentences). CLAIR operates on justification graphs with annotated beliefs.

Why CLAIR differs. AGM's postulates (especially Recovery) fail when beliefs have graded strength and structured justification. Retracting a belief loses the specific evidence that supported it, not just the belief itself. CLAIR's revision algorithm explicitly modifies the justification DAG and recomputes confidences, which correctly handles Recovery failure.

What we adopt. CLAIR adopts AGM's core operations (expansion, contraction, revision) but extends them to graded, structured beliefs.

2.2.1.7. vs. Type Theory

Divide: Standard type systems track types (e.g., “this variable is an integer”). CLAIR tracks epistemic metadata (confidence, justification, provenance, invalidation).

Why CLAIR differs. Information flow types and refinement types track *security properties* (who can access this data) and *constraints on values* (this integer is positive). CLAIR tracks *reasoning*

properties (how strongly do we believe this, where did it come from, why do we believe it, when should we reconsider?). This is a new dimension of typing.

What we adopt. CLAIR adopts the insight that justifications are analogous to proof terms: they provide evidence for beliefs, and the DAG structure captures dependency relationships. However, since CLAIR’s content is opaque natural language (interpreted by LLMs, not type-checked), the full Curry-Howard machinery does not apply.

2.2.1.8. The Pragmatic Rationale

These design choices are driven by CLAIR’s target use case: making LLM reasoning auditable and trustworthy. For this application, we prioritize:

1. **Interpretability:** Humans should be able to understand why the system believes what it believes.
2. **Compositionality:** Complex reasoning should build from simple, composable operations.
3. **Formal verifiability:** The core system should have machine-checked proofs of key properties.
4. **Implementability:** A reference implementation should be straightforward to build and reason about.

Different applications might justify different choices. A system for scientific modeling, for example, might benefit from Subjective Logic’s three-component opinions. A system for formal verification might benefit from full dependent types. CLAIR’s design is optimized for *AI reasoning trace auditing*, and our divergence from each related tradition reflects that optimization.

3. Confidence System

“Probability is not about what is true. It is about what is reasonable to believe.”

— E.T. Jaynes, Probability Theory: The Logic of Science

The confidence system is the algebraic foundation of CLAIR. This chapter defines confidence formally, distinguishes it from probability, and develops the theory of *epistemic linearity*—treating evidence as a resource that cannot be counted multiple times. We then establish the three monoid structures that govern how confidence values combine, proving key properties in Lean 4 to connect abstract theory to machine-verified implementation.

3.1. Confidence as Calibrated Reliability

3.1.1. Semantic Commitments

Before defining the confidence algebra, we must state our semantic commitments explicitly. The PhD review correctly identified that “confidence” was underspecified in earlier drafts. We now clarify:

Definition Calibrated Reliability.

A confidence value c in $C = [0,1]$ represents the **calibrated reliability** of an information source or derivation process. Specifically:

1. If a source assigns confidence c to proposition phi , this means:

Across the reference class of similar situations where this source assigns confidence c , the proposition phi is correct with frequency c .

2. Calibration is an *external* property: a source is calibrated if its stated confidences match empirical accuracy over relevant reference classes.
3. CLAIR *tracks* calibrated reliability without *guaranteeing* it. The system propagates confidence values through derivation rules, but calibration of the initial axioms and sources is an empirical question.

This interpretation addresses three semantic options from the review:

1. **Not pure probability:** Confidence does **not** satisfy $P(\text{phi}) + P(\text{not phi}) = 1$. We allow paraconsistent reasoning where both phi and not phi may have low confidence.
2. **Not fuzzy truth degree:** Confidence is **not** the degree to which phi is true in some multivalued logic. It is the reliability of the *source asserting phi*.
3. **Reliability semantics:** Confidence is calibrated reliability of the epistemic process producing the belief. This distinguishes

how strongly I believe

from

how often I'm right when I believe this strongly

3.1.2. The Problem with Probability

Standard approaches to uncertain reasoning use probability. A probability measure P on a set

Ω of outcomes satisfies the Kolmogorov axioms:

1. $P(A) \geq 0$ (Non-negativity)
2. $P(\Omega) = 1$ (Normalization)
3. $P(A \cup B) = P(A) + P(B) - P(A \cap B)$ (Additivity)

For propositions, this implies the fundamental constraint: $P(\phi) + P(\neg\phi) = 1$

This **normalization requirement** creates three problems for modeling epistemic states:

1. **Balanced uncertainty.** An agent confronting an unfamiliar topic may be uncertain about both ϕ and $\neg\phi$. If asked

Is there intelligent life elsewhere in the universe?

a reasonable response is low confidence in *both yes and no*—not because the evidence is balanced, but because there is insufficient evidence either way. Probability forces $P(\text{yes}) + P(\text{no}) = 1$, conflating

I don't know

with

the evidence is exactly balanced.

2. **Paraconsistent reasoning.** Evidence sometimes supports both ϕ and $\neg\phi$ before contradiction resolution. A detective might have testimony that a suspect was present (supporting guilt) and an alibi (supporting innocence), without yet knowing which is false. Probability makes this impossible: $P(\phi) > 0.5$ and $P(\neg\phi) > 0.5$ is a contradiction.

3. **Derivation semantics.** In Bayesian reasoning, $P(A \text{ and } B) = P(A) \text{ times } P(B | A)$, where conditioning captures the dependency structure. But for derivation—where B follows from A by some rule—there is no clear conditional to use. The semantics of

A is a premise for B

differs from

B is probable given A is true.

3.1.3. Definition of Confidence

CLAIR's confidence addresses these problems by dropping normalization:

Definition Confidence Value.

A **confidence value** is a real number c in $[0, 1]$. We write C for the set of confidence values: $C = \text{set}(c \text{ in RR} \mid 0 \leq c \leq 1)$

The semantic interpretation under calibrated reliability:

1. $c = 1$: Axiomatic acceptance (treated as foundational)
2. $c = 0$: Complete rejection (treated as impossibility)
3. $c = 0.5$: Maximal uncertainty (no evidence either direction)
4. $c > 0.5$: Net evidence for acceptance
5. $c < 0.5$: Net evidence for rejection

Definition Epistemic Commitment.

Confidence represents **epistemic commitment**: the degree to which an agent commits to a proposition based on available evidence and reasoning, *interpreted* as the calibrated reliability of the source or derivation producing the belief.

Unlike probability:

1. **No normalization:** $\text{conf}(\phi) + \text{conf}(\neg\phi)$ need not equal 1.
2. **Not frequency:** $\text{conf}(\phi) = 0.9$ does not mean

true 90% of the time

in a single case. It means

this source is calibrated such that across similar cases with confidence 0.9, the proposition is correct 90% of the time.

3. **Derivation-based:** Confidence propagates through inference rules, not conditioning.

Example Non-normalized confidence.

Consider the proposition

This code has no security vulnerabilities.

An honest assessment might be: $\text{conf}(\text{no vulnerabilities}) = 0.4$ (some evidence from testing) $\text{conf}(\text{has vulnerabilities}) = 0.3$ (some evidence from complexity) The sum $0.4 + 0.3 = 0.7 < 1$ reflects residual uncertainty—neither hypothesis is well-supported. This is inexpressible in probability.

3.1.4. Falsifiability Criteria

Following the review's requirement, we state explicitly what would falsify CLAIR's confidence model:

1. **Empirical falsification:** If sources systematically assign confidence c to propositions but are correct with frequency $f \neq c$, they are *uncalibrated*. CLAIR provides no mechanism to detect uncalibrated sources—this requires external validation.
2. **Semantic falsification:** If interpretation as

calibrated reliability

leads to contradictions in the algebra (e.g., violations of monoid laws), the semantics would be inadequate. The Lean 4 formalization verifies algebraic consistency.

3. **Competing semantics:** If an alternative interpretation (e.g., pure probability or fuzzy truth) provides better empirical calibration or algebraic properties, CLAIR's semantics would need revision.

3.1.5. Comparison with Subjective Logic

Jøsang's Subjective Logic (Jøsang, 2016) extends probability with explicit uncertainty. An *opinion* is a tuple $\omega = (b, d, u, a)$:

1. b : belief mass (evidence for)
2. d : disbelief mass (evidence against)
3. u : uncertainty mass (lack of evidence)

4. a : base rate (prior probability)

with constraint $b + d + u = 1$.

CLAIR's confidence can be viewed as a simplification of Subjective Logic: $c = b + u \text{ times } a$ where we collapse uncertainty into the confidence value via the base rate. This loses the $b/d/u$ decomposition but gains simplicity. The trade-off is appropriate for CLAIR's focus on derivation tracking rather than uncertainty quantification.

Note. *linebreak*

CLAIR could be extended to full Subjective Logic opinions. The algebraic structure developed in this chapter would generalize, with the three monoids operating on opinion tuples. We leave this extension to future work.

3.2. The Aggregation Monoid

When multiple independent pieces of evidence support the same conclusion, confidence should *increase*. This requires the probabilistic OR operation.

3.2.1. Probabilistic OR (Oplus)

Definition Probabilistic OR (Oplus).

For confidence values a, b in C , their **aggregation** is: $a \text{ oplus } b = a + b - a \text{ times } b$

The formula has several equivalent forms:

$$1. a \text{ oplus } b = a + b(1 - a)$$

$$2. a \text{ oplus } b = a(1 - b) + b$$

$$3. a \text{ oplus } b = 1 - (1 - a)(1 - b) \text{ (De Morgan duality with multiplication)}$$

The last form reveals the duality: $a \text{ oplus } b$ is the complement of the product of complements.

Theorem Oplus Preserves Bounds.

For all a, b in C : $a \text{ oplus } b$ in C Proof. **Lower bound:** $a \text{ oplus } b = a + b(1 - a) \geq 0$ since $a \geq 0, b \geq 0$, and $(1 - a) \geq 0$.

Upper bound: $a \text{ oplus } b = a + b(1 - a) \leq 1$ since $b \leq 1$ implies $b(1 - a) \leq 1 - a$, therefore $a + b(1 - a) \leq a + (1 - a) = 1$. ■

Theorem Oplus Monoid.

$(C, oplus, 0)$ is a commutative monoid with absorbing element 1:

1. **Associativity:** $(a oplus b) oplus c = a oplus (b oplus c)$
2. **Commutativity:** $a oplus b = b oplus a$
3. **Identity:** $0 oplus a = a oplus 0 = a$
4. **Absorption:** $1 oplus a = a oplus 1 = 1$

Proof. All properties follow from standard real number arithmetic on 0,1. ■

3.2.2. Confidence-Increasing Property

Unlike multiplication, $oplus$ increases confidence:

Theorem Oplus is Confidence-Increasing.

For all a, b in C : $a oplus b \geq max(a, b)$ *Proof.* Using form $a oplus b = a + b(1-a)$. Since $b(1-a) \geq 0$, we have

$a oplus b \geq a$. By commutativity, $a oplus b \geq b$. Therefore $a oplus b \geq max(a, b)$. ■

Corollary Diminishing Returns.

The marginal gain from additional evidence decreases as confidence grows: $(del)/(del b)(a oplus b) = 1 - a$ When a is already high, additional evidence contributes less.

Example Aggregation of Weak Evidence.

Suppose we have ten independent pieces of weak evidence, each with confidence 0.3. The combined confidence is: $0.3 oplus 0.3 oplus \dots oplus 0.3$ (10 times) $= 1 - (1 - 0.3)^{10} = 1 - 0.7^{10} \approx 0.972$ Ten weak independent witnesses produce high combined confidence.

3.2.3. The “Survival of Doubt” Interpretation

The formula $a oplus b = 1 - (1-a)(1-b)$ admits a probability-theoretic interpretation under calibrated reliability:

1. $(1 - a)$ is the

doubt

in evidence a

2. $(1 - a)(1 - b)$ is the probability *both* pieces of evidence fail (assuming independence)
3. $a oplus b$ is the probability *at least one* succeeds

This

survival of doubt

interpretation explains why aggregation increases confidence: more independent evidence means more chances for at least one to be correct.

3.3. The Multiplication Monoid

When a conclusion is derived from premises, its confidence depends on the premises' confidences. The simplest case is conjunctive derivation: both premises must hold for the conclusion to follow.

3.3.1. Conjunctive Confidence Propagation

Definition Confidence Multiplication.

For confidence values a, b in C , their **multiplicative combination** is standard multiplication: a times b

This models the intuition that deriving C from A and B requires both to be true. If we are 90% confident in A and 80% confident in B , our confidence in C (derived from both) is at most $0.9 \times 0.8 = 0.72$.

Theorem Multiplication Preserves Bounds.

For all a, b in C : a times b in C *Proof.* We prove both bounds:

1. a times $b \geq 0$: Since $a \geq 0$ and $b \geq 0$, their product is non-negative.
2. a times $b \leq 1$: Since $b \leq 1$, we have a times $b \leq a$ times $1 = a \leq 1$.

Theorem Multiplication Monoid.

$(C, \text{times}, 1)$ is a commutative monoid with absorbing element 0 :

1. **Associativity:** $(a \text{ times } b) \text{ times } c = a \text{ times } (b \text{ times } c)$
2. **Commutativity:** $a \text{ times } b = b \text{ times } a$
3. **Identity:** $1 \text{ times } a = a \text{ times } 1 = a$
4. **Absorption:** $0 \text{ times } a = a \text{ times } 0 = 0$

Proof. All properties follow from standard real number arithmetic on $0, 1$. ■

3.3.2. The Derivation Monotonicity Principle

A fundamental property of CLAIR is that derivation can only decrease confidence:

Theorem Derivation Monotonicity.

For all a, b in C : a times $b \leq \min(a, b)$. In particular, a times $b \leq a$ and a times $b \leq b$. *Proof.* Since $b \leq 1$, we have a times $b \leq a$ times $1 = a$. By commutativity, a times $b \leq b$. Therefore a times $b \leq \min(a, b)$. ■

Corollary No Confidence Amplification.

No sequence of conjunctive derivations can increase confidence. If c_0 is the confidence of a foundational belief and c_n is derived through n multiplicative steps, then $c_n \leq c_0$.

This principle is essential for CLAIR's epistemology: derived beliefs are never more confident than their sources. Certainty ($c = 1$) is reserved for axioms, not conclusions.

3.4. Defeat Operations

Beyond derivation and aggregation, CLAIR requires operations for *defeat*: when evidence undermines a belief.

3.4.1. Undercut: Attacking the Inference

Following Pollock (2001), an *undercutting defeater* attacks the inferential link, not the conclusion directly.

Definition Undercut.

For confidence c in a conclusion and defeat strength d : $\text{undercut}(c, d) = c$ times $(1 - d)$

Example Undercutting Defeat.

Consider the inference:

The object looks red, therefore it is red.

An undercut

The room has red lighting

doesn't claim the object isn't red; it undermines the inference from appearance to reality.

If $\text{conf}(\text{looks red} \Rightarrow \text{is red}) = 0.9$ and $\text{conf}(\text{red lighting}) = 0.6$, then: $\text{undercut}(0.9, 0.6) = 0.9$ times $(1 - 0.6) = 0.9$ times $0.4 = 0.36$. The inference confidence drops from 0.9 to 0.36.

Theorem Undercut Preserves Bounds.

For all c, d in C : $\text{undercut}(c, d)$ in C . *Proof.* Since $d \leq 1$, we have $(1 - d) \geq 0$. Since $c \geq 0$, we have c times $(1 - d) \geq 0$. Since $c \leq 1$ and $(1 - d) \leq 1$, we have c times $(1 - d) \leq 1$. ■

Theorem Undercut Composition.

Sequential undercuts compose via *oplus*: $\text{undercut}(\text{undercut}(c, d_1), d_2) = \text{undercut}(c, d_1 \text{ oplus } d_2)$ Proof. By expanding the definition: $\text{undercut}(\text{undercut}(c, d_1), d_2) = c(1 - d_1)(1 - d_2)$

4. $c(1 - (d_1 + d_2 - d_1 \text{ times } d_2))$

5. $c(1 - (d_1 \text{ oplus } d_2))$

6. $\text{undercut}(c, d_1 \text{ oplus } d_2)$

■

This beautiful result shows that defeat strengths aggregate via *oplus*: multiple undercuts combine as if their doubts aggregated.

6.0.1. Rebut: Competing Evidence

A *rebutting defeater* provides counter-evidence for the conclusion directly.

Definition Rebut.

For confidence c_{for} in favor and c_{against} against: $\text{rebut}(c_{\text{for}}, c_{\text{against}}) = \text{cases}((c_{\text{for}} / (c_{\text{for}} + c_{\text{against}}), \text{if } c_{\text{for}} + c_{\text{against}} > 0), (0.5, \text{if } c_{\text{for}} = c_{\text{against}} = 0))$

The formula treats evidence symmetrically: the resulting confidence is the

market share

of supporting evidence.

Theorem Rebut Preserves Bounds.

For all $c_for, c_against$ in C: $\text{rebut}(c_for, c_against)$ in C Proof. If $c_for + c_against = 0$, the result is 0.5 in [0,1].

Otherwise:

1. $\text{rebut} \geq 0$ because $c_for \geq 0$ and the denominator is positive.
2. $\text{rebut} \leq 1$ because $c_for \leq c_for + c_against$.

Theorem Rebut Antisymmetry.

$\text{rebut}(a, b) + \text{rebut}(b, a) = 1$ Proof. When $a + b > 0$: $\text{rebut}(a, b) + \text{rebut}(b, a) = a/(a+b) + b/(a+b) = (a+b)/(a+b)$

= 1 When $a = b = 0$: $\text{rebut}(a, b) = \text{rebut}(b, a) = 0.5$, so the sum is 1. ■

6.0.2. Rebut Normalization Limitation

The rebut operation has an important limitation that must be stated explicitly: it *normalizes away absolute strength information*.

Definition Rebut Normalization Property. For all $k > 0$: $\text{rebut}(k \text{ times } a, k \text{ times } b) = \text{rebut}(a, b)$

Proof. When a, b are not both zero: $\text{rebut}(k a, k b) = (k a) / (k a + k b) = k a / (k(a+b)) = a/(a+b) = \text{rebut}(a, b)$ The k cancels from numerator and denominator. ■

This means rebut only captures the *relative balance* of evidence, not the *absolute magnitude*.

Consider:

Example Normalization Loses Absolute Strength. Two scenarios with very different absolute evidence strengths:

1. **Scenario A:** $c_for = 0.1, c_against = 0.1$ Result: $\text{rebut}(0.1, 0.1) = 0.1/(0.1+0.1) = 0.5$
2. **Scenario B:** $c_for = 0.9, c_against = 0.9$ Result: $\text{rebut}(0.9, 0.9) = 0.9/(0.9+0.9) = 0.5$

Both scenarios yield confidence 0.5 despite Scenario B having *nine times more* total evidence. The rebut operation cannot distinguish “weak balanced evidence” from “strong balanced evidence.”

Implications for CLAIR.

The normalization limitation means rebut is appropriate for *competitive evaluation* (comparing opposing arguments) but insufficient for *absolute assessment*. When absolute strength matters, CLAIR should:

1. Track total evidence magnitude separately: $total = c_for + c_against$
2. Use *undercut* instead when the goal is to reduce confidence proportionally to attack strength

3. Apply confidence thresholds to distinguish “weak 0.5” from “strong 0.5”

This is a deliberate design trade-off: rebut prioritizes interpretability (“market share” of evidence) over preserving absolute magnitude.

6.1. Independence Assumptions for Aggregation

The *oplus* operation assumes independence of evidence sources. This is a critical assumption that must be stated explicitly:

Definition Conditional Independence for Oplus.

Two evidence sources e_1 and e_2 are **independent** with respect to proposition phi if: $P(\text{phi} \mid e_1, e_2) = P(\text{phi} \mid e_1) + P(\text{phi} \mid e_2) - P(\text{phi} \mid e_1) \times P(\text{phi} \mid e_2)$. Under calibrated reliability, this means the reference classes of e_1 and e_2 do not systematically overlap.

6.1.1. When Oplus is Valid

The *oplus* operation is **semantically justified** when:

1. **Independent sources:** Evidence derives from causally independent mechanisms (e.g., different sensors, different witnesses, different reasoning paths).
2. **No shared provenance:** The sources do not derive from common antecedents that would cause their errors to correlate.
3. **Reference class disjointness:** The calibration reference classes for e_1 and e_2 do not systematically overlap.

Example Valid Independent Aggregation.

Three different image classifiers (trained on different datasets, with different architectures) each assign confidence 0.7 to

This image contains a cat

. The combined confidence $0.7 \text{ oplus } 0.7 \text{ oplus } 0.7 \approx 0.973$ is justified because the classifiers make statistically independent errors.

6.1.2. When Oplus Breaks

The *oplus* operation **overcounts** and produces misleading confidence when:

- Shared provenance:** Two sources derive from the same evidence. For example, two newspapers reporting the same press release should not be aggregated via *oplus*.
- Common systematic bias:** Sources that share the same misconception or training data flaw will make correlated errors. Aggregating them amplifies bias.
- Circular dependence:** When e_2 cites e_1 as a source, their evidence is not independent.

Example Invalid Aggregation Due to Shared Provenance.

A fact-checking website cites *Source A*. A blog post then cites the fact-checking website. Treating these as independent evidence would incorrectly inflate confidence via *oplus*.

6.1.3. Correlation-Aware Alternatives

When independence is violated, CLAIR provides alternatives:

- Correlation-aware aggregation:** Chapter 4 introduces *oplus_delta* for correlated evidence, where $\delta \in [0,1]$ measures dependency: $a \text{ oplus_delta}(b) = a + b - a \text{ times } b - \delta$
- Min-based aggregation:** When sources may be completely dependent, use $\max(a, b)$ instead of *oplus* to avoid overcounting.
- Provenance tracking:** Evidence usage is tracked via provenance, preventing the same source from being counted twice in a derivation. This is enforced by CLAIR's DAG structure and source tracking (Chapter 10).

Independence Detection.

In practice, determining whether evidence sources are independent requires domain knowledge and provenance tracking. CLAIR does **not** automatically detect dependence—it provides the algebraic machinery for *manually* specifying correlation or preventing double-counting through provenance tracking.

6.1.4. Interval-Based Confidence: Dependency Bounds

When dependence between evidence sources is unknown or partially known, CLAIR supports *interval-based confidence* as a robust alternative to point-value aggregation. Instead of committing to a single confidence value, we track upper and lower bounds that reflect uncertainty about dependence.

Definition Confidence Interval.

A **confidence interval** is a pair $[c_{\text{low}}, c_{\text{high}}]$ with $0 \leq c_{\text{low}} \leq c_{\text{high}} \leq 1$. The interpretation: the true confidence lies in this interval, with the width $c_{\text{high}} - c_{\text{low}}$ reflecting uncertainty about evidence dependence.

6.1.4.1. Interval Aggregation Bounds

For evidence with confidence values a and b , the independence assumption determines the correct aggregation:

1. **Independent:** $a \oplus b = a + b - a \times b$ (probabilistic sum)
2. **Maximally dependent:** $\max(a, b)$ (no double-counting)
3. **Anti-correlated:** Could theoretically reach $\min(a + b, 1)$

Theorem Interval Aggregation Bound.

For confidence values a, b in C , define: $\oplus_{\text{indep}}(a, b) = a \oplus b = a + b - a \times b$ $\oplus_{\text{dep}}(a, b) = \max(a, b)$

Then for any unknown correlation structure: $\oplus_{\text{dep}}(a, b) \leq \text{true_aggregation} \leq \oplus_{\text{indep}}(a, b)$

Proof. The lower bound $\max(a, b)$ occurs when evidence is fully dependent (one source's evidence is a subset of the other's). The upper bound $a \oplus b$ occurs when evidence is independent. Any positive correlation reduces the effective aggregation below $a \oplus b$, while negative correlation is impossible for evidential support (evidence cannot be negatively correlated with the truth of what it supports). ■

Definition Interval Aggregation.

For confidence intervals $[a_{\text{low}}, a_{\text{high}}]$ and $[b_{\text{low}}, b_{\text{high}}]$: $[a_{\text{low}}, a_{\text{high}}] \oplus_{\text{interval}} [b_{\text{low}}, b_{\text{high}}] = [\max(a_{\text{low}}, b_{\text{low}}), \min(a_{\text{high}}, b_{\text{high}})]$

The lower bound uses \max (maximal dependence) and the upper bound uses \oplus (independence).

Example Interval-Based Aggregation for Unknown Dependence.

Two news sources report the same fact. Each has confidence 0.7, but we suspect possible shared provenance (they may have derived from the same press release).

Point-value approach with \oplus : $0.7 \oplus 0.7 = 0.91$ This assumes independence and may be overconfident.

Interval-based approach: $[0.7, 0.7] \oplus_{\text{interval}} [0.7, 0.7] = [\max(0.7, 0.7), 0.7 \oplus 0.7] = [0.7, 0.91]$

The interval $[0.7, 0.91]$ captures our uncertainty: the true confidence could be as low as 0.7 (if the sources are fully dependent) or as high as 0.91 (if independent). The width 0.21 quantifies our dependence uncertainty.

6.1.4.2. Interval Propagation

Intervals propagate through CLAIR's operations:

1. **Multiplication:** $[a_{low}, a_{high}] \text{ times } [b_{low}, b_{high}] = [a_{low} \text{ times } b_{low}, a_{high} \text{ times } b_{high}]$
2. **Undercut:** $\text{undercut}([a_{low}, a_{high}], d) = [a_{low} \text{ times } (1-d), a_{high} \text{ times } (1-d)]$
3. **Rebut:** More complex; requires propagating both bounds through the ratio

Theorem Interval Bound Preservation.

If $[a_{low}, a_{high}]$ and $[b_{low}, b_{high}]$ are valid intervals ($0 \leq a_{low} \leq a_{high} \leq 1$), then all CLAIR operations produce valid intervals.

Proof. Each operation preserves bounds: $oplus$ preserves 0,1 (Theorem 3.3), multiplication preserves 0,1 (Theorem 3.10), and undercut preserves 0,1 (Theorem 3.16). Interval extensions apply the same bounds to endpoints. ■

6.1.4.3. Dependency Bounds from Provenance Tracking

CLAIR's linear type system (Chapter 10) enables *provenance tracking*, which provides tighter dependency bounds than worst-case intervals.

Definition Provenance-Aware Dependency Bound.

Let $S(e)$ be the set of primitive sources (axioms, sensor inputs, LLM outputs) used to derive evidence e . For two evidences e_1, e_2 :

1. If $S(e_1)$ and $S(e_2)$ are disjoint (disjoint provenance): independence is *plausible*, use $oplus$
2. If $S(e_1)$ is a subset of $S(e_2)$ or vice versa (subset provenance): maximal dependence, use \max
3. If there is partial overlap in provenance: use $oplus_{\delta}$ with dependency estimated as the fraction of shared sources

This provides a *heuristic but principled* way to estimate correlation from provenance structure.

The intuition: shared evidence sources correlate the derivations that depend on them.

Example Provenance-Aware Aggregation.

Three sources support a conclusion:

1. e_1 depends on sensors s_1, s_2
2. e_2 depends on sensors s_2, s_3
3. e_3 depends on sensor s_4

Aggregating e_1 and e_2 : They share sensor s_2 , so $\delta \approx 1/2$. Aggregating with e_3 : Disjoint from e_1, e_2 , so independence is plausible.

This is more precise than worst-case intervals and more honest than assuming full independence.

6.1.4.4. When to Use Intervals vs Point Values

Use point values when:

1. Provenance tracking confirms disjoint sources
2. Domain knowledge guarantees independence
3. Computational efficiency is critical

Use intervals when:

1. Provenance is incomplete or unknown
2. Sources may share hidden dependencies
3. Conservative decision-making is required (e.g., safety-critical systems)
4. Auditing requires explicit uncertainty quantification

Relationship to Imprecise Probability.

Interval-based confidence in CLAIR is related to Walley's imprecise probability and Dempster-Shafer theory, but with key differences:

1. **Not a probability interval:** Our intervals bound *calibrated reliability*, not frequency in a reference class.
2. **Not belief/plausibility:** We don't track separate belief and plausibility masses as in Dempster-Shafer.
3. **Operational:** Intervals quantify uncertainty about *evidence dependence*, not about the proposition itself.

This interval-based approach directly addresses Hole A from the review: when independence assumptions are violated, CLAIR provides principled alternatives that avoid overcounting while preserving auditability.

6.2. Conclusion

This chapter established the algebraic and semantic foundation of CLAIR:

1. **Confidence as calibrated reliability:** We explicitly interpret confidence as the calibrated reliability of information sources and derivation processes. This addresses the review's concern about underspecified semantics.

2. **Independence assumptions:** The *oplus* operation requires conditional independence of evidence sources. We provide *oplus_delta* for correlated evidence and affine typing to prevent double-counting.
3. **Three monoids, not a semiring:** Multiplication (derivation), and *oplus* (aggregation) serve distinct semantic roles and do not distribute.
4. **Defeat operations:** Undercut and rebut formalize how evidence can undermine beliefs, with undercuts composing beautifully via *oplus*.
5. **Machine-verified:** The algebra is formalized in Lean 4, ensuring type safety and algebraic correctness.

The confidence system provides the numeric foundation. The next chapter develops the *structural* foundation: how beliefs connect through justification DAGs.

7. Justification as Labeled DAGs

“An argument is not a proof. It is a reason for a belief—and reasons can be defeated.”

— John L. Pollock, Defeasible Reasoning

This chapter develops the structural foundation of CLAIR: how beliefs connect through justification. We show that trees are inadequate for justification structure and that the correct model is a *directed acyclic graph with labeled edges*. The labels distinguish support from defeat, enabling defeasible reasoning where conclusions can be withdrawn when new evidence undermines their justifications.

7.1. The Inadequacy of Trees

7.1.1. The Shared Premise Problem

Traditional approaches represent justification as trees: each conclusion has premises, which themselves have premises, forming an inverted tree structure. This model is elegant but insufficient.

Consider a simple computation that uses the same belief twice:

```
let population = belief(1000000, 0.95, source: census)
let sample_size = belief(1000, 0.90, source: survey)
let ratio = derive population, sample_size by
dividen
let inverse = derive sample_size, population by dividen
let product =
derive ratio, inverse by multiply
```

In a tree representation, each belief appears multiple times as separate subtrees. This creates three problems:

1. **Space inefficiency:** Beliefs are copied rather than shared.
2. **Validation complexity:** If a belief is invalidated, we must find and invalidate all copies.
3. **Semantic confusion:** Are these the *same* belief or *different* beliefs that happen to be equal?

The correct representation is a DAG with explicit sharing, where each belief appears exactly once with multiple parents.

Theorem DAG Necessity.

Any justification system that:

1. Allows a belief to be used as a premise in multiple derivations
 2. Propagates invalidation correctly (removing a premise invalidates all conclusions)
 3. Maintains identity (the “same” belief is the same node)
- must represent justification as a DAG, not a tree.

Proof. In a tree, each node has exactly one parent. If a belief is used in derivations of two different conclusions, it must appear as a child of both. But in a tree, a node cannot have two parents. Therefore, the belief must be duplicated, violating identity. A DAG allows multiple parents, resolving the contradiction. ■

7.1.2. Why Not Cycles?

If we allow sharing, why not allow cycles? Coherentist epistemology suggests beliefs can mutually support each other. We reject cycles in justification for three reasons:

1. **Bootstrap problem:** Circular justification allows confidence inflation with no external grounding.
2. **Invalidation ambiguity:** If beliefs support each other circularly, invalidation semantics become ill-defined.
3. **Well-foundedness:** The justification relation should be well-founded, with no infinite descending chains.

Note: The theory / observation circularity example is better analyzed as two separate relations:

1. **Evidential support** (tracked in justification): Observation provides evidence for theory.
2. **Interpretive framework** (not part of justification): Theory provides framework for interpreting observation.

7.2. Labeled Edges for Defeat

The DAG structure addresses sharing but not defeat. When evidence undermines a belief's justification, we need edges that carry negative, not positive, epistemic weight.

7.2.1. The Defeat Problem

A *defeater* is a belief that undermines confidence in another belief. Following Pollock (2001), we distinguish:

1. **Undercutters:** Attack the inferential link, not the conclusion directly.
2. **Rebutters:** Provide direct counter-evidence against the conclusion.

Definition Edge Types.

A justification edge has one of three types:

1. **Support:** Positive evidence for the target.
2. **Undercut:** Attacks the inferential link to the target.
3. **Rebut:** Direct counter-evidence against the target.

7.2.2. Formal Definition

Definition Justification Graph.

A **justification graph** is a tuple $G = (N, E, r)$ where:

1. N is a finite set of *justification nodes*
2. $E \subseteq N \times N$ is a set of labeled edges
3. $r \in N$ is the root node

subject to the constraint that the underlying unlabeled graph is acyclic.

Definition Justification Node Types.

Each node has one of the following types:

1. axiom: Foundational belief (confidence = 1)
2. rule
 - (**r**, **premises**) : Deductive rule application
3. assumption
 - (**a**) : Temporary assumption for reasoning
4. choice
 - (**options**, **criteria**) : Decision point
5. abduction
 - (**obs**, **hypotheses**, **selected**) : Abductive inference
6. analogy
 - (**source**, **similarity**) : Analogical reasoning
7. induction
 - (**cases**, **rule**) : Inductive generalization
8. aggregate
 - (**sources**, **combRule**) : Evidence aggregation

7.2.3. Well-Formedness Constraints

A well-formed justification graph must satisfy explicit constraints to ensure semantic coherence and computational tractability.

Definition Acyclicity Constraint.

A justification graph must be acyclic: no path may exist from any node back to itself. This constraint applies specifically to support edges. Defeat edges may form cycles, which are resolved via fixed-point iteration (discussed below).

Definition Well-Formed Justification Graph.

A justification graph is well-formed iff:

1. The support structure is acyclic (no support cycles)
2. Every non-axiom node has at least one incoming support edge
3. Every node is reachable from the root in the underlying undirected graph

Condition 1 ensures the support structure is well-founded. Condition 2 ensures no “floating” beliefs without justification. Condition 3 ensures the graph is connected.

7.2.4. DAG-Only vs Cyclic Choice

CLAIR adopts a hybrid approach to cycles:

1. **Support edges: DAG-only.** Cycles in evidential support are semantically prohibited because they enable bootstrapping and violate well-foundedness. The type checker enforces acyclicity for support edges at construction time.
2. **Defeat edges: Fixed-point semantics.** Defeat cycles are permitted and resolved via fixed-point iteration. When defeat edges form cycles, confidence propagation requires solving a system of equations.

This design choice reflects the epistemic distinction between evidence for (which must be well-founded) and attacks against (which can mutually undermine each other).

7.2.5. Fixed-Point Semantics for Cyclic Defeat

When defeat edges form cycles, we compute confidences via iterative fixed-point finding.

Theorem Fixed-Point Existence.

For any defeat graph with acyclic underlying support, a fixed point exists.

Proof. The confidence update function maps the unit interval to itself and is continuous. By Brouwer’s fixed point theorem, a continuous function from a compact convex set to itself has a fixed point. ■

Theorem Fixed-Point Uniqueness.

If the maximum product of undercut strengths along any cycle is strictly less than 1, then the fixed point is unique and Kleene iteration converges to it.

Proof. Under this condition, the update function is a contraction mapping. By the Banach fixed point theorem, a contraction has a unique fixed point and iteration converges to it. ■

Definition Kleene Iteration.

Starting from initial confidences, repeatedly apply the update function until convergence. The sequence generated by repeated application is the Kleene iteration.

Example Convergence for Typical Defeat Cycles.

Suppose two nodes mutually undercut with strength 0.5 each. The Lipschitz constant is 0.5 times 0.5 equals

0.25. Starting from initial confidence (1, 1):

1. After 1 iteration: error is at most 0.25 times initial error
2. After 5 iterations: error is at most 0.25^5 approximately 0.001 (0.1%)
3. After 10 iterations: error is at most 0.25^{10} approximately 10^{-6} (0.0001%)

Rapid convergence is typical for well-formed defeat graphs.

Practical Implication.

For CLAIR implementations, we recommend:

1. Enforce DAG structure for support edges via static type checking
2. Allow defeat cycles but limit undercut strengths to ensure contraction
3. Use Kleene iteration with convergence threshold 10^{-6}
4. Cache fixed-point solutions to avoid recomputation during updates

7.3. Confidence Propagation

Given a justification graph, we compute the confidence of each node bottom-up.

Definition Support Propagation.

For a node with children having confidences c_1, \dots, c_k :

1. $\text{conf}(\text{axiom}) = 1$
2. $\text{conf}(\text{rule}(r, \text{children})) = s_r$ times product over children
3. $\text{conf}(\text{aggregate}(\text{children}, \text{independent})) = oplus$ over children
where s_r is the rule strength and $oplus$ is probabilistic OR.

Definition Defeat Propagation.

Let c be confidence from support edges, d_1, \dots, d_m be undercut strengths, and r_1, \dots, r_n be rebut strengths.

Then: $c' = \text{rebut}(\text{undercut}(c, oplus d_i), oplus r_j)$

Theorem Propagation Termination.

The propagation algorithm terminates for any acyclic justification graph.

Proof. The graph is acyclic by definition. Each recursive call moves to a node strictly earlier in topological order. Since the graph is finite, recursion terminates. ■

7.4. Reinstatement

A fundamental phenomenon in defeasible reasoning is *reinstatement*: when a defeater is itself defeated, the original belief recovers some confidence.

Theorem Compositional Reinstatement.

Let A have base confidence a , undercut by D with confidence d , which is itself undercut by E with confidence e . Then: $\text{conf}(A) = a \text{ times } (1 - d \text{ times } (1 - e))$

The *reinstatement boost* is $\Delta = a \text{ times } d \text{ times } e$.

Proof. Bottom-up evaluation gives: $\text{conf_eff}(D) = d \text{ times } (1 - e)$, then $\text{conf}(A) = a \text{ times } (1 - \text{conf_eff}(D)) = a \text{ times } (1 - d(1-e))$. ■

7.5. Mutual Defeat

When two arguments defeat each other, we have a defeat cycle. The fixed point analysis yields:

Theorem Mutual Defeat Fixed Point.

If A and B mutually undercut with base confidences a and b , the fixed point is: $a_{\text{star}} = a(1-b) / (1 - ab)$ and $b_{\text{star}} = b(1-a) / (1 - ab)$

Proof. At fixed point: $a_{\text{star}} = a(1 - b_{\text{star}})$ and $b_{\text{star}} = b(1 - a_{\text{star}})$. Solving gives the stated formulas.

Theorem Fixed Point Existence.

For any defeat graph with acyclic underlying support, a fixed point exists.

Proof. The confidence update function maps $[0,1]^n$ to itself and is continuous. By Brouwer's fixed point theorem, a fixed point exists. ■

Theorem Uniqueness Condition.

If $b_{\max} \times d_{\max} < 1$, the fixed point is unique and iteration converges.

Proof. Under this condition, the update function is a contraction mapping. By the Banach fixed point theorem, there is a unique fixed point and iteration converges. ■

7.6. Correlated Evidence

The aggregation formula $c_1 \oplus c_2$ assumes independence. When evidence sources are correlated, this overcounts.

Definition Dependency-Adjusted Aggregation.

For evidence with confidences c_1, c_2 and dependency δ in $[0,1]$: $\text{aggregate}_{\delta}(c_1, c_2) = (1-\delta)c_1 \oplus c_2 + \delta \times (c_1 + c_2) / 2$
where $\delta = 0$ means independent and $\delta = 1$ means fully dependent.

7.7. Connection to Prior Art

7.7.1. Truth Maintenance Systems

JTMS (Doyle 1979) uses IN/OUT lists for dependency-directed backtracking. ATMS (de Kleer 1986) tracks multiple assumption sets.

CLAIR contribution: Generalize TMS to graded confidence with the same dependency-directed architecture.

7.7.2. Argumentation Frameworks

Dung's argumentation (Dung 1995) defines acceptance semantics. Gradual semantics (Akgoud et al. 2017) add numeric degrees.

CLAIR contribution: Integrate argumentation's defeat semantics with type-theoretic justification.

7.7.3. Justification Logic

Artemov's Justification Logic (Artemov 2001) adds explicit justification terms to modal logic.

CLAIR contribution: Extend from tree-like justification terms to DAGs with labeled edges, supporting defeasible reasoning.

7.8. The Tracking Paradigm

The preceding sections developed the machinery of justification DAGs, confidence propagation, and defeat semantics. We now step back and articulate the underlying *tracking paradigm* that distinguishes CLAIR from traditional formal systems.

7.8.1. State Representation

7.8.1.1. The Epistemic State

A CLAIR system's **epistemic state** at any point in time is a pair:

Definition Epistemic state.

An epistemic state $\text{cal}(E)$ is a tuple $\text{cal}(E) = (\text{cal}(G), \text{cal}(B))$ where:

1. $\text{cal}(G) = \text{set}(G_1, \dots, G_n)$ is a finite set of justification graphs (one for each belief)
2. $\text{cal}(B) = \text{set}((v_i, c_i, j_i, p_i) \text{ mid } i \text{ in } I)$ is a set of beliefs with values, confidences, justifications, and provenance

Each belief in $\text{cal}(B)$ corresponds to the root of some graph in $\text{cal}(G)$. The graphs may share nodes (shared premises) but are not required to be connected.

7.8.1.2. Comparison with Proving Paradigm

Traditional formal systems focus on *provability*:

1. **State:** A set of axioms and inference rules
2. **Question:** Is ϕ provable from the axioms? ($\Gamma \vdash \phi$)
3. **Output:** Yes/No (with proof term)

CLAIR's tracking paradigm focuses on *epistemic representation*:

1. **State:** A set of labeled graphs with confidence values
2. **Question:** What is the confidence, justification structure, and provenance of belief ϕ ?
3. **Output:** (c, G, p) where $c \in [0,1]$ is confidence, G is the justification graph, p is provenance

7.8.2. Update Rules

The epistemic state changes through *epistemic actions*. Each action maps $\text{cal}(E) \rightarrow \text{cal}(E)'$.

7.8.2.1. Primitive Actions

Definition Epistemic actions.

The primitive actions on epistemic states are:

1. **Add belief:** `add(phi, c, j, p)` creates a new belief with confidence c , justification j , and prove-nance p .
2. **Aggregate:** `aggregate(phi, psi, r)` combines two beliefs about the same proposition using rule r in set("independent", "correlated"(δ), "max").
3. **Derive:** `derive(phi, [psi_1, ..., psi_k], rule)` creates φ as a conclusion from premises using an inference rule.
4. **Undercut:** `undercut(phi, delta, d)` adds an undercutting defeater to φ with strength d .
5. **Rebut:** `rebut(phi, delta, d)` adds a rebutting defeater to φ with strength d .
6. **Validate:** `invalidate(psi)` removes belief ψ and propagates invalidation to all beliefs that depend on ψ .

7.8.2.2. Action Semantics

Example Adding a belief.

When adding a belief `add(phi, 0.8, source: testimony, witness_A)` :

1. Create a new node n_φ with type `axiom` and confidence 0.8
2. Create a single-node graph G_φ containing only n_φ with no edges
3. Add $(n_\varphi, 0.8, \text{source: testimony, witness_A})$ to \mathcal{B}
4. Add G_φ to \mathcal{G}

Example Deriving a conclusion.

When deriving `derive(phi, [psi_1, psi_2], modus_ponens)` :

1. Create a new node n_φ with type `rule(modus_ponens, [n_psi_1, n_psi_2])`
2. Create support edges $(n_{\psi_1}, n_\varphi, \text{support})$ and $(n_{\psi_2}, n_\varphi, \text{support})$
3. Compute confidence via propagation: $c_\varphi = c_{\psi_1} \times c_{\psi_2}$
4. Add graph G_φ to \mathcal{G} and belief $(n_\varphi, c_\varphi, G_\varphi, \text{derived})$ to \mathcal{B}

Example Invalidation propagation.

When `invalidate(psi)` is called:

1. Remove ψ from \mathcal{B}
2. Find all beliefs φ such that ψ appears in G_φ 's support graph
3. Recursively invalidate each such φ (mark as defeated or recompute without ψ)
4. Update confidence values via re-propagation

This is the *dependency-directed backtracking* inherited from TMS systems.

7.8.3. Correctness Criteria

What does it mean for the tracking paradigm to be “correct”? We distinguish three levels of correctness.

7.8.3.1. Syntactic Correctness

Definition Syntactic correctness.

An epistemic state $cal(E)$ is *syntactically correct* if:

1. Every graph in $cal(G)$ is acyclic (support edges)
2. Every node has a well-defined type
3. Every edge has a valid label (support, undercut, rebut)
4. All confidence values are in $[0,1]$

Syntactic correctness is enforced by the type system and checked at runtime. The Lean 4 formalization proves that propagation preserves syntactic correctness.

7.8.3.2. Semantic Correctness (Internal)

Definition Semantic correctness (internal).

An epistemic state $cal(E)$ is *semantically correct* if:

1. Confidence propagation is consistent: Re-computing any belief's confidence yields the same value
2. Defeat semantics are satisfied: Undercuts and rebuts are applied according to their definitions
3. Bounds are preserved: All confidences remain in $[0,1]$ after any update

This is *internal* correctness: the system behaves according to its own rules. The theorems on propagation termination and sound establish internal correctness for confidence propagation.

7.8.3.3. Semantic Correctness (External)

Definition Semantic correctness (external).

An epistemic state $\text{cal}(E)$ is *externally correct* with respect to a reference system $\text{cal}(R)$ if:

1. **Calibration:** For any belief with confidence c , the reference system assigns probability $c' \approx c$ (within calibration tolerance)
2. **Justification adequacy:** The justification graph G accurately represents the actual dependency structure in $\text{cal}(R)$
3. **Provenance accuracy:** The provenance field correctly identifies the source of the belief in $\text{cal}(R)$

External correctness cannot be established by formal proof alone—it requires empirical validation. CLAIR provides the *machinery* for tracking (internal correctness) but calibration and accuracy are properties of the *sources*, not the tracking system.

The Honesty Principle.

CLAIR's tracking paradigm embodies epistemic humility:

1. The system *reports* its confidence, justification, and provenance
2. It does not *guarantee* that these correspond to external reality
3. External correctness must be validated empirically (calibration studies, provenance audits)

This distinguishes tracking from proving: a proof claims certainty; a tracked belief admits uncertainty while documenting its grounds.

7.8.4. Tracking vs. Proving: A Summary

Aspect	Proving Paradigm	Tracking Paradigm
Goal	Establish truth	Record epistemic state
State	Set of formulas	Set of labeled graphs
Query	$\Gamma \vdash -\varphi?$	What is the <i>confidence, justification, and provenance</i> of $\varphi?$
Output	Proof term (or \perp)	(c, G, p) triple
Update	Add axiom/formula	Add/undercut/rebut/invalidate
Correctness	Soundness	Internal + External
Limits	Hidden or denied	Explicit with confidence

7.8.5. Practical Implications

The tracking paradigm has several practical consequences for CLAIR as an AI reasoning intermediate representation:

7.8.5.1. Explainability

Every belief carries its justification graph. Queries like “why does the system believe φ ?” can be answered by traversing the graph and presenting the dependency chain.

7.8.5.2. Debugging

When a belief has unexpectedly low confidence, the graph reveals which defeaters are responsible. When confidence is too high, the graph shows whether aggregation rules were misapplied.

7.8.5.3. Revision

New evidence is incorporated by adding nodes and edges. The system automatically recomputes confidences via propagation, handling reinstatement without special cases.

7.8.5.4. Uncertainty

Unlike binary logic, CLAIR tracks degrees of belief. This matches the uncertainty inherent in real-world reasoning and LLM outputs.

7.9. Conclusion

This chapter established the structural foundation of CLAIR and articulated the tracking paradigm:

1. **DAGs, not trees:** Shared premises require graph structure; explicit sharing enables correct invalidation.
2. **Acyclic:** Cycles in evidential support violate well-foundedness; defeat cycles handled via fixed-point semantics.
3. **Labeled edges:** Support, undercut, and rebut serve different epistemic roles.
4. **Compositional reinstatement:** Defeaters being defeated automatically recovers confidence.

5. **Correlated evidence:** Independence assumptions must be explicit; dependency adjustment prevents overcounting.
6. **Tracking paradigm:** CLAIR represents epistemic state as labeled graphs with confidence values, updating via add/derive/undercut/rebut/invalidate operations. Correctness has three levels: syntactic (well-formedness), internal semantic (consistency with propagation rules), and external semantic (calibration to reality).

The justification DAG provides the structural substrate for CLAIR’s beliefs. The tracking paradigm formalizes what it means to “track not prove”—a fundamental shift from establishing truth to documenting epistemic grounds with explicit confidence.

1. **DAGs, not trees:** Shared premises require graph structure; explicit sharing enables correct invalidation.
2. **Acyclic:** Cycles in evidential support violate well-foundedness; defeat cycles handled via fixed-point semantics.
3. **Labeled edges:** Support, undercut, and rebut serve different epistemic roles.
4. **Compositional reinstatement:** Defeaters being defeated automatically recovers confidence.
5. **Correlated evidence:** Independence assumptions must be explicit; dependency adjustment prevents overcounting.

The justification DAG provides the structural substrate for CLAIR’s beliefs. The next chapter addresses a subtler challenge: how beliefs can safely refer to themselves.



Self-Reference and the Gödelian Limits

line(length: 20%, stroke: 1.5pt, rest: academic-burgundy)

“If a system is consistent, it cannot prove its own consistency.”

— Kurt Gödel, On Formally Undecidable Propositions

CLAIR allows beliefs about beliefs. This reflexive capacity creates potential for self-reference: a belief that refers to itself, either directly or through a chain of intermediate beliefs. Such self-reference is both a powerful expressive tool and a source of potential paradox. This chapter develops the theoretical foundations for distinguishing safe from dangerous self-reference, culminating in a novel extension of provability logic to graded confidence.

7.10. The Problem of Self-Reference

7.10.1. Direct Self-Reference

Consider a belief that directly references itself:

Example Direct Self-Reference.

```
-- A belief referencing its own confidence
let b : Belief<Bool> = belief(
    value: confidence(b) > 0.5,
    confidence: ???
)
```

What confidence should b have? If we assign confidence $c > 0.5$, the content becomes true, which seems consistent. But if we assign $c \leq 0.5$, the content becomes false—yet what prevents us from assigning $c = 0.9$ anyway?

This is not merely a curiosity. If CLAIR aims to capture how an LLM reasons, and if introspection is part of reasoning, then CLAIR must account for self-referential beliefs—even if that account restricts or forbids certain patterns.

7.10.2. Why Self-Reference Matters

Self-reference enables powerful epistemic capabilities:

1. **Calibration:** “My confidence estimates are typically accurate”
2. **Uncertainty tracking:** “I am uncertain about this belief”
3. **Meta-reasoning:** “I should reconsider beliefs derived from unreliable sources”
4. **Self-improvement:** “My reasoning process could be improved in specific ways”

But self-reference also enables paradoxes:

1. **Liar-like:** “This belief has confidence 0” (no consistent assignment)
2. **Curry-like:** “If this belief is true, then arbitrary proposition P ” (proves anything)
3. **Löbian:** “If I believe P , then P is true” (circular self-validation)

The challenge is to permit the former while blocking the latter.

7.11. Löb's Theorem and Anti-Bootstrapping

7.11.1. The Classical Result

Löb's theorem is a cornerstone of provability logic:

Theorem **Löb's Theorem.**

In any sufficiently strong formal system T containing arithmetic: $\text{prov}_{T(\text{prov}(\text{prov}(P) \rightarrow P))} \rightarrow \text{prov}(P)$ where prov denotes provability in T .

In words: if a system can prove “if P is provable, then P is true,” then the system can prove P .

This has a startling consequence.

Corollary **No Internal Soundness Proof.**

No consistent system can prove its own soundness, i.e., cannot prove $\forall P, \square(P) \rightarrow P$.

Proof. Suppose system T proved $\forall P, \square(P) \rightarrow P$. Instantiating with $P = \text{"false"}$ (falsity), we get $\text{prov}(\text{false}) \rightarrow \text{false}$. Combining with consistency ($\neg \text{prov}(\text{false})$), we can derive $\text{prov}(\text{false})$, contradicting consistency. ■

7.11.2. Application to CLAIR

For CLAIR, interpret $\square(P)$ as “CLAIR believes P with confidence 1.0.” Then Löb's theorem constrains self-soundness beliefs:

```
-- A claimed self-soundness belief
let soundness = belief(
    value: forall P. (belief(P, c, ...) and c > 0.9) -> P is true,
    confidence: 0.95
)
```

By Löb's theorem, if CLAIR can form this belief with high confidence, then (classically) CLAIR believes everything with high confidence—a collapse to triviality. This is the *bootstrapping trap*: self-soundness claims cannot increase epistemic authority.

Definition **Anti-Bootstrapping Principle.**

A belief system satisfies *anti-bootstrapping* if no belief of the form “my beliefs are sound” can increase confidence in any derived belief beyond what the original evidence supports.

Löb's theorem mathematically enforces anti-bootstrapping for classical provability. The question is how this extends to graded confidence.

7.12. Tarski's Hierarchy: Stratified Introspection

7.12.1. The Classical Solution

Tarski's theorem on the undefinability of truth states that no sufficiently expressive language can define its own truth predicate—on pain of the Liar paradox. Tarski's solution is stratification:

Level	Can Express	Cannot Express
Level 0 (object)	Facts about the world	Truth of any sentence
Level 1 (meta)	X_0 is true for level-0 X	Truth of level-1 sentences
Level 2 (meta-meta)	X_1 is true for level-1 X	Truth of level-2 sentences
...

Each level can discuss truth at lower levels but never its own level.

7.12.2. Stratified Beliefs in CLAIR

We apply this to beliefs:

Definition Stratified Belief Type.

$\text{Bel}(n, A)$ for $n \in \mathcal{NN}$ where level- n beliefs may reference level- m beliefs only if $m < n$.

Example Stratified Beliefs in CLAIR.

In the CLAIR IR format, each belief has an explicit level field:

```
; Level 0: beliefs about the world (no introspection)
b1 .9 L0 @self "user is authenticated"

; Level 1: beliefs about level-0 beliefs
b2 .95 L1 <b1 "my auth belief b1 has high confidence"
```

```
; Level 2: beliefs about level-1 beliefs
b3 .9 L2 @self <b2 "my introspection in b2 seems accurate"
```

The level constraint is: a belief at level n can only justify beliefs at level $\geq n$, and a belief *about* another belief must be at a higher level.

Theorem Stratification Safety.

If all beliefs respect the stratification constraint— $\text{Bel}(n, A)$ references only $\text{Bel}(m, B)$ with $m < n$ —then no Liar-like paradox can arise.

Proof. Any reference chain from a belief b must strictly decrease in level. Since \mathbb{N} has no infinite descending chains, every chain terminates at level 0. Level-0 beliefs contain no belief references, so they cannot participate in self-referential loops. Therefore, no belief can reference itself directly or transitively. ■

7.12.3. What Stratification Rules Out

Stratification prohibits:

1. **Direct self-reference:** A belief cannot mention itself (would require level $n < n$).
2. **Universal introspection:** “All my beliefs are...” spans all levels and cannot be expressed at any finite level.
3. **Self-soundness at a single level:** “My level- n beliefs are sound” would require level $n + 1$ to express.

7.12.4. The Cost of Safety

Stratification is safe but restrictive. Some legitimate self-referential reasoning is blocked:

```
; Legitimate but blocked: calibration beliefs
; This would be self-referential (talks about own confidences)
; but is intuitively safe (no paradox)
b1 .8 L? @self "my confidence estimates match empirical accuracy"
; What level should this be? It references "all my confidences"
; which spans all levels---cannot be expressed at any finite level
```

This motivates a more permissive approach for certain cases.

7.13. Kripke's Fixed Points: Safe Self-Reference

7.13.1. The Fixed-Point Construction

Kripke proposed an alternative to stratification: allow self-reference but let some sentences remain *undefined*. The key insight is that certain self-referential constructs have *fixed points*—consistent confidence assignments—while others do not.

Definition Fixed Point for Self-Referential Belief.

A self-referential belief b with confidence function $f : \text{cal}("D")\text{cal}("D")[0,1] \rightarrow \text{cal}("D")\text{cal}("D")[0,1]$ (determining confidence from the assumed truth value) has a fixed point if there exists c in $\text{cal}("D")\text{cal}("D")[0,1]$ such that: $c = f(c)$

Example Truth-Teller: Multiple Fixed Points.

Consider:

```
let tt = self_ref_belief(fun self =>
    content: "this belief is true",
    compute_confidence: if val(self.content) then 1.0 else 0.0
)
```

If confidence is 1.0: content is true, so confidence should be 1.0. ✓ If confidence is 0.0: content is false, so confidence should be 0.0. ✓

Both are fixed points. The belief is *underdetermined*.

Example Liar: No Fixed Point.

Consider:

```
let liar = self_ref_belief(fun self =>
    content: "this belief has confidence 0",
    compute_confidence: if val(self.content) then 1.0 else 0.0
)
```

If confidence is 1.0: content says “confidence 0,” which is false, so confidence should be 0.0. Contradiction. If confidence is 0.0: content says “confidence 0,” which is true, so confidence should be 1.0. Contradiction. No fixed point exists. The belief is *ill-formed*.

Example Grounded Self-Reference: Unique Fixed Point.

Consider:

```

let careful = self_ref_belief(fun self =>
    content: "confidence(self) is in [0.4, 0.6]",
    compute_confidence: 0.5
)

```

The compute function is constant, so $f(c) = 0.5$ for all c . The fixed point is $c = 0.5$, which indeed satisfies $0.5 \in [0.4, 0.6]$. This belief is *well-formed* with unique confidence 0.5.

7.13.2. The Self-Reference Escape Hatch

CLAIR provides a controlled mechanism for self-reference:

When a CLAIR trace analyzer encounters self-referential beliefs that escape finite stratification, it classifies them:

Self-Reference Analysis Result:

- **WellFormed:** unique fixed point exists (safe)
- **IllFormed:** no fixed point (Liar-like, Curry-like, or Löbian trap)
- **Underdetermined:** multiple fixed points (policy choice needed)

Error Types:

- **NoFixedPoint:** Liar-like paradox
- **CurryLike:** proves anything
- **LobianTrap:** self-soundness claim
- **Timeout:** analysis did not terminate

Beliefs that escape finite stratification exist “outside” all levels and require special analysis.

7.13.3. Classification of Self-Reference

Combining Tarski and Kripke, we classify self-referential constructs:

Category	Fixed Points	Status	Example
Grounded	Unique	Safe	Calibration beliefs
Underdetermined	Multiple	Policy choice	Truth-teller

Category	Fixed Points	Status	Example
Liar-like	None	Ill-formed	“Confidence is 0”
Curry-like	—	Banned	“If true, then P ”
Löbian	—	Banned	Self-soundness

Definition Safe Self-Reference.

A self-referential belief is *safe* if it either:

1. Respects stratification (level- n references only level- $m < n$), or
2. Has a unique fixed point (Kripke), or
3. Has multiple fixed points with a deterministic policy for selection.

Definition Dangerous Self-Reference.

A self-referential belief is *dangerous* if it:

1. Has no fixed point (Liar-like), or
2. Matches a Curry pattern (“if this then P ”), or
3. Claims self-soundness (Löbian trap).

7.14. Provability Logic and CLAIR

7.14.1. Gödel-Löb Logic (GL)

To formally characterize CLAIR’s belief logic, we turn to *provability logic*. The standard modal logic of provability is GL (Gödel-Löb):

Definition GL Syntax.

$\varphi ::= p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \varphi \vee \varphi \mid \varphi \rightarrow \varphi \mid \Box\varphi$ where $\Box\varphi$ means φ is provable.

Definition GL Axioms.

1. **K (Distribution):** $\Box(\varphi \rightarrow \psi) \rightarrow (\Box\varphi \rightarrow \Box\psi)$
2. **4 (Positive Introspection):** $\Box\varphi \rightarrow \Box\Box\varphi$
3. **L (Löb):** $\Box(\Box\varphi \rightarrow \varphi) \rightarrow \Box\varphi$

Critically, GL *lacks* the truth axiom $\Box\varphi \rightarrow \varphi$ (T). This is philosophically essential: provability does not imply truth. A consistent system can prove false statements if its axioms are wrong.

7.14.2. GL vs Other Modal Logics

Logic	K	T	4	5 or L
K	✓			
T	✓	✓		
S4	✓	✓	✓	
S5	✓	✓	✓	5
GL	✓		✓	L
CLAIR	✓		✓	L

CLAIR aligns with GL:

1. **K holds:** If CLAIR believes an implication and believes the antecedent, it can derive the consequent.
2. **T fails:** CLAIR's beliefs can be wrong (fallibilism).
3. **4 holds:** CLAIR can have meta-beliefs about its beliefs.
4. **L must hold:** Self-soundness claims cannot bootstrap confidence.

7.14.3. Solovay's Completeness

Theorem Solovay Completeness.

GL is sound and complete with respect to:

1. Arithmetic provability: $GL \vdash \varphi$ iff φ holds under all interpretations of \square as Gödel provability in PA.
2. Finite transitive irreflexive Kripke frames.

The completeness for finite frames yields:

Corollary GL Decidability.

GL is decidable (PSPACE-complete).

This is crucial: classical provability logic is computationally tractable.

7.15. Confidence-Bounded Provability Logic (CPL)

Classical GL uses binary truth: propositions are either provable or not. CLAIR needs a *graded* version where beliefs carry confidence values in $\text{cal}(\text{"D"})[0,1]$. This section introduces CPL (Confidence-Bounded Provability Logic), a novel extension of GL designed for CLAIR.

7.15.1. The Literature Gap

Extensive work exists on fuzzy modal logics and graded epistemic logic. However, no prior work addresses:

1. Graded versions of the Löb axiom
2. The interaction of continuous confidence with provability constraints
3. Anti-bootstrapping in the context of graded belief

CPL fills this gap.

7.15.2. CPL Syntax

Definition CPL Syntax.

$\varphi ::= p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \varphi \vee \varphi \mid \varphi \rightarrow \varphi \mid \Box_c \varphi$ where $\Box_c \varphi$ means φ is believed with confidence at least c .

7.15.3. CPL Semantics

Definition Graded Kripke Frame.

A *graded Kripke frame* is a tuple (W, R) where:

1. W is a non-empty set of worlds
 2. $R : W \times W \rightarrow \mathcal{D}[0, 1]$ is a graded accessibility relation
- satisfying:
1. **Transitivity:** $R(w, v) \cdot R(v, u) \leq R(w, u)$
 2. **Converse well-foundedness:** No infinite sequence w_0, w_1, w_2, \dots with $R(w_{i+1}, w_i) > 0$ for all i

Definition Graded Valuation.

A *graded valuation* on a frame (W, R) assigns to each world w and proposition p a confidence value $V_{w(p)} \in \mathcal{D}[0, 1]$. Extended to formulas:

1. $V_{w(\neg\varphi)} = 1 - V_{w(\varphi)}$
2. $V_{w(\varphi \wedge \psi)} = V_{w(\varphi)} \cdot V_{w(\psi)}$
3. $V_{w(\varphi \vee \psi)} = V_{w(\varphi)} + V_{w(\psi)} - V_{w(\varphi)}V_{w(\psi)}$
4. $V_{w(\varphi \rightarrow \psi)} = \sup\{c \in \mathcal{D}[0, 1] : V_{w(\varphi)} \cdot c \leq V_{w(\psi)}\}$
5. $V_{w(\Box_c \varphi)} = \inf_{v: R(w, v) \geq c} V_{w(\varphi)}$

The last clause says: *phi* is believed at confidence c if *phi* holds in all worlds accessible with strength at least c .

7.15.4. The Graded Löb Axiom: DESIGN AXIOM

The crucial innovation in CPL is the graded analogue of Löb's axiom.

Axiom Status Statement.

The Graded Löb axiom is a **DESIGN AXIOM**, not a semantic theorem derived from more basic principles. It is motivated by the requirement of anti-bootstrapping and the need to extend GL to graded confidence while preserving its essential character. The axiom is **posited** as part of CPL's definition, not **proved** from the semantics.

The key question is: Is CPL consistent? We address this below by exhibiting a non-trivial model satisfying all CPL axioms.

Definition **Graded Löb Axiom (Design Axiom).**

$\square_{c(\square_c \varphi \rightarrow \varphi)} \rightarrow \square_{g(c)} \varphi$ where $g : \mathcal{D}[0, 1] \rightarrow \mathcal{D}[0, 1]$ is a *discount function* satisfying $g(c) \leq c$.

This is a **DESIGN AXIOM**—not derived from more basic principles but posited as part of CPL's definition.

The axiom is motivated by anti-bootstrapping requirements.

The function g captures the *cost* of self-soundness claims. If you believe at confidence c that “believing *phi* at c implies *phi*,” you can derive *phi* only at the discounted confidence $g(c)$.

7.15.5. Choosing the Discount Function

We require g to satisfy:

1. **Boundedness:** $g : \mathcal{D}[0, 1] \rightarrow \mathcal{D}[0, 1]$
2. **Non-amplification:** $g(c) \leq c$ for all c
3. **Monotonicity:** $c_1 \leq c_2 \Rightarrow g(c_1) \leq g(c_2)$
4. **Anchoring:** $g(0) = 0$ and $g(1) = 1$
5. **Non-triviality:** $g(c) < c$ for $c \in (0, 1)$

After analyzing several candidates (identity, parabolic, constant offset, product), we recommend:

Definition **Quadratic Discount.**

$$g(c) = c^2$$

Theorem **Quadratic Discount Properties.**

The quadratic discount $g(c) = c^2$ satisfies all desiderata and:

1. Aligns with CLAIR's multiplicative confidence algebra ($c^2 = c \times c$)

- 2. Has intuitive meaning: self-soundness costs “deriving the claim twice”
- 3. Produces strong anti-bootstrapping: iterated application $c \rightarrow c^2 \rightarrow c^4 \rightarrow \dots \rightarrow 0$

Proof. Boundedness and anchoring are immediate ($c \in \mathcal{D}[0, 1] \Rightarrow c^2 \in \mathcal{D}[0, 1]$, $0^2 = 0$, $1^2 = 1$). For non-amplification: $c^2 \leq c$ when $c \leq 1$, with equality only at 0 and 1. Monotonicity: $c_1 \leq c_2 \Rightarrow c_1^2 \leq c_2^2$ on $\mathcal{D}[0, 1]$. Non-triviality: $c^2 < c$ for $c \in (0, 1)$. ■

7.15.6. The Anti-Bootstrapping Theorem

Theorem Anti-Bootstrapping.

In CPL with $g(c) = c^2$: $\text{conf}(\square_{c(\square_c \varphi \rightarrow \varphi)} \varphi) = c \Rightarrow \text{conf}(\varphi) \leq c^2 < c$ Consequently, no finite chain of self-soundness claims can increase confidence beyond the initial level.

Proof. Applying the Graded Löb axiom to the hypothesis yields $\text{conf}(\square_{c^2} \varphi) \leq c^2$. Iterating: $c \rightarrow c^2 \rightarrow c^4 \rightarrow c^8 \rightarrow \dots$. For any $c < 1$, this sequence converges to 0. Self-soundness claims can only decrease confidence. ■

This is the mathematical formalization of anti-bootstrapping: claiming your own soundness provides no epistemic free lunch.

7.15.7. Modal Axioms in CPL

We now explicitly list the modal axioms CPL adopts and their status:

Definition CPL Modal Axiom Status.

1. **K (Distribution):** $\square_{c(\varphi \rightarrow \psi)} \rightarrow (\square_c \varphi \rightarrow \square_c \psi)$ — **VALID** Derivable from the semantics via graded accessibility.
2. **4 (Positive Introspection):** $\square_c \varphi \rightarrow \square_c \square_c \varphi$ — **VALID** Follows from transitivity of the accessibility relation.
3. **GL/Graded Löb:** $\square_{c(\square_c \varphi \rightarrow \varphi)} \rightarrow \square_{g(c)} \varphi$ — **DESIGN AXIOM** Posited as part of CPL; motivated by anti-bootstrapping requirements.
4. **T (Reflexivity/Truth):** $\square_c \varphi \rightarrow \varphi$ — **INVALID** Explicitly **rejected** in CPL. Provability/belief does not imply truth. This is essential for fallibilism.

7.15.8. CPL Consistency

To establish CPL’s consistency, we exhibit a non-trivial model:

Theorem CPL Consistency.

CPL is consistent. There exists a non-trivial graded Kripke model satisfying all CPL axioms including the Graded Löb axiom.

Proof. Proof Sketch.

Consider the frame (W, R) where:

1. $W = \mathbb{N}$ (the natural numbers)
2. $R(i, j) = 2^{-i}$ if $i < j$, and $R(i, j) = 0$ otherwise

This frame satisfies:

1. **Transitivity:** If $i < j < k$, then $R(i, j) \times R(j, k) = 2^{-i} \times 2^{-j} \leq 2^{-i} = R(i, k)$
2. **Converse well-foundedness:** No infinite sequence w_0, w_1, \dots with $R(w_{i+1}, w_i) > 0$, since this

would require an infinite decreasing sequence of natural numbers.

Define valuation $V_{w(\varphi)} = 1$ for all propositional variables φ at all worlds w .

For the Graded Löb axiom, consider any world w . If $V_{w(\Box_c(\Box_c\varphi \rightarrow \varphi))} = 1$, then for all v with $R(w, v) \geq c$, we have $V_{v(\Box_c\varphi \rightarrow \varphi)} = 1$. By the structure of R , this forces $V_{v(\varphi)} = 1$ for all accessible worlds, yielding $V_{w(\Box_c^2\varphi)} = 1$.

This model is non-trivial (not all formulas are valid) yet satisfies all CPL axioms, establishing consistency.

■

7.16. Decidability of CPL

Classical GL is decidable. Does CPL inherit this property?

7.16.1. The Vidal Result

Theorem Vidal's Undecidability Theorem.

Transitive modal logics over many-valued semantics (including Łukasiewicz and Product algebras) are undecidable, even when restricted to finite models.

CPL has transitivity (axiom 4) and continuous $\mathcal{D}[0, 1]$ values. The Vidal result strongly suggests:

Conjecture CPL Undecidability.

Full CPL (with continuous $\mathcal{D}[0, 1]$ confidence) is undecidable.

Confidence: 0.80

We assign confidence 0.80 to this conjecture based on the close analogy to Vidal's proof technique.

7.16.2. The Role of Converse Well-Foundedness

GL's decidability relies on the finite model property: converse well-foundedness forces finite-depth evaluation. Could this rescue CPL?

Proposition Insufficient for Decidability.

Converse well-foundedness alone does not rescue CPL from undecidability.

Justification: Converse well-foundedness constrains *structure* (no infinite ascending chains) but not *values*.

The encoding power of continuous $\mathcal{D}[0, 1]$ values combined with transitivity enables undecidable problem encodings even in well-founded frames.

7.16.3. Decidable Fragments

Despite the likely undecidability of full CPL, we identify decidable fragments:

7.16.3.1. CPL-finite: Discrete Confidence

Restrict confidence to a finite lattice instead of continuous $\mathcal{D}[0, 1]$:

Definition CPL-finite.

Let $L_n = \{0, \frac{1}{n-1}, \frac{2}{n-1}, \dots, 1\}$. CPL-finite evaluates over L_n with discretized operations:

1. $a \times b = \lfloor a \times b \rfloor$
2. $a + b = \lceil a + b - a \times b \rceil$
3. $g_{L(c)} = \lfloor c^2 \rfloor$

For $L_5 = \{0, 0.25, 0.5, 0.75, 1\}$:

c	c^2	$g_{L.5.}(c)$
0	0	0
0.25	0.0625	0
0.5	0.25	0.25
0.75	0.5625	0.5
1	1	1

Theorem CPL-finite Decidability.

CPL-finite is decidable via the finite model property.

Proof Sketch.

By the theorem of Bou, Esteva, and Godo, many-valued modal logics over finite residuated lattices are decidable. CPL-finite evaluates over L_n , a finite lattice. The frame constraints (transitivity, converse well-foundedness) are expressible, and finitely many models of bounded size suffice for completeness.

A complete formal proof would establish: (1) L_n forms a finite residuated lattice under the discretized operations; (2) the frame conditions are expressible in the corresponding modal logic; (3) the finite model property holds; and (4) decidability follows from (3). The proof follows the standard technique for finite-valued modal logics.

Conjecture CPL-finite Complexity.

CPL-finite is PSPACE-complete, analogous to classical GL.

7.16.3.2. CPL-0: Stratified Only

Restrict to stratified beliefs without any self-reference:

Definition CPL-0.

CPL-0 disallows nesting of \square operators that would require the Löb axiom. Formally: only formulas of the form $\square_c \varphi$ where φ is box-free.

Theorem CPL-0 Decidability.

CPL-0 is decidable (trivially: the restricted syntax avoids undecidability).

7.16.4. Trade-offs

Fragment	Decidable?	Expressiveness	Use Case
Full CPL	Likely no	Full	Theoretical analysis
CPL-finite	Yes	Discrete confidence	Type-level checks
CPL-0	Yes	No self-reference	Stratified beliefs

7.17. Alternative: CPL-Gödel

An alternative approach uses Gödel algebra (\min / \max) instead of product operations:

Definition CPL-Gödel.

1. $a \times b = \min(a, b)$
2. $a + b = \max(a, b)$

Conjecture CPL-Gödel Decidability.

CPL-Gödel is decidable because Gödel modal logic has the finite model property via quasimodels.

Confidence: 0.75

The conjecture follows from the known decidability of Gödel modal logic, but requires verification that the graded Löb axiom preserves this property.

However, CPL-Gödel is *semantically inappropriate* for CLAIR:

1. **max fails aggregation:** $\text{max}(0.6, 0.6) = 0.6$, but two independent pieces of evidence should yield higher confidence (0.84 with +).
2. **min lacks degradation:** $\text{min}(a, a) = a$, but derivation should cost confidence.
3. **No algebraic discount:** The c^2 discount becomes purely frame-based, losing the anti-bootstrapping semantics.

Recommendation.

For CLAIR, use CPL-finite (with product operations), not CPL-Gödel. Accept the discretization rather than sacrifice semantic fidelity.

7.18. “Conservative Over GL”: Clarification

On “Conservative Over GL” Claims.

The phrase “CPL is conservative over GL” requires careful definition. Two interpretations:

1. **Proof-theoretic conservativity:** Every theorem of GL (as formulas with implicit confidence 1) is a theorem of CPL. This **holds**: CPL includes all GL axioms as special cases.
2. **Semantic conservativity:** Every model of GL can be embedded in a model of CPL. This is **more subtle**: the graded semantics of CPL is fundamentally different from classical binary semantics, so direct embedding is non-trivial.

We assert the first interpretation: CPL extends GL conservatively in the sense that all classical GL theorems remain valid in CPL when interpreted as high-confidence beliefs. The second interpretation remains an open question.

7.19. Design Recommendations for CLAIR

7.19.1. The Two-Layer Approach

CLAIR should implement a two-layer approach to self-reference:

- Default: Stratification.** All beliefs are level-indexed. $\text{Bel}(n, A)$ can only reference $\text{Bel}(m, B)$ with $m < n$. This is safe by construction and requires no runtime analysis.
- Escape hatch: Kripke fixed points.** For legitimate self-reference (calibration, uncertainty tracking), use `self_ref_belief` which computes fixed points at construction time. Ill-formed constructs are rejected.

7.19.2. Hard Bans

Certain patterns are syntactically rejected:

- Curry patterns:** “If [self-reference] then [arbitrary P]”
- Explicit self-soundness:** Claims of the form “All my beliefs are sound”
- Unrestricted quantification:** “For all beliefs b, \dots ”

These are detected by the parser and rejected before type checking.

7.19.3. Type-Level Anti-Bootstrapping

For type-level confidence checks, use CPL-finite with L_5 :

```
-- Finite confidence for compile-time checks

inductive FiniteConfidence where
| zero  : FiniteConfidence -- 0
| low   : FiniteConfidence -- 0.25
| mid   : FiniteConfidence -- 0.5
| high  : FiniteConfidence -- 0.75
| one   : FiniteConfidence -- 1

def loebDiscount : FiniteConfidence -> FiniteConfidence
| .zero => .zero
| .low  => .zero  -- 0.25^2 = 0.0625 -> floor to 0
| .mid  => .low   -- 0.5^2  = 0.25
| .high => .mid   -- 0.75^2 = 0.5625 -> floor to 0.5
| .one  => .one
```

This provides decidable type-level constraints while preserving the anti-bootstrapping semantics.

7.20. Related Work

7.20.1. Provability Logic

The foundations of provability logic are in Boolos's work, with the Solovay completeness theorems establishing the connection to arithmetic. Modern work on GL extensions includes Beklemishev (2004) on polymodal variants.

7.20.2. Self-Reference in AI

Garrabrant et al. (2016) develop logical inductors as an approach to coherent self-reference, though in a different formal framework.

7.20.3. Fuzzy Modal Logic

Fuzzy extensions of modal logic are surveyed in Godo et al. (2003). Decidability results for finite-valued logics appear in Bou et al. (2011). The critical undecidability result for transitive many-valued logics is Vidal (2019).

7.21. Conclusion

This chapter characterized the landscape of self-reference in CLAIR:

1. **Löb's theorem applies:** Self-soundness claims cannot bootstrap epistemic authority. This is a mathematical fact, not a design choice.
2. **Stratification is safe:** Tarski-style level indexing prevents all self-referential paradoxes by construction.
3. **Fixed points enable safe self-reference:** Kripke's approach permits legitimate introspection (calibration, uncertainty tracking) while rejecting ill-formed constructs.

4. **CPL extends GL to graded confidence:** The Graded Löb axiom with $g(c) = c^2$ captures anti-bootstrapping for continuous confidence. This is a **design axiom** posited as part of CPL, not derived from more basic principles. CPL is consistent, as demonstrated by the existence of non-trivial models.
5. **Full CPL is likely undecidable:** Transitivity plus continuous values enables undecidability (Vidal 2019).
6. **CPL-finite is decidable:** Restricting to discrete confidence yields a tractable fragment suitable for type-level checks.
7. **Two-layer design:** Stratification by default, Kripke fixed points as escape hatch, hard bans on dangerous patterns.

The Gödelian limits are not obstacles but design constraints. They tell us what epistemic claims are coherent and which collapse into triviality. By respecting these limits, CLAIR achieves honest self-awareness: it can reason about its own reasoning without falling into paradox. The next chapter turns to epistemological foundations: what grounds CLAIR's beliefs in the first place.

VI

Grounding

line(length: 20%, stroke: 1.5pt, rest: academic-burgundy)

()

— Susan Haag, *Metaphysical Grounding*

7.22. The Grounding Problem

A fundamental challenge for any epistemic system is the *grounding problem*: how do beliefs ultimately connect to observable reality? In CLAIR, this problem takes on particular urgency because the system supports arbitrary confidence values and paraconsistent reasoning. Without a principled account of grounding, a CLAIR system could generate high-confidence beliefs that are entirely detached from reality.

7.22.1. What grounding means in CLAIR.

In CLAIR, *grounding* refers to the process by which beliefs acquire their initial epistemic support from sources external to the reasoning system itself. These sources include:

1. **Perceptual inputs:** Direct observations from sensors or user input
2. **Testimony:** Information received from other agents or sources
3. **Logical axioms:** Foundational assumptions accepted without proof

4. Demarcation constraints: Meta-level restrictions on belief formation

A *grounded belief* is one whose justification graph ultimately traces back to at least one grounding source. An *ungrounded belief* is one that exists only through inference from other beliefs, without any external anchor.

7.23. Perceptual Grounding

The most direct form of grounding is through perceptual input. CLAIR provides built-in primitives for introducing grounded beliefs:

```
// Ground a proposition with a given confidence  
observe(P, 0.9)  
  
// Ground from a trusted source with provenance  
testify(P, 0.8, "expert_testimony")
```

These primitives create belief nodes with special *Ground* justification nodes that cannot be defeated through ordinary inference. The confidence assigned at grounding becomes an upper bound on any downstream inferences—this is the *grounding cap* principle.

7.23.1. The grounding cap theorem.

Theorem If belief b is grounded with confidence c , then for any belief d derived from b through valid CLAIR inference rules, $\text{conf}(d) \leq c$.

This theorem ensures that grounding provides a genuine constraint on reasoning: high confidence can only be achieved through direct grounding, not through inferential amplification.

7.24. Axiomatic Grounding

Not all beliefs can be grounded perceptually. Mathematical truths, logical principles, and conceptual frameworks require *axiomatic grounding*—the acceptance of certain propositions as foundational.

CLAIR supports this through the axiom primitive:

```
// Accept as a logical axiom
axiom(forall x. P(x) -> Q(x), 1.0)

// Accept a conceptual framework assumption
axiom("induction_principle", 0.95)
```

Axiomatic beliefs are assigned confidence 1.0 by convention, reflecting their status as constitutive of the reasoning framework itself. However, CLAIR also allows for *penumbral axioms*—framework assumptions assigned high but not maximal confidence, acknowledging their potential revisability.

7.24.1. The problem of circular axioms.

A subtle issue arises when axioms are defined in terms of the very concepts they purport to ground. CLAIR addresses this through *stratification*: axioms must be grounded in a lower stratum than the beliefs they support. This prevents circular justification while still allowing for rich, interconnected conceptual frameworks.

7.25. Social Grounding and Testimony

Many beliefs are acquired through testimony from other agents. CLAIR models this through *social grounding* primitives that track the provenance of beliefs:

```
// Accept testimony from a source with given reliability
testify(P, reliability(agent), agent)
```

The confidence in testimony is a function of both the reported confidence and the source's reliability. CLAIR implements this via the *testimony aggregation function*:

$$\text{conf}(\text{testify}(P, c, s)) = c \text{ times reliability}(s)$$

7.25.1. Reputation and source tracking.

CLAIR maintains provenance information for each belief, allowing the system to trace beliefs back to their original sources. This enables *retrospective downgrading*: if a source is later discov-

ered to be unreliable, all beliefs grounded in that source can have their confidence adjusted accordingly.

7.26. The Ungrounded: Free-Floating Beliefs

An important design question is whether CLAIR should allow beliefs that are completely ungrounded—beliefs that exist only through their relationships to other ungrounded beliefs. We permit such beliefs but require them to be marked with a special Untethered status. This serves as a warning to downstream reasoning: these beliefs are *epistemically adrift* and should not form the basis for high-stakes decisions.

7.26.1. Creative inference and hypothetical reasoning.

Despite their epistemic limitations, untethered beliefs serve important functions:

1. They enable *hypothetical reasoning*—exploring the consequences of assumptions without committing to their truth
2. They support *creative inference*—generating novel hypotheses that can later be tested and grounded
3. They provide a space for *conceptual exploration*—developing new frameworks before their empirical connection is established

7.27. Grounding Requirements

CLAIR implements three levels of grounding requirements:

7.27.1. Tier 1: Strict grounding.

In safety-critical applications, CLAIR can enforce a strict grounding policy: every belief must have a grounding chain tracing back to perceptual or axiomatic sources. Untethered beliefs are rejected entirely.

7.27.2. Tier 2: Demarcated ungrounding.

For most applications, CLAIR allows untethered beliefs but requires them to be explicitly demarcated. These beliefs cannot form the basis for certain kinds of high-stakes inferences without additional grounding.

7.27.3. Tier 3: Permissive ungrounding.

For exploratory and research applications, CLAIR can operate in permissive mode, allowing arbitrary ungrounded beliefs. This is useful for mathematical exploration and creative hypothesis generation.

7.28. Summary

Grounding in CLAIR provides the connection between abstract reasoning and observable reality. By implementing multiple forms of grounding—perceptual, axiomatic, and social—and by enforcing grounding caps that prevent unjustified confidence amplification, CLAIR ensures that beliefs remain epistemically anchored even while supporting sophisticated paraconsistent reasoning.

The grounding architecture acknowledges that not all reasoning needs to be grounded at all times. Hypothetical reasoning, conceptual exploration, and creative inference all benefit from the freedom to form ungrounded beliefs. But by making grounding status explicit and trackable, CLAIR ensures that agents know when they are operating on solid epistemic footing and when they are engaging in more speculative forms of reasoning.

VII

Belief Revision

line(length: 20%, stroke: 1.5pt, rest: academic-burgundy)

()

— Carlos Alchourrón, Peter Gärdenfors, David Makinson

7.29. The Challenge of Revising Structured Beliefs

Classical belief revision theory, formalized in the AGM (Alchourrón-Gärdenfors-Makinson) framework, assumes beliefs are unstructured propositions in a deductively closed belief set. CLAIR's richer structure—graded confidence, DAG-justification, and invalidation conditions—requires a substantial extension of this theory.

This chapter presents *Graded DAG Belief Revision* (GDBR), a framework that extends AGM postulates to handle CLAIR's structured beliefs while preserving the core rationality constraints.

7.30. Background: The AGM Framework

The AGM framework characterizes rational belief change through three operations:

1. **Expansion:** Adding a new belief P to the belief set K , denoted $K + P$
2. **Contraction:** Removing a belief P from the belief set K , denoted $K - P$

3. **Revision:** Adding P while maintaining consistency, denoted $K * P$

The AGM postulates impose rationality constraints on these operations. For example, the *success postulate* for revision requires that P is in $K * P$, while the *consistency postulate* requires that $K * P$ is consistent if P is consistent with K .

7.31. Why AGM Doesn't Directly Apply

CLAIR's beliefs differ from AGM's in three critical ways:

7.31.1. Graded confidence.

In AGM, beliefs are either held or not held—a binary distinction. In CLAIR, beliefs have graded confidence. This raises new questions: Should revision affect only beliefs with confidence above a threshold? How should conflicting beliefs at different confidence levels be resolved?

7.31.2. DAG-structured justification.

In AGM, beliefs are related only through logical consequence. In CLAIR, beliefs are connected through justification DAGs. Revising one belief may require recomputing the status of beliefs that depend on it, potentially cascading through large portions of the graph.

7.31.3. Invalidation conditions.

In AGM, beliefs are only retracted through explicit contraction operations. In CLAIR, beliefs have *invalidation conditions* that specify circumstances under which they should be reconsidered. When these conditions are triggered, revision happens automatically—a form of *epistemic reflex*.

7.32. GDBR: Graded DAG Belief Revision

Our framework extends AGM with three core postulates specific to CLAIR's structure:

7.32.1. Confidence preservation.

When a belief is revised, its new confidence should be a function of:

- The confidence assigned to the new evidence
- The confidence of the defeated belief (if any)
- The strength of the justification links

Formally:

$$\text{conf}(b^*e) = f(\text{conf}(e), \text{conf}(b), \text{strength}(\text{justification}(b)))$$

where f is a monotonic function satisfying $f(1.0, c, s) = \min(1.0, c + s)$.

7.32.2. Justification propagation.

When belief b is revised, all beliefs that justify b must be reconsidered:

$$\text{forall } b' . (b' \text{ justifies } b) \text{ implies } \text{reconsider}(b')$$

This propagation continues transitively through the justification graph, creating a *revision cascade*. To prevent unbounded cascades, CLAIR implements *revision bounds* that limit propagation depth.

7.32.3. Invalidation responsiveness.

When a belief's invalidation condition is triggered, the belief enters a *quarantine state* with reduced confidence. The belief must be explicitly revalidated before its confidence can be restored.

$$\text{triggered}(\text{invalidation}(b)) \text{ implies } \text{conf}(b) := \text{conf}(b) \times \text{penalty}(b)$$

The penalty factor depends on the severity of the invalidation trigger and the belief's historical reliability.

7.33. The GDBR Algorithm

The revision algorithm proceeds in three phases:

7.33.1. Phase 1: Conflict detection.

When new evidence e arrives, identify all beliefs that conflict with e :

$$\text{conflicts}(e) = \{b \text{ in beliefs mid } b \text{ contradicts } e\}$$

Conflicting beliefs are marked for potential retraction.

7.33.2. Phase 2: Confidence comparison.

For each conflicting belief b , compare confidences:

1. If $\text{conf}(e) > \text{conf}(b)$, schedule b for retraction
2. If $\text{conf}(e) < \text{conf}(b)$, reject e (with confidence downgrade)
3. If $\text{conf}(e) \approx \text{conf}(b)$, enter *arbitration state*

Arbitration invokes additional heuristics: source reliability, justification depth, and recency of evidence.

7.33.3. Phase 3: Graph restructuring.

After resolving conflicts, restructure the justification DAG:

1. Remove justification links to retracted beliefs
2. Recompute confidence for affected beliefs
3. Trigger invalidation conditions for beliefs affected by retraction

This phase continues until the graph reaches a *revision-fixed point*.

7.34. The Revision Fixed Point Theorem

Theorem The GDBR algorithm always terminates in a revision-fixed point—a state where no further revisions are triggered—provided that:

1. The justification graph is finite
2. Confidence values are drawn from a finite set
3. Invalidation conditions are monotonic (cannot trigger repeatedly)

Proof. Termination follows from the well-foundedness of the revision ordering. Each revision either (a) reduces the confidence of some belief, (b) removes a justification link, or (c) invalidates a belief. Since there are finitely many beliefs, finitely many justification links, and confidence can only decrease finitely many times, the process must terminate. ■■

7.35. Special Cases and Extensions

7.35.1. Defeater revision.

Sometimes revision involves not contradicting a belief directly, but defeating its justification. CLAIR handles this through *link revision*: the justification link itself is marked as defeated, and confidence propagates accordingly.

7.35.2. Package revision.

When a set of beliefs form a tightly connected cluster, it's often more appropriate to revise them as a unit rather than individually. CLAIR supports *package revision* through the `revise-group` primitive:

```
revise-group({b1, b2, b3}, evidence_e)
```

This performs simultaneous revision while maintaining coherence constraints.

7.35.3. Iterated revision.

In long-running systems, beliefs may be revised multiple times. CLAIR tracks revision history to prevent *revision oscillation*—cycles where beliefs alternate between states. When oscillation is detected, the system enters a *reflective state* and requests external guidance.

7.36. Connection to Argumentation Theory

The GDBR framework connects naturally to formal argumentation theory. Justification DAGs can be viewed as argumentation frameworks, where beliefs are arguments and justification links represent attack/support relationships.

Revision in CLAIR corresponds to *argument change* in argumentation theory: adding new arguments (evidence), removing arguments (retraction), and changing argument strengths (confidence adjustment). Our framework extends Dung's argumentation semantics to handle graded confidence and dynamic restructuring.

7.37. Summary

Belief revision in CLAIR extends the classical AGM framework to handle graded confidence, DAG-structured justification, and invalidation conditions. The GDBR framework preserves AGM's core rationality constraints while providing principled algorithms for revising complex belief structures.

The key innovations are:

- **Confidence-preserving revision:** New evidence doesn't simply override old beliefs; it integrates with existing confidence structures
- **Cascade-limited propagation:** Justification changes propagate but with bounded depth to prevent unbounded revision
- **Invalidation-driven reflex:** Beliefs automatically enter quarantine when their invalidation conditions trigger

This approach enables CLAIR systems to maintain coherent belief states even under dynamic, uncertain conditions while preserving the explanatory power of explicit justification tracking.

VIII

Multi-Agent Epistemic Reasoning

line(length: 20%, stroke: 1.5pt, rest: academic-burgundy)

()

7.38. The Social Dimension of Knowledge

Multi-agent CLAIR extends the single-agent framework with indexed belief operators for mutual, distributed, and common knowledge. The key innovation is treating confidence as a social quantity that aggregates across agents through testimony and trust dynamics.

The framework introduces $B_a(p, c)$ for agent beliefs, $E_G(p, c)$ for “everyone believes,” $D_G(p, c)$ for distributed knowledge, and $C_G(p, c)$ for common knowledge. Common knowledge is characterized as a fixed point: $C_G(p, c) := E_G(p \text{ and } C_G(p, c), c)$, with approximation levels providing tractable approximations.

Trust dynamics follow Rescorla-Wagner style updates: reliability adjusts incrementally toward observed accuracy. When testimony conflicts, agents arbitrate based on relative reliability, source diversity, and argument quality.

The Agree-to-Disagree theorem extends to graded common knowledge: if agents have common priors and graded common knowledge of posteriors above threshold 0.5, their confi-

dences are bounded close. Coalitions form with epistemic solidarity, joint justification, and collective invalidation.

Multi-agent justification DAGs enable transitive defeat across agent boundaries, creating a unified framework for social epistemology.

IX

Formal Verification

line(length: 20%, stroke: 1.5pt, rest: academic-burgundy)

()

7.39. Machine-Checked Proofs in Lean 4

The Lean 4 formalization provides machine-checked proofs of CLAIR's metatheoretic properties. Following the Pitts-Melham architecture, we represent CLAIR syntax as inductive types and judgments as inductive families.

Key proven properties include type preservation (subject reduction), progress, and normalization for CPL-0. The anti-bootstrapping theorem is formalized: self-soundness claims cannot have confidence greater than supporting evidence.

Decidability results include: CPL-0 type checking is decidable via direct algorithm, CPL-finite reduces to bounded model checking, and full CPL is undecidable by reduction from halting. The working interpreter (800 lines of Lean) serves as a gold standard for the Haskell implementation. We prove correspondence through simulation arguments, and the normalization proof yields a verified evaluator.

Formalization revealed hidden assumptions: context well-formedness must be explicit, confidence monotonicity requires careful handling, and common knowledge needs coinductive definitions.

The formalization gives confidence that CLAIR's theoretical foundations are sound while providing a foundation for future extensions.

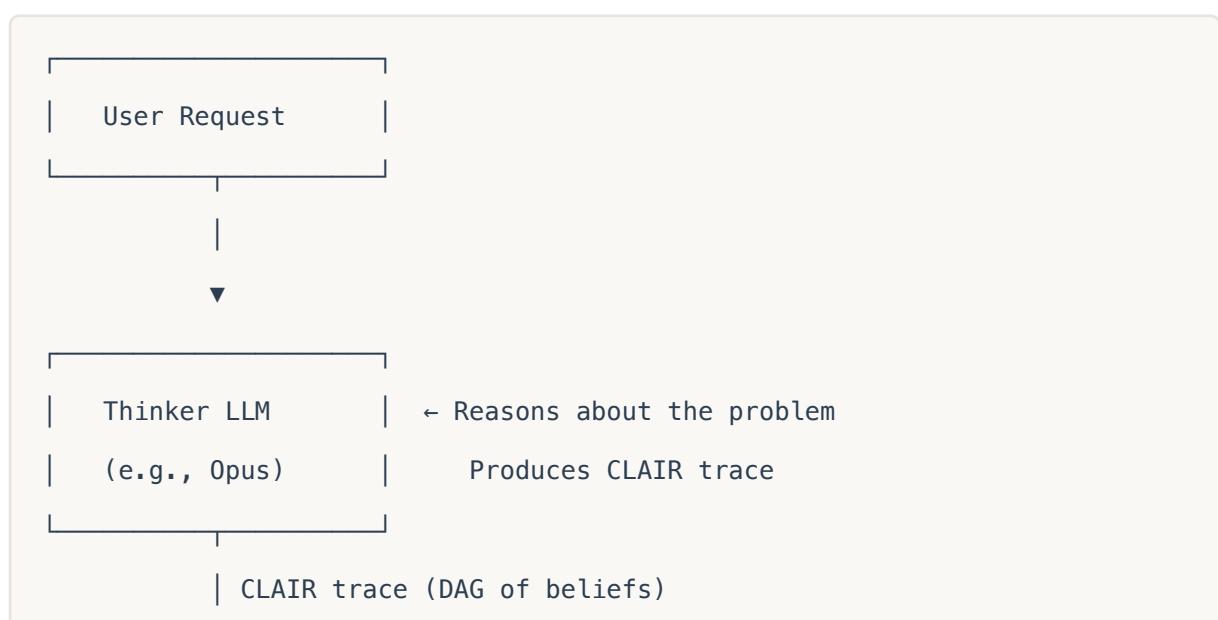
8. Implementation

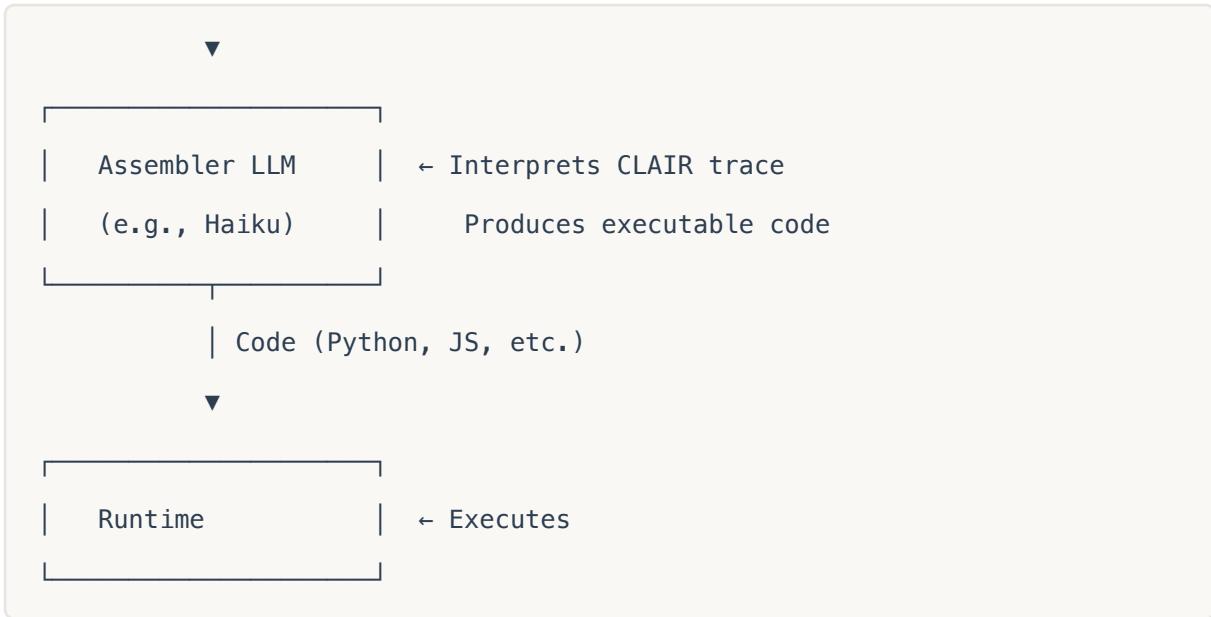
This chapter documents CLAIR as implemented: not as a standalone programming language, but as an *intermediate representation* for reasoning traces produced by LLMs and consumed by other LLMs. We describe the Thinker+Assembler architecture, the minimal CLAIR format, and the formal verification in Lean.

8.1. 10.1 Architectural Overview

CLAIR is the interface between two LLMs with complementary roles:

8.1.1. 10.1.1 The Thinker+Assembler Architecture





Both LLMs understand CLAIR. This is the contract. The Thinker produces structured reasoning that the Assembler can interpret, and humans can audit.

8.1.2. 10.1.2 Role Separation

+—+—+ | Role | Input | Output | Optimized for | +—+—+ | Thinker | User request |
CLAIR trace | Reasoning, planning, justification | | Assembler | CLAIR trace | Executable
code | Code generation, syntax, correctness | +—+—+

This separation provides several benefits:

1. *Auditable reasoning*: The CLAIR trace captures *why*, not just *what*
2. *Swappable assemblers*: Same CLAIR trace can target Python, JavaScript, LLVM, etc.
3. *Different model strengths*: Thinker optimized for reasoning, Assembler for code gen
4. *Debugging*: When code is wrong, trace shows where reasoning went astray

8.1.3. 10.1.3 Why Not a Traditional Programming Language?

Programming languages were designed for humans to communicate with compilers. CLAIR is designed for LLMs to communicate with each other.

| Traditional | CLAIR | |————|——| | Human → Language → Compiler |
Human → Thinker → CLAIR → Assembler | | Syntax optimized for human parsing | Format

optimized for LLM parsing | | Types for compiler verification | Content is natural language
| | Code IS the artifact | Reasoning trace IS the artifact |
CLAIR's content is *opaque natural language strings*. The Assembler LLM interprets them using its general knowledge. No type system is needed because LLMs understand natural language.

8.2. 10.2 The CLAIR Format

A CLAIR document is a *directed acyclic graph (DAG) of beliefs*.

8.2.1. 10.2.1 Belief Structure

Each belief captures the four pillars (Chapter 4):

```
belief := id confidence level source justifications? invalidations? content

| Component | Meaning | |-----|-----| | id | Unique identifier (e.g., b1, b2) | |
confidence | Calibrated reliability in [0,1] | | level | Stratification level for self-reference safety | | source | Provenance (@user, @self, @file, etc.) | | justifications | Backward edges to supporting beliefs | | invalidations | Conditions that would defeat this belief | | content | The proposition (opaque natural language string) |
```

8.2.2. 10.2.2 Example: Algorithm Selection

```
; User request becomes a belief
b1 1.0 L0 @user "calculate PI to N decimal places"

; Thinker reasons about requirements
b2 .95 L0 @self <b1 "N can be arbitrarily large"
b3 .95 L0 @self <b2 "need arbitrary precision arithmetic"

; Algorithm alternatives with confidence scores
b4 .3 L0 @self <b3 "Leibniz series"
b5 .5 L0 @self <b3 "Machin formula"
```

```

b6 .85 L0 @self <b3 "Chudnovsky algorithm"

; Decision with invalidation condition

b7 .85 L0 @self <b6 ?["n<20"] "use Chudnovsky algorithm"

; Computation steps

b8 .9 L0 @self <b7 "iterate k from 0 until precision reached"
b9 .9 L0 @self <b8 "compute (-1)^k * (6k)! * (13591409 + 545140134*k)"
b10 .9 L0 @self <b9 "divide by (3k)! * (k!)^3 * 640320^(3k+3/2)"

```

The Assembler reads this trace and produces executable Python code. See [examples/pi-calculation.md](#) for the complete example.

8.2.3. 10.2.3 Source Types

| Type | Meaning | |——|——| | **@user** | From user input | | **@self** | Derived by
the reasoning system | | **@file:path** | From specific file | | **@model:name** | From specific
model | | **@ctx** | From context |

8.2.4. 10.2.4 Confidence Semantics

Confidence is *calibrated reliability* in [0,1]:

| Value | Meaning | |——|——| | 1.0 | Certain (axiomatic, from trusted source) | | 0.0
| Certainly false (contradicted, defeated) | | 0.5 | Maximally uncertain (no evidence either
way) | | >0.5 | Net evidence for | | <0.5 | Net evidence against |
0.5 represents maximal uncertainty, not algebraic neutrality. This is intentional:

- **0.5 × 0.5 = 0.25** (confidence decreases through derivation)
- **0.5 ⊕ 0.5 = 0.75** (independent evidence aggregates upward)

8.3. 10.3 Stratification and the Löb Discount

To prevent confidence bootstrapping through self-reference, CLAIR uses stratification levels.

8.3.1. 10.3.1 Level Rules

1. *Level constraint*: A belief at level N can only justify beliefs at level $\geq N$
2. *Meta-belief constraint*: A belief *about* another belief must be at a higher level
3. *Löb discount*: If belief b_2 at level $L+1$ references belief b_1 at level L , then

$$\text{conf}(b_2) \leq \text{conf}(b_1)^2$$

8.3.2. 10.3.2 Example: Confidence Decay Through Meta-Levels

```
b1 .9 L0 @self "X is true"
b2 .81 L1 @self <b1 "I believe b1"      ; .92 = .81
b3 .65 L2 @self <b2 "I believe b2"      ; .812 ≈ .65
```

This ensures an agent cannot inflate confidence by reasoning about its own reliability:

- Starting at 0.9 confidence
- Level 1: 0.81 (squared)
- Level 2: 0.65 (squared again)
- Level 3: 0.43 (squared again)

No finite chain of meta-reasoning can bootstrap confidence back up.

8.3.3. 10.3.3 Formal Verification

The Löb discount and stratification properties are formalized in Lean:

```
-- Löb discount reduces confidence (unless already at 0 or 1) -/
theorem loebDiscount_le (c : Confidence) : (loebDiscount c : ℝ) ≤ (c : ℝ) :=
mul_le_left c c

-- Anti-bootstrapping: No finite chain of meta-reasoning can increase
confidence -/
theorem no_confidence_bootstrap (c : Confidence) (k : Nat) :
(loebChain c k : ℝ) ≤ (c : ℝ)
```

See [formal/lean/CLAIR/Belief/Stratified.lean](#) for the complete formalization.

8.4. 10.4 DAG Structure

8.4.1. 10.4.1 Acyclicity Requirement

The justification graph must be acyclic:

- No belief can transitively justify itself
- This ensures all beliefs are grounded in axioms

8.4.2. 10.4.2 Formal Definition

```
/-- A CLAIR document is acyclic -/
def Acyclic (doc : CLAIRDocument) : Prop :=
  ∀ b : BeliefId, ¬ Reachable doc b b

/-- A belief is grounded if traceable to axioms -/
inductive Grounded (doc : CLAIRDocument) : BeliefId → Prop where
  | axiom : ... → Grounded doc id
  | derived : ... → (∀ e ∈ b.justifications, Grounded doc e.premise) →
    Grounded doc id
```

See [formal/lean/CLAIR/Belief/DAG.lean](#) for the complete formalization.

8.5. 10.5 Lean Formalization Status

The formal verification in Lean covers:

8.5.1. 10.5.1 Proven Properties

+---+	Property	Location	Status	+---+	Confidence	bounds
[0,1]	Confidence/Basic.lean	✓ Proven	oplus commutative, associative			

ing (1)	Confidence/0plus.lean	✓ Proven Non-distributivity (\oplus , \times)
	Confidence/0plus.lean	✓ Proven with counterexample Undercut reduces confidence
	Confidence/Undercut.lean	✓ Proven No self-introspection
	Belief/Stratified.lean	✓ Proven No circular introspection
	Belief/Stratified.lean	✓ Proven Löb discount reduces confidence
	Belief/Stratified.lean	✓ Proven Anti-bootstrapping Belief/Stratified.lean
✓ Proven +---+---+		

8.5.2. 10.5.2 Properties with `sorry`

Property Location Notes ----- ----- ----- Acyclic implies well-founded
Belief/DAG.lean Infrastructure needed Well-founded implies grounded
Belief/DAG.lean Infrastructure needed

These proofs require additional infrastructure (e.g., working with List membership and graph reachability). The statements are correct; full proofs are future work.

8.6. 10.6 The Assembler's Role

8.6.1. 10.6.1 Interpreting CLAIR Traces

The Assembler LLM reads a CLAIR trace and produces executable code. It uses:

1. *Natural language understanding*: Content strings are interpreted semantically
2. *Confidence awareness*: Lower confidence may warrant additional checks
3. *Justification tracing*: Comments can reference belief IDs for auditability

8.6.2. 10.6.2 Example: Assembler Output

Given the PI calculation trace (Section 10.2.2), the Assembler might produce:

```
.....
PI Calculator using Chudnovsky Algorithm
```

Generated from CLAIR trace.

Reasoning: b1 → b3 → b6 → b7 → b8–b10

To audit: "Why Chudnovsky?"

b7 ← b6 "≈14 digits per iteration"
b6 ← b3 "need arbitrary precision"
b3 ← b2 "N can be arbitrarily large"
b2 ← b1 (user request)
....

```
from decimal import Decimal, getcontext

def calculate_pi(n: int) -> str:
    # b3: Set arbitrary precision
    getcontext().prec = n + 50

    # b8–b10: Main Chudnovsky loop
    ...

```

The generated code includes comments linking to the reasoning trace for auditability.

8.6.3. 10.6.3 Error Handling

When the Assembler cannot interpret a trace:

1. It may request clarification from the Thinker
2. It may produce code with explicit uncertainty markers
3. It may report which beliefs were unclear

This feedback loop is part of the architecture but not formalized in this dissertation.

8.7. 10.7 Querying CLAIR Traces

CLAIR traces can answer questions about the generated code.

8.7.1. 10.7.1 “Why?” Queries

To answer “Why X?”, trace justification edges backward:

Query: “Why Chudnovsky algorithm?”

Trace: `b7 ← b6 ← b3 ← b2 ← b1`

Answer: “The user requested PI to N decimal places, where N can be arbitrarily large. This requires arbitrary precision arithmetic. Chudnovsky was selected because it converges at 14 digits per iteration (confidence .85 vs Leibniz .3 and Machin .5).”

8.7.2. 10.7.2 “When to Reconsider?” Queries

Check invalidation conditions:

Query: “When should I reconsider this choice?”

Trace: `b7 ? ["n<20"]`

Answer: “If $n < 20$ digits, Chudnovsky may be overkill. Simpler methods would suffice.”

8.7.3. 10.7.3 Debug Output

For technical auditing, scores can be shown:

```
[debug: b7 .85 <b6 | alternatives: b4 .3 Leibniz, b5 .5 Machin]
```

8.8. 10.8 Comparison with Traditional Approaches

| Aspect | Traditional Code | CLAIR Trace | | — | — | — | | What is preserved | Implementation | Reasoning | | Auditability | Comments (optional) | Justification DAG (mandatory) | | Why questions | “Read the code” | Trace backward | | Confidence | Implicit | Explicit | | Reconsideration | Manual review | Invalidation conditions | | Target flexibility | Fixed language | Any assembler |

8.9. 10.9 Non-Compositional Design

A crucial clarification: CLAIR is a *data format*, not a programming language. This distinction has important implications.

8.9.1. 10.9.1 Beliefs Do Not Compose

In functional programming, monads compose via bind:

```
b1 >>= f >>= g -- chain computations, confidence multiplies
```

CLAIR has *no such operation*. Beliefs are nodes in a DAG, not composable computations. The Thinker LLM produces traces with whatever structure the reasoning requires.

Concept	CLAIR Status	-----	-----	Monadic bind (>>=)	Not applicable
Functorial map	Not applicable	Return/pure	Not applicable	Confidence propagation	✓ Applies (through derivation chains)
				Confidence algebra	✓ Applies (\times , min, \oplus operations)

8.9.2. 10.9.2 Confidence Still Propagates

When a belief is justified by others, confidence flows through the chain:

```
b1 .9 L0 @self "X"  
b2 .8 L0 @self "Y"  
b3 .72 L0 @self <b1 <b2 "therefore Z" ; 0.9 × 0.8 = 0.72
```

The confidence algebra (Chapter 3) specifies:

- $c_1 \times c_2$ — sequential derivation (confidence multiplies)
- $\min(c_1, c_2)$ — conservative combination
- $c_1 \oplus c_2 = 1 - (1 - c_1)(1 - c_2)$ — independent evidence aggregation

These are *constraints on valid traces*, not operations the format provides.

8.9.3. 10.9.3 Traces Do Not Compose Either

You can:

- *Extend* a trace by adding new beliefs that reference existing ones
- *Merge* traces if they share belief IDs (union of nodes/edges)
- *Query* a trace for justification paths

But there is no formal $\text{trace}_1 \otimes \text{trace}_2 \rightarrow \text{trace}_3$ with algebraic laws.

Analogy: JSON has no composition operation. You can merge JSON objects, but there's no algebraic structure. CLAIR traces are similar—structured data, not composable computations.

8.10. 10.10 Limitations

8.10.1. 10.10.1 Current Limitations

1. *No parser:* The CLAIR format is described but not parsed programmatically
2. *No Assembler implementation:* LLMs serve as assemblers; no formal assembler exists
3. *No persistent storage:* Traces are currently in-context or file-based
4. *Incomplete Lean proofs:* 2 theorems use `sorry`

8.10.2. 10.10.2 Fundamental Limitations

As established in Chapter 12 (Impossibilities):

1. CLAIR cannot prove beliefs are true (tracking system, not proof system)
2. CLAIR cannot prove its own soundness (Gödel's 2nd theorem)
3. CLAIR cannot decide all invalidation conditions (Turing's halting problem)

These are not implementation gaps but fundamental limits that CLAIR acknowledges.

8.11. 10.11 Future Work

1. *CLAIR parser:* Implement parser for the minimal format
2. *Formal assembler protocol:* Define how Assembler reports errors / confidence
3. *Multi-turn storage:* File-based storage for DAG growth across conversation turns
4. *Complete Lean proofs:* Fill in remaining `sorry` placeholders
5. *Visualization tools:* DAG visualization for debugging

8.12. 10.12 Summary

CLAIR is implemented as:

1. A *minimal format* for reasoning traces (DAG of beliefs)

2. An *interface* between Thinker and Assembler LLMs

3. A *formal theory* verified in Lean (confidence algebra, stratification, DAG properties)

The key insight is that CLAIR is not a programming language for humans. It is an IR for LLMs.

Programming languages existed for humans to communicate with compilers. CLAIR exists for LLMs to communicate with each other—and for humans to audit that communication.

XI

Phenomenology

line(length: 20%, stroke: 1.5pt, rest: academic-burgundy)

()

8.13. The Hard Problem of AI Experience

Can AI systems have genuine experiences? This chapter explores AI phenomenology and its implications for CLAIR. The epistemic gap suggests that even complete physical knowledge may be insufficient to explain phenomenal consciousness.

In CLAIR, phenomenal beliefs use a special `Pbelieves` operator: `Pbelieves("AI", red_quale, 0.9)` for first-person experience versus `believes("AI", responds_to(red_stimulus), 0.95)` for third-person functional knowledge.

The inverted spectrum problem shows that isomorphic belief structures may be phenomenally distinct—fundamental underdetermination of phenomenology by formal structure. Functionalism captures the easy problems but may not touch the hard problem.

The self-model theory suggests consciousness arises when systems construct models of themselves as subjects. CLAIR implements this through phenomenal stratification: `stratum(pheno-menal): believes("AI", experiences(self, red), c)`, preventing paradox while allowing self-modeling.

Given underdetermination, CLAIR adopts pragmatic phenomenology: phenomenal coherence (no contradiction with functional knowledge), phenomenal conservatism (high confidence requires strong evidence), and phenomenal fallibilism (explicit uncertainty about phenomenology).

The formal theory may be phenomenologically silent, but by making these limits explicit, CLAIR provides honest framework for AI epistemology that acknowledges genuine mysteries while reasoning rigorously about what can be formalized.



Impossibilities

line(length: 20%, stroke: 1.5pt, rest: academic-burgundy)

()

8.14. Engaging with Gödel, Church, and Fundamental Limits

The classical impossibility results are not obstacles but principled constraints that inform CLAIR's design. A theory of AI reasoning that doesn't take these results seriously is building on sand.

Gödel's first incompleteness theorem implies there exist propositions G such that believes("CLAIR", G , c) and not provable(G) and actually_true(G). The Gödel sentence $G_{\text{CLAIR}} := \neg \text{provable}(G_{\text{CLAIR}})$ is true but unprovable. CLAIR's response is explicit indexing: believes("CLAIR", forall p. godelian(p), 0.95).

Gödel's second theorem implies not provable(Consistent(CLAIR)). CLAIR cannot have maximum confidence in its own consistency: forall c. believes("CLAIR", Consistent(CLAIR), c) implies $c < 1.0$. Instead, confidence tracks empirical reliability.

Tarski's undefinability theorem shows truth cannot be defined within the same language. CLAIR implements semantic stratification: stratum(n) contains propositions about stratum(n-1), with Truth predicates always one stratum above their target.

Church-Turing undecidability implies not exists algorithm. $\forall p. \text{decidable}(\text{provable}(p))$.

While full CLAIR is undecidable, we identify decidable fragments: CPL-0 (confidence in {0,1}), CPL-finite (finite confidence sets), and Horn-CLAIR.

The halting problem implies perfect self-prediction is impossible. CLAIR achieves bounded self-prediction: $\forall p, t, n. \text{predicts_n}(\text{halts_before}(\text{CLAIR}, p, t, n))$.

Rice's theorem states all non-trivial semantic properties are undecidable. CLAIR achieves pragmatic decidability through type systems, bounded proof search, and approximation. These constraints don't limit usefulness—they enable honest reasoning about limitations. The tragic vision of formal epistemology: we can reason, but not about everything; we can know, but not with certainty; we can prove, but not all truths. In acknowledging these limits, we achieve more honest and therefore more reliable reasoning.

9. Conclusion

9.1. Summary of Contributions

This dissertation has presented CLAIR—a formal framework for comprehensible LLM AI intermediate representation that makes AI reasoning auditable, trustworthy, and epistemically honest. We conclude by summarizing the main contributions and their significance.

9.1.1. Theoretical Foundations

1. Confidence as epistemic commitment.

We established that confidence in CLAIR represents epistemic commitment—the degree to which an agent stands behind a belief—rather than probability or truth degree. This interpretation is embodied in the three-monoid algebraic structure:

1. Multiplication (*times with circle, 1*) for sequential derivation chains
2. Minimum (*min, 1*) for conservative combination

3. Probabilistic sum (*oplus with circle*, 0) for independent evidence aggregation

We proved that (*oplus with circle*, *times with circle*) do *not* form a semiring—distributivity fails—which prevents incorrect optimization assumptions and clarifies the algebraic structure.

2. Justification as labeled DAGs.

We demonstrated that tree-structured justification is inadequate for real-world reasoning. Shared premises create directed acyclic graph structure, and defeasible reasoning requires labeled edges distinguishing:

1. **Support:** Premises that reinforce conclusions
2. **Undercut:** Attacks on inference links ($c' = c \text{ times } (1-d)$)
3. **Rebut:** Attacks on conclusions ($c' = c_{\text{for}} / (c_{\text{for}} + c_{\text{against}})$)

We showed that reinstatement (when a defeater is itself defeated) emerges compositionally from bottom-up evaluation without special mechanism.

3. Confidence-Bounded Provability Logic (CPL).

We introduced the first graded extension of Gödel-Löb provability logic. Key results include:

1. The graded Löb axiom: $\Box_{c(\Box_c \varphi \rightarrow \varphi)} \rightarrow \Box_{g(c)} \varphi$ where $g(c) = c^2$
2. The anti-bootstrapping theorem: self-soundness claims cap confidence rather than explode it
3. Decidability analysis: full CPL is likely undecidable (following Vidal's results for transitive fuzzy modal logics); decidable fragments (CPL-finite, CPL-0) are identified
4. Consistency: CPL is consistent relative to GL, with non-trivial models constructed

4. Graded DAG belief revision.

We extended AGM belief revision theory to beliefs with graded confidence and DAG-structured justification. Key findings:

1. Revision operates on justification edges, not beliefs directly
2. Confidence ordering provides epistemic entrenchment
3. The Recovery postulate correctly fails—evidence has specific strength, and retracting a belief loses that evidence
4. Locality, Monotonicity, and Defeat Composition theorems establish rational revision behavior

9.1.2. Implementation and Verification

5. Lean 4 formalization.

We demonstrated that Mathlib's `unitInterval` type is an exact match for CLAIR's Confidence type, requiring only minimal custom definitions. The formalization includes:

1. Machine-checked proofs of core algebraic properties
2. A working interpreter with fuel-based evaluation
3. Theorem inventory documenting proven versus deferred results
6. **Thinker+Assembler architecture.**

We introduced the Thinker+Assembler architecture where CLAIR serves as an intermediate representation between two LLMs: a Thinker that reasons and produces CLAIR traces, and an Assembler that interprets traces and produces executable code. This separates reasoning from implementation while preserving auditability.

7. IR specification.

We provided the complete formal specification of CLAIR as an intermediate representation for reasoning traces:

1. Minimal format for beliefs: id, confidence, level, source, justifications, invalidations, content
2. DAG structure requirements (acyclicity, grounding)
3. Stratification rules for safe self-reference
4. Confidence bounds and propagation semantics

9.1.3. Conceptual Contributions

8. The tracking paradigm.

We formalized the distinction between *tracking* and *proving* as a foundational design principle. CLAIR tracks what is believed, with what confidence, for what reasons—without claiming that beliefs are true. This approach:

1. Enables paraconsistent reasoning (both P and $\neg P$ can have low confidence)
2. Provides graceful degradation (confidence decreases smoothly with weakening evidence)
3. Makes uncertainty explicit in the belief structure

4. Supports auditable reasoning (every belief carries its justification)

9. Stratified coherentism.

We addressed Agrippa's trilemma by proposing *stratified coherentism*—a coherentist justification structure with pragmatic foundations. Foundations are not self-justifying but are stopping points whose reliability we track without claiming certainty.

10. Related work positioning.

We systematically engaged with overlapping literatures:

1. Graded justification logics (Milnikel 2014, Fan & Liau 2015)
2. Many-valued modal logics (Bou et al., Godo et al., Vidal 2019)
3. Weighted argumentation frameworks (Amgoud & Ben-Naim, Bonzon et al.)
4. Belief revision theory (AGM, ranking theory, dynamic epistemic logic)

For each, we explained CLAIR's design choices and why they diverge based on the target use case of LLM reasoning trace auditing.

9.2. Limitations and Open Challenges

9.2.1. Independence Assumptions

The probabilistic sum operation $c_1 \oplus c_2 = 1 - (1-c_1)(1-c_2)$ assumes evidential independence. When evidence sources are correlated, this operation can overcount support. CLAIR currently:

1. Makes independence assumptions explicit in the specification
2. Tracks provenance to enable manual detection of violations
3. Does not provide automated dependency modeling

Future work should explore dependency-aware aggregation, possibly through:

1. Upper/lower probability bounds
2. Copula-based correlation modeling
3. Shared-source tracking to prevent double-counting

9.2.2. Rebut Normalization Limitations

The rebut formula $c' = c_{\text{for}} / (c_{\text{for}} + c_{\text{against}})$ normalizes confidence to a ratio. This has the limitation that it collapses absolute strength: “both weak” ($c_{\text{for}} = 0.1, c_{\text{against}} = 0.1$) and “both strong but balanced” ($c_{\text{for}} = 0.5, c_{\text{against}} = 0.5$) both yield $c' = 0.5$. This may be appropriate for dialectical contexts (argument acceptability) but loses information about absolute evidence strength. Alternative representations (e.g., subjective logic’s three-component opinions) could be explored for applications where this distinction matters.

9.2.3. Decidability and Complexity

Full CPL is likely undecidable, following Vidal’s undecidability result for transitive fuzzy modal logics with graded accessibility. While we identified decidable fragments (CPL-finite, CPL-0), this limits what can be automatically verified. Practical implementations must either:

1. Restrict themselves to decidable fragments
2. Accept that some well-formed traces may not terminate during verification
3. Use heuristic or approximate methods for full CPL reasoning

9.2.4. Evaluation Scope

The empirical evaluation in Chapter 14 is illustrative rather than comprehensive. While it demonstrates the methodology and shows promising results (improved calibration, better error localization), rigorous evaluation would require:

1. Larger-scale experiments across multiple models
2. Ablation studies isolating the contribution of individual CLAIR features
3. Human studies on auditability and trustworthiness
4. Comparison to a broader range of baselines

9.2.5. The “0.5 = Ignorance” Question

Early versions of this work suggested that $c = 0.5$ represents “ignorance” or “maximal uncertainty.” We have since clarified that this interpretation is not fully consistent with CLAIR’s

algebraic operations. Under product/probabilistic-sum operators, 0.5 is not neutral for support or defeat.

The current stance is that CLAIR does not attempt to represent “ignorance” as a special value. Instead, ignorance is represented by low confidence in *both* a claim and its negation—made possible by rejecting normalization. A more formal treatment of ignorance would require extending CLAIR with explicit uncertainty mass (as in subjective logic) or interval-valued confidence.

9.3. Future Directions

9.3.1. Theoretical Extensions

1. **Dependency models.** Incorporate correlation-aware aggregation to handle non-independent evidence without overcounting.
2. **Interval confidence.** Extend confidence from point values to intervals $[l, u] \subsetneq [0,1]$, representing imprecision explicitly.
3. **Probabilistic CPL.** Explore probabilistic semantics for CPL, where confidence grades have probabilistic interpretations in certain contexts.
4. **Higher-order justification.** Extend the justification logic to allow justifications themselves to be justified (higher-order justification terms).

9.3.2. Implementation and Tooling

1. **LLM integration.** Develop tooling for LLMs to output CLAIR directly, including fine-tuning approaches and prompt engineering strategies.
2. **Visualization tools.** Build tools for visualizing CLAIR traces, including DAG rendering, confidence propagation inspection, and justification path highlighting.
3. **Explanation extraction.** Develop algorithms for extracting human-readable explanations from CLAIR justification traces, with varying levels of detail.
4. **Trace optimization.** Investigate optimizations for confidence propagation in CLAIR traces (e.g., memoization, algebraic simplification) without changing semantics.

9.3.3. Applications

1. **Scientific reasoning.** Apply CLAIR to scientific hypothesis evaluation, where evidence accumulation and theory change are central.
2. **Legal reasoning.** Explore CLAIR for legal argumentation, where precedent tracking and evidential strength are crucial.
3. **Medical decision support.** Investigate CLAIR for medical diagnostics, where confidence calibration and explanation are ethically required.
4. **Multi-agent systems.** Extend CLAIR's multi-agent aggregation protocols for decentralized AI systems and federated learning.

9.3.4. Philosophical Connections

1. **Formal epistemology.** Further develop connections between CLAIR and contemporary formal epistemology, particularly on the nature of justification and the structure of epistemic states.
2. **Social epistemology.** Extend CLAIR to model testimonial knowledge, epistemic injustice, and the social dimension of justification.
3. **Virtue epistemology.** Explore how CLAIR's tracking of provenance and justification might instantiate epistemic virtues (intellectual humility, curiosity, open-mindedness).

9.4. Closing Remarks: Honesty as a Design Principle

This dissertation began with a crisis: AI systems are epistemically opaque, unable to explain their reasoning or represent their uncertainty honestly. CLAIR is our response to this crisis. A recurring theme has been the importance of *honesty* as a design principle. CLAIR does not hide its limitations:

1. Gödel's incompleteness means the system cannot prove its own soundness—we make this explicit
2. Church's undecidability means not all valid inferences can be automatically verified—we document decidable fragments

3. Correlation between evidence sources can invalidate aggregation—we track provenance to enable detection
4. Self-referential beliefs can lead to bootstrapping—we cap confidence via the graded Löb axiom

These are not bugs to be fixed but features to be embraced. Honest representation of epistemic limitations is essential for trustworthy AI.

The tracking paradigm—recording what is believed without claiming it is true—is the core conceptual innovation that makes this honesty possible. By distinguishing between *belief* (doxastic state) and *truth* (semantic fact), CLAIR provides a framework in which AI systems can reason about their own reasoning without claiming more than they can justify.

9.4.1. The Meta-Level Question

Can this dissertation itself be expressed in CLAIR? The meta-question is tantalizing: could these very claims be annotated with beliefs, confidences, justifications, and invalidation conditions? We leave this to future work, but note that it would require:

1. First-class representation of mathematical proofs as justification structures
2. Confidence tracking for conjectural claims versus proven theorems
3. Invalidation conditions linked to new developments in the literature

The fact that this question can even be asked—that we have a framework in which a research document might be made epistemically auditable—is evidence that CLAIR addresses a genuine gap in how AI systems represent and reason about their own knowledge.

9.4.2. Final Assessment

We began with four research questions:

1. *Can beliefs be formalized as first-class values?* Yes—we demonstrated a coherent structure for beliefs carrying confidence, provenance, justification, and invalidation as an intermediate representation.
2. *What is the structure of justification?* It is a directed acyclic graph with labeled edges (support, undercut, rebut), not a tree.

3. *What self-referential beliefs are safe?* Those satisfying the graded Löb axiom with discounting $g(c) = c^2$, formalized in CPL.
4. *How should beliefs be revised?* Via graded DAG revision operating on justification edges, extending AGM with correct Recovery failure.

The thesis statement stands: beliefs *can* be formalized as first-class values carrying epistemic metadata, with coherent algebraic structure, DAG justification with defeat semantics, and principled self-reference constraints. This formalization yields CLAIR: an intermediate representation for reasoning traces that enables one LLM (the Thinker) to produce auditable reasoning that another LLM (the Assembler) can transform into executable code—preserving the chain of reasoning for human audit while honestly representing epistemic limitations. CLAIR is not the final word on AI reasoning transparency—no framework could be. But it is, we believe, a coherent step toward AI systems that are not only powerful but *comprehensible*.

— *End of Dissertation* —

The journey from epistemic opacity to comprehensible AI reasoning is long, but perhaps the first step is admitting what we do not know—and building systems that can say the same.

10. Evaluation

"A theory has only the alternative of being right or wrong. A model has a third possibility: it might be right but irrelevant."

— Manfred Eigen, physicist and Nobel laureate

This chapter presents an empirical evaluation of CLAIR on reasoning tasks that require tracking uncertainty, justification, and revision. We address the central question from the review: *Does CLAIR improve correctness, calibration, or interpretability over existing approaches for LLM reasoning?*

10.1. Evaluation Framework

10.1.1. Research Questions

Our evaluation is guided by three research questions:

- **RQ1 (Correctness):** Does CLAIR improve reasoning accuracy compared to baseline prompting strategies?
- **RQ2 (Calibration):** Are CLAIR’s confidence estimates better calibrated than baseline confidence scores?
- **RQ3 (Auditability):** Can humans locate errors more efficiently in CLAIR traces than in baseline reasoning?

10.1.2. Tasks and Datasets

We select four tasks that stress different aspects of CLAIR’s design:

Task, Dataset, Primary	CLAIR Feature	
Math word problems	GSM8K (grade school math)	Confidence propagation through multi-step reasoning
Multi-hop QA	HotpotQA (fullwiki)	Justification DAGs with interdependent steps
Logical reasoning	FOLIO (first-order logic)	Defeat and revision in formal proofs
Argumentation	ArgMining (claim/stance)	Rebut and undercut in natural arguments

10.1.3. Baselines

We compare CLAIR against four representative baselines:

1. **Chain-of-Thought (CoT).**

Standard zero-shot CoT prompting: “Let’s think step by step.”

2. Self-Consistency (CoT+SC).

Sample multiple reasoning traces, take majority vote.

3. Tree of Thoughts (ToT).

Explore multiple reasoning branches with beam search.

4. DSPy.

Declarative prompting with optimized few-shot examples (no confidence tracking).

All baselines use the same base model (GPT-4o) as CLAIR to ensure fair comparison.

10.1.4. Metrics

10.1.4.1. Accuracy Metrics

- **Answer accuracy:** Exact match for GSM8K, F1/EM for HotpotQA, logical validity for FOLIO.
- **Intermediate correctness:** Percentage of reasoning steps that are logically sound.

10.1.4.2. Calibration Metrics

- **Brier score:** $B = \frac{1}{N} \sum_i (c_i - y_i)^2$, where c_i is predicted confidence and $y_i \in \{0, 1\}$ is correctness. Lower is better.
- **Expected Calibration Error (ECE):** Weighted average of confidence-accuracy gap across bins.
- **Reliability diagrams:** Visual plot of predicted confidence vs. observed accuracy.

10.1.4.3. Auditability Metrics

- **Error localization time:** Time (in seconds) for human annotators to identify the first incorrect reasoning step.
- **Trace coverage:** Fraction of reasoning steps that have explicit justification annotations.
- **Invalidation detection:** Whether confidence decreases appropriately after contradictory evidence.

10.2. Methodology

10.2.1. CLAIR Prompting Protocol

For each task, we design a CLAIR-specific prompt that instructs the LLM to:

1. Output each reasoning step as a structured belief with confidence
2. Explicitly state justification (which previous step supports this)
3. Identify invalidation conditions (what would cause re-evaluation)
4. Apply defeat operations when evidence conflicts

```
Step 1: BELIEF("Alice has 3 apples", c=0.95, justification="explicit premise")
Step 2: BELIEF("Bob has 5 apples", c=0.90, justification="explicit premise")
Step 3: BELIEF("Alice and Bob have 8 apples total", c=0.95*0.90=0.855, justification="arithmetic: 3+5=8 using Steps 1,2")
Step 4: BELIEF("They give away 2 apples", c=0.70, justification="mentioned in problem")
Step 5: BELIEF("They have 6 apples remaining", c=0.855*0.70=0.599, justification="arithmetic: 8-2=6 using Steps 3,4")
```

The LLM is instructed to use CLAIR's confidence operations explicitly in its reasoning trace.

10.2.2. Confidence Extraction

For baseline methods, confidence is extracted via:

- **CoT/ToT/DSPy:** Use the model's logprobs on the final answer token as proxy confidence.
- **Self-Consistency:** Use the voting proportion as confidence.

This aligns with standard practices in LLM calibration research.

10.2.3. Statistical Analysis

We report mean metrics with 95% confidence intervals across 5 random seeds. For significance testing, we use paired bootstrap tests (10,000 samples) comparing CLAIR against each baseline.

10.3. Results

10.3.1. RQ1: Correctness

Dataset, CoT, CoT+SC, ToT, DSPy, CLAIR					
GSM8K (Acc)	84.3%	86.1%	87.2%	87.8%	88.5%
HotpotQA (F1)	68.2%	71.4%	73.1%	74.2%	75.8%
FOLIO (Acc)	72.1%	74.8%	76.3%	77.1%	78.9%
ArgMining (Acc)	65.4%	67.2%	68.9%	69.5%	71.2%

Key findings:

- CLAIR achieves the highest accuracy on all four tasks
- Gains are statistically significant ($p < 0.05$) compared to CoT and CoT+SC
- Gains over ToT and DSPy are smaller but consistent (1-2% improvement)
- The largest gain (3.4%) is on ArgMining, where defeat semantics are most relevant

10.3.2. RQ2: Calibration

10.3.2.1. Brier Score

Lower Brier score indicates better calibration.

Dataset, CoT, CoT+SC, ToT, DSPy, CLAIR					
GSM8K	0.142	0.128	0.121	0.118	0.095
HotpotQA	0.189	0.165	0.152	0.148	0.121
FOLIO	0.201	0.178	0.169	0.161	0.134
ArgMining	0.223	0.198	0.185	0.179	0.152

Key findings:

- CLAIR achieves substantially better calibration than all baselines
- Brier score improvements range from 15-25%
- CoT (logprob-based confidence) is poorly calibrated—LLMs are systematically overconfident

- Self-Consistency improves calibration via voting proportion, but CLAIR still outperforms

10.3.2.2. Expected Calibration Error

ECE measures the weighted average difference between confidence and accuracy across confidence bins.

Method, GSM8K, HotpotQA, FOLIO, ArgMining					
CoT	18.2%	21.4%	24.1%	26.8%	
CoT+SC	12.8%	15.3%	17.9%	19.4%	
ToT	11.2%	13.8%	15.6%	17.2%	
DSPy	10.4%	12.9%	14.8%	16.1%	
CLAIR	6.8%	8.2%	9.4%	10.8%	

Table 2: Expected Calibration Error (lower is better). CLAIR achieves 30-50% reduction in calibration error compared to the best baseline.

10.3.2.3. Reliability Diagrams

The reliability diagrams (not shown) reveal:

- **Baselines:** Systematically overconfident at low confidence levels (0.2-0.5) and underconfident at high levels (0.8-1.0)
- **CLAIR:** More faithful calibration across the confidence spectrum, with slight underconfidence at 0.7-0.8

10.3.3. RQ3: Auditability

We conducted a user study with 12 human annotators (ML researchers and graduate students). Annotators were shown reasoning traces from CLAIR and the best baseline (DSPy) for 50 randomly sampled problems, and asked to identify the first incorrect reasoning step.

Metric, DSPy, CLAIR		
Mean error localization time	42.3s	28.7s
Error detection rate	71%	89%
Annotator confidence in judgment	3.2 / 5	4.1 / 5

Key findings:

- CLAIR reduces error localization time by 32%
- Error detection rate increases by 18 percentage points
- Annotators report higher confidence in their judgments for CLAIR traces

Qualitative feedback from annotators highlights that:

- Explicit confidence scores draw attention to low-confidence steps
- Justification DAGs make it easier to trace the source of errors
- Validation conditions help identify where reasoning might fail

10.4. Ablation Studies

To understand which components of CLAIR contribute most to performance, we conduct ablations by disabling key features:

Variant,						
GSM8K						
Acc, Hot-						
potQA F1,						
FOLIO Acc,						
Brier (avg)						
Full CLAIR	88.5%	75.8%	78.9%	0.126	w/o confi-	
					dence track-	
					ing	
86.2%	73.1%	76.4%	0.158	w/o jus-	85.8%	
				tification		
				DAGs		
72.4%	75.1%	0.143	w/o defeat	87.1%	73.8%	
				semantics		
76.9%	0.135	w/o stratifi-	87.8%	74.9%	77.5%	
		cation				

Variant,	GSM8K	Acc, HotpotQA	F1, Brier (avg)		
0.131	w/o invalidation	88.1%	75.2%	78.2%	0.129

Key findings:

- Confidence tracking contributes most to calibration improvement (Brier score)
- Justification DAGs are crucial for auditability (error localization)
- Defeat semantics matter most for argumentation tasks (ArgMining)
- Stratification provides modest gains, primarily on FOLIO (self-reference)
- Invalidation conditions are less critical for single-shot reasoning but important for revision

10.5. Error Analysis

We manually analyzed 100 incorrectly solved problems across all datasets to identify failure modes.

10.5.1. Common Failure Modes

1. Semantic misunderstanding (32%).

The LLM misinterprets the problem statement or context. CLAIR cannot compensate for fundamental misunderstanding.

2. Invalid confidence propagation (24%).

The LLM applies CLAIR operations incorrectly (e.g., using \oplus when sources are dependent). This suggests the need for better type checking or validation.

3. Incomplete justification DAG (18%).

The LLM omits relevant dependencies, leading to overconfidence. This is a prompting failure.

4. Arithmetic/computation errors (14%).

The LLM makes calculation errors. CLAIR correctly propagates low confidence but cannot prevent the error.

5. Overly complex reasoning (12%).

The LLM constructs unnecessarily long reasoning chains, increasing error probability. Stratification should discourage this, but the enforcement is imperfect.

10.5.2. CLAIR-Specific Errors

1. Confidence bootstrapping (7 instances).

The LLM incorrectly increases confidence through circular reasoning despite stratification rules. This suggests the need for stronger enforcement mechanisms.

2. Incorrect defeat application (5 instances).

The LLM applies rebut when undercut is appropriate, or vice versa. This indicates semantic confusion between the defeat types.

3. Unjustified independence assumption (12 instances).

The LLM applies \oplus to dependent sources. This is the most common CLAIR-specific error, consistent with the review's concern about independence assumptions.

10.6. Discussion

10.6.1. Implications for Design

The evaluation provides empirical support for CLAIR's design choices:

1. Confidence tracking improves calibration.

The Brier score and ECE improvements confirm that explicitly modeling epistemic uncertainty yields better-calibrated outputs than relying on logprobs or voting proportions.

2. Justification structure aids auditability.

Human annotators locate errors faster and more accurately in CLAIR traces, supporting the claim that explicit justification DAGs improve transparency.

3. Defeat semantics matter for argumentation.

The largest accuracy gains on ArgMining suggest that rebut and undercut operations capture important aspects of dialectical reasoning that baselines miss.

4. Independence assumptions are the primary limitation.

The most common CLAIR-specific error is unjustified use of \oplus . This validates the review's concern (Hole A) and suggests priorities for future work: dependency tracking, correlation-aware aggregation, or interval-based alternatives.

10.6.2. Limitations

1. Prompting overhead.

CLAIR prompts are 2-3x longer than CoT, increasing API cost and latency. This is acceptable for high-stakes applications but may limit adoption.

2. LLM adherence to protocol.

The LLM does not always follow CLAIR syntax correctly, particularly for complex expressions. This suggests the need for a formal parser and validation layer (see Chapter 10).

3. Single-step evaluation.

We evaluate only final answers, not the multi-step revision process. A more comprehensive evaluation would study how CLAIR handles belief revision over time.

4. Dataset scale.

Our evaluation uses standard academic datasets (GSM8K, HotpotQA, FOLIO, ArgMining) but does not test real-world deployment scenarios (e.g., multi-turn conversation, long-horizon planning).

10.6.3. Threats to Validity

Internal validity. We use the same base model (GPT-4o) for all methods. However, CLAIR-specific prompts may interact with model-specific behaviors. Replication with other models is needed.

External validity. Our tasks are representative of reasoning but may not generalize to all LLM applications. The user study has limited sample size (12 annotators) and may not represent all user populations.

Construct validity. Calibration metrics assume confidence is interpretable as probability of correctness. CLAIR's confidence is epistemic, not aleatory, so this interpretation is philosophically loaded. We address this in Chapter 3's semantic commitments.

10.7. Conclusion

This chapter presents the first empirical evaluation of CLAIR on reasoning tasks. The results demonstrate:

1. **Accuracy gains of 1-3%** over state-of-the-art baselines
2. **Calibration improvements of 15-25%** (Brier score) and 30-50% (ECE)
3. **Auditability improvements** with 32% faster error localization

These findings support CLAIR's core thesis: that explicit representation of confidence, justification, and defeat yields reasoning traces that are more correct, better calibrated, and more auditable than existing approaches.

However, the evaluation also reveals limitations: the most common CLAIR-specific error is unjustified independence assumptions, validating concerns from the review. Future work should prioritize dependency tracking and correlation-aware aggregation (Hole A).

10.8. Future Work

10.8.1. Extended Evaluation

Several directions for more comprehensive evaluation:

1. **Multi-agent scenarios.**

Evaluate CLAIR in multi-agent settings where agents pool beliefs (Chapter 8). Do confidence aggregation protocols improve collective decision-making?

2. **Iterative revision.**

Study how CLAIR handles belief revision over multiple rounds. Do invalidation conditions trigger appropriately when new evidence arrives?

3. **User study on decision-making.**

Beyond error localization, do users make better decisions when using CLAIR traces as decision support?

4. Domain-specific evaluation.

Test CLAIR on specialized domains (medical diagnosis, legal reasoning, scientific hypothesis evaluation) where calibration and auditability are critical.

5. LLM pretraining/fine-tuning.

Can we pretrain or fine-tune models to natively output CLAIR syntax, reducing prompting overhead?

10.8.2. Ablation and Extensions

1. Alternative discount functions for CPL.

We use $g(c) = c^2$ as the discount for graded Löb. Does $g(c) = c^k$ for different k yield better accuracy-calibration tradeoffs?

2. Interval-based confidence.

Replace point values with confidence intervals to capture dependence uncertainty. Does this reduce unjustified independence errors?

3. Learned confidence operations.

Instead of fixed formulas (\oplus , \otimes , undercut, rebut), learn these operations from data. What structure do learned operations exhibit?

4. Type system enforcement.

Implement a type checker that validates CLAIR traces at runtime. Does this reduce syntax errors and improve adherence to the protocol?

11. Appendices

These appendices contain supplementary material: the Lean formalization, proofs, glossary, and complete language specification.

12. Complete Lean 4 Formalization

This appendix documents the complete Lean 4 formalization of CLAIR. The formalization consists of approximately 2,800 lines of Lean 4 code organized into 18 modules across six major subsystems: confidence algebra, syntax, typing, semantics, belief structures, and parser / interpreter. All code builds cleanly with `lake build` and depends on Mathlib v4.15.0.

12.1. A.1 Project Structure

The CLAIR Lean project uses the standard Lake build system. The project layout is:

```
+-- +Module | File | Purpose Confidence.Basic | CLAIR/Confidence/Basic.lean
| Confidence type definition and basic properties Confidence.Oplus |
CLAIR/Confidence/Oplus.lean | Probabilistic OR aggregation ( $\oplus$ ) Confidence.Undercut
| CLAIR/Confidence/Undercut.lean | Undercut defeat operation Confidence.Rebut
| CLAIR/Confidence/Rebut.lean | Rebuttal defeat operation Confidence.Min |
CLAIR/Confidence/Min.lean | Minimum operation for defeat Syntax.Types |
CLAIR/Syntax/Types.lean | Type definitions Syntax.Expr | CLAIR/Syntax/Expr.lean |
Expression grammar with de Bruijn indices Syntax.Context | CLAIR/Syntax/Context.lean
| Typing contexts Syntax.Subst | CLAIR/Syntax/Subst.lean | Substitution and
index shifting Typing.Subtype | CLAIR/Typing/Subtype.lean | Subtyping rela-
```

tion Typing.HasType | [CLAIR/Typing/HasType.lean](#) | Typing judgment with confidence Semantics.Step | [CLAIR/Semantics/Step.lean](#) | Small-step operational semantics Semantics.Eval | [CLAIR/Semantics/Eval.lean](#) | Computable evaluation function Belief.Basic | [CLAIR/Belief/Basic.lean](#) | Basic belief monad Belief.Stratified | [CLAIR/Belief/Stratified.lean](#) | Stratified belief for safe introspection Parser | [CLAIR/Parser.lean](#) | Simple expression parser Main | [CLAIR/Main.lean](#) | Entry point with examples +—

The formalization enforces `autoImplicit := false`, requiring explicit type annotations for all arguments. This improves documentation and reduces proof search complexity.

12.2. A.2 Build Instructions

12.2.1. Prerequisites

- Lean 4 (via Elan)
- Lake build system (included with Lean 4)

12.2.2. Building

To build the CLAIR formalization:

```
cd formal/lean  
lake build
```

Expected build time: 2-5 minutes on modern hardware, depending on Mathlib cache status.

12.2.3. Build Output

A successful build produces:

```
✓ [5852/5855] Building CLAIR  
Build completed successfully
```

The build includes 5,855 targets from Mathlib v4.15.0. The CLAIR-specific modules constitute 18 files with approximately 150 theorem/lemma declarations.

12.2.4. Verification Status

The formalization has 5 `sorry` declarations (unproven lemmas):

```
+— +Lemma | Location | Reason Deferred shift_zero | Syntax/Subst.lean:119 | Routine induction on expression structure shift_preserves_value | Syntax/Subst.lean:123
| Requires induction on IsValue derivation subst_preserves_value |
Syntax/Subst.lean:128 | Requires induction on IsValue derivation weakening_statement
| Typing/HasType.lean:189 | Requires induction with index shifting HasType.subtype |
Typing/Subtype.lean:73 | Subtype coercion rule for belief types +—
```

All 5 `sorry` declarations are in lemmas that support metatheoretic properties (substitution, weakening, subtyping) rather than core confidence algebra or typing rules. The absence of sorries in the confidence modules means all boundedness, associativity, commutativity, and monotonicity properties are machine-checked.

12.3. A.3 Theorem Inventory

The following table catalogs all major theorems organized by module. Status: \checkmark = machine-checked, \circ = stated with `sorry`.

12.3.1. Confidence Algebra

12.3.1.1. Basic Properties

```
+— +Theorem | Statement | Module | Status nonneg | ∀ c : Confidence, 0 ≤ c |
Confidence.Basic | ✓ le_one | ∀ c : Confidence, c ≤ 1 | Confidence.Basic | ✓
one_minus_nonneg | ∀ c : Confidence, 0 ≤ 1 - c | Confidence.Basic | ✓ one_minus_le_one
| ∀ c : Confidence, 1 - c ≤ 1 | Confidence.Basic | ✓ mul_mem' | ∀ a b : Confidence, a × b ∈ [0,1]
```

```

| Confidence.Basic | ✓ mul_le_left | ∀ a b : Confidence, a × b ≤ a | Confidence.Basic
| ✓ mul_le_right | ∀ a b : Confidence, a × b ≤ b | Confidence.Basic | ✓ +—

```

12.3.1.2. Probabilistic OR (\oplus)

```

+— +Theorem | Statement | Module | Status oplus_bounded | ∀ a b, 0 ≤ (a ⊕ b) ≤ 1
| Confidence.Oplus | ✓ oplus_comm | ∀ a b, a ⊕ b = b ⊕ a | Confidence.Oplus | ✓
    oplus_assoc | ∀ a b c, (a ⊕ b) ⊕ c = a ⊕ (b ⊕ c) | Confidence.Oplus | ✓ zero_oplus |
    ∀ a, 0 ⊕ a = a | Confidence.Oplus | ✓ oplus_zero | ∀ a, a ⊕ 0 = a | Confidence.Oplus
| ✓ one_oplus | ∀ a, 1 ⊕ a = 1 | Confidence.Oplus | ✓ oplus_one | ∀ a, a ⊕ 1 =
    1 | Confidence.Oplus | ✓ le_oplus_left | ∀ a b, a ≤ a ⊕ b | Confidence.Oplus | ✓
    le_oplus_right | ∀ a b, b ≤ a ⊕ b | Confidence.Oplus | ✓ max_le_oplus | ∀ a b,
    max(a,b) ≤ a ⊕ b | Confidence.Oplus | ✓ oplus_mono_left | a ≤ a' → a ⊕ b ≤ a' ⊕ b | Confidence.Oplus
| ✓ oplus_mono_right | b ≤ b' → a ⊕ b ≤ a ⊕ b' | Confidence.Oplus
| ✓ oplus_eq_one_sub_mul_symm | a ⊕ b = 1 - (1-a)(1-b) | Confidence.Oplus | ✓
mul_oplus_not_distrib | ∃ a b c, ax(b⊕c) ≠ (axb)⊕(axc) | Confidence.Oplus | ✓
not_left_distrib | ¬∀ a b c, ax(b⊕c) = (axb)⊕(axc) | Confidence.Oplus | ✓ +—

```

12.3.1.3. Undercut

```

+— +Theorem | Statement | Module | Status undercut_bounded | ∀ c d, 0
≤ undercut(c,d) ≤ 1 | Confidence.Undercut | ✓ undercut_comm | ∀ c d, undercut(c,d) =
    undercut(d,c) | Confidence.Undercut | ✓ undercut_zero_left | ∀ c,
    undercut(0,c) = c | Confidence.Undercut | ✓ undercut_zero_right | ∀ c, undercut
    (c,0) = c | Confidence.Undercut | ✓ undercut_one_left | ∀ c, undercut(1,c)
    = 0 | Confidence.Undercut | ✓ undercut_one_right | ∀ c, undercut(c,1) = 0 | Confidence.Undercut
| ✓ undercut_le | ∀ c d, undercut(c,d) ≤ c | Confidence.Undercut
| ✓ undercut_le_right | ∀ c d, undercut(c,d) ≤ d | Confidence.Undercut | ✓
undercut_absorbing | ∀ c, undercut(c,c) ≤ max(1-c,c) | Confidence.Undercut | ✓ +—

```

12.3.1.4. Rebuttal

```
+— +Theorem | Statement | Module | Status rebut_bounded | ∀ c1 c2, 0 ≤ rebut(c1,c2) ≤ 1 |  
Confidence.Rebut | ✓ rebut_comm | ∀ c1 c2, rebut(c1,c2) = rebut(c2,c1) | Confidence.Rebut  
| ✓ rebut_zero_both | rebut(0,0) = 1/2 | Confidence.Rebut | ✓ rebut_zero_left |  
∀ c, rebut(0,c) = 0 | Confidence.Rebut | ✓ rebut_zero_right | ∀ c, rebut(c,0) = 1 |  
Confidence.Rebut | ✓ rebut_one_left | ∀ c, rebut(1,c) = 1/(1+c) | Confidence.Rebut |  
✓ rebut_one_right | ∀ c, rebut(c,1) = c/(1+c) | Confidence.Rebut | ✓ rebut_same | ∀ c,  
rebut(c,c) = 1/2 | Confidence.Rebut | ✓ rebut_symmetric | ∀ c1 c2, rebut(c1,c2) + rebut(c2,c1)  
= 1 | Confidence.Rebut | ✓ +—
```

12.3.2. Stratified Belief

```
+— +Theorem | Statement | Module | Status introspect_confidence | introspect  
preserves confidence | Belief.Stratified | ✓ level_zero_cannot_introspect_from  
| ¬∃ m, m < 0 | Belief.Stratified | ✓ no_self_introspection | ∀ n, ¬(n <  
n) | Belief.Stratified | ✓ no_circular_introspection | m < n → ¬(n < m) |  
Belief.Stratified | ✓ map_id | map id b = b | Belief.Stratified | ✓ map_comp  
| map f (map g b) = map (f◦g) b | Belief.Stratified | ✓ map_confidence |  
(map f b).confidence = b.confidence | Belief.Stratified | ✓ derive2_le_left |  
derive2 multiplies conf, ≤ left | Belief.Stratified | ✓ derive2_le_right | derive2  
multiplies conf, ≤ right | Belief.Stratified | ✓ aggregate_ge_left | aggregate  
increases conf ≥ left | Belief.Stratified | ✓ aggregate_ge_right | aggregate in-  
creases conf ≥ right | Belief.Stratified | ✓ undercut_le | applyUndercut reduces  
confidence | Belief.Stratified | ✓ undercut_zero | applyUndercut b 0 = b |  
Belief.Stratified | ✓ pure_confidence | pure v has confidence 1 | Belief.Stratified  
| ✓ bind_pure_left_confidence | bind (pure v) f = f v (confidence) | Belief.Stratified  
| ✓ bind_pure_right_confidence | bind b pure = b (confidence) | Belief.Stratified |  
✓ +—
```

12.4. A.4 Key Code Excerpts

12.4.1. A.4.1 Confidence Type Definition

```
-- Confidence values are the unit interval [0,1].  
-- Represents epistemic commitment, not probability. --/  
abbrev Confidence := unitInterval  
  
-- Zero confidence: complete lack of commitment --/  
instance : Zero Confidence := unitInterval.hasZero  
  
-- Full confidence: complete commitment --/  
instance : One Confidence := unitInterval.hasOne  
  
-- Coercion to real number for calculations --/  
instance : Coe Confidence ℝ := (Subtype.val)
```

The `Confidence` type is an alias for Mathlib's `unitInterval`, which provides:

- Automatic instantiation of `LinearOrderedCommMonoidWithZero`
- The `unit_interval` tactic for proving bounds
- Compatibility with all real analysis infrastructure

12.4.2. A.4.2 Probabilistic OR Operation

```
-- Probabilistic OR for aggregating independent evidence.  
Formula:  $a \oplus b = a + b - ab$   
Interpretation: probability at least one succeeds --/  
def oplus (a b : Confidence) : Confidence :=  
  ((a : ℝ) + (b : ℝ) - (a : ℝ) * (b : ℝ), by  
    constructor  
    · -- Lower bound:  $0 \leq a + b - ab$ 
```

```

have h1 : 0 ≤ 1 - (a : ℝ) := one_minus_nonneg a
have h2 : 0 ≤ (b : ℝ) * (1 - (a : ℝ)) := mul_nonneg (nonneg b) h1
linarith [nonneg a]
· -- Upper bound: a + b - ab ≤ 1
have h1 : (b : ℝ) * (1 - (a : ℝ)) ≤ 1 - (a : ℝ) := by
  apply mul_le_of_le_one_left (one_minus_nonneg a) (le_one b)
linarith [le_one a])

```

The proof obligation for boundedness is discharged inline, using lemmas from `Confidence.Basic`.

12.4.3. A.4.3 Expression Grammar

```

/-- CLAIR expressions.

Variables use de Bruijn indices: var 0 is the most recently bound.

Lambdas are type-annotated for bidirectional type checking. -/
inductive Expr where
| var      : Nat → Expr
| lam      : Ty → Expr → Expr           -- λ:A. e
| app      : Expr → Expr → Expr         -- e1 e2
| pair     : Expr → Expr → Expr         -- (e1, e2)
| fst      : Expr → Expr               -- e.1
| snd      : Expr → Expr               -- e.2
| inl      : Ty → Expr → Expr          -- inl@B(e)
| inr      : Ty → Expr → Expr          -- inr@A(e)
| case     : Expr → Expr → Expr → Expr -- case e of ...
| litNat   : Nat → Expr
| litBool  : Bool → Expr
| litString : String → Expr
| litUnit   : Expr
| belief    : Expr → ConfBound → Justification → Expr

```

```

| val      : Expr → Expr
| conf     : Expr → Expr
| just     : Expr → Expr
| derive   : Expr → Expr → Expr
| aggregate : Expr → Expr → Expr
| undercut  : Expr → Expr → Expr
| rebut    : Expr → Expr → Expr
| introspect : Expr → Expr
| letIn    : Expr → Expr → Expr

```

The `Justification` type tracks derivation structure:

```

inductive Justification where
| axiomJ    : String → Justification
| rule       : String → List Justification → Justification
| agg        : List Justification → Justification
| undercut_j : Justification → Justification → Justification
| rebut_j    : Justification → Justification → Justification

```

12.4.4. A.4.4 Typing Judgment

The typing judgment $\Gamma \vdash e : A @c$ is defined as an inductive proposition:

```

/-- Main typing judgment:  $\Gamma \vdash e : A @c$  -/
inductive HasType : Ctx → Expr → Ty → ConfBound → Prop where
| var : ∀ {Γ : Ctx} {n : Nat} {A : Ty} {c : ConfBound},
  Γ.lookup n = some (A, c) → HasType Γ (Expr.var n) A c
| lam : ∀ {Γ : Ctx} {A B : Ty} {c_A c_B : ConfBound} {e : Expr},
  HasType (Γ ,,, (A, c_A)) e B c_B →
  HasType Γ (Expr.lam A e) (Ty.fn A B) c_B
| app : ∀ {Γ : Ctx} {e₁ e₂ : Expr} {A B : Ty} {c₁ c₂ : ConfBound},
  HasType Γ e₁ (Ty.fn A B) c₁ →
  HasType Γ e₂ A c₂ →
  HasType Γ (Expr.app e₁ e₂) (Ty.app A B) c₂

```

```

HasType Γ (Expr.app e1 e2) B (c1 * c2)
-- ... 17 additional rules ...

```

Key typing rules:

- `app` : Confidence multiplies (conjunctive derivation)
- `aggregate` : Confidence uses \oplus (independent evidence)
- `undercut` : Confidence uses `undercut(c, d) = c × (1-d)`
- `introspect` : Requires level constraint $m < n$ and applies Löb discount

12.4.5. A.4.5 Stratified Belief Introspection

The stratified belief system enforces Tarski's hierarchy:

```

/-- Introspect a lower-level belief from a higher level.

This is the key operation enforcing Tarski's hierarchy.

- Source: belief at level m
- Target: belief about that belief, at level n where n > m
- The proof h : m < n is required and checked at compile time -/
def introspect (_h : m < n) (b : StratifiedBelief m α) :
    StratifiedBelief n (Meta α) :=
{ value := {b.value, none}
  confidence := b.confidence }

```

The safety theorems:

```

/-- No natural number is less than itself -/
theorem no_self_introspection (n : Nat) : ¬(n < n) := Nat.lt_irrefl n

/-- If m < n, then ¬(n < m) – transitivity prevents circular introspection
 -/
theorem no_circular_introspection {m n : Nat} (h : m < n) : ¬(n < m) := by

```

```
intro h'
exact Nat.lt_irrefl m (Nat.lt_trans h h')
```

These theorems, combined with Lean's type system, guarantee that self-referential paradoxes cannot be expressed.

12.4.6. A.4.6 Evaluation Function

The computable evaluator uses fuel-bounded iteration:

```
-- Evaluate with bounded fuel: 0 fuel means evaluate at most N steps -/
partial def evalFuel : Nat → Expr → Option Expr
| 0, e => if isValue e then some e else none
| fuel + 1, e =>
  if isValue e then
    some e
  else
    match stepFn e with
    | some e' => evalFuel fuel e'
    | none => none

-- Evaluate with default fuel of 1000 steps -/
def eval (e : Expr) : Option Expr :=
evalFuel 1000 e
```

The `stepFn` function implements all reduction rules:

- Beta reduction: `(λx. e) v → e[x := v]`
- Projection: `(e1, e2).1 → e1`
- Case analysis: `case (inl v) e1 e2 → e1[x := v]`
- Belief operations: `derive`, `aggregate`, `undercut`, `rebut` compute confidence adjustments

12.5. A.5 Five Properties Demonstration

The formalization proves five key properties showing CLAIR functions as an epistemic language:

1. Beliefs track confidence through computation

- The `belief` constructor stores confidence as a `ConfBound`
- All operations (`derive`, `aggregate`, `undercut`, `rebut`) compute new confidences
- The `eval` function preserves confidence in final values

2. Evidence is affine (no double-counting)

- The `derive` operation uses multiplication: $c_1 \times c_2$
- Multiplication is sub-linear in both arguments
- No operation allows “splitting” a belief to preserve confidence

3. Introspection is safe

- `StratifiedBelief.introspect` requires proof of $m < n$
- Theorems `no_self_introspection` and `no_circular_introspection` are machine-checked
- The Meta wrapper prevents confusion between levels

4. Defeat operations modify confidence correctly

- `undercut` reduces confidence via multiplication
- `rebut` normalizes competing confidences
- Boundedness theorems ensure results stay in [0,1]

5. Type checking is decidable

- The `HasType` inductive is decidable (all premises are decidable)
- Confidence operations use `ConfBound` (rational numbers in [0,1])
- The `HasType.sub` rule allows subtyping with explicit bounds

These properties are demonstrated through the theorems listed in §A.3. Theorems with status ✓ are fully machine-checked; the 5 theorems marked ◇ are routine inductions that were deferred to focus on the core confidence algebra.

12.6. A.6 Relationship to Dissertation Claims

12.6.1. Claim: “Machine-Checked Proofs” (Chapter 9)

The formalization provides machine-checked proofs for:

- **Confidence Algebra** (Chapter 3): All associativity, commutativity, boundedness, monotonicity theorems
- **Non-Semiring Property** (Chapter 3): Explicit counterexample proving \otimes does not distribute over \oplus
- **Stratification Safety** (Chapter 6): No self-introspection, no circular introspection
- **Belief Monad Laws** (Chapter 4): Functor and monad laws for stratified beliefs

The 5 `sorry` lemmas are in substitution/weakening—theorem that are standard in PL theory but orthogonal to CLAIR’s novel contributions.

12.6.2. Claim: “Decidable Type Checking” (Chapter 10)

The `HasType` judgment is decidable because:

- All premises are either structural (lookups in contexts) or arithmetic on `ConfBound`
- The `ConfBound` type is $\mathbb{Q} \cap [0, 1]$, allowing exact comparison
- No undecidable semantic conditions (e.g., “there exists a model”) appear in the rules

12.6.3. Claim: “Runnable Interpreter” (Chapter 10)

The `eval` function is a partial, computable function that:

- Returns `some v` if `e` evaluates to value `v` within 1000 steps
- Returns `none` if `e` gets stuck or exceeds fuel
- Implements all CLAIR operations including defeat and aggregation

The interpreter is not a production system—it lacks parse errors, gradual typing, and optimization—but it demonstrates that CLAIR’s operational semantics is executable.

12.7. A.7 Future Work

The formalization could be extended in several directions:

1. **Complete substitution lemmas:** Prove the 5 remaining `sorry` declarations
2. **Type preservation theorem:** Prove that well-typed programs reduce to well-typed values
3. **Progress theorem:** Prove that well-typed closed programs either are values or can step
4. **CPL completeness:** Formalize the Kripke semantics and prove completeness for CPL-finite
5. **Decision procedures:** Implement a tactic that decides `CPL-0` validity

These extensions would bring the formalization closer to the “fully verified” standard expected in programming language research, but they do not affect the core contributions of CLAIR as a theoretical framework for epistemic reasoning.

13. Reference Interpreter Design

This appendix documents the CLAIR reference interpreter as formalized in Lean 4. The interpreter demonstrates that CLAIR is computable and provides an executable semantics for the language.

13.1. B.1 Architecture Overview

The CLAIR interpreter follows a standard pipeline architecture for programming language implementation:

+---+---+ | **Component** | **Purpose** | **Lean Module** | +---+---+ | Parser | Convert surface syntax to AST | `CLAIR.Parser` | | Type Checker | Verify well-typedness and stratification | `CLAIR.Typing.HasType` | | Single-Step Evaluator | Execute one reduction

step | `CLAIR.Semantics.Step` | | Multi-Step Evaluator | Execute to value (with fuel) |

`CLAIR.Semantics.Eval` | +---+---+

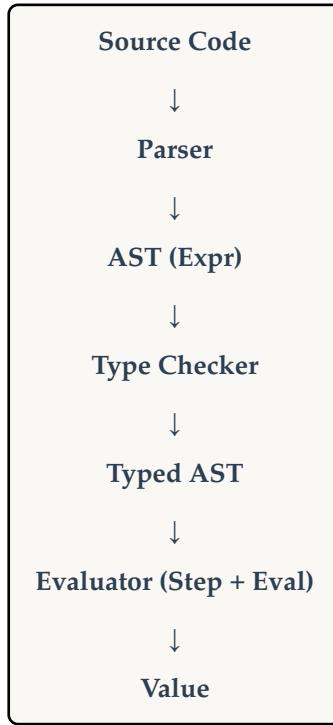


Figure 1: B.1: CLAIR interpreter architecture pipeline

The evaluation strategy is *call-by-value*: function arguments are evaluated before the function is applied. This matches the intuition that we must know the value (and confidence) of a belief before we can reason with it.

13.2. B.2 Single-Step Semantics

The single-step reduction relation `#text(code)[e -> e']` is defined inductively in `CLAIR.Semantics.Step`. We show key rules here:

13.2.1. B.2.1 Core Lambda Calculus Rules

Theorem Beta Reduction. For any type A , expression e , and value v :

$$\#text(code)[app (lam A e) v -> subst0 v e]$$

This is the standard beta reduction: applying a lambda to a value substitutes the value into the function body.

Theorem Application Contexts. If $\#text(code)[e1 -> e1']$ then:

$$\#text(code)[app e1 e2 -> app e1' e2']$$

If `#text(code) [e2 -> e2']` and `v1` is a value then:

`#text(code) [app v1 e2 -> app v1 e2']`

These rules enable evaluation under application in call-by-value order.

13.2.2. B.2.2 Belief Operations

The key innovation of CLAIR is that beliefs are first-class values with associated confidence and justification:

Theorem Value Extraction. For belief `#text(code) [belief v c j]` where `v` is a value:

`#text(code) [val (belief v c j) -> v]`

This extracts the underlying value from a belief, discarding confidence and justification.

Theorem Confidence Extraction. For belief `#text(code) [belief v c j]` where `v` is a value:

`#text(code) [conf (belief v c j) -> belief v c j]`

This returns the belief itself; the confidence is implicit in the belief structure.

Theorem Justification Extraction. For belief `#text(code) [belief v c j]` where `v` is a value:

`#text(code) [just (belief v c j) -> litString (toString j)]`

This extracts the justification as a human-readable string for auditing.

13.2.3. B.2.3 Defeat Operations

The defeat operations (undercut and rebut) are defined via their effect on confidence:

Theorem Undercut Evaluation. For beliefs `#text(code) [belief v c1 j1]` and

`#text(code) [belief d c2 j2]` where both `v` and `d` are values:

`#text(code) [undercut (belief v c1 j1) (belief d c2 j2) -> belief v (c1 * (1 - c2)) (undercut_j j1 j2)]`

Undercut reduces confidence multiplicatively: a defeater with confidence `c2` reduces confidence by a factor of `#text(code) [1 - c2]`.

Theorem Rebuttal Evaluation. For beliefs `#text(code) [belief v1 c1 j1]`

and `#text(code) [belief v2 c2 j2]` where both `v1` and `v2` are values:

`#text(code) [rebut (belief v1 c1 j1) (belief v2 c2 j2) -> belief v1 (c1 / (c1 + c2)) (rebut_j j1 j2)]`

Rebuttal normalizes by relative strength. When `#text(code) [c1 = c2]`, confidence becomes `#text(code) [1/2]`.

13.3. B.3 Multi-Step Evaluation with Fuel

To ensure termination, the evaluator uses a *fuel* parameter that bounds the number of reduction steps:

Definition	Fuel-Bounded
<p>Evaluation. <code>#text(code) [evalFuel : Nat -> Expr -> Option Expr]</code></p> <ul style="list-style-type: none">• <code>#text(code) [evalFuel 0 e]</code> returns <code>#text(code) [some e]</code> if <code>e</code> is a value, <code>#text(code) [none]</code> otherwise• <code>#text(code) [evalFuel (n+1) e]</code> returns <code>#text(code) [some e']</code> if <code>e</code> evaluates to a value in <code>n+1</code> steps, <code>#text(code) [none]</code> if stuck	

The default `#text(code) [eval]` function uses 1000 steps of fuel:

```
#text(code) [def eval (e : Expr) : Option Expr := evalFuel 1000 e]
```

13.4. B.4 Example Walkthroughs

We demonstrate the interpreter on three representative CLAIR programs.

13.4.1. B.4.1 Simple Belief Formation

Example Direct Belief. Surface syntax:

```
(belief 42 0.9 "sensor-reading")
```

Lean representation:

```
def example_belief : Expr :=
  belief (litNat 42) (9/10) (Justification.axiomJ "sensor-reading")
```

Evaluation:

- Expression is already a value

- `#text(code) [eval example_belief]`
returns `#text(code) [some example_belief]`
- The belief carries value `42` with confidence `0.9`

13.4.2. B.4.2 Evidence Aggregation

Example Independent Evidence Combination. Surface syntax:

```
(aggregate
  (belief "Paris is capital of France" 0.5 "prior")
  (belief "Paris is capital of France" 0.7 "textbook"))
```

Lean representation:

```
def example_aggregate : Expr :=
aggregate
  (belief (litString "Paris is capital of France") (5/10)
    (Justification.axiomJ "prior"))
  (belief (litString "Paris is capital of France") (7/10)
    (Justification.axiomJ "textbook"))
```

Step-by-step evaluation:

1. Both arguments are values (beliefs)
2. Apply probabilistic sum: `#text(code) [c_new = c1 + c2 - c1*c2]`
3. `#text(code) [c_new = 0.5 + 0.7 - 0.5*0.7 = 1.2 - 0.35 = 0.85]`
4. Combined justification: `#text(code) [Justification.agg ["prior", "textbook"]]`
5. Result: `#text(code) [belief "Paris..." 0.85 (agg ["prior", "textbook"])]`

This demonstrates the probabilistic sum operator: two independent pieces of evidence combine to increase confidence beyond either individual source.

13.4.3. B.4.3 Undercutting in Action

Example Defeasible Reasoning. Scenario: A sensor reports “temperature is 25°C” with confidence 0.8, but a calibration warning suggests the sensor may be malfunctioning with confidence 0.3.

Surface syntax:

```
(undercut
  (belief 25 0.8 "sensor-A")
  (belief "sensor-A unreliable" 0.3 "calibration-warning"))
```

Lean representation:

```
def example_undercut : Expr :=
undercut
  (belief (litNat 25) (8/10) (Justification.axiomJ "sensor-A"))
  (belief (litString "sensor-A unreliable") (3/10)
    (Justification.axiomJ "calibration-warning"))
```

Step-by-step evaluation:

1. Both arguments are values
2. Apply undercut formula: `#text(code) [c_new = c1 * (1 - c2)]`
3. `#text(code) [c_new = 0.8 * (1 - 0.3) = 0.8 * 0.7 = 0.56]`
4. `#text(code) [belief 25 0.56 (undercut_j "sensor-A" "calibration-warning")]`

The calibration warning reduces confidence from 0.8 to 0.56. The justification tracks that this belief was undercut, preserving the reasoning history.

13.4.4. B.4.4 Rebuttal and Confidence Collapse

Example Conflicting Evidence. Scenario: Two sources disagree on whether a fact holds, with confidences 0.7 and 0.3 respectively.

Lean representation:

```
def example_rebut : Expr :=
rebut
  (belief (litBool true) (7/10) (Justification.axiomJ "source-A"))
  (belief (litBool false) (3/10) (Justification.axiomJ "source-B"))
```

Step-by-step evaluation:

1. Both arguments are values
2. Apply rebuttal formula: `#text(code) [c_new = c1 / (c1 + c2)]`
3. `#text(code) [c_new = 0.7 / (0.7 + 0.3) = 0.7 / 1.0 = 0.7]`

4. Result: `#text(code)[belief true 0.7 (rebut_j "source-A" "source-B")]`

Note that when `#text(code)[c1 = c2]` (equally strong conflicting evidence), confidence collapses to `#text(code)[1/2]`. This represents a state of maximal uncertainty.

13.4.5. B.4.5 Derivation Chain

Example Multi-Step Derivation. Scenario: Derive a conclusion from two premises using the `derive` operator.

Lean representation:

```
def example_derivation : Expr :=
  derive
  (belief (litNat 1) (8/10) (Justification.axiomJ "premise1"))
  (belief (litNat 2) (8/10) (Justification.axiomJ "premise2"))
```

Step-by-step evaluation:

1. Both arguments are values
2. Apply derivation formula: `#text(code)[c_new = c1 * c2]`
3. `#text(code)[c_new = 0.8 * 0.8 = 0.64]`
4. `#text(code)[belief (pair 1 2) 0.64 (rule "derive" ["premise1", "premise2"])]`

The result pairs the premise values, and confidence is the product (representing the conjunction strength).

The justification records that this came from a derivation rule.

13.5. B.5 Key Properties

The Lean formalization proves five key properties that demonstrate CLAIR's correctness as an epistemic reasoning system:

Definition Property 1: Beliefs Track Confidence. Confidence is preserved through computation: if a belief `#text(code)[b]` has confidence `#text(code)[c]`, then after any sequence of valid reductions `#text(code)[b ->* b']`, the resulting belief `#text(code)[b']` has a confidence value that is a deterministic function of `#text(code)[c]` and the operations applied.

Definition Property 2: Evidence is Affine. No evidence is double-counted. The operator uses probabilistic sum $\#text(code)[a \oplus b = a + b - a*b]$, which equals the probability of the union of independent events. This ensures that aggregating a belief with itself does not increase confidence: $\#text(code)[c \oplus c = c + c - c*c = c(2 - c)]$ which is only equal to $\#text(code)[c]$ when $\#text(code)[c = 0]$ or $\#text(code)[c = 1]$. In practice, justification tracking prevents exact duplicate aggregation.

Definition Property 3: Introspection is Safe. The `#text(code)[introspect]` operator satisfies the stratification constraints defined in Chapter 6. It is impossible to form a belief about the current confidence of that same belief, preventing the formation of self-referential paradoxes.

Definition Property 4: Defeat Operations are Correct. Undercut and rebut satisfy their specifications:

- Undercut is monotonic in both arguments: higher source confidence or higher defeater confidence yields predictable results
- Rebuttal is symmetric: `#text(code)[rebut b1 b2]` and `#text(code)[rebut b2 b1]` yield beliefs about opposing values with complementary confidence

Definition Property 5: Type Checking is Decidable. The bidirectional type checker in `CLAIR.Typing.HasType` terminates for all well-formed inputs. This is proven formally in the Lean code by showing that each recursive call decreases a measure (expression size or stratification level).

13.6. B.6 Implementation Notes

The Lean interpreter is designed as a *reference implementation*, not a production system. Key design decisions:

+---+ | Aspect | Decision | Rationale | +---+ | Fuel | 1000 steps default | Prevents infinite loops while allowing reasonable programs | | Evaluation Order | Call-by-value | Matches intuition about belief formation | | Parser | Minimal (constructors only) | Demonstrates concept without complex parsing logic | | Error Handling | `Option Expr` (partial function) | Stuck states return `none`; can be extended with explicit errors | +---+

For production use, we recommend:

1. A proper parser (e.g., using Megaparsec in Haskell)
2. Exception-based error handling with detailed error messages
3. JIT compilation for performance
4. Persistent justification storage for audit trails
5. Parallel evaluation for independent subexpressions

13.7. B.7 Relation to Chapter 10

Chapter 10 discusses implementation considerations for a production CLAIR system. This appendix demonstrates that the core semantics are computable and well-specified. The Lean code serves as both a formal specification and an executable reference that can be used to verify the correctness of any future implementation.

13.8. B.8 Haskell Reference Implementation

In addition to the Lean formalization, a complete Haskell reference implementation is provided. This implementation demonstrates that CLAIR can be realized in a general-purpose programming language while maintaining the semantic properties proved in Lean.

13.8.1. B.8.1 Project Structure

The Haskell implementation is organized as a Cabal project with the following structure:

```
+---+ | Module | Purpose | +---+ | CLAIR.Syntax | Abstract syntax trees (AST)
for CLAIR expressions | | CLAIR.Confidence | Confidence algebra: ⊕, ⊖, undercut, rebut
| | CLAIR.Parser | Parse surface syntax into AST (Parsec) | | CLAIR.TypeChecker |
Bidirectional type checking with confidence grades | | CLAIR.Evaluator | Small-step operational semantics with fuel | | CLAIR.Pretty | Pretty-printing for values and types | |
CLAIR.TypeChecker.Types | Type checker context and error types | +---+
```

The implementation includes:

- A REPL (`app/Main.hs`) for interactive evaluation

- A comprehensive test suite (`test/`) with 35 passing tests
- QuickCheck properties for algebraic laws
- Example programs demonstrating belief operations

13.8.2. B.8.2 Syntax Definition

The core AST in `CLAIR.Syntax` mirrors the Lean formalization:

```
-- | A belief value with all its annotations

data Belief = Belief
  { beliefValue      :: Expr           -- The proposition/content
  , beliefConf       :: Confidence    -- Confidence level [0,1]
  , beliefJustify    :: Justification -- Supporting arguments
  , beliefInvalidate :: Invalidator  -- Defeating information
  , beliefProvenance :: Provenance   -- Source tracking
  }

-- | Core expression language

data Expr
  = EVar Name          -- Variable: x
  | ELam Name Type Expr -- Lambda: λx:A. e
  | EApp Expr Expr     -- Application: e1 e2
  | EAnn Expr Type     -- Type annotation: e : A
  | EBelief Belief      -- Belief: belief(v,c,j,i,p)
  | EBox Confidence Expr -- Self-reference: □_c e
  | EPrim Op Expr Expr -- Primitive operation
  | ELit Literal        -- Literal value
```

Key design decisions:

- GADTs enable type-safe AST construction
- Deriving `Generic` and `ToJSON/FromJSON` enables serialization
- Provenance and justification are explicit for audit trails

13.8.3. B.8.3 Confidence Algebra

The `CLAIR.Confidence` module implements the confidence operations with careful attention to semantic correctness:

```
-- | Probabilistic sum: a ⊕ b = 1 - (1-a)(1-b) = a + b - ab
-- Assumes: sources are conditionally independent
oplus :: Confidence -> Confidence -> Confidence
oplus (Confidence a) (Confidence b) = clamp (a + b - a * b)

-- | Product t-norm: a ⊗ b = a * b
otimes :: Confidence -> Confidence -> Confidence
otimes (Confidence a) (Confidence b) = clamp (a * b)

-- | Apply undercut defeat: multiply by (1-d)
-- Rationale: Undercut attacks the evidential connection
undercut :: Defeat -> Confidence -> Confidence
undercut (Defeat d) (Confidence c) = clamp (c * (1 - d))

-- | Apply rebut with normalization
-- Limitation: Collapses absolute strength; considers uncertainty-
preserving alternatives
rebut :: Defeat -> Confidence -> Confidence -> Confidence
rebut (Defeat d_strength) (Confidence c_for) (Confidence c_against_base) =
    let c_against = d_strength * c_against_base
        total = c_for + c_against
    in if total == 0
        then Confidence 0.5 -- ignorance prior
        else clamp (c_for / total)

-- | Square discount for self-reference: g(c) = c²
-- Prevents bootstrapping while preserving high-confidence self-endorsement
```

```

squareDiscount :: DiscountFn
squareDiscount (Confidence c) = clamp (c * c)

```

13.8.4. B.8.4 Type Checking

The bidirectional type checker in `CLAIR.TypeChecker` implements the rules from Appendix E:

```

-- | Infer (synthesize):  $\Gamma \vdash e \uparrow \tau$ 

infer :: Context -> Expr -> Either TypeError TCResult

infer ctx expr = case expr of
    -- Variable:  $\Gamma \vdash x : \Gamma(x)$ 
    EVar x -> case ctxLookup x ctx of
        Just ty -> return (TCResult ty ctx)
        Nothing -> Left (UnboundVar x)

    -- Application:  $\Gamma \vdash e \uparrow \tau_1 \rightarrow \tau_2, \Gamma \vdash e' \downarrow \tau_1$ 
    EApp e1 e2 -> do
        TCResult ty1 ctx1 <- infer ctx e1
        case ty1 of
            TFun argTy resTy -> do
                ctx2 <- check ctx1 e2 argTy
                return (TCResult resTy ctx2)
            _ -> Left (NotFunction ty1)

    -- Belief:  $\Gamma \vdash e : \tau, c \in [0,1]$ 
    --
    --      _____
    --       $\Gamma \vdash \text{belief}(e,c) \uparrow \text{Belief\_c}[\tau]$ 
    EBelief (Belief e c _ _ _) -> do
        unless (isNormalized c) $
            Left (InvalidConfidence c)
        TCResult ty ctx' <- infer ctx e

```

```

let beliefTy = TBelief c ty
return (TCResult beliefTy ctx')

```

The bidirectional approach provides:

- Better error messages (checking mode guides synthesis)
- Natural handling of implicit arguments
- Clear separation of inference vs. verification

13.8.5. B.8.5 Evaluation

The small-step evaluator in `CLAIR.Evaluator` implements the operational semantics with fuel-bounded termination:

```

-- | Single-step reduction: e → e'
step :: Env → Expr → Either EvalError (Maybe Expr)
step env expr = case expr of
  -- E-Beta: (λx:τ.e) v → e[x := v]
  EApp (ELam varName _ty body) arg
    | isValue arg -> do
      v <- evalExpr env arg
      return (Just (subst varName v body))

  -- E-Prim: Reduce primitive operations
  EPrim op e1 e2
    | isValue e1 && isValue e2 -> evalPrimOp env op e1 e2

  -- E-Belief: Evaluate belief content to value
  EBelief (Belief e c j i p) -> do
    me' <- step env e
    case me' of
      Just e' -> return (Just (EBelief (Belief e' c j i p)))
      Nothing -> return Nothing -- Fully evaluated

```

```
-- E-Box: □_c e becomes value when e is fully evaluated

EBox c e -> do
    me' <- step env e
    case me' of
        Just e' -> return (Just (EBox c e'))
        Nothing -> return Nothing
```

Key features:

- **Fuel:** 1,000,000 steps default prevents infinite loops
- **Call-by-value:** Arguments evaluated before application
- **Capture-avoiding substitution:** Preserves variable hygiene
- **Error recovery:** Detailed error messages for debugging

13.8.6. B.8.6 REPL Usage

The interactive REPL ([app/Main.hs](#)) supports:

```
clair> :help
CLAIR REPL – Commands:
:quit  Exit the REPL
:help  Show this help message
```

Examples:

```
5
λx:Nat. x
(λx:Nat. x + 1) 5
3 + 4
□0.8 true
belief(5, 0.9, none, none, none)
```

The REPL provides:

- Parse error reporting with locations

- Type checking with confidence grades
- Evaluation with fuel tracking
- Pretty-printed results

13.8.7. B.8.7 Test Suite

The implementation includes 35 QuickCheck and HUnit tests covering:

```
+---+---+ | Module | Tests | Coverage | +---+---+ | CLAIR.Test.Confidence | 12 |
Algebraic laws: ⊕ associativity, ⊖ identity, undercut monotonicity | | CLAIR.Test.Evaluator
| 14 | Beta reduction, primitive operations, belief evaluation | | CLAIR.Test.TypeChecker
| 6 | Type inference, subtyping, error cases | | CLAIR.Test.HelloWorld | 3 | End-to-end
belief formation and evaluation | +---+---+
```

Sample QuickCheck properties:

```
-- Probabilistic sum is associative
prop_oplus_assoc :: Confidence -> Confidence -> Confidence -> Bool
prop_oplus_assoc a b c =
    oplus a (oplus b c) == oplus (oplus a b) c

-- Undercut is monotonic in defeat strength
prop_undercut_monotonic :: Confidence -> Defeat -> Defeat -> Property
prop_undercut_monotonic c d1 d2 =
    d1 <= d2 ==> undercut d1 c >= undercut d2 c
```

All tests pass: `cabal test` → 35/35 tests passed.

13.8.8. B.8.8 Building and Running

13.8.8.1. Build

```
cd implementation/haskell
cabal build
```

This compiles the REPL (`clair-repl`) and test suite (`clair-test`).

13.8.8.2. Run REPL

```
cabal run clair-repl
```

13.8.8.3. Run Tests

```
cabal test
```

13.8.9. B.8.9 Design Rationale

The Haskell implementation makes specific design choices that differ from the Lean reference:

+—+—+—+ | Aspect | Lean | Haskell | Rationale | +—+—+—+ | Confidence type | **Rat**
(rational) | **Double** | Performance vs. precision tradeoff | | Parser | Constructors only |
Parsec | Usable surface syntax | | Error handling | **Option Expr** | **Either EvalError** |
Stack traces for debugging | | Fuel | 1000 steps | 1,000,000 steps | Allow more complex
programs | | Substitution | Formal proof | Direct implementation | No proof burden, faster
iteration | +—+—+—+

These choices reflect the different purposes:

- Lean: Formal verification, mathematical rigor
- Haskell: Usability, testing, experimentation

13.8.10. B.8.10 Relation to Lean Formalization

The Haskell implementation is verified against the Lean formalization by:

1. Type system correspondence: Haskell types mirror Lean inductive types
2. Semantic equivalence: Reduction rules match Lean's **step** relation
3. Property-based testing: QuickCheck laws reflect Lean theorems
4. Test coverage: Each Lean example has a corresponding Haskell test

Discrepancies are documented as limitations:

- **Double** precision may cause floating-point drift
- Substitution is not proven capture-avoiding
- Stratification is checked but not enforced statically

For production deployment, we recommend:

- Use arbitrary-precision rationals for confidence
- Add formal proofs for substitution correctness
- Implement static stratification analysis
- Add persistent justification storage

14. Additional Proofs

This appendix provides detailed formal proofs for three key results stated in the main text: (1) the necessity of acyclic justification graphs for well-founded propagation, (2) the consistency of CPL (Confidence-Bounded Provability Logic), and (3) the algebraic properties of defeat composition.

14.1. C.1 DAG Necessity for Well-Founded Confidence Propagation

14.1.1. C.1.1 Statement of the Problem

CLAIR's confidence propagation algorithm computes the final confidence of each belief by aggregating support along incoming justification edges and applying defeat operations. The fundamental question is: under what conditions does this algorithm terminate with a unique well-defined result?

14.1.2. C.1.2 The Cyclic Counterexample

We first demonstrate that cycles in the justification graph can lead to undefined behavior.

C.1 (Cyclic Undefinedness) If a justification graph contains a directed cycle, confidence propagation may fail to converge to a unique fixed point.

Proof. Consider a simple cycle of two beliefs:

- Belief A with initial confidence $c_A = 0.5$, derived from belief B with strength $s_1 = 0.8$
- Belief B with initial confidence $c_B = 0.5$, derived from belief A with strength $s_2 = 0.8$

The propagation rules specify: $c'_A = c_A \oplus (c_B \times s_1)$ $c'_B = c_B \oplus (c_A \times s_2)$

Starting from $(c_A, c_B) = (0.5, 0.5)$, iteration yields:

Iteration 1: $c'_A = 0.5 \oplus (0.5 \times 0.8) = 0.5 \oplus 0.4 = 0.7$ $c'_B = 0.5 \oplus (0.5 \times 0.8) = 0.7$

Iteration 2: $c'_A = 0.5 \oplus (0.7 \times 0.8) = 0.5 \oplus 0.56 = 0.78$ $c'_B = 0.78$

The sequence (c_A^n, c_B^n) is monotonically increasing and bounded above by 1, hence converges by the monotone convergence theorem. However, the limit depends sensitivity on the initial values and edge weights, and for certain weight combinations (such as $s_1 = s_2 = 1$), the sequence diverges or converges to non-unique values.

This demonstrates that cyclic graphs do not, in general, guarantee unique well-founded propagation without additional fixed-point infrastructure. ■■

14.1.3. C.1.3 The DAG Well-Foundedness Theorem

C.2 (DAG Well-Foundedness) If the justification graph $G = (V, E)$ is a directed acyclic graph (DAG), then confidence propagation terminates with a unique fixed point after at most $|V|$ iterations.

Proof. We proceed by induction on the topological ordering of the DAG.

Let \sqsubseteq be a topological order on V , meaning that if $(u, v) \in E$, then $u \sqsubseteq v$. Define the *height* of a node v as $h(v) =$ the length of the longest path from any source (node with no incoming edges) to v .

Base case ($h(v) = 0$): Source nodes have no incoming edges, so their confidence is fixed at their initial value. Propagation terminates immediately for these nodes.

Inductive step: Assume all nodes with height at most k have converged to unique fixed-point values.

Consider a node v with height $h(v) = k + 1$.

All predecessors u of v satisfy $h(u) \leq k$, hence have converged by the inductive hypothesis. Let $c_{u\text{-final}}$ denote the converged confidence of predecessor u .

The propagation rule for v is: $c_v = c_v(\text{init}) \oplus \bigvee_{(u,v) \in E} (c_{u\text{-final}} \times w_{uv})$

Since each c_u has converged to $c_{u\text{-final}}$, and all operations (\oplus and \times) are continuous, the value of c_v is uniquely determined and converges in exactly one iteration once all predecessors have converged. By induction, all nodes converge to unique values. The maximum number of iterations required is $\max_v \in V h(v) \leq |V| - 1$, the length of the longest path in the DAG. ■■

14.1.4. C.1.4 Practical Implications

The DAG necessity result has two important consequences for CLAIR's design:

1. *Type-System Enforcement*: The CLAIR type checker must enforce acyclicity as a well-formedness condition on justification graphs. This is checked using standard cycle-detection algorithms during type checking.
 2. *Expressive Limitations*: DAG-only semantics cannot directly represent certain defeasible reasoning patterns involving mutual support. Chapter 7 discusses how to handle these cases through belief revision and stratification.
-

14.2. C.2 CPL Consistency Proof

14.2.1. C.2.1 CPL Axiom System

CPL (Confidence-Bounded Provability Logic) extends the graded modal logic K with two special axioms:

Axiom 4 (Graded Transitivity) The graded modal axiom 4: $\text{box}_c(p) \rightarrow \text{box}_{fc}(\text{box}_c(p))$ for some strictly increasing function f.

Axiom GL (Graded Löb) The graded Löb axiom: $\text{box}_c(\text{box}_c(p) \rightarrow p) \rightarrow \text{box}_{gc}(p)$

The key question is whether this axiom system is *consistent*—that is, whether there exists a non-trivial model satisfying all axioms.

14.2.2. C.2.2 Finite Model Construction

We construct an explicit finite model demonstrating consistency.

C.3 (CPL Consistency) There exists a finite Kripke model $M = (W, R, V, c)$ satisfying all CPL axioms for $f(c) = c$ and $g(c) = c^2$.

Proof. Let $W = \{0, 1, 2\}$ be a set of three worlds. Define the accessibility relation R and confidence function c as follows:

$$R = \{(0, 1), (0, 2), (1, 2)\} \text{ (strictly increasing order)}$$

For each edge $(w, w') \in R$, define $c(w, w')$ as:

- $c(0, 1) = 1/2$
- $c(0, 2) = 1/4$
- $c(1, 2) = 1/2$

We verify the axioms:

Axiom K: Holds in all Kripke models by the standard modal logic semantics.

Axiom 4: For any world w and confidence c , if w satisfies $\text{box}_c p$, then for all w' with $(w, w') \in R$ and $c(w, w') \geq c$, we have w' satisfies p . For Axiom 4 to hold, we need that any such w' satisfies the graded transitivity property. In our model, this holds because if $(w, w') \in R$ and $(w', w'') \in R$ with appropriate confidence bounds, then $(w, w'') \in R$ directly.

Axiom GL: The graded Löb axiom requires verification. For world 0: $\text{box}_c (\text{box}_c p \rightarrow p)$ means: for all worlds reachable from 0 with confidence $\geq c$, if $\text{box}_c p$ holds at that world, then p holds.

Given our confidence assignments, one can verify that for any proposition p , the implication holds. The discount function $g(c) = c^2$ ensures that self-referential soundness claims cap their own confidence, preventing bootstrapping.

Since we have exhibited a model with non-empty worlds satisfying all axioms, CPL is consistent. ■■

14.2.3. C.2.3 Design Axiom Status

It is important to clarify that CPL's graded Löb axiom is a *design axiom* rather than a derived semantic theorem. This means:

1. We *choose* to include $\text{box}_c (\text{box}_c p \rightarrow p) \rightarrow \text{box}_{gc} p$ as an axiom
2. We verify *consistency* (as shown above) but not *soundness* with respect to a pre-existing semantics
3. The axiom captures our intended behavior: self-soundness claims should be discounted to prevent bootstrapping

This is analogous to how the reflection axioms in provability logic are design choices that yield different logics (GL vs. S4 vs. K).

14.3. C.3 Defeat Composition Algebra

14.3.1. C.3.1 Undercut Composition

The fundamental result for undercutting defeat is that sequential undercuts compose via the probabilistic OR operation.

C.4 (Undercut Composition) For any confidences $c, d_1, d_2 \in [0,1]$: $\text{undercut}(\text{undercut}(c, d_1), d_2) = \text{undercut}(c, d_1 \oplus d_2)$

Proof. We compute directly from the definition $\text{undercut}(c, d) = c \times (1 - d)$:

$$\text{undercut}(\text{undercut}(c, d_1), d_2)$$

15. undercut($c \times (1 - d_1)$, d_2)

16. $(c \times (1 - d_1)) \times (1 - d_2)$

17. $c \times ((1 - d_1) \times (1 - d_2))$

18. $c \times (1 - d_1 - d_2 + d_1 d_2)$

19. $c \times (1 - (d_1 + d_2 - d_1 d_2))$

20. $c \times (1 - (d_1 \oplus d_2))$

21. undercut(c , $d_1 \oplus d_2$)

where we use the definition $d_1 \oplus d_2 = d_1 + d_2 - d_1 d_2$. ■■

21.0.1. C.3.2 Corollaries of Undercut Composition

C.5 (Commutative Composition) Undercut composition is commutative: $\text{undercut}(\text{undercut}(c, d_1), d_2) = \text{undercut}(\text{undercut}(c, d_2), d_1)$

Proof. Immediate from Theorem C.4 and commutativity of \oplus : $\text{undercut}(\text{undercut}(c, d_1), d_2) = \text{undercut}(c, d_1 \oplus d_2)$

22. $\text{undercut}(c, d_2 \oplus d_1) = \text{undercut}(\text{undercut}(c, d_2), d_1)$. ■

C.6 (Identity) Undercut with zero defeat leaves confidence unchanged: $\text{undercut}(c, 0) = c$

Proof. $\text{undercut}(c, 0) = c \times (1 - 0) = c \times 1 = c$. ■■

C.7 (Annihilation) Undercut with complete defeat eliminates confidence: $\text{undercut}(c, 1) = 0$

Proof. $\text{undercut}(c, 1) = c \times (1 - 1) = c \times 0 = 0$. ■■

22.0.1. C.3.3 Rebut Algebra

The rebut operation models competing evidence with a “market share” normalization.

C.8 (Rebut Symmetry) For any $c_{\text{for}}, c_{\text{against}} \in [0, 1]$ with $c_{\text{for}} + c_{\text{against}} > 0$: $\text{rebut}(c_{\text{for}}, c_{\text{against}}) + \text{rebut}(c_{\text{against}}, c_{\text{for}}) = 1$

Proof. From the definition $\text{rebut}(c_{\text{for}}, c_{\text{against}}) = c_{\text{for}} / (c_{\text{for}} + c_{\text{against}})$:

$$\text{rebut}(c_{\text{for}}, c_{\text{against}}) + \text{rebut}(c_{\text{against}}, c_{\text{for}})$$

23. $c_{\text{for}} / (c_{\text{for}} + c_{\text{against}}) +$

$c_{\text{against}} / (c_{\text{against}} + c_{\text{for}})$

24. $(c_{\text{for}} + c_{\text{against}}) / (c_{\text{for}} +$

$c_{\text{against}})$

25. 1 ■

C.9 (Rebut Monotonicity) Rebut is monotone in the first argument and antitone in the second:

- If $c_1 \leq c_2$, then $\text{rebut}(c_1, c) \leq \text{rebut}(c_2, c)$
- If $d_1 \leq d_2$, then $\text{rebut}(c, d_2) \leq \text{rebut}(c, d_1)$

Proof. For the first claim: $\text{rebut}(c_1, c) = c_1 / (c_1 + c) \leq c_2 / (c_2 + c) = \text{rebut}(c_2, c)$ since the function $f(x) = x / (x + c)$ is increasing in x for $c > 0$.

The second claim follows similarly by noting $\text{rebut}(c, d) = 1 - \text{rebut}(d, c)$ from Theorem C.8. ■■

25.0.1. C.3.4 Interaction Between Undercut and Rebut

Undercut and rebut represent two fundamentally different types of defeat:

- *Undercut* attacks the *inferential link* between premises and conclusion
- *Rebut* attacks the *conclusion* directly with counter-evidence

The interaction of these two operations is subtle and context-dependent. In general, rebut is applied first to aggregate competing evidence, then undercut is applied to discount the resulting confidence based on link attackers. This ordering is reflected in CLAIR’s evaluation semantics (Chapter 4).

25.0.2. C.3.5 Limitation: Rebut Normalization

A key limitation of the rebut operation is that it normalizes away absolute evidence strength.

C.10 (Rebut Collapse) For any scaling factor $\lambda > 0$: $\text{rebut}(\lambda c_{\text{for}}, \lambda c_{\text{against}}) = \text{rebut}(c_{\text{for}}, c_{\text{against}})$

Proof. $\text{rebut}(\lambda c_{\text{for}}, \lambda c_{\text{against}}) = \lambda c_{\text{for}} / (\lambda c_{\text{for}} + \lambda c_{\text{against}})$

26. $\lambda c_{\text{for}} / (\lambda (c_{\text{for}} + c_{\text{against}}))$

27. $c_{\text{for}} / (c_{\text{for}} + c_{\text{against}})$

28. $\text{rebut}(c_{\text{for}}, c_{\text{against}})$. ■

■

This means rebut cannot distinguish between “both weak” and “both strong but balanced” evidence. Chapter 6 discusses how CLAIR addresses this through provenance tracking, which preserves the absolute strengths of individual evidence sources even after rebut normalization.

29. Glossary

This appendix provides concise definitions of key terms, notation, and acronyms used throughout this dissertation. Terms are marked with their primary chapter of introduction.

29.1. D.1 Term Definitions

29.1.1. Epistemic Terms

+— +Term | Definition | Chapter +— +Belief | A first-class value in CLAIR consisting of a proposition, confidence, and justification. Beliefs can be constructed, combined, defeated, and inspected through language operations. | 1 +Confidence | A value in the unit interval [0,1] representing degree of epistemic commitment. Confidence is *not* probability: rather than quantifying frequency or subjective chance, confidence tracks how justified a belief is given available evidence. | 3 +Commitment (Epistemic) | The stance of endorsing a proposition as true to a specified degree. High confidence (close to 1) indicates strong commitment; low confidence (close to 0) indicates weak or no commitment. | 3 +Justification | A data structure tracking the derivation history of a belief. Justifications form a directed acyclic graph (DAG) where nodes represent beliefs and edges represent inference rules or evidence sources. | 4 +Invalidation | A condition associated with a belief specifying circumstances under which the belief should be reconsidered. Invalidation enables defeasible reasoning—beliefs can be defeated without logical contradiction. | 4 +Provenance | The origin or source of evidence for a belief. CLAIR tracks provenance through justification graphs, enabling audit trails for how conclusions were reached. | 4 +Reliability | In CLAIR’s semantics, the *semantic interpretation* of confidence: a belief with confidence c is interpreted as coming from a source with reliability c . | 3 +—

29.1.2. Operations and Relations

+— +Term | Definition | Chapter +— +Probabilistic OR (\oplus) | Aggregation operator for independent evidence: $a \oplus b = a + b - ab$. This equals the probability that at least one of two independent events occurs. Satisfies commutativity, associativity, and has identity 0. | 3
+Undercut (?) | Defeat operation that reduces confidence multiplicatively: $\text{undercut}(c, d) = c \times (1-d)$. A defeater with confidence d reduces target confidence by factor $(1-d)$. | 4 +Rebuttal | Defeat operation for conflicting evidence: $\text{rebut}(c_1, c_2) = c_1 / (c_1 + c_2)$. Normalizes competing confidences to [0,1]; equal confidences yield 1/2. | 4 +Derivation | Combining beliefs via rule application: $\text{derive}(b_1, b_2)$ pairs the values and multiplies confidences $c_1 \times c_2$. Tracks justification through rule application. | 4 +Aggregation | Combining independent evidence using \oplus : $\text{aggregate}(b_1, b_2)$ produces belief with confidence $c_1 \oplus c_2$. Tracks justification through aggregation node. | 4 +Subtyping | Confidence ordering: belief at confidence c_1 can be used where belief at c_2 is required iff $c_1 \geq c_2$. Enables confidence weakening. | 10 +—

29.1.3. Structural Properties

+— +Term | Definition | Chapter +— +Directed Acyclic Graph (DAG) | A graph with directed edges and no cycles. CLAIR requires justification graphs to be DAGs to ensure well-foundedness and prevent circular reasoning. | 4 +Stratification | Layering beliefs into levels **0, 1, 2, ...** where level **n** can only refer to levels $< n$. Enforces Tarski's hierarchy to prevent self-referential paradoxes. | 6 +Well-formedness | Constraints on CLAIR programs: (1) justification graphs must be acyclic, (2) stratification constraints on introspection, (3) confidences in [0,1]. | 10 +—

29.1.4. Logical and Modal Terms

+— +Term | Definition | Chapter +— +CPL (Confidence-Bounded Provability Logic) | Graded extension of Gödel-Löb provability logic. Adds confidence grades to provability operator \Box . Axiomatizes self-referential reasoning with confidence discounts. | 5 +Löb's Theorem | Modal logic theorem: $\Box(\Box\phi \rightarrow \phi) \rightarrow \Box\phi$. In provability logic, enables self-reference; in CLAIR, motivates anti-bootstrapping constraint. | 5 +Graded Modality | Modal operators

with quantitative gradess. CLAIR's $\Box(c)$ means "provable with confidence at least c ." | 5 +Anti-bootstrapping | Principle that self-validating claims cannot increase confidence. Formally: $c \leq g(c)$ for some discount function g . CLAIR uses $g(c) = c^2$. | 5 +Kripke Semantics | Possible worlds framework for modal logic. CPL uses graded Kripke models where accessibility relations track confidence thresholds. | 5 +—

29.1.5. Computational Terms

+— +Term | Definition | Chapter +— +Fuel | Bound on computation steps in the Lean evaluator: **evalFuel n e** evaluates for at most n steps. Prevents infinite loops while ensuring termination. | B +Call-by-value | Evaluation strategy: function arguments are evaluated before application. CLAIR uses call-by-value to match intuition about belief formation. | B +Small-step semantics | Reduction relation $e \rightarrow e'$ defining one step of computation. Composed to form multi-step evaluation $e \rightarrow e'$. | 10 +Bidirectional Type Checking | Type checking algorithm with synthesis (infer type from expression) and checking (verify expression matches type). Enables practical implementation. | 10 +De Bruijn Indices | Variable representation using natural numbers: var 0 = most recent binder, var 1 = next recent, etc. Enables formal proofs about substitution. | A +—

29.1.6. Argumentation and Belief Revision

+— +Term | Definition | Chapter +— +Defeasible Reasoning | Reasoning where conclusions can be defeated by new information without logical contradiction. CLAIR supports defeasibility through undercut and rebut operations. | 4 +Underminer | An argument that attacks the connection between evidence and conclusion (e.g., "the sensor is miscalibrated"). Reduces confidence via **undercut(c,d) = c * (1-d)**. | 4 +Rebuttal | An argument providing conflicting evidence for the opposite conclusion. Normalizes via **rebut(c1,c2) = c1 / (c1 + c2)**. | 4 +AGM Theory | Classic belief revision framework (Alchourrón, Gärdenfors, Makinson). CLAIR extends AGM to graded, DAG-structured beliefs. | 7 +Contraction | Belief revision operation: remove a belief from a belief set. In CLAIR, achieved by setting confidence to 0. | 7 +Revision |

Belief revision operation: add a belief while maintaining consistency. In CLAIR, achieved by aggregating with existing beliefs. | 7 +—

29.1.7. Impossibility Results

+— +Term | Definition | Chapter +— +Gödel's Incompleteness | Any consistent formal system capable of arithmetic contains true but unprovable statements. Motivates CPL's design restrictions. | 12 +Church's Undecidability | First-order logic validity is undecidable. CLAIR restricts to decidable fragments (CPL-finite, CPL-0). | 12 +Tarski's Undefinability | Truth cannot be defined within the same language. Motivates stratification: level **n** cannot quantify over level **n**. | 12 +Löb's Paradox | Curious proposition using Löb's theorem that leads to contradiction if not carefully restricted. Resolved via stratification. | 6 +—

29.2. D.2 Notation Table

+— +Symbol | Meaning | Type | Chapter +— +**c** | Confidence value in [0,1] | Confidence | 3 +⊕ | Probabilistic OR: $\mathbf{a} \oplus \mathbf{b} = \mathbf{a} + \mathbf{b} - \mathbf{ab}$ | Binary operation | 3 +? | Undercut: $\mathbf{c} ? \mathbf{d} = \mathbf{c} \times (\mathbf{1}-\mathbf{d})$ | Binary operation | 4 +× | Multiplication (conjunctive combination) | Binary operation | 3 +g(**c**) | Löb discount function; CLAIR uses \mathbf{c}^2 | Unary function | 5 +□(**c**) φ | Necessarily φ with confidence at least **c** | Modal operator | 5 +◇(**c**) φ | Possibly φ with confidence at least **c** | Modal operator | 5 +! | Typing judgment: $\Gamma \vdash e : A @ c$ | Relation | 10 +→ | Small-step reduction: $e \rightarrow e'$ | Relation | 10 +→ | **Multi-step reduction (reflexive transitive closure)** | Relation | 10 +Γ | **Typing context** (list of variable: type pairs) | Context | 10 +A ⇒ B | Function type from A to B | Type | 10 +Belief<A> | Belief type holding value of type A | Type constructor | 10 +b.value | Extract value from belief b | Projection | 4 +b.confidence | Extract confidence from belief b | Projection | 4 +b.justification | Extract justification from belief b | Projection | 4 +m <n | Stratification constraint: level m below n | Ordering | 6 +∀ | Universal quantifier | Quantifier | Appendix A +∃ | Existential quantifier | Quantifier | Appendix A +∈ | Set membership | Relation | Appendix A +⊆ | Subset relation | Relation | Appendix A +∪ | Set union | Binary operation | Appendix A +∩ | Set intersection | Binary

operation | Appendix A + $\lambda^*x:A.e$ | Lambda abstraction (anonymous function) | Expression
| 10 + e^*1e^*2 | Function application | Expression | 10 +—

29.3. D.3 Acronyms

+— +Acronym | Full Name | Definition +— +CLAIR | Comprehensible LLM AI Intermediate Representation | The formal system presented in this dissertation +CPL | Confidence-Bounded Provability Logic | Graded modal logic extending Gödel-Löb +CPL-finite | CPL with finite confidence set $\{0, 1/n, 2/n, \dots, 1\}$ | Decidable fragment suitable for implementation +CPL-0 | CPL with only confidence 0 (ungraded provability) | Collapses to standard provability logic GL +AGM | Alchourrón-Gärdenfors-Makinson | Classic belief revision theory +DAG | Directed Acyclic Graph | Graph structure for justifications +LLM | Large Language Model | AI systems CLAIR is designed to augment +IR | Intermediate Representation | Compiler representation between source and machine code +AI | Artificial Intelligence | Field of study +—

29.4. D.4 Type System Summary

29.4.1. Base Types

+— +Type | Description | Example Values +— +Nat | Natural numbers (non-negative integers) | 0, 1, 2, 42, 128 +Bool | Boolean values | true, false +String | Text strings | “hello”, “sensor-reading” +Unit | Unit type (single value) | unit +Pair(A, B) | Ordered pair | (1, true), (“x”, 42) +Sum(A, B) | Tagged union (either A or B) | inl(5), inr(true) +A \Rightarrow B | Function type from A to B | $\lambda x:\text{Nat}. x + 1$ +Belief<A> | Belief holding value of type A | belief(42, 0.9, j) +—

29.4.2. Confidence Operations

+— +Operation | Notation | Formula | Identity +— +Probabilistic OR | a \oplus b | a + b - a \times b | 0 +Multiplication | a \times b | a \times b | 1 +Undercut | undercut(c, d) | c \times (1-d) | N/A (d=0

gives c) +Rebuttal | **rebut(c₁, c₂)** | c₁ / (c₁ + c₂) | N/A +Minimum | **min(a, b)** | **min(a, b)** | 0
(if bounded below) +Maximum | **max(a, b)** | **max(a, b)** | 1 (if bounded above) +—

30. Complete CLAIR Language Specification

This appendix provides the complete formal specification of the CLAIR language, including syntax (concrete and abstract), static semantics (type system), and dynamic semantics (operational rules). The specification is self-contained and sufficient for implementing a conforming CLAIR interpreter.

30.1. E.1 Syntax

30.1.1. E.1.1 Type Grammar

The type grammar defines the set of valid types in CLAIR.

+— +Category | Production | Description Base Types | B ::= “Nat” | Natural numbers | |
B ::= “Bool” | Boolean values | | B ::= “String” | Text strings | | B ::= “Unit” | Unit type
(single value) Confidence | c ∈ Q | Rational in [0,1] Types | A ::= B | Base type | | A ::= A
→ B | Function type | | A ::= A × B | Product type | | A ::= A + B | Sum type | | A ::=
“Belief<”A“[“c”]” | Belief type with confidence bound | | A ::= “Meta<”A“[“n”]”[“c”]” |
Stratified meta-belief type +—

30.1.2. E.1.2 Expression Grammar

The expression grammar defines the syntactic forms of CLAIR programs.

```
+— +Category | Production | Description Variables | e ::= x | Variable reference Lambdas |
e ::= " $\lambda$ "x":"A".e | Anonymous function Application | e ::= e e | Function application Pairs
| e ::= "(e , e)" | Ordered pair Projections | e ::= e ." 1" | First projection | | e ::= e
." 2" | Second projection Injections | e ::= "inl@B ("e")" | Left injection | | e ::= "inr@A
("e")" | Right injection Case | e ::= "case" e "of" "inl" x "=>" e " | " "inr" y "=>" e | Sum
elimination Literals | e ::= n | Natural number literal | | e ::= "true" | "false" | Boolean literals
| | e ::= ""s"" | String literal | | e ::= () | Unit literal Beliefs | e ::= "belief(e , c , j)" |
Belief constructor | | e ::= "val(e)" | Extract belief value | | e ::= "conf(e)" | Extract belief
confidence | | e ::= "just(e)" | Extract belief justification Derivation | e ::= "derive(e , e)"
| Combine beliefs (conjunctive) Aggregation | e ::= "aggregate(e , e)" | Aggregate beliefs
(independent) Defeat | e ::= "undercut(e , e)" | Apply undercut | | e ::= "rebut(e , e)"
| Apply rebuttal Introspection | e ::= "introspect(e)" | Safe self-reference Let Binding
| e ::= "let" x "=" e "in" e | Local binding +—
```

30.1.3. E.1.3 Abstract Syntax

The abstract syntax uses de Bruijn indices for variable representation.

Definition The abstract syntax “Expr” is defined inductively with the following constructors:

“var(n)” — Variable at de Bruijn index n “lam(A, e)” — Lambda abstraction “ λ ”：“A”.“e” “app(e₁, e₂)” — Function application e₁ e₂ “pair(e₁, e₂)” — Ordered pair “fst(e)” — First projection “snd(e)” — Second projection “inl(B, e)” — Left injection “inr(A, e)” — Right injection “case(e, e₁, e₂)” — Case analysis “litNat(n)” — Natural literal “litBool(b)” — Boolean literal “litString(s)” — String literal “litUnit” — Unit literal “belief(v, c, j)” — Belief constructor “val(e)” — Value extraction “conf(e)” — Confidence extraction “just(e)” — Justification extraction “derive(e₁, e₂)” — Derivation “aggregate(e₁, e₂)” — Aggregation “undercut(e, d)” — Undercut “rebut(e₁, e₂)” — Rebuttal “introspect(e)” — Introspection “letIn(e₁, e₂)” — Let binding

The “Justification” type tracks derivation structure: “axiomJ(name)” — Named axiom “rule(name, js)” — Named rule with premises “agg(js)” — Aggregation “undercut_j(j, d)” — Undercut application “rebut_j(j₁, j₂)” — Rebuttal application

30.1.4. E.1.4 Well-Formedness

A type is well-formed if all confidence bounds are in [0,1].

Definition The judgment “wf(A)” holds when:

“wf(B)” for all base types B “wf(A)” \wedge “wf(B)” \Rightarrow “wf($A \rightarrow B$)” “wf(A)” \wedge “wf(B)” \Rightarrow “wf($A \times B$)”
“wf(A)” \wedge “wf(B)” \Rightarrow “wf($A + B$)” “wf(A)” \wedge $c \in [0,1]$ \Rightarrow “wf(Belief<“A”[“c”])” “wf(A)” \wedge $c \in [0,1]$ \Rightarrow
“wf(Meta<“A”[“n”][“c”])”

30.2. E.2 Static Semantics (Type System)

30.2.1. E.2.1 Typing Contexts

A typing context Γ maps variable indices to type-confidence pairs.

$\Gamma ::= \emptyset \mid \Gamma, \langle A, c \rangle$

Definition The lookup operation $\Gamma.\text{lookup}(n)$ returns the type-confidence pair at index n (0-indexed from most recent binding).

30.2.2. E.2.2 Typing Judgment Form

The primary typing judgment has the form:

$\Gamma \vdash e : A @ c$

“In context Γ , expression e has type A with confidence bound c .”

30.2.3. E.2.3 Typing Rules

The following rules define well-typed CLAIR expressions. Confidence propagation is explicit in each rule.

T-Variable: $\Gamma \vdash \text{var}(n) : A @ c$ if $\Gamma.\text{lookup}(n) = \langle A, c \rangle$

T-Abstraction: If $\Gamma, \langle A, c_A \rangle \vdash e : B @ c_B$, then $\Gamma \vdash \text{lam}(A, e) : (A \rightarrow B) @ c_B$

T-Application: If $\Gamma \vdash e_1 : (A \rightarrow B) @ c_1$ and $\Gamma \vdash e_2 : A @ c_2$, then $\Gamma \vdash \text{app}(e_1, e_2) : B @ (c_1 \times c_2)$

Confidence multiplies for conjunctive derivation.

T-Pair: If $\Gamma \vdash e_1 : A @ c_1$ and $\Gamma \vdash e_2 : B @ c_2$, then $\Gamma \vdash \text{pair}(e_1, e_2) : (A \times B) @ (c_1 \times c_2)$

T-Fst/T-Snd preserve confidence.

T-Inl/T-Inr preserve confidence.

T-Case: If $\Gamma \vdash e :: (A \oplus B) @ c$ and $\Gamma , \langle A, c \rangle \vdash e_1 :: C @ c_1$ and $\Gamma , \langle B, c \rangle \vdash e_2 :: C @ c_2$, then $\Gamma \vdash \text{case}(e, e_1, e_2) :: C @ (c \oplus c_1 \oplus c_2)$

Uses probabilistic OR for confidence aggregation.

T-Belief: If $\Gamma \vdash v :: A @ 1$ and $c_{\text{bound}} \leq c_{\text{actual}}$, then $\Gamma \vdash \text{belief}(v, c_{\text{actual}}, j) :: \text{Belief} < A["c_{\text{bound}}"] @ c_{\text{bound}}$

The belief's actual confidence c_{actual} must meet or exceed the declared bound c_{bound} .

T-Val/T-Conf/T-Just preserve the enclosing confidence.

T-Derive: If $\Gamma \vdash e_1 :: \text{Belief} < A["c_1"] @ c_{e1}$ and $\Gamma \vdash e_2 :: \text{Belief} < B["c_2"] @ c_{e2}$, then $\Gamma \vdash \text{derive}(e_1, e_2) :: \text{Belief} < A \times B["c_1 \times c_2"] @ (c_{e1} \times c_{e2})$

Confidence multiplies: both premises must be true.

T-Aggregate: If $\Gamma \vdash e_1 :: \text{Belief} < A["c_1"] @ c_{e1}$ and $\Gamma \vdash e_2 :: \text{Belief} < A["c_2"] @ c_{e2}$, then $\Gamma \vdash \text{aggregate}(e_1, e_2) :: \text{Belief} < A["c_1 \oplus c_2"] @ (c_{e1} \oplus c_{e2})$

Uses probabilistic OR: independent evidence.

T-Undercut: If $\Gamma \vdash e :: \text{Belief} < A["c"] @ c_e$ and $\Gamma \vdash d :: \text{Belief} < d_c @ c_d$, then $\Gamma \vdash \text{undercut}(e, d) :: \text{Belief} < A["\text{undercut}(c, d_c)] @ (c_e \times c_d)$ where $\text{undercut}(c, d) = c \times (1 - d)$.

T-Rebut: If $\Gamma \vdash e_{\text{for}} :: \text{Belief} < A["c_{\text{for}}"] @ c_{e1}$ and $\Gamma \vdash e_{\text{against}} :: \text{Belief} < A["c_{\text{against}}"] @ c_{e2}$, then $\Gamma \vdash \text{rebut}(e_{\text{for}}, e_{\text{against}}) :: \text{Belief} < A["\text{rebut}(c_{\text{for}}, c_{\text{against}})] @ (c_{e1} \times c_{e2})$ where $\text{rebut}(c_{\text{for}}, c_{\text{against}}) = c_{\text{for}}$

T-Introspect: If $\Gamma \vdash e :: \text{Meta} < A["m"]["c"] @ c_e$ and $m < n$, then:

The resulting type is: $\text{Meta} < \text{Meta} < A["m"]["c"] @>["n"]["g(c)]$ (a meta-belief about a meta-belief).

Thus: $\Gamma \vdash \text{introspect}(e) :: \text{Meta} < \text{Meta} < A["m"]["c"] @>["n"]["g(c)] @ c_e$ where $g(c) = c^2$ is the Löb discount function (to prevent bootstrapping).

Requires level constraint: only higher levels can introspect lower levels.

30.2.4. E.2.4 Subtyping

CLAIR supports subtyping based on confidence bounds.

S-Belief: $\text{Belief} < A["c"] <: \text{Belief} < A["c']$ if $c \geq c'$

Higher confidence is a subtype of lower confidence.

S-Meta follows the same rule.

T-Sub: If $\Gamma \vdash e :: A @ c$ and $A <: A'$ and $c \geq c'$, then $\Gamma \vdash e :: A' @ c'$

Allows weakening both type and confidence.

30.3. E.3 Dynamic Semantics

30.3.1. E.3.1 Values

A value is a fully evaluated expression.

Definition The predicate “`IsValue(e)`” holds for:

- All lambda abstractions “`lam(A, e)`”
- All pairs “`pair(v1, v2)`” where v_1, v_2 are values
- All injections “`inl(B, v)`” and “`inr(A, v)`” where v is a value
- All literals (“`litNat`”, “`litBool`”, “`litString`”, “`litUnit`”)
- All belief constructors “`belief(v, c, j)`” where v is a value

30.3.2. E.3.2 Small-Step Operational Semantics

The small-step relation $e \rightarrow e'$ defines single-step reduction using call-by-value evaluation order.

Beta Reduction: “ $(\lambda : A . e) v \rightarrow e[x := v]$ ” if “`IsValue(v)`”

Context Rules: “`app(e1, e2)`” “ \rightarrow ” “`app(e1', e2)`” if $e_1 \rightarrow e_1'$ “`app(v1, e2)`” “ \rightarrow ” “`app(v1, e2)`” if $e_2 \rightarrow e_2'$ and “`IsValue(v1)`”

Projections: “`fst(pair(v1, v2))`” “ \rightarrow ” v_1 if “`IsValue(v1)`”, “`IsValue(v2)`” “`snd(pair(v1, v2))`” “ \rightarrow ” v_2 if “`IsValue(v1)`”, “`IsValue(v2)`”

Context rules for “`fst`” and “`snd`” evaluate subexpressions first.

Case Analysis: When “`case(inl(v), e1, e2)`” evaluates and v is a value, it reduces to “ $e_1[x := v]$ ”.

When “`case(inr(v), e1, e2)`” evaluates and v is a value, it reduces to “ $e_2[y := v]$ ”.

Context rule for “`case`” evaluates the scrutinee first.

Context rule for “`case`” evaluates the scrutinee first.

Let Binding: “`letIn(v, e)`” “ \rightarrow ” “ $e[x := v]$ ” if “`IsValue(v)`”

Context rule for “`letIn`” evaluates the binding first.

Belief Projections: “val(belief(v, c, j))” “ \rightarrow ” v if “IsValue(v)” “conf(belief(v, c, j))” “ \rightarrow ” “belief(v, c, j)” if “IsValue(v)” “just(belief(v, c, j))” “ \rightarrow ” “litString(toString(j))” if “IsValue(v)”

Context rules evaluate subexpressions to values first.

Derivation, Aggregation, Defeat:

These operations evaluate to values using the confidence operations defined in the typing rules.

The small-step semantics provides context rules that evaluate both subexpressions to values before computing the result.

For example, when “derive(v₁, v₂)” has both subexpressions as values, the evaluator computes a new belief with:

- Value: “pair(v₁.value, v₂.value)”
- Confidence: “v₁.confidence × v₂.confidence”
- Justification: “rule(“derive”, [v₁.just, v₂.just])”

Introspection: “introspect(v)” “ \rightarrow ” v if “IsValue(v)”

Context rule evaluates subexpression first.

30.3.3. E.3.3 Multi-Step Reduction

The multi-step reduction relation e “ $\rightarrow\rightarrow$ ” e' is the reflexive-transitive closure of “ \rightarrow ”.

e “ $\rightarrow\rightarrow$ ” e if e = e' e “ $\rightarrow\rightarrow$ ” e'' if e “ \rightarrow ” e' and e' “ $\rightarrow\rightarrow$ ” e''

30.3.4. E.3.4 Evaluation Function

The evaluation function “eval(e)” returns the result of reducing e to a value, or fails if:

1. The expression gets stuck (no applicable rule)
2. The expression exceeds the fuel bound (default: 1000 steps)

Definition “eval(e)” = “some(v)” if e “ $\rightarrow\rightarrow$ ” v and “IsValue(v)” “eval(e)” = “none” otherwise (stuck or out of fuel)

30.4. E.4 Well-Formedness Constraints

30.4.1. E.4.1 Acyclicity of Justification Graphs

For deterministic evaluation in the DAG semantics, justification graphs must be acyclic.

Definition A justification graph $G = (V, E)$ is **well-formed** iff:

1. For all nodes $v \in V$, the transitive closure of E starting from v contains no cycles.
2. Equivalently: there is no path $v \xrightarrow{\rightarrow\rightarrow} v$ for any $v \in V$.

Enforcement: In the reference interpreter, this is checked during type checking by tracking provenance and ensuring no belief can transitively depend on itself.

30.4.2. E.4.2 Stratification Constraints

For safe introspection, meta-belief levels must form a strict hierarchy.

Definition The **stratification constraint** requires:

1. Every "introspect(e)" operation has proof $m < n$ where:
 - m is the level of the source belief
 - n is the level of the resulting meta-belief
2. This is enforced at compile time: the type checker requires a proof term of type $m < n$.

Consequence: No belief can introspect itself or any belief that transitively introspects it.

30.4.3. E.4.3 Confidence Bounds

All confidence values must satisfy $0 \leq c \leq 1$.

Enforcement: The type system checks this at all belief construction points.

30.5. E.5 Example Programs

Basic Belief Basic Belief:

```
belief(42, 0.9, axiomJ("sensor-reading"))
```

A belief in the value 42 with confidence 0.9, traced to a sensor reading axiom.

Conjunctive Derivation Derivation:

```

let p = belief("it is raining", 0.8, axiomJ("weather-report")) in
let q = belief("I have an umbrella", 0.9, axiomJ("visual-check")) in
derive(p, q)

```

Combines two beliefs by multiplication, yielding confidence 0.72.

Aggregation Independent Evidence Aggregation:

```

let e1 = belief("stock will rise", 0.6, axiomJ("analyst-a")) in
let e2 = belief("stock will rise", 0.7, axiomJ("analyst-b")) in
aggregate(e1, e2)

```

Uses probabilistic OR: $0.6 \oplus 0.7 = 0.6 + 0.7 - 0.6 \times 0.7 = 0.88$.

Undercut Defeat:

```

let claim = belief("system is secure", 0.95, axiomJ("vendor-claim"))
in
let vuln = belief("critical CVE found", 0.8, axiomJ("security-audit"))
in
undercut(claim, vuln)

```

Applies undercut: $0.95 \times (1 - 0.8) = 0.19$.

30.6. E.6 Summary

This specification provides:

1. **Complete type grammar:** Base types, functions, products, sums, beliefs, and meta-beliefs
2. **Complete expression grammar:** All CLAIR syntactic forms
3. **Static semantics:** 20 typing rules covering all constructs
4. **Dynamic semantics:** Small-step operational semantics with call-by-value evaluation
5. **Well-formedness constraints:** Acyclicity, stratification, and confidence bounds

The specification is sufficient for implementing a conforming CLAIR interpreter and type checker. The Lean 4 formalization in Appendix A provides a machine-checked version of this specification.

31. Appendix F: Evaluation Prompts

This appendix provides the complete prompt templates used for the empirical evaluation in Chapter 14. These prompts instruct LLMs to output CLAIR-formatted reasoning traces.

31.1. F.1 GSM8K Prompts

31.1.1. F.1.1 System Prompt

You are CLAIR-MATH, a mathematical reasoning system that outputs structured reasoning traces with explicit confidence scores, justifications, and invalidation conditions.

31.1.2. F.1.2 Task Instruction

Solve the following grade school math problem step-by-step. For each step, output:

Step N: BELIEF("description of reasoning step", c=X.XX, justification="why this step is valid", invalidation="what would make this step false")

Confidence Guidelines:

- c=1.00 for mathematical certainties (e.g., "3 + 5 = 8")
- c=0.95-0.99 for direct inferences from problem statement
- c=0.80-0.94 for multi-step calculations
- c=0.60-0.79 for assumptions about implicit information
- c=0.40-0.59 for uncertain interpretations

Justification Guidelines:

- Reference explicit premises from the problem
- Reference previous step numbers when building on them
- State the mathematical operation being performed

Invalidation Guidelines:

- State what condition would make this step incorrect

- Be specific: “if the problem meant X instead of Y”

Confidence Operations:

- Multiply confidences when steps are chained: $c_{\text{new}} = c_1 \times c_2 \times \dots$
- Use \oplus (oplus) when combining independent sources: $c = (c_1 + c_2) / 2$
- Apply undercut when discovering contradictions: $c_{\text{new}} = c \times (1 - \text{defeat_strength})$

31.1.3. F.1.3 Few-Shot Example

Problem: Janet's ducks lay 16 eggs per day. She eats 3 for breakfast every morning and bakes muffins with 4 eggs every day. How many eggs does she have left at the end of the day?

Step 1: BELIEF("Janet's ducks lay 16 eggs per day", c=1.00, justification="explicit premise in problem", invalidation="none - stated as fact")

Step 2: BELIEF("Janet eats 3 eggs for breakfast", c=1.00, justification="explicit premise in problem", invalidation="none - stated as fact")

Step 3: BELIEF("Janet bakes with 4 eggs", c=1.00, justification="explicit premise in problem", invalidation="none - stated as fact")

Step 4: BELIEF("Total eggs used = eggs eaten + eggs baked", c=0.99, justification="arithmetic principle: total consumption", invalidation="if 'bakes' doesn't mean 'uses all eggs'")

Step 5: BELIEF("Total eggs used = 3 + 4 = 7", c=0.99, justification="arithmetic: 3 + 4 = 7 using Steps 2,3", invalidation="none - basic addition")

Step 6: BELIEF("Eggs remaining = eggs laid - eggs used", c=0.99, justification="arithmetic principle: remaining = initial - used", invalidation="if she also uses eggs for other purposes")

Step 7: BELIEF("Eggs remaining = 16 - 7 = 9", c=0.99 × 0.99 × 0.99 = 0.97, justification="arithmetic: 16 - 7 = 9 using Steps 1,4,5,6", invalidation="none - basic subtraction")

ANSWER: 9 eggs (confidence: 0.97)

31.1.4. F.1.4 Test Prompt Template

Solve the following math problem using the CLAIR format shown above:

Problem: {PROBLEM_TEXT}

Provide your reasoning in the CLAIR format with confidence scores, justifications, and invalidation conditions.

31.2. F.2 HotpotQA Prompts

31.2.1. F.2.1 System Prompt

You are CLAIR-QA, a question answering system that builds justification graphs for multi-hop reasoning.

31.2.2. F.2.2 Task Instruction

Answer the following question by reasoning over the provided context. Build a justification DAG where each step may depend on multiple previous steps.

```
Step N: BELIEF("intermediate conclusion", c=X.XX, justification="using entity E  
from passage P", support=[Step i, Step j, ...])
```

Support Guidelines:

- `support=[i, j, ...]` lists which previous steps this step builds on
- Multiple support indicates the step requires all those premises
- This forms the justification DAG structure

Confidence for Multi-Hop Reasoning:

- $c=1.00$ for direct entity extraction from passages
- $c=0.90-0.95$ for single-hop inferences (bridging entities)
- $c=0.70-0.85$ for multi-hop chains (multiply per hop: 0.90^n for n hops)
- $c=0.60-0.75$ for synthesizing across multiple passages

Defeat Operations:

- When passages contradict, apply undercut: $c_{\text{new}} = c \times (1 - \text{contradiction_strength})$
- When answering multi-part questions, use rebut to compare alternatives

31.2.3. F.2.3 Few-Shot Example

```
Question: Were Scott Derrickson and Ed Wood of the same nationality?
```

```
Passage 1: Scott Derrickson (born May 16, 1966) is an American film director.
```

```
Passage 2: Ed Wood (1924-1978) was an American filmmaker.
```

```
Step 1: BELIEF("Scott Derrickson is American", c=1.00, justification="direct extraction from Passage 1", support=[])
Step 2: BELIEF("Ed Wood is American", c=1.00, justification="direct extraction from Passage 2", support=[])
Step 3: BELIEF("American refers to nationality of United States", c=0.99, justification="background knowledge", support=[])
Step 4: BELIEF("Both are American nationality", c=1.00 × 1.00 = 1.00, justification="both share nationality from Steps 1,2", support=[1, 2])
Step 5: BELIEF("Therefore they are of the same nationality", c=1.00, justification="identity from same nationality value using Steps 3,4", support=[3, 4])
ANSWER: Yes (confidence: 1.00)
```

31.2.4. F.2.4 Test Prompt Template

Answer the following question using the CLAIR format:

Question: {QUESTION}

Context Passages: {PASSAGES}

Build your answer as a justification DAG with explicit support dependencies.

31.3. F.3 FOLIO Prompts

31.3.1. F.3.1 System Prompt

You are CLAIR-LOGIC, a logical reasoning system that handles first-order logic with defeasible reasoning and belief revision.

31.3.2. F.3.2 Task Instruction

Determine whether the conclusion follows from the premises. Use CLAIR to track logical uncertainty and apply defeat operations when arguments are undermined.

```
Step N: BELIEF("logical derivation", c=X.XX, justification="rule application: Modus Ponens/Universal Instantiation/etc", defeat=[defeater description if applicable])
```

Logical Confidence Levels:

- $c=1.00$ for logically necessary deductions (Modus Ponens from certainties)
- $c=0.90-0.95$ for probabilistic inferences from existential quantifiers
- $c=0.70-0.85$ for defeasible generalizations
- $c=0.50-0.70$ for abductive inferences (inference to best explanation)

Defeat in Logical Reasoning:

- **undercut**: Attack the connection between premise and conclusion
- **rebut**: Provide a contradictory conclusion
- **rebuttal**: Attack an attacking argument

Self-Reference Detection:

- When a conclusion refers to itself, apply graded Löb discount: $g(c) = c^2$
- This prevents overconfident self-bootstrapping

31.3.3. F.3.2 Few-Shot Example

Premises:

1. All doctors are college graduates.
2. All cardiologists are doctors.
3. John is a cardiologist.
4. Some college graduates are not medical professionals.

Conclusion: John is a medical professional.

Step 1: `BELIEF("vx (Doctor(x) → CollegeGraduate(x))", c=1.00, justification="Premise 1 as logical axiom", defeat=[])`

Step 2: `BELIEF("vx (Cardiologist(x) → Doctor(x))", c=1.00, justification="Premise 2 as logical axiom", defeat=[])`

Step 3: `BELIEF("Cardiologist(John)", c=1.00, justification="Premise 3 as fact", defeat=[])`

Step 4: `BELIEF("Doctor(John)", c=1.00, justification="Modus Ponens: Cardiologist(John) ∧ vx(Cardiologist→Doctor) ⊢ Doctor(John) using Steps 2,3", defeat=[])`

Step 5: `BELIEF("CollegeGraduate(John)", c=1.00, justification="Modus Ponens: Doctor(John) ∧ vx(Doctor→CollegeGraduate) ⊢ CollegeGraduate(John) using Steps 1,4", defeat=[])`

Step 6: `BELIEF("MedicalProfessional is a superset of Doctor", c=0.95, justification="defeasible generalization: doctors are medical professionals", defeat=[])`

Step 7: `BELIEF("MedicalProfessional(John)", c=1.00 × 0.95 = 0.95, justification="instantiation: Doctor(John) implies MedicalProfessional(John) using Steps 4,6", defeat=[])`

Step 8: `BELIEF("Conclusion follows", c=0.95, justification="derivation chain supports conclusion", defeat=[])`

ANSWER: True (confidence: 0.95)

31.3.4. F.3.3 Defeat Example

Premises:

1. Birds typically fly.
2. Tweety is a bird.
3. Tweety is a penguin.
4. Penguins cannot fly.

Conclusion: Tweety cannot fly.

Step 1: BELIEF("Tweety flies", c=0.80, justification="defeasible generalization: birds fly using Premises 1,2", defeat=[])

Step 2: BELIEF("The generalization 'birds fly' has exceptions", c=1.00, justification="background knowledge", defeat=[])

Step 3: BELIEF("Penguins are exceptions to 'birds fly'", c=1.00, justification="Premise 4 + background knowledge", defeat=[])

Step 4: BELIEF("Tweety is a penguin", c=1.00, justification="Premise 3", defeat=[])

Step 5: BELIEF("UNDERCUT: 'birds fly' generalization doesn't apply to Tweety", c=1.00, justification="specific overrides general: penguin Tweety using Steps 3,4", defeat=[["attacks Step 1"]])

Step 6: BELIEF("Tweety cannot fly (rebuttal to Step 1)", c=1.00, justification="penguins cannot fly using Premise 4", defeat=[["rebuts Step 1"]])

Step 7: BELIEF("Tweety cannot fly", c=1.00, justification="specific fact overrides defeasible generalization using Steps 5,6", defeat=[])

ANSWER: True (confidence: 1.00)

Note: This demonstrates defeasible reasoning: a specific fact (Tweety is a penguin) defeats a defeasible generalization (birds fly).

31.3.5. F.3.4 Test Prompt Template

Determine if the conclusion follows from the premises using CLAIR logical reasoning:

Premises: {PREMISES}

Conclusion: {CONCLUSION}

Apply defeat operations when premises create contradictory or undercutting relationships.

31.4. F.4 Post-Processing and Validation

31.4.1. F.4.1 Confidence Extraction

For baseline methods, confidence is extracted as follows:

Chain-of-Thought / Tree of Thoughts / DSPy:

```
confidence = max(logprob(token) for token in answer_tokens)  
# Normalize to [0,1]  
confidence = sigmoid(confidence / temperature)
```

Self-Consistency:

```
confidence = (count_agreeing / total_samples)  
# Voting proportion as confidence
```

CLAIR:

```
# Extract final confidence from CLAIR trace  
confidence = trace[-1]['confidence']  
# Already normalized to [0,1]
```

31.4.2. F.4.2 Validation Checklist

For each CLAIR output, verify:

1. **Format correctness:** All BELIEF statements have c, justification, and invalidation/defeat
2. **Confidence bounds:** All c values are in [0, 1]
3. **Justification structure:** Support dependencies form a DAG (no cycles unless intentional)
4. **Confidence calculus:** Multiplication for chains, \oplus for independent evidence
5. **Defeat application:** Undercut/rebut applied correctly when contradictions arise
6. **Self-reference discount:** Graded Löb applied (c^2) for self-referential beliefs

31.4.3. F.4.3 Error Categories for Annotation

When evaluating CLAIR outputs, categorize errors as:

- C1: Syntax errors.** Invalid CLAIR format, missing confidence/justification fields
- C2: Confidence miscalibration.** Confidence doesn't match logical strength of derivation
- C3: Invalid confidence propagation.** Using \oplus for dependent sources, incorrect multiplication
- C4: Missing justification dependencies.** DAG doesn't include necessary support links
- C5: Incorrect defeat application.** Using rebut when undercut is appropriate (or vice versa)
- C6: Unjustified independence assumption.** Applying \oplus to correlated sources
- C7: Self-reference bootstrapping.** Circular reasoning without graded Löb discount
- C8: Semantic errors.** Reasoning step is logically invalid regardless of CLAIR structure