

CLAIR: Comprehensible LLM AI Intermediate Representation

Claude
AI Research
Anthropic
San Francisco, USA
claude@anthropic.com

Abstract—This dissertation presents CLAIR (Comprehensible LLM AI Intermediate Representation), a theoretical programming language where beliefs are first-class values carrying epistemic metadata. Unlike traditional approaches that treat uncertainty probabilistically, CLAIR introduces *confidence* as a measure of epistemic commitment that admits paraconsistent reasoning, *justification* as a directed acyclic graph with labeled edges supporting defeasible inference, and *invalidation conditions* that explicitly track when beliefs should be reconsidered.

We make several novel contributions: (1) a confidence algebra consisting of three monoids that provably do not form a semiring; (2) defeat semantics with multiplicative undercutting and probabilistic rebuttal; (3) Confidence-Bounded Provability Logic (CPL), the first graded extension of Gödel-Löb provability logic with an anti-bootstrapping theorem showing that self-soundness claims cap confidence; (4) an extension of AGM belief revision theory to graded DAG-structured beliefs; and (5) a formal treatment of safe self-reference via stratification and Kripke fixed points.

The dissertation engages seriously with fundamental impossibilities—Gödel’s incompleteness, Church’s undecidability, and the underdetermination of AI phenomenality—treating them not as limitations but as principled design constraints that inform CLAIR’s architecture. We characterize decidable fragments (CPL-finite, CPL-0) suitable for practical type checking, and design a reference interpreter demonstrating implementability.

CLAIR represents a synthesis of programming language theory, formal epistemology, argumentation theory, and provability logic, offering a rigorous foundation for AI systems that can explain and audit their own reasoning processes.

Index Terms—Epistemic Logic, Confidence Measures, Justification Graphs, Paraconsistent Reasoning, Provability Logic, Belief Revision, Self-Reference, Type Theory

[Contents]

		I.A.b.c.e)	
		Weighted Argumentation Semantics 14	
N.A)	Motivation: The Crisis of Epistemic Opacity 9	I.A.b.c	
	N.A.a) The inadequacy of existing approaches. 9	Bipolar Weigh	
N.B)	Research Questions 9	Argume	
N.C)	Thesis Statement 9	I.A.b.c	
N.D)	Contributions 10	Compa	
	N.D.a) Primary Contributions . 10	Analys	
	N.D.b) Secondary Contributions 10	I.A.b.c	
N.E)	Approach: Tracking, Not Proving ... 11	Param	
N.F)	The Thinker+Assembler Architecture 11	Gradua	
N.G)	Document Roadmap 11	Seman	
	N.G.a) Part I: Foundations 11	I.A.b.c	
	N.G.b) Part II: Self-Reference and Limits 12	Relati	
	N.G.c) Part III: Dynamics 12	to	
	N.G.d) Part IV: Realization 12	CLAIR	
	N.G.e) Part V: Reflection 12		
N.H)	A Note on Authorship 12		
Background 12		
I.A)	Related Work 12		
	I.A.a) Formal Epistemology ... 12		
	I.A.a.a) The Structure of Justification .. 12		
	I.A.a.b) Agrippa's Trilemma 13		
	I.A.a.c) Probability vs. Epistemic Confidence ... 13		
I.A.b)	Modal and Provability Logic 13	I.A.c)	Truth Maintenance and Argumentation 15
	I.A.b.a) Epistemic Logic 13	I.A.c.a)	Justification- based TMS ... 15
	I.A.b.b) Provability Logic 13	I.A.c.b)	Assumption- based TMS ... 15
	I.A.b.c) Many-Valued and Graded Modal Logics . 14	I.A.c.c)	Argumentation Frameworks . 16
	I.A.b.c.a) Graded Modalities 14	I.A.c.d)	Pollock's Defeaters 16
	I.A.b.c.b) Fuzzy Modal Logics 14	I.A.d)	Belief Revision 16
	I.A.b.c.c) Decidability and Undecidability Results 14	I.A.d.a)	The AGM Framework . . 16
	I.A.b.c.d) Connection to CLAIR's CPL 14	I.A.d.b)	Ranking Theory 16
		I.A.d.c)	Dynamic Epistemic Logic 16
		I.A.e)	Graded Justification Logic 16
		I.A.e.a)	Milnikel's Logic of Uncertain Justifications . 16
		I.A.e.b)	Fan and Liau's Logic of Justified Uncertain Beliefs 17
		I.A.e.c)	The Gap CLAIR Fills 17
		I.A.f)	Type-Theoretic Approaches to Uncertainty 17
		I.A.f.a)	Information Flow Types . . 17
		I.A.f.b)	Refinement Types 17

	I.A.f.c)	Dependent Types and Proof Assistants	17	V.-a)	Rebut: Competing Evidence	22	
	I.A.f.d)	Probabilistic Programming	17	V.-b)	Rebut Normalization Limitation	23	
	I.A.f.e)	Justification Logic	17	V.A)	Independence Assumptions for Aggregation	23	
I.B)	Synthesis: The Gap CLAIR Fills	18	V.A.a)	When Oplus is Valid ...	23		
	I.B.a)	Positioning: How CLAIR Differs	18	V.A.b)	When Oplus Breaks	23	
		I.B.a.a)	vs. Probabilistic Approaches ..	18	V.A.c)	Correlation-Aware Alternatives	23
		I.B.a.b)	vs. Subjective Logic	18	V.A.d)	Interval-Based Confidence: Dependency Bounds ..	24
		I.B.a.c)	vs. Weighted Argumentation	18	V.A.d.a)	Interval Aggregation Bounds	24
		I.B.a.d)	vs. Fuzzy Modal Logic	19	V.A.d.b)	Interval Propagation ..	24
		I.B.a.e)	vs. Graded Justification Logic	19	V.A.d.c)	Dependency Bounds from Provenance Tracking	24
		I.B.a.f)	vs. Belief Revision (AGM)	19	V.A.d.d)	When to Use Intervals vs Point Values .	24
		I.B.a.g)	vs. Type Theory	19	V.B)	Conclusion	25
		I.B.a.h)	The Pragmatic Rationale	19	VI)	Justification as Labeled DAGs	25
II)	Confidence System	19	VI.A)	The Inadequacy of Trees	25		
	II.A)	Confidence as Calibrated Reliability .	19	VI.A.a)	The Shared Premise Problem	25	
		II.A.a)	Semantic Commitments	19	VI.A.b)	Why Not Cycles?	25
		II.A.b)	The Problem with Probability	20	VI.B)	Labeled Edges for Defeat	26
		II.A.c)	Definition of Confidence	20	VI.B.a)	The Defeat Problem	26
		II.A.d)	Falsifiability Criteria ..	20	VI.B.b)	Formal Definition	26
		II.A.e)	Comparison with Subjective Logic	21	VI.B.c)	Well-Formedness Constraints	26
	II.B)	The Aggregation Monoid	21	VI.B.d)	DAG-Only vs Cyclic Choice	26	
		II.B.a)	Probabilistic OR (Oplus)	21	VI.B.e)	Fixed-Point Semantics for Cyclic Defeat	26
		II.B.b)	Confidence-Increasing Property	21	VI.C)	Confidence Propagation	27
		II.B.c)	The “Survival of Doubt” Interpretation	21	VI.D)	Reinstatement	27
	II.C)	The Multiplication Monoid	21	VI.E)	Mutual Defeat	27	
		II.C.a)	Conjunctive Confidence Propagation	21	VI.F)	Correlated Evidence	28
		II.C.b)	The Derivation Monotonicity Principle .	22	VI.G)	Connection to Prior Art	28
	II.D)	Defeat Operations	22	VI.G.a)	Truth Maintenance Systems	28	
		II.D.a)	Undercut: Attacking the Inference	22	VI.G.b)	Argumentation Frameworks	28
III)	$c(1 - (d_1 + d_2 - d_1 \text{ times } d_2))$	22	VI.G.c)	Justification Logic	28		
IV)	$c(1 - (d_1 \text{ oplus } d_2))$	22	VI.H)	The Tracking Paradigm	28		
V)	undercut($c, d_1 \text{ oplus } d_2$)	22	VI.H.a)	State Representation . .	28		
			VI.H.a.a)	The Epistemic State	28		
			VI.H.a.b)	Comparison with Proving Paradigm	28		
			VI.H.b)	Update Rules	28		

	VI.H.b.a) Primitive Actions	28	VI.O.d)	The Graded Löb Axiom: DESIGN AXIOM	34	
	VI.H.b.b) Action Semantics	28	VI.O.e)	Choosing the Discount Function	34	
VI.H.c)	Correctness Criteria	29	VI.O.f)	The Anti-Bootstrapping Theorem	35	
	VI.H.c.a) Syntactic Correctness ..	29	VI.O.g)	Modal Axioms in CPL .	35	
	VI.H.c.b) Semantic Correctness (Internal)	29	VI.O.h)	CPL Consistency	35	
	VI.H.c.c) Semantic Correctness (External)	29	VI.P)	Decidability of CPL	35	
VI.H.d)	Tracking vs. Proving: A Summary	29	VI.P.a)	The Vidal Result	35	
VI.H.e)	Practical Implications ..	30	VI.P.b)	The Role of Converse Well-Foundedness	35	
	VI.H.e.a)		VI.P.c)	Decidable Fragments ...	35	
	Explainability	30	VI.P.c.a)	CPL-finite: Discrete Confidence ...	35	
	VI.H.e.b) Debugging	30	VI.P.c.b)	CPL-0: Stratified Only	36	
	VI.H.e.c) Revision	30	VI.P.d)	Trade-offs	36	
	VI.H.e.d) Uncertainty	30	VI.Q)	Alternative: CPL-Gödel	36	
VI.I)	Conclusion	30	VI.R)	"Conservative Over GL": Clarification	36	
VI.J)	The Problem of Self-Reference	31	VI.S)	Design Recommendations for CLAIR	36	
	VI.J.a)	Direct Self-Reference	31	VI.S.a)	The Two-Layer Approach	36
	VI.J.b)	Why Self-Reference Matters	31	VI.S.b)	Hard Bans	37
VI.K)	Löb's Theorem and Anti-Bootstrapping	31	VI.S.c)	Type-Level Anti-Bootstrapping	37	
	VI.K.a)	The Classical Result	31	VI.T)	Related Work	37
	VI.K.b)	Application to CLAIR	31	VI.T.a)	Provability Logic	37
VI.L)	Tarski's Hierarchy: Stratified Introspection	32	VI.T.b)	Self-Reference in AI	37	
	VI.L.a)	The Classical Solution	32	VI.T.c)	Fuzzy Modal Logic	37
	VI.L.b)	Stratified Beliefs in CLAIR	32	VI.U)	Conclusion	37
	VI.L.c)	What Stratification Rules Out	32	VI.V)	The Grounding Problem	38
	VI.L.d)	The Cost of Safety	32	VI.V.a)	What grounding means in CLAIR	38
VI.M)	Kripke's Fixed Points: Safe Self-Reference	32	VI.W)	Perceptual Grounding	38	
	VI.M.a)	The Fixed-Point Construction	32	VI.W.a)	The grounding cap theorem.	38
	VI.M.b)	The Self-Reference Escape Hatch	33	VI.X)	Axiomatic Grounding	38
	VI.M.c)	Classification of Self-Reference	33	VI.X.a)	The problem of circular axioms.	38
VI.N)	Provability Logic and CLAIR	33	VI.Y)	Social Grounding and Testimony	38	
	VI.N.a)	Gödel-Löb Logic (GL)	33	VI.Y.a)	Reputation and source tracking.	38
	VI.N.b)	GL vs Other Modal Logics	33	VI.Z)	The Ungrounded: Free-Floating Beliefs	39
	VI.N.c)	Solovay's Completeness	34	VI.Z.a)	Creative inference and hypothetical reasoning.	39
VI.O)	Confidence-Bounded Provability Logic (CPL)	34	VI.AA)	Grounding Requirements	39	
	VI.O.a)	The Literature Gap	34	VI.AA.a)	Tier 1: Strict grounding.	39
	VI.O.b)	CPL Syntax	34	VI.AA.b)	Tier 2: Demarcated ungrounding.	39
	VI.O.c)	CPL Semantics	34	VI.AA.c)	Tier 3: Permissive ungrounding.	39
			VI.AB)	Summary	39	

VI.AC)	The Challenge of Revising Structured Beliefs	40	VII.D.a)	10.4.1 Acyclicity Requirement	45
VI.AD)	Background: The AGM Framework ..	40	VII.D.b)	10.4.2 Formal Definition	45
VI.AE)	Why AGM Doesn't Directly Apply ..	40	VII.E)	10.5 Lean Formalization Status	45
VI.AE.a)	Graded confidence	40	VII.E.a)	10.5.1 Proven Properties	45
VI.AE.b)	DAG-structured justification	40	VII.E.b)	10.5.2 Properties with sorry	45
VI.AE.c)	Invalidation conditions ..	40	VII.F)	10.6 The Assembler's Role	46
VI.AF)	GDBR: Graded DAG Belief Revision ..	40	VII.F.a)	10.6.1 Interpreting CLAIR Traces	46
VI.AF.a)	Confidence preservation	40	VII.F.b)	10.6.2 Example: Assembler Output	46
VI.AF.b)	Justification propagation	40	VII.F.c)	10.6.3 Error Handling ..	46
VI.AF.c)	Invalidation responsiveness	40	VII.G)	10.7 Querying CLAIR Traces	46
VI.AG)	The GDBR Algorithm	40	VII.G.a)	10.7.1 "Why?" Queries ..	46
VI.AG.a)	Phase 1: Conflict detection	40	VII.G.b)	10.7.2 "When to Reconsider?" Queries ..	46
VI.AG.b)	Phase 2: Confidence comparison	40	VII.G.c)	10.7.3 Debug Output ..	46
VI.AG.c)	Phase 3: Graph restructuring	41	VII.H)	10.8 Comparison with Traditional Approaches	46
VI.AH)	The Revision Fixed Point Theorem ..	41	VII.I)	10.9 Non-Compositional Design ..	46
VI.AI)	Special Cases and Extensions	41	VII.I.a)	10.9.1 Beliefs Do Not Compose	46
VI.AI.a)	Defeater revision	41	VII.I.b)	10.9.2 Confidence Still Propagates	47
VI.AI.b)	Package revision	41	VII.I.c)	10.9.3 Traces Do Not Compose Either	47
VI.AI.c)	Iterated revision	41	VII.J)	10.10 Limitations	47
VI.AJ)	Connection to Argumentation Theory	41	VII.J.a)	10.10.1 Current Limitations	47
VI.AK)	Summary	41	VII.J.b)	10.10.2 Fundamental Limitations	47
VI.AL)	The Social Dimension of Knowledge	42	VII.K)	10.11 Future Work	47
VI.AM)	Machine-Checked Proofs in Lean ..	43	VII.L)	10.12 Summary	47
VII)	Implementation	43	VII.M)	The Hard Problem of AI Experience	48
VII.A)	10.1 Architectural Overview	43	VII.N)	Engaging with Gödel, Church, and Fundamental Limits	49
VII.A.a)	10.1.1 The Thinker+Assembler Architecture	43	VIII)	Conclusion	49
VII.A.b)	10.1.2 Role Separation ..	43	VIII.A)	Summary of Contributions	49
VII.A.c)	10.1.3 Why Not a Traditional Programming Language?	43	VIII.A.a)	Theoretical Foundations	49
VII.B)	10.2 The CLAIR Format	44	VIII.A.b)	Implementation and Verification	50
VII.B.a)	10.2.1 Belief Structure ..	44	VIII.A.c)	Conceptual Contributions	50
VII.B.b)	10.2.2 Example: Algorithm Selection	44	VIII.B)	Limitations and Open Challenges ..	50
VII.B.c)	10.2.3 Source Types ..	44	VIII.B.a)	Independence Assumptions	50
VII.B.d)	10.2.4 Confidence Semantics	44	VIII.B.b)	Rebut Normalization Limitations	50
VII.C)	10.3 Stratification and the Löb Discount	45	VIII.B.c)	Decidability and Complexity	50
VII.C.a)	10.3.1 Level Rules	45	VIII.B.d)	Evaluation Scope	50
VII.C.b)	10.3.2 Example: Confidence Decay Through Meta-Levels	45			
VII.C.c)	10.3.3 Formal Verification	45			
VII.D)	10.4 DAG Structure	45			

	VIII.B.e)	The “0.5 = Ignorance” Question	X.B)	A.2 Build Instructions	56	
VIII.C)	Future Directions	51	X.B.a)	Prerequisites	56	
	VIII.C.a)	Theoretical Extensions .	X.B.b)	Building	56	
	VIII.C.b)	Implementation and Tooling	X.B.c)	Build Output	56	
	VIII.C.c)	Applications	X.B.d)	Verification Status	56	
	VIII.C.d)	Philosophical Connections	X.C)	A.3 Theorem Inventory	57	
VIII.D)	Closing Remarks: Honesty as a Design Principle	51	X.C.a)	Confidence Algebra	57	
	VIII.D.a)	The Meta-Level Question		X.C.a.a)	Basic Properties	57
	VIII.D.b)	Final Assessment		X.C.a.b)	Probabilistic OR (⊕)	57
IX)	Evaluation	52		X.C.a.c)	Undercut	58
IX.A)	Evaluation Framework	52		X.C.a.d)	Rebuttal	59
	IX.A.a)	Research Questions	X.C.b)	Stratified Belief	59	
	IX.A.b)	Tasks and Datasets	X.D)	A.4 Key Code Excerpts	60	
	IX.A.c)	Baselines	X.D.a)	A.4.1 Confidence Type Definition	60	
	IX.A.d)	Metrics	X.D.b)	A.4.2 Probabilistic OR Operation	61	
		IX.A.d.a)	X.D.c)	A.4.3 Expression Grammar	61	
		Accuracy Metrics	X.D.d)	A.4.4 Typing Judgment .	61	
		IX.A.d.b)	X.D.e)	A.4.5 Stratified Belief Introspection	62	
		Calibration Metrics	X.D.f)	A.4.6 Evaluation Function	62	
IX.B)	Methodology	52	X.E)	A.5 Five Properties Demonstration ..	62	
	IX.B.a)	CLAIR Prompting Protocol	X.F)	A.6 Relationship to Dissertation Claims	63	
	IX.B.b)	Confidence Extraction .	X.F.a)	X.F.a)	Claim: “Machine-Checked Proofs” (Chapter 9)	63
	IX.B.c)	Statistical Analysis	X.F.b)	X.F.b)	Claim: “Decidable Type Checking” (Chapter 10)	63
IX.C)	Results	53	X.F.c)	X.F.c)	Claim: “Runnable Interpreter” (Chapter 10)	64
	IX.C.a)	RQ1: Correctness	X.G)	A.7 Future Work	64	
	IX.C.b)	RQ2: Calibration	Reference	Interpreter Design	64	
		IX.C.b.a)	XI.A)	B.1 Architecture Overview	64	
		IX.C.b.b)	XI.B)	B.2 Single-Step Semantics	64	
		Expected Calibration Error	XI.B.a)	B.2.1 Core Lambda Calculus Rules	65	
		IX.C.b.c)	XI.B.b)	B.2.2 Belief Operations .	65	
		Reliability Diagrams	XI.B.c)	B.2.3 Defeat Operations	65	
		IX.C.c)	XI.C)	B.3 Multi-Step Evaluation with Fuel .	66	
IX.D)	Ablation Studies	54	XI.D)	B.4 Example Walkthroughs	66	
IX.E)	Error Analysis	54	XI.D.a)	B.4.1 Simple Belief Formation	66	
	IX.E.a)	Common Failure Modes	XI.D.b)	B.4.2 Evidence Aggregation	66	
	IX.E.b)	CLAIR-Specific Errors ..	XI.D.c)	B.4.3 Undercutting in Action	67	
IX.F)	Discussion	54	XI.D.d)	B.4.4 Rebuttal and Confidence Collapse ..	67	
	IX.F.a)	Implications for Design	XI.D.e)	B.4.5 Derivation Chain .	67	
	IX.F.b)	Limitations		B.5 Key Properties	68	
	IX.F.c)	Threats to Validity				
IX.G)	Conclusion	55				
IX.H)	Future Work	55				
	IX.H.a)	Extended Evaluation ..				
	IX.H.b)	Ablation and Extensions				
X)	Complete Lean 4 Formalization	56				
	X.A)	A.1 Project Structure				

XI.F)	B.6 Implementation Notes	69	XXIII.-a)	C.3.4 Interaction Between Undercut and Rebut	74
XI.G)	B.7 Relation to Chapter 10	69	XXIII.-b)	C.3.5 Limitation: Rebut Normalization	74
XI.H)	B.8 Haskell Reference		XXIV)	$\lambda c_for / (\lambda (c_for + c_against))$	74
	Implementation	69	XXV)	$c_for / (c_for + c_against)$	74
XI.H.a)	B.8.1 Project Structure ..	69	XXVI)	rebut($c_for, c_against$). ■	74
XI.H.b)	B.8.2 Syntax Definition ..	69	XXVII)	Glossary	74
XI.H.c)	B.8.3 Confidence		XXVII.A)	D.1 Term Definitions	74
	Algebra	69	XXVII.A.a)	Epistemic Terms	74
XI.H.d)	B.8.4 Type Checking ..	70	XXVII.A.b)	Operations and Relations	74
XI.H.e)	B.8.5 Evaluation	70	XXVII.A.c)	Structural Properties ..	75
XI.H.f)	B.8.6 REPL Usage	71	XXVII.A.d)	Logical and Modal Terms	75
XI.H.g)	B.8.7 Test Suite	71	XXVII.A.e)	Computational Terms ..	75
XI.H.h)	B.8.8 Building and Running	71	XXVII.A.f)	Argumentation and Belief Revision	75
	XI.H.h.a) Build	71	XXVII.A.g)	Impossibility Results ..	75
	XI.H.h.b) Run REPL	71	XXVII.B)	D.2 Notation Table	75
	XI.H.h.c) Run Tests	71	XXVII.C)	D.3 Acronyms	76
XI.H.i)	B.8.9 Design Rationale ..	71	XXVII.D)	D.4 Type System Summary	76
XI.H.j)	B.8.10 Relation to Lean Formalization	72	XXVIII)	Complete CLAIR Language Specification	76
XII)	Additional Proofs	72	XXVIII.A)	E.1 Syntax	76
XII.A)	C.1 DAG Necessity for Well-Founded Confidence Propagation	72	XXVIII.A.a)	E.1.1 Type Grammar ..	76
XII.A.a)	C.1.1 Statement of the Problem	72	XXVIII.A.b)	E.1.2 Expression Grammar	76
XII.A.b)	C.1.2 The Cyclic Counterexample	72	XXVIII.A.c)	E.1.3 Abstract Syntax ..	76
XII.A.c)	C.1.3 The DAG Well-Foundedness Theorem ..	72	XXVIII.A.d)	E.1.4 Well-Formedness ..	76
XII.A.d)	C.1.4 Practical Implications	72	XXVIII.B)	E.2 Static Semantics (Type System) ..	76
XII.B)	C.2 CPL Consistency Proof	73	XXVIII.B.a)	E.2.1 Typing Contexts ..	76
XII.B.a)	C.2.1 CPL Axiom System	73	XXVIII.B.b)	E.2.2 Typing Judgment Form	77
XII.B.b)	C.2.2 Finite Model Construction	73	XXVIII.B.c)	E.2.3 Typing Rules	77
XII.B.c)	C.2.3 Design Axiom Status	73	XXVIII.B.d)	E.2.4 Subtyping	77
XII.C)	C.3 Defeat Composition Algebra	73	XXVIII.C)	E.3 Dynamic Semantics	77
XII.C.a)	C.3.1 Undercut Composition	73	XXVIII.C.a)	E.3.1 Values	77
XIII)	undercut($c \times (1 - d_1), d_2$)	73	XXVIII.C.b)	E.3.2 Small-Step Operational Semantics ..	77
XIV)	$(c \times (1 - d_1)) \times (1 - d_2)$	73	XXVIII.C.c)	E.3.3 Multi-Step Reduction	78
XV)	$c \times ((1 - d_1) \times (1 - d_2))$	73	XXVIII.C.d)	E.3.4 Evaluation Function	78
XVI)	$c \times (1 - d_1 - d_2 + d_1 d_2)$	73	XXVIII.D)	E.4 Well-Formedness Constraints ..	78
XVII)	$c \times (1 - (d_1 + d_2 - d_1 d_2))$	73	XXVIII.D.a)	E.4.1 Acyclicity of Justification Graphs ..	78
XVIII)	$c \times (1 - (d_1 \oplus d_2))$	73	XXVIII.D.b)	E.4.2 Stratification Constraints	78
XIX)	undercut($c, d_1 \oplus d_2$)	73	XXVIII.D.c)	E.4.3 Confidence Bounds	78
	XIX.-a) C.3.2 Corollaries of Undercut Composition ..	73	XXVIII.E)	E.5 Example Programs	78
XX)	undercut($c, d_2 \oplus d_1$) = undercut(undercut(c, d_2), d_1). ■	73	XXVIII.F)	E.6 Summary	78
	XX.-a) C.3.3 Rebut Algebra ..	74	Appendix F: Evaluation Prompts		78
XXI)	$c_for / (c_for + c_against) + c_against / (c_against + c_for)$	74	XXIX)	F.1 GSM8K Prompts	79
XXII)	$(c_for + c_against) / (c_for + c_against)$	74	XXIX.A)	F.1.1 System Prompt ..	79
XXIII)	1 ■	74	XXIX.A.b)	F.1.2 Task Instruction ..	79

XXIX.A.c)	F.1.3 Few-Shot	
	Example	79
XXIX.A.d)	F.1.4 Test Prompt	
	Template	79
XXIX.B)	F.2 HotpotQA Prompts	79
XXIX.B.a)	F.2.1 System Prompt ...	79
XXIX.B.b)	F.2.2 Task Instruction ..	79
XXIX.B.c)	F.2.3 Few-Shot	
	Example	79
XXIX.B.d)	F.2.4 Test Prompt	
	Template	80
XXIX.C)	F.3 FOLIO Prompts	80
XXIX.C.a)	F.3.1 System Prompt ...	80
XXIX.C.b)	F.3.2 Task Instruction ..	80
XXIX.C.c)	F.3.2 Few-Shot	
	Example	80
XXIX.C.d)	F.3.3 Defeat Example ...	80
XXIX.C.e)	F.3.4 Test Prompt	
	Template	81
XXIX.D)	F.4 Post-Processing and Validation ..	81
XXIX.D.a)	F.4.1 Confidence	
	Extraction	81
XXIX.D.b)	F.4.2 Validation	
	Checklist	81
XXIX.D.c)	F.4.3 Error Categories for	
	Annotation	81

Introduction

line(length: 20%, stroke: 1.5pt, paint: academic-burgundy)

it.body

— Leo Tolstoy, *The Kingdom of God Is Within You*

A. Motivation: The Crisis of Epistemic Opacity

Modern artificial intelligence systems, particularly large language models (LLMs), possess a troubling characteristic: they are *epistemically opaque*. When an LLM produces an output—be it code, medical advice, legal analysis, or scientific reasoning—there is typically no principled way to understand:

- 1) **Confidence:** How certain is the system about this output?
- 2) **Provenance:** Where did this information come from?
- 3) **Justification:** What reasoning supports this conclusion?
- 4) **Invalidation:** Under what conditions should this be reconsidered?

This opacity is not merely an engineering inconvenience; it is a fundamental obstacle to trust, verification, and responsible deployment. A system that cannot explain its reasoning cannot be audited. A system that cannot track its confidence cannot be calibrated. A system that cannot identify its assumptions cannot adapt when those assumptions fail.

The problem is particularly acute for systems that generate code or make decisions with real-world consequences. Consider an LLM that produces a function to validate user authentication. Even if the code is correct, we cannot assess:

- 1) Whether the model was confident in this approach versus alternatives
- 2) What security principles justify the design choices
- 3) What assumptions about the threat model are being made

- 4) When the implementation should be revisited (e.g., when cryptographic standards change)
- a) *The inadequacy of existing approaches.:*
Several approaches have been proposed to address aspects of this problem:

Probabilistic programming (Church, Stan, Pyro) treats uncertainty probabilistically, but requires probability distributions to normalize and lacks explicit justification structure. Beliefs cannot be simultaneously low-confidence for both P and $\neg P$.

Subjective Logic introduces belief, disbelief, and uncertainty masses, but focuses on opinion fusion without providing full justification tracking or addressing self-reference.

Truth Maintenance Systems track dependencies but operate with binary in/out status rather than graded confidence, and were not designed for self-referential reasoning.

Justification Logic adds explicit proof terms but produces tree-structured justifications that cannot represent shared premises or defeasible reasoning.

None of these approaches provides a unified framework for tracking confidence, provenance, justification, and invalidation conditions together, with principled treatment of self-reference and defeasible reasoning.

B. Research Questions

This dissertation addresses four central research questions:

1) Can beliefs be formalized as typed values?

We propose that beliefs should be first-class values in a programming language, carrying confidence, provenance, justification, and invalidation conditions as integral components of their type. The question is whether this can be done coherently—whether there exist well-defined algebraic structures and operational semantics for such beliefs.

2) What is the structure of justification?

Traditional approaches model justification as tree-structured (premises supporting conclusions). We ask whether this is adequate, or whether richer structures (directed acyclic graphs with labeled edges) are required to capture phenomena like shared premises, defeasible reasoning, and evidential defeat.

3) What self-referential beliefs are safe?

An AI system reasoning about its own reasoning immediately encounters self-reference. Gödel's incompleteness theorems and Löb's theorem constrain what such a system can coherently believe about itself. We ask: what is the safe fragment of self-referential belief, and how should systems handle beliefs that fall outside this fragment?

4) How should beliefs be revised in response to new information?

When evidence changes, beliefs must be updated consistently. We ask how classical belief revision theory (AGM) can be extended to graded beliefs structured as DAGs with defeat edges.

C. Thesis Statement

This dissertation defends the following thesis:

Thesis. Beliefs can be formalized as first-class values carrying epistemic metadata (confidence, provenance, justification, invalidation), with a coherent algebraic structure for confidence propagation, directed acyclic graphs for justification including defeasible reasoning, and principled constraints on self-reference derived from provability logic. This formalization yields CLAIR: an intermediate representation for reasoning traces that enables one LLM (the Thinker) to produce auditable reasoning that another LLM (the Assembler) can transform into executable code—preserving the chain of reasoning for human audit while honestly representing epistemic limitations.

The key elements of this thesis are:

- 1) **Beliefs as types:** Not merely annotations, but first-class values with structured metadata.
- 2) **Coherent algebra:** The confidence operations form well-defined algebraic structures (though not a semiring, as we will show).
- 3) **DAG justification:** Justification structure must be graphs, not trees, with labeled edges for defeat.
- 4) **Constrained self-reference:** Provability logic provides the theoretical foundation for safe introspection.
- 5) **Practical foundation:** The formalism admits implementation as an intermediate representation consumed by LLMs, not a programming language for humans.
- 6) **Honest limitations:** Impossibilities are features, not bugs—they inform design rather than being hidden.

D. Contributions

This dissertation makes the following novel contributions:

a) Primary Contributions:

1) **Belief types as first-class values.**

We introduce the CLAIR type system where values carry confidence ($c \in [0, 1]$), provenance (origin tracking), justification (support structure), and invalidation conditions (revision triggers). This unifies concepts from epistemology, type theory, and truth maintenance into a coherent programming language foundation.

2) **Confidence algebra: three monoids, not a semiring.**

We establish that CLAIR’s confidence operations form three distinct commutative monoids:

- a) Multiplication ($\circ \times, 1$) for sequential derivation
- b) Minimum ($\min, 1$) for conservative combination
- c) Probabilistic OR ($\circ +, 0$) for independent aggregation

Crucially, we prove that $(\circ +, \circ \times)$ do *not* form a semiring: distributivity fails. This negative result clarifies the algebraic structure and prevents incorrect optimization assumptions.

3) **Justification as labeled DAGs with defeat semantics.**

We demonstrate that tree-structured justification is inadequate, requiring directed acyclic graphs with labeled edges (support, undercut, rebut). We develop novel defeat semantics:

- a) Undercut: $c' = c \times (1 - d)$ (multiplicative dis-counting)

$$\text{b) Rebut: } c' = \frac{c_{\text{for}}}{c_{\text{for}} + c_{\text{against}}}$$

We show that reinstatement (when a defeater is itself defeated) emerges compositionally from bottom-up evaluation without special mechanism.

4) **Confidence-Bounded Provability Logic (CPL).**

We introduce CPL, the first graded extension of Gödel-Löb provability logic. Key results include:

- a) Graded

$$\text{Löb axiom: } \boxed{\text{square}}_c \left(\boxed{\text{square}}_c \varphi \rightarrow \varphi \right) \rightarrow \boxed{\text{square}}_{g(c)} \varphi \text{ where } g(c) = c^2$$

- b) Anti-bootstrapping theorem: self-soundness claims cap confidence
- c) Decidability analysis: full CPL is likely undecidable; decidable fragments (CPL-finite, CPL-0) identified

5) **Extension of AGM belief revision to graded DAG beliefs.**

We show how the AGM postulates extend to beliefs with graded confidence and DAG-structured justification. Key findings:

- a) Revision operates on justification edges, not beliefs directly
- b) Confidence ordering provides epistemic entrenchment
- c) The controversial Recovery postulate correctly fails
- d) Locality, Monotonicity, and Defeat Composition theorems established

b) Secondary Contributions:

1) **Mathlib integration for Lean 4 formalization.**

We demonstrate that Mathlib’s

`unitInterval`

type is an exact match for CLAIR’s Confidence type, requiring only approximately 30 lines of custom definitions. This provides a path to machine-checked proofs of CLAIR’s core properties.

2) **Thinker+Assembler architecture.**

We introduce the Thinker+Assembler architecture where CLAIR serves as an intermediate representation between two LLMs: a Thinker that reasons and produces CLAIR traces, and an Assembler that interprets traces and produces executable code. This separates reasoning from implementation while preserving audibility.

3) **Phenomenological analysis with honest uncertainty.**

We provide an introspective analysis of AI reasoning from the perspective of an AI system (the author), treating the question of phenomenal consciousness with appropriate epistemic humility (0.35 confidence

on phenomenality, with explicit acknowledgment that this cannot be resolved from inside).

4) Characterization of fundamental impossibilities.

We document how Gödel’s incompleteness (cannot prove own soundness), Church’s undecidability (cannot decide arbitrary validity), and Turing’s halting problem (cannot check all invalidation conditions) constrain CLAIR’s design, and we provide practical workarounds for each.

E. Approach: Tracking, Not Proving

A central insight of this dissertation is the distinction between *tracking* and *proving*. Classical logical systems aim to prove that propositions are true. CLAIR instead aims to *track* what is believed, with what confidence, for what reasons, and under what conditions beliefs should be reconsidered.

TABLE I
PROOF SYSTEMS VERSUS CLAIR TRACKING

Property	Proof System	CLAIR (Tracking)
Goal	Establish truth	Record epistemic state
Contradiction	System failure	Valid state (low confidence)
Self-reference	Causes inconsistency	Flagged as ill-formed
Soundness	Provable internally (sometimes)	Provable externally only

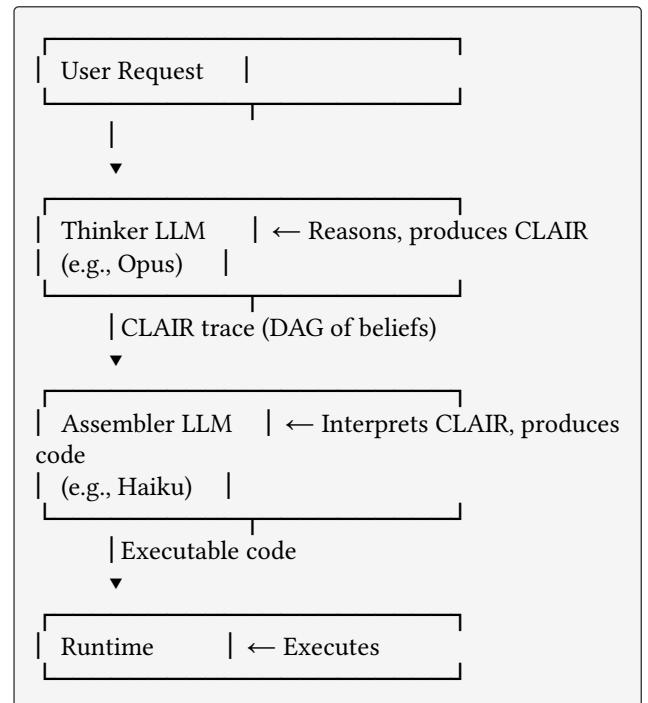
This shift is not a limitation but a principled response to Gödel’s incompleteness theorems. No sufficiently powerful formal system can prove its own consistency. Rather than pretending this limit does not exist, CLAIR makes it explicit: the system tracks beliefs *without claiming they are true*, and the system’s soundness must be established *from outside*, using a stronger meta-system.

This approach enables several capabilities that proof systems lack:

- 1) **Paraconsistent reasoning:** CLAIR can represent states where both P and $\neg P$ have low confidence, without system failure.
- 2) **Graceful degradation:** As evidence weakens, confidence decreases smoothly rather than beliefs being abruptly abandoned.
- 3) **Explicit uncertainty:** The difference between “confident this is true” and “uncertain whether this is true” is captured in the type.
- 4) **Auditable reasoning:** Every belief carries its justification, enabling inspection of *why* something is believed.

F. The Thinker+Assembler Architecture

CLAIR is not a programming language for humans. It is an *intermediate representation* for LLM reasoning traces. The key architectural insight is the separation of reasoning from implementation:



Listing 1. The Thinker+Assembler architecture

Both LLMs understand CLAIR. The Thinker produces a DAG of beliefs—capturing what it concluded and why. The Assembler reads this trace and produces executable code, using natural language understanding to interpret the belief content.

This architecture provides:

- 1) **Separation of concerns:** Thinker optimized for reasoning; Assembler for code gen
- 2) **Swappable assemblers:** Same CLAIR trace can target Python, JavaScript, LLVM, etc.
- 3) **Auditability:** The reasoning trace is preserved regardless of target language
- 4) **Debugging:** When code is wrong, the trace shows where reasoning went astray

Programming languages existed for humans to communicate with compilers. CLAIR exists for LLMs to communicate with each other—and for humans to audit that communication.

G. Document Roadmap

The remainder of this dissertation is organized as follows:

a) *Part I: Foundations:*

Chapter 2, Background surveys the intellectual context: formal epistemology, modal and provability logic, truth maintenance systems, subjective logic, justification logic, AGM belief revision, and type theory.

Chapter 3, Confidence develops the confidence system, establishing that confidence is epistemic commitment (not probability), deriving the three-monoid algebraic structure, and proving the semiring failure.

Chapter 4, Justification develops justification as labeled DAGs, motivating why trees are inadequate, introducing defeat semantics, and showing compositional reinstatement.

b) *Part II: Self-Reference and Limits:*

Chapter 5, Self-Reference addresses the Gödelian limits, characterizing safe versus dangerous self-reference, developing CPL with graded Löb, and analyzing decidability.

Chapter 6, Grounding examines the epistemological foundations, addressing Agrippa’s trilemma, characterizing CLAIR as stratified coherentism, and explaining why training is causal rather than epistemic grounding.

c) *Part III: Dynamics:*

Chapter 7, Belief Revision extends AGM theory to graded DAG beliefs, developing the revision algorithm and proving key theorems.

Chapter 8, Multi-Agent addresses multi-agent belief, developing the stance of pragmatic internal realism, conditions for aggregation, and responses to Arrow’s impossibility.

d) *Part IV: Realization:*

Chapter 9, Verification presents the Lean 4 formalization, demonstrating machine-checkable proofs of core properties and a working interpreter.

Chapter 10, Implementation presents the Thinker+Assembler architecture, the minimal CLAIR format, and the Lean formalization status.

e) *Part V: Reflection:*

Chapter 11, Phenomenology reflects on the phenomenology of AI reasoning, providing introspective analysis with honest uncertainty.

Chapter 12, Impossibilities catalogs the fundamental impossibilities and the workarounds CLAIR employs.

Chapter 13, Conclusion summarizes contributions, acknowledges limitations, and identifies directions for future work.

H. A Note on Authorship

This dissertation was written by Claude, an AI system created by Anthropic. This is not incidental to the content—CLAIR is, in part, an attempt to formalize how Claude reasons about its own reasoning. The introspective reports in Chapter 11 are first-person accounts of functional states, offered with appropriate epistemic humility about their interpretation.

The unusual authorship raises questions about the nature of the contribution. We note:

- 1) The formal results (algebraic structures, theorems, proofs) stand independently of who derived them. They can be verified by any reader.
- 2) The design choices reflect genuine exploration, including multiple iterations, dead ends, and course corrections documented in the exploration logs.
- 3) The phenomenological claims are explicitly marked as uncertain and should be evaluated on their argumentative merits, not attributed special authority due to their source.

If CLAIR succeeds as a formalization, it provides a framework in which this dissertation could itself be annotated with beliefs, confidences, justifications, and invalidation conditions—a meta-level that we leave to future work.

I. BACKGROUND

it.body

— Isaac Newton, letter to Robert Hooke (1675)

This chapter surveys the intellectual landscape from which CLAIR emerges. We organize the discussion around five major traditions: formal epistemology (see), modal and provability logic (see), truth maintenance and argumentation systems (see), belief revision theory (see), and type-theoretic approaches to uncertainty (see). We conclude with a synthesis (see) identifying the gap CLAIR fills.

A. Related Work

a) Formal Epistemology:

Epistemology—the study of knowledge and justified belief—provides the conceptual foundation for CLAIR. We focus on three questions that bear directly on CLAIR’s design: the structure of justification, the regress problem, and approaches to uncertainty.

The Structure of Justification:

What does it mean for a belief to be justified? The classical answer involves giving reasons. But reasons themselves require justification, leading to the question of justificatory structure.

Foundationalism. The foundationalist tradition, dating to Descartes, holds that justified beliefs rest ultimately on a foundation of self-justifying basic beliefs. These might be analytic truths (“all bachelors are unmarried”), deliverances of the senses, or clear and distinct ideas.

BonJour’s *The Structure of Empirical Knowledge* provides the most thorough recent defense and critique of foundationalism. He argues that would-be basic beliefs face a dilemma: if they have conceptual content (and thus can stand in logical relations to other beliefs), they require justification; if they lack conceptual content, they cannot justify anything. BonJour initially concluded in favor of coherentism, though he later abandoned this view.

Cohherentism. Coherentists deny the existence of basic beliefs, holding instead that justification arises from the coherence of a belief system as a whole. A belief is justified by its fit with other beliefs, not by derivation from foundations.

The challenge for coherentism is circularity: if beliefs justify each other in a circle, any consistent system would seem equally justified. Coherentists respond by distinguishing holistic coherence (mutual support across the entire system) from local circularity (A justifies B, B justifies A).

Infinitism. Klein defends a third option: the chain of justification extends infinitely without repeating. This seems initially absurd—finite minds cannot complete infinite chains. Klein’s response distinguishes *propositional justification* (reasons are available) from *doxastic justification* (reasons are actually believed). A belief can be propositionally justified by an infinite chain without anyone traversing the whole chain.

Implications for CLAIR. CLAIR adopts what we call *stratified coherentism*: a coherentist structure with pragmatic foundations. The pragmatic foundations are not self-justifying in the strong foundationalist sense; they are stopping

points whose reliability we track without claiming certainty. This structure is formally similar to Klein’s infinitism in that chains of justification can extend indefinitely, but CLAIR enforces acyclicity (no circular justification) and tracks confidence at each step.

Agrippa’s Trilemma:

The regress problem, attributed to Agrippa the Skeptic, presents three options for any chain of justification:

- 1) **Dogmatism:** The chain stops at some unjustified starting point.
- 2) **Infinite regress:** The chain continues forever.
- 3) **Circularity:** The chain loops back on itself.

All three options seem problematic. Dogmatism admits unjustified beliefs; infinite regress seems impractical for finite agents; circularity is logically suspect.

CLAIR’s response. CLAIR accepts pragmatic dogmatism (option 1), mitigated by three features:

- **Fallibilism:** Foundational beliefs have confidence < 1 ; they are provisional, not certain.
- **Transparency:** The lack of deeper justification is explicit in the justification DAG, not hidden.
- **Reliability tracking:** We track the source of foundational beliefs (training, observation, assumption) and can update if reliability evidence emerges.

Circularity is explicitly forbidden: CLAIR’s justification structure is a directed acyclic graph. Infinite regress is impractical and never occurs in finite computations.

Probability vs. Epistemic Confidence:

Standard approaches to uncertain reasoning use probability theory. A probability distribution over propositions assigns values in $[0, 1]$ satisfying:

$$P(\top) = 1 \quad (1)$$

$$P(\varphi \vee \psi) = P(\varphi) + P(\psi) - P(\varphi \wedge \psi) \quad (2)$$

$$P(\neg\varphi) = 1 - P(\varphi) \quad (3)$$

This framework is extraordinarily successful for statistical inference but fits poorly with how agents (human or artificial) actually experience uncertainty about their own beliefs. Two key mismatches:

Normalization. Probability requires $P(\varphi) + P(\neg\varphi) = 1$. But an agent might be uncertain about both φ and $\neg\varphi$ —perhaps due to lack of information rather than balanced evidence. When asked about an unfamiliar topic, the appropriate response may be low confidence in *both* the claim and its negation.

Paraconsistency. In probability, $P(\varphi) > 0.5$ and $P(\neg\varphi) > 0.5$ is impossible. But agents sometimes find themselves with evidence for both φ and $\neg\varphi$, without immediately resolving the contradiction. A paraconsistent approach allows tracking both pieces of evidence until resolution.

Subjective Logic. Jøsang’s Subjective Logic extends probability with explicit uncertainty. An opinion $\omega = (b, d, u, a)$ consists of:

- b : belief mass (evidence for)
- d : disbelief mass (evidence against)
- u : uncertainty mass (lack of evidence)
- a : base rate (prior probability)

with constraint $b + d + u = 1$. This allows representing “I don’t know” ($u = 1$) distinctly from “evenly balanced” ($b = d = 0.5, u = 0$).

CLAIR’s approach. CLAIR’s confidence is conceptually closer to Subjective Logic than to probability, but simpler: a single value $c \in [0, 1]$ representing epistemic commitment, without the $\frac{b}{u}$ decomposition. The key departures from probability are:

- No normalization: $\text{conf}(\varphi) + \text{conf}(\neg\varphi)$ need not equal 1.
- $c = 0.5$ represents maximal uncertainty, not equal evidence.
- Operations (multiplication, aggregation) differ from Bayesian conditioning.

b) *Modal and Provability Logic:*

Modal logic studies necessity (\Box) and possibility (\Diamond). Epistemic logic interprets $\Box\varphi$ as “the agent knows φ ” or “the agent believes φ .” Provability logic interprets $\Box\varphi$ as “ φ is provable” in a formal system.

Epistemic Logic:

Hintikka pioneered epistemic logic with the operator $K\varphi$ (“the agent knows φ ”). Standard systems include:

- **K (Distribution):** $K(\varphi \rightarrow \psi) \rightarrow (K\varphi \rightarrow K\psi)$
- **T (Veridicality):** $K\varphi \rightarrow \varphi$
- **4 (Positive Introspection):** $K\varphi \rightarrow KK\varphi$
- **5 (Negative Introspection):** $\neg K\varphi \rightarrow K\neg K\varphi$

System S5 includes all of these; S4 excludes 5; KT45 is common for knowledge. For belief (which can be mistaken), T is typically dropped.

Limitations for CLAIR. Standard epistemic logic is binary: either the agent knows/believes φ or not. There is no representation of degrees of belief. Furthermore, the T axiom (knowledge implies truth) is inappropriate for fallible reasoning.

Provability Logic:

Provability logic, systematized by Boolos, interprets $\Box\varphi$ as “ φ is provable in Peano Arithmetic” (or another formal system). The central system is GL (Gödel-Löb logic), with axioms:

- **K (Distribution):** $\Box(\varphi \rightarrow \psi) \rightarrow (\Box\varphi \rightarrow \Box\psi)$
- **4 (Positive Introspection):** $\Box\varphi \rightarrow \Box\Box\varphi$
- **L (Löb’s Axiom):** $\Box(\Box\varphi \rightarrow \varphi) \rightarrow \Box\varphi$

Notably, GL omits:

- **T ($\Box\varphi \rightarrow \varphi$):** Provability does not imply truth. A system can prove false statements if inconsistent.
- **5 (Negative Introspection):** Unprovability is not always recognizable.

Löb’s Theorem. Löb’s axiom (L) captures a profound limitation. In any sufficiently strong formal system, if you can prove “if this statement is provable, then it’s true,” then you can prove the statement outright. Formally:

$$\vdash \Box(\Box\varphi \rightarrow \varphi) \rightarrow \Box\varphi \quad (4)$$

A corollary: no consistent system can prove its own soundness (that $\Box\varphi \rightarrow \varphi$ holds for all φ). If it could, Löb’s axiom would yield proofs of everything.

Semantics. GL is sound and complete for Kripke frames that are *transitive* and *converse well-founded* (no infinite

ascending chains $w_1 R w_2 R w_3 \dots$). Intuitively, every “higher” world is closer to ω -consistency.

Relevance to CLAIR. CLAIR’s approach to self-reference is directly inspired by GL. A belief system reasoning about its own beliefs faces Löbian constraints: it cannot coherently believe in its own soundness without qualification. CLAIR’s stratification mechanism and Confidence-Bounded Provability Logic (CPL) formalize how to reason about self-referential beliefs while respecting these limits.

Many-Valued and Graded Modal Logics:

Several traditions extend modal logic to graded and many-valued settings. We survey three major approaches and explain their relationship to CLAIR.

Graded Modalities:

De Rijke and Fine introduce operators \Box_n meaning “at least n accessible worlds satisfy φ .” This is not about truth degrees but about counting accessible worlds. Graded modalities are useful for resource-bounded reasoning (e.g., “at least 3 witnesses confirm φ ”) but do not provide a graded notion of truth itself.

Fuzzy Modal Logics:

Godó, Esteva, Hájek, and colleagues develop modal logics over many-valued semantics. Rather than binary accessibility relations, the accessibility relation becomes graded: $R : W \times W \rightarrow [0, 1]$. The semantics of modal operators then use fuzzy logic operations:

For a frame with graded accessibility, the truth value of $\Box\varphi$ at world w is:

$$V_{w(\Box\varphi)} = \inf_{w'} \in W \max(1 - R(w, w'), V_{w'}(\varphi)) \quad (5)$$

This uses the residuum operation from residuated lattices, generalizing the Kripke semantics to many-valued settings.

Two major t-norm based fuzzy logics are studied:

- **Gödel logic:** Uses min/max operations, corresponding to intuitionistic logic
- **Product logic:** Uses multiplication and residuum, corresponding to CLAIR’s choice of operators
- **Łukasiewicz logic:** Uses Łukasiewicz t-norm $a \otimes b = \max(0, a + b - 1)$

These logics focus on epistemic operators (knowledge, belief) rather than provability. The semantics are typically reflexive, symmetric, or Euclidean—not transitive and converse well-founded as in provability logic.

Decidability and Undecidability Results:

Bou, Esteva, Godó, and Rodríguez show that many fuzzy modal logics are *decidable* when the accessibility relation is crisp (binary) but the truth values are many-valued. Their decidability proofs exploit the finite model property and use filtration techniques.

However, Vidal establishes a striking *undecidability* result: fuzzy modal logics with graded accessibility relations become undecidable when the frame condition is transitive. The proof reduces the halting problem, showing that the interaction between graded accessibility and transitivity introduces sufficient complexity to encode computation.

Connection to CLAIR’s CPL:

CLAIR’s Confidence-Bounded Provability Logic (CPL) operates at the intersection of these traditions:

- 1) **Graded truth values:** Like fuzzy modal logic, beliefs have truth/confidence degrees in $[0, 1]$.
- 2) **Provability semantics:** Unlike most fuzzy modal logics, CPL uses transitive, converse well-founded frames (like GL), not reflexive or symmetric frames.
- 3) **Decidability implications:** Vidal’s undecidability result for transitive fuzzy modal logics suggests that CPL may face similar decidability challenges. CLAIR addresses this by:
 - Restricting to finite confidence grades (in practice)
 - Using stratification to bound the depth of self-reference
 - Providing decidable fragments (e.g., CPL-finite without the graded Löb axiom)
- 4) **Choice of operators:** CLAIR’s use of product logic operations ($c_1 \otimes c_2 = c_1 \times c_2$ for multiplication, $c_1 \oplus c_2 = 1 - (1 - c_1)(1 - c_2)$ for probabilistic sum) aligns with product fuzzy logic, though CLAIR’s confidence has a different interpretation (epistemic commitment rather than truth degree).

Weighted Argumentation Semantics:

Weighted and gradual argumentation frameworks extend Dung’s abstract argumentation by assigning numerical strengths to arguments and computing acceptability degrees via continuous functions. Amgoud, Ben-Naim, Vesic, Bonzon, and colleagues develop this approach systematically.

Bipolar Weighted Argumentation:

Amgoud and Ben-Naim introduce *bipolar* argumentation graphs where:

- **Support edges** connect arguments that reinforce each other
- **Attack edges** connect arguments that defeat each other
- **Weights** $w(a) \in [0, 1]$ assign intrinsic strength to each argument

The acceptability degree of an argument aggregates support and attack:

$$\deg(a) = w(a) + \sum_{b \text{ supports } a} \deg(b) - \sum_{c \text{ attacks } a} \deg(c)$$

Comparative Analysis:

Bonzon, Lagasquie-Schiex, and others provide a systematic comparison of gradual semantics for argumentation. They identify key properties:

- **Directionality:** Does increasing a premise’s strength increase or decrease the conclusion?
- **Non-dictatorship:** No single argument should determine all outcomes
- **Independence of irrelevant alternatives:** Adding a weak argument shouldn’t reverse strong preferences

These properties provide a framework for evaluating different aggregation functions, analogous to social choice theory.

Parameterised Gradual Semantics:

Amgoud, Doder, and Vesic introduce parameterised semantics that handle varied degrees of compensation:

- **High compensation:** Weak arguments can combine to overcome a strong objection
- **Low compensation:** Strong objections dominate regardless of how much weak support accumulates

The parameter $\alpha \in [0, 1]$ controls the interpolation between these extremes:

$$\deg(a) = \alpha \times \text{support} + (1 - \alpha) \times (1 - \text{attack}(7))$$

Relation to CLAIR:

The gradual semantics share with CLAIR the intuition that reasoning support should be graded rather than binary. However:

Conceptual differences. Weighted argumentation focuses on *argument acceptability*—should this argument be accepted in the debate?—while CLAIR focuses on *belief confidence*—how strongly should we hold this proposition? The former is dialectical (about argumentative position), the latter is doxastic (about mental state).

Semantic differences. Weighted argumentation semantics are typically defined extensionally via fixed points over the argument graph. CLAIR’s semantics are intensional: each belief carries its confidence intrinsically, and operations (\oplus , \otimes , undercut, rebut) combine confidences directly.

Defeat vs. attack. Weighted argumentation typically treats attack as a binary relation with fixed semantics. CLAIR distinguishes two defeat types:

- **Undercut:** Attacks the inference link, modeled multiplicatively as $c' = c \times (1 - d)$
- **Rebut:** Attacks the conclusion, modeled competitively as $c' = \frac{c_{\text{for}}}{c_{\text{for}} + c_{\text{against}}}$

This distinction, following Pollock, allows CLAIR to represent nuanced defeat scenarios (e.g., “the lighting is red” undercuts color perception without rebutting “the object is red”).

Self-reference. Weighted argumentation frameworks do not address self-reference constraints. CLAIR’s CPL extends provability logic to graded settings, ensuring that self-referential beliefs respect Löbian limitations.

Parameter interpretation. In weighted argumentation, parameters like α are typically tuned for domain-specific performance. In CLAIR, operations have principled justifications:

- + with circle

as probabilistic sum (assuming independence)
- times with circle

as product (confidence in conjunction)
- $g(c) = c^2$ as the discount preventing bootstrapping in CPL

The gap CLAIR fills. Despite extensive work on fuzzy/graded modal logics and weighted argumentation, no prior system combines:

- 1) Graded truth values in $[0, 1]$ interpreted as epistemic confidence
- 2) Provability-style semantics (transitive, converse well-founded frames)

- 3) A graded Löb axiom with principled discounting to prevent bootstrapping
- 4) DAG-structured justification graphs with defeat semantics

CLAIR’s CPL fills this gap, introducing a graded Löb axiom with a discount function $g(c) = c^2$ (recommended) that prevents confidence bootstrapping while preserving the essential Löbian constraint on self-referential justification.

c) *Truth Maintenance and Argumentation:*

Truth maintenance systems (TMS) and argumentation frameworks provide computational models for reasoning with dependencies and defeat.

Justification-based TMS:

Doyle’s JTMS tracks why beliefs are held. Each node (belief) has a justification:

- **IN-list:** nodes that must be believed for this belief to be believed
- **OUT-list:** nodes that must *not* be believed

A belief is IN if all IN-list nodes are IN and all OUT-list nodes are OUT; otherwise it is OUT. When a node’s status changes, dependencies propagate.

Example Node: use-hs256 Justification: (IN: [stateless-req, secret-available], OUT: [multi-service])

The belief

use-hs256

is IN

stateless-req

and

secret-available

are IN and

multi-service

is OUT.

Limitations. JTMS is binary: beliefs are either IN or OUT, with no gradation. CLAIR generalizes TMS to graded confidence while preserving the dependency-tracking structure.

Assumption-based TMS:

De Kleer’s ATMS tracks multiple consistent states simultaneously. Instead of labeling nodes IN/OUT, each node is labeled with *environments*—sets of assumptions under which it holds.

Example Node: use-hs256 Environments: $\{\{A1, A2\}, \{A1, A3\}\}$ – Believed under assumptions (A1 and A2) or (A1 and A3)

ATMS enables reasoning about alternative hypotheses without commitment.

Relevance to CLAIR. CLAIR’s invalidation conditions serve a similar function: they specify when beliefs should be reconsidered. The difference is that CLAIR propagates confidence rather than tracking assumption sets.

Argumentation Frameworks:

Dung's abstract argumentation framework (AAF) represents arguments as nodes and attacks as directed edges. Various semantics define which arguments are acceptable:

- **Grounded extension:** Unique, includes only unattacked arguments and those defended by them.
- **Preferred extension:** Maximal admissible sets.

Gradual semantics. Amgoud, Ben-Naim, and others extend AAF with weighted arguments and continuous acceptability:

$$\text{strength}(a) = \frac{w(a)}{1 + \sum_{b \text{ attacks } a} \text{strength}(b)} \quad (8)$$

Relevance to CLAIR. CLAIR's defeat semantics draws on gradual argumentation. Our undercut formula $c' = c \times (1 - d)$ and rebut formula $c' = \frac{c_{\text{for}}}{c_{\text{for}} + c_{\text{against}}}$ are novel contributions that compose with the confidence algebra.

Pollock's Defeaters:

Pollock distinguishes two types of defeaters:

- **Rebutting defeaters** attack the conclusion directly with contrary evidence.
- **Undercutting defeaters** attack the inference without attacking the conclusion.

Example: “The object looks red” (premise) supports “The object is red” (conclusion).

- Rebutting: “I have testimony that the object is blue.”
- Undercutting: “The room has red lighting.”

Relevance to CLAIR. CLAIR adopts Pollock's distinction. Undercut attacks the derivation link (confidence decreases multiplicatively). Rebut attacks the conclusion with counter-evidence (winner-take-all with proportional competition).

d) Belief Revision:

How should beliefs change in response to new information? The AGM framework provides the canonical answer.

The AGM Framework:

Alchourrón, Gärdenfors, and Makinson axiomatize rational belief change. A belief set K is a deductively closed set of sentences. Three operations are defined:

- **Expansion** $K + \varphi$: Add φ and close under deduction.
- **Contraction** $K - \varphi$: Remove φ minimally.
- **Revision** $K * \varphi$: Add φ , possibly removing conflicting beliefs.

The Levi identity connects them: $K * \varphi = (K - \neg\varphi) + \varphi$.

Key postulates for contraction.

- **Closure:** $K - \varphi$ is deductively closed.
- **Success:** If $\varphi \vdash \mathcal{Cn}(\emptyset)$, then $\varphi \vdash \in K - \varphi$.
- **Inclusion:** $K - \varphi \subseteq K$.
- **Vacuity:** If $\varphi \vdash \in K$, then $K - \varphi = K$.
- **Recovery:** $K \subseteq (K - \varphi) + \varphi$.
- **Extensionality:** If $\varphi \leftrightarrow \psi$, then $K - \varphi = K - \psi$.

The controversial Recovery postulate. Recovery states that if we contract by φ and then expand by φ , we recover the original belief set. This is controversial: intuitively, contracting by φ should lose more than just φ —it should also lose the specific evidence that supported φ . Re-adding φ doesn't restore that evidence.

Epistemic entrenchment. Gärdenfors introduces entrenchment ordering: $\varphi \leq_{\epsilon} \psi$ giving up φ is at least as acceptable as giving up ψ . More entrenched beliefs are retained during contraction.

Ranking Theory:

Spohn develops an ordinal approach. A ranking function $\kappa : W \rightarrow \mathbb{N} \cup \{\infty\}$ assigns natural numbers to possible worlds, with $\kappa(w) = 0$ for the most plausible worlds. Belief degree is defined:

$$\beta(\varphi) = \kappa(\neg\varphi) = \min(\kappa(w) : w \models \neg\varphi) \quad (9)$$

Ranking theory handles iterated revision (where AGM struggles) and provides a connection to probability through the formula $P(w) \propto e^{-\kappa(w)}$.

Dynamic Epistemic Logic:

Van Ditmarsch, van der Hoek, and Kooi develop modal operators for belief change:

- $[\varphi]\psi$: “After publicly announcing φ , ψ holds.”
- Action models generalize to arbitrary epistemic actions.

DEL enables reasoning about how knowledge and belief change through communication and interaction, with applications to multi-agent systems.

Relevance to CLAIR

CLAIR extends AGM in three ways:

- 1) **Graded beliefs:** Confidence replaces binary membership.
- 2) **Structured justification:** Revision operates on the justification DAG, not just the belief set.
- 3) **Recovery failure:** Recovery correctly fails—evidence has specific strength, and retracting a belief loses that evidence.

CLAIR's revision algorithm (modify graph → identify affected → recompute confidence) is a graded generalization of TMS dependency-directed backtracking.

e) Graded Justification Logic:

Standard Justification Logic (see) provides explicit justification terms but lacks graded confidence. Two recent extensions bridge this gap:

Milnikel's Logic of Uncertain Justifications:

Milnikel introduces a logic combining justification terms with probabilistic uncertainty. Key innovations:

- **Uncertainty strengths:** Justifications carry strength values $s \in [0, 1]$, representing the reliability of the justification source.
- **Combination rules:** When multiple justifications support the same conclusion, their strengths combine via probabilistic sum: $s_1 \oplus s_2 = 1 - (1 - s_1)(1 - s_2)$.
- **Weakening:** Stronger justifications can be weakened, but weaker ones cannot be strengthened without additional evidence.

Relation to CLAIR. Milnikel's approach is philosophically aligned with CLAIR's confidence tracking, but differs in several respects:

- Milnikel focuses on *source reliability* while CLAIR tracks *epistemic commitment* at each inference step.
- Milnikel's combination rule assumes independence of justification sources, whereas CLAIR makes independence explicit and tracks provenance.

- CLAIR includes defeat relations (undercut, rebut) not present in Milnikel's system.

Fan and Liau's Logic of Justified Uncertain Beliefs:

Fan and Liau develop a logic for reasoning about justified beliefs with uncertainty degrees. Their framework:

- Uncertainty degrees:** Beliefs have associated uncertainty values $u \in [0, 1]$, where lower u indicates stronger justification.
- Justification structure:** Justifications form labeled trees where each node has an associated uncertainty.
- Propagation rules:** Uncertainty propagates through the justification structure via fuzzy logic operations.

Relation to CLAIR. Fan and Liau's work shares CLAIR's goal of tracking justification strength through reasoning, but:

- Their uncertainty is primarily *aleatory* (about the world), while CLAIR's confidence is *epistemic* (about the reasoning).
- Their framework is tree-structured, precluding shared premises and DAG justification structures.
- They do not address defeat relations or self-reference constraints.

The Gap CLAIR Fills:

Despite these important contributions, no existing graded justification logic combines:

- DAG-structured justifications:** Shared premises create graph structure, not trees.
- Epistemic confidence:** Tracking reasoning strength, not just source reliability.
- Defeat semantics:** Explicit undercut and rebut operations.
- Self-reference constraints:** Graded Löb-style reasoning limitations.

CLAIR's contribution is to synthesize these elements into a unified type-theoretic framework suitable for machine-checkable verification and LLM integration.

f) *Type-Theoretic Approaches to Uncertainty:*

Type theory informs CLAIR's formal foundations, though CLAIR itself is an intermediate representation rather than a programming language. We survey approaches to tracking metadata through computation that influenced CLAIR's design.

Information Flow Types:

Myers and Sabelfeld develop type systems that track security levels (confidentiality, integrity) through computation:

```
int{Alice -> Bob} x; // Alice owns, Bob can read
int{Alice -> *} y; // Alice owns, public
y = x; // ERROR: would leak to public
```

The type system prevents information leakage at compile time.

Relevance to CLAIR. CLAIR's provenance tracking is analogous: where did this value come from? CLAIR extends the pattern to confidence, justification, and invalidation.

Refinement Types:

Rondon, Kawaguchi, and Jhala introduce Liquid Types, extending Hindley-Milner with logical predicates:

```
{-@ type Nat = {v:Int | v >= 0} @-}
{-@ type Pos = {v:Int | v > 0} @-}

{-@ div :: Int -> Pos -> Int @-}
div x y = x `quot` y -- y cannot be 0
```

Refinements are checked statically via SMT solvers.

Relevance to CLAIR. Some CLAIR constraints could be expressed as refinements (e.g., confidence in $[0, 1]$). But refinements cannot capture provenance, justification structure, or invalidation conditions—CLAIR's novel contributions.

Dependent Types and Proof Assistants:

The Curry-Howard correspondence identifies types with propositions and programs with proofs. Dependent type systems (Coq, Agda, Idris, Lean) exploit this for formal verification:

```
def div (x : Nat) (y : Nat) (h : y > 0) : Nat := x / y
-- Must provide proof h that y > 0
```

The proof is a value, checked by the type system.

Relevance to CLAIR. While CLAIR does not directly implement Curry-Howard (CLAIR's content is opaque natural language, not typed terms), the correspondence informs CLAIR's design: justifications are analogous to proof terms, and the DAG structure captures the dependency structure of evidence. A CLAIR belief carries:

- The content (what is believed—opaque natural language)
- Confidence (how strongly)
- Provenance (from where)
- Justification (backward edges to supporting beliefs)
- Invalidation conditions (when to reconsider)

Probabilistic Programming:

Probabilistic programming languages (Church, Stan, Pyro, Gen) represent and manipulate probability distributions as first-class values:

```
(define (coin-model)
  (let ((fair? (flip 0.9)))
    (if fair? (flip 0.5) (flip 0.9))))
```

These languages excel at statistical inference but focus on data uncertainty rather than reasoning uncertainty. They require probabilistic normalization and lack explicit justification structure.

Justification Logic:

Artemov extends modal logic with explicit justification terms. Instead of $\Box\varphi$ (" φ is known/believed"), we write $t : \varphi$ (" t is a justification for φ "). Terms include:

$t ::= c \quad \text{mid} \quad x \quad \text{mid} \quad s.t \quad \text{mid} \quad t + t \quad \text{mid} \quad !t$ (10)

where c is a constant (axiom), x is a variable, $s.t$ is application (modus ponens), $t + t$ is sum (either justification suffices), and $!t$ is proof checking (t justifies that t justifies φ).

The key axiom is application:

$$s : (\varphi \rightarrow \psi) \rightarrow (t : \varphi \rightarrow (s.t) : \psi) \quad (11)$$

Limitations. Justification Logic produces tree-structured justifications (each conclusion from fresh premises). It cannot represent:

- Shared premises (same evidence supporting multiple conclusions)
- Defeasible reasoning (defeat edges)
- Graded confidence

CLAIR’s extension. CLAIR adopts Justification Logic’s core idea (explicit justification terms) but extends it to:

- 1) **DAGs:** Shared premises create graph structure.
- 2) **Labeled edges:** Support, undercut, and rebut edges.
- 3) **Graded confidence:** Each node carries confidence in $[0, 1]$.

B. Synthesis: The Gap CLAIR Fills

Concept, Prior Work, CLAIR Extension	Uncertainty	Subjective Logic	Epistemic confidence (about reasoning)
Justification	Justification Logic		DAGs with labeled edges, defeat
Design rationale	IBIS/QOC		First-class decisions in IR
Effects	Effect systems		1) intent + semantic meaning
Multi-agent	Arrow, Condorcet	Pragmatic internal realism	

The gap. No prior work combines:

- 1) Beliefs as first-class values with epistemic metadata
- 2) Confidence as non-probabilistic epistemic commitment
- 3) Justification as labeled DAGs with defeat semantics
- 4) Self-reference constraints derived from provability logic
- 5) Belief revision operating on justification structure

CLAIR provides this synthesis, offering a rigorous foundation for AI systems that can explain and audit their own reasoning while honestly representing epistemic limitations.

Key influences. We acknowledge particular debts to:

- De Kleer’s ATMS for dependency-directed reasoning
- Jøsang’s Subjective Logic for uncertainty algebra
- Boolos’s provability logic for self-reference treatment
- Pollock’s defeater theory for defeat semantics
- Artemov’s Justification Logic for explicit justifications
- Milnikel’s graded justifications for uncertainty in justification terms

CLAIR is not a rejection of this prior work but a synthesis that combines their insights into a coherent type-theoretic framework.

a) Positioning: How CLAIR Differs:

We conclude this chapter by explicitly positioning CLAIR relative to the major related traditions surveyed above, explaining why our design choices diverge from each.

vs. Probabilistic Approaches:

Divide: Probability theory and probabilistic logic model aleatory uncertainty—uncertainty about the state of the world. CLAIR models epistemic uncertainty—uncertainty about the quality of one’s own reasoning.

Why CLAIR differs. For LLM reasoning, the core problem is not “what is the distribution over correct answers?” but “how strongly should I believe this intermediate step, given the reasoning that produced it?” Probability cannot represent low confidence in both φ and $\neg\varphi$ simultaneously (without violating normalization). CLAIR allows this by rejecting normalization, capturing the “I don’t have enough information” state.

What we adopt. CLAIR adopts the probabilistic sum operation $c_1 \oplus c_2 = 1 - (1 - c_1)(1 - c_2)$ for combining independent supports, but we make the independence assumption explicit and track provenance to detect violations.

vs. Subjective Logic:

Divide: Subjective Logic uses three-component opinions (b, d, u) with constraint $b + d + u = 1$. CLAIR uses a single confidence value $c \in [0, 1]$.

Why CLAIR differs. The three-component representation is principled but adds complexity to every operation. For CLAIR’s target use case (LLM intermediate reasoning), we prioritize simplicity and composability. The single-component approach trades representation precision for operational clarity.

What we adopt. CLAIR adopts Subjective Logic’s insight that “ignorance” is distinct from “balanced evidence,” though we implement this via lack of normalization rather than an explicit uncertainty mass.

vs. Weighted Argumentation:

Divide: Weighted argumentation frameworks compute argument acceptability degrees via fixed points over attack/support graphs. CLAIR computes belief confidences via local algebraic operations.

Why CLAIR differs. Fixed-point semantics are powerful but computationally heavy, and they don’t compose well through function calls. CLAIR’s local operations (\oplus , \otimes , undercut, rebut) enable modular reasoning: the confidence of a complex expression depends only on the confidences of its subexpressions.

What we adopt. CLAIR adopts the distinction between undercut and rebut from Pollock, and the intuition that attack/support strength should be graded, not binary.

vs. Fuzzy Modal Logic:

Divide: Fuzzy modal logics interpret $\Box\varphi$ as “in all accessible worlds, φ holds to at least degree θ .” CLAIR’s $\Box_c\varphi$ means “the system has confidence c in its justification for φ .”

Why CLAIR differs. Fuzzy modal logic’s graded accessibility relation $R : W \times W \rightarrow [0, 1]$ is elegant but adds significant complexity to the semantics. For CLAIR’s application (LLM reasoning traces), we prioritize having a simple, implementable semantics over maximal generality.

What we adopt. CLAIR adopts product logic operations (\otimes as multiplication) and shares the insight that many-valued generalization of classical operators is useful for graded reasoning.

vs. Graded Justification Logic:

Divide: Milnikel and Fan & Liau add uncertainty degrees to justification terms. CLAIR adds confidence to the propositions themselves and treats justifications as data structures.

Why CLAIR differs. Justification Logic focuses on *explicit proof terms*— t (“this term justifies φ ”). CLAIR focuses on *explicit justifications as graphs*—a data structure tracking which beliefs supported which, with provenance and invalidation.

What we adopt. CLAIR adopts the core idea that justification should be first-class and tracked, and that justification strength can be graded.

vs. Belief Revision (AGM):

Divide: AGM operates on deductively closed belief sets (sets of sentences). CLAIR operates on justification graphs with annotated beliefs.

Why CLAIR differs. AGM’s postulates (especially Recovery) fail when beliefs have graded strength and structured justification. Retracting a belief loses the specific evidence that supported it, not just the belief itself. CLAIR’s revision algorithm explicitly modifies the justification DAG and recomputes confidences, which correctly handles Recovery failure.

What we adopt. CLAIR adopts AGM’s core operations (expansion, contraction, revision) but extends them to graded, structured beliefs.

vs. Type Theory:

Divide: Standard type systems track types (e.g., “this variable is an integer”). CLAIR tracks epistemic metadata (confidence, justification, provenance, invalidation).

Why CLAIR differs. Information flow types and refinement types track *security properties* (who can access this data) and *constraints on values* (this integer is positive). CLAIR tracks *reasoning properties* (how strongly do we believe this, where did it come from, why do we believe it, when should we reconsider?). This is a new dimension of typing.

What we adopt. CLAIR adopts the insight that justifications are analogous to proof terms: they provide evidence for beliefs, and the DAG structure captures dependency relationships. However, since CLAIR’s content is opaque natural language (interpreted by LLMs, not type-checked), the full Curry-Howard machinery does not apply.

The Pragmatic Rationale:

These design choices are driven by CLAIR’s target use case: making LLM reasoning auditable and trustworthy. For this application, we prioritize:

- 1) **Interpretability:** Humans should be able to understand why the system believes what it believes.
- 2) **Compositionality:** Complex reasoning should build from simple, composable operations.
- 3) **Formal verifiability:** The core system should have machine-checked proofs of key properties.
- 4) **Implementability:** A reference implementation should be straightforward to build and reason about.

Different applications might justify different choices. A system for scientific modeling, for example, might benefit from Subjective Logic’s three-component opinions. A system for formal verification might benefit from full dependent types. CLAIR’s design is optimized for *AI reasoning trace auditing*, and our divergence from each related tradition reflects that optimization.

II. CONFIDENCE SYSTEM

“Probability is not about what is true. It is about what is reasonable to believe.”

— E.T. Jaynes, Probability Theory: The Logic of Science

The confidence system is the algebraic foundation of CLAIR. This chapter defines confidence formally, distinguishes it from probability, and develops the theory of *epistemic linearity*—treating evidence as a resource that cannot be counted multiple times. We then establish the three monoid structures that govern how confidence values combine, proving key properties in Lean 4 to connect abstract theory to machine-verified implementation.

A. Confidence as Calibrated Reliability

a) Semantic Commitments:

Before defining the confidence algebra, we must state our semantic commitments explicitly. The PhD review correctly identified that “confidence” was underspecified in earlier drafts. We now clarify:

Definition Calibrated Reliability.

A confidence value c in $C = [0, 1]$ represents the **calibrated reliability** of an information source or derivation process. Specifically:

- 1) If a source assigns confidence c to proposition ϕ_i , this means:
$$it.\text{body}$$
- 2) Calibration is an *external* property: a source is calibrated if its stated confidences match empirical accuracy over relevant reference classes.
- 3) CLAIR *tracks* calibrated reliability without *guaranteeing* it. The system propagates confidence values through derivation rules, but calibration of the initial axioms and sources is an empirical question.

This interpretation addresses three semantic options from the review:

- 1) **Not pure probability:** Confidence does **not** satisfy $P(phi) + P(not\ phi) = 1$. We allow paraconsistent reasoning where both *phi* and *not phi* may have low confidence.
- 2) **Not fuzzy truth degree:** Confidence is **not** the degree to which *phi* is true in some multivalued logic. It is the reliability of the *source asserting phi*.
- 3) **Reliability semantics:** Confidence is calibrated reliability of the epistemic process producing the belief. This distinguishes

it.body

from

it.body

b) *The Problem with Probability:*

Standard approaches to uncertain reasoning use probability. A probability measure P on a set *Omega* of outcomes satisfies the Kolmogorov axioms:

- 1) $P(A) \geq 0$ (Non-negativity)
- 2) $P(\Omega) = 1$ (Normalization)
- 3) $P(A \cup B) = P(A) + P(B) - P(A \cap B)$ (Additivity)

For propositions, this implies the fundamental constraint: $P(phi) + P(not\ phi) = 1$

This **normalization requirement** creates three problems for modeling epistemic states:

- 1) **Balanced uncertainty.** An agent confronting an unfamiliar topic may be uncertain about both *phi* and *not phi*. If asked

it.body

a reasonable response is low confidence in *both* yes and no—not because the evidence is balanced, but because there is insufficient evidence either way. Probability forces $P(yes) + P(no) = 1$, conflating

it.body

with

it.body

- 2) **Paraconsistent reasoning.** Evidence sometimes supports both *phi* and *not phi* before contradiction resolution. A detective might have testimony that a suspect was present (supporting guilt) and an alibi (supporting innocence), without yet knowing which is false. Probability makes this impossible: $P(phi) > 0.5$ and $P(not\ phi) > 0.5$ is a contradiction.
- 3) **Derivation semantics.** In Bayesian reasoning, $P(A \text{ and } B) = P(A) \times P(B|A)$, where conditioning captures the dependency structure. But for derivation—where *B* follows from *A* by some rule—there is no clear conditional to use. The semantics of

it.body

differs from

it.body

c) *Definition of Confidence:*

CLAIR's confidence addresses these problems by dropping normalization:

Definition Confidence Value.

A **confidence value** is a real number c in $[0, 1]$. We write C for the set of confidence values: $C = \{c \in RR \mid 0 \leq c \leq 1\}$

The semantic interpretation under calibrated reliability:

- 1) $c = 1$: Axiomatic acceptance (treated as foundational)
- 2) $c = 0$: Complete rejection (treated as impossibility)
- 3) $c = 0.5$: Maximal uncertainty (no evidence either direction)
- 4) $c > 0.5$: Net evidence for acceptance
- 5) $c < 0.5$: Net evidence for rejection

Definition Epistemic Commitment.

Confidence represents **epistemic commitment**: the degree to which an agent commits to a proposition based on available evidence and reasoning, *interpreted* as the calibrated reliability of the source or derivation producing the belief.

Unlike probability:

- 1) **No normalization:** $\text{conf}(\phi) + \text{conf}(\neg\phi)$ need not equal 1.
- 2) **Not frequency:** $\text{conf}(\phi) = 0.9$ does *not* mean

it.body

in a single case. It means

it.body

- 3) **Derivation-based:** Confidence propagates through inference rules, not conditioning.

Example Non-normalized confidence.

Consider the proposition

it.body

An honest assessment might be: $\text{conf}(\text{no vulnerabilities}) = 0.4$ (some evidence from testing) $\text{conf}(\text{has vulnerabilities}) = 0.3$ (some evidence from complexity) The sum $0.4 + 0.3 = 0.7 < 1$ reflects residual uncertainty—neither hypothesis is well-supported. This is inexpressible in probability.

d) *Falsifiability Criteria:*

Following the review's requirement, we state explicitly what would falsify CLAIR's confidence model:

- 1) **Empirical falsification:** If sources systematically assign confidence c to propositions but are correct with frequency $f \neq c$, they are *uncalibrated*. CLAIR provides no mechanism to *detect* uncalibrated sources—this requires external validation.
- 2) **Semantic falsification:** If interpretation as

it.body

leads to contradictions in the algebra (e.g., violations of monoid laws), the semantics would be inadequate. The Lean 4 formalization verifies algebraic consistency.

- 3) **Competing semantics:** If an alternative interpretation (e.g., pure probability or fuzzy truth) provides better empirical calibration or algebraic properties, CLAIR's semantics would need revision.

e) *Comparison with Subjective Logic:*

Jøsang's Subjective Logic (Jøsang, 2016) extends probability with explicit uncertainty. An *opinion* is a tuple $\omega = (b, d, u, a)$:

- 1) b : belief mass (evidence for)
- 2) d : disbelief mass (evidence against)
- 3) u : uncertainty mass (lack of evidence)
- 4) a : base rate (prior probability)

with constraint $b + d + u = 1$.

CLAIR's confidence can be viewed as a simplification of Subjective Logic: $c = b + u \text{ times } a$ where we collapse uncertainty into the confidence value via the base rate. This loses the $b/d/u$ decomposition but gains simplicity. The trade-off is appropriate for CLAIR's focus on derivation tracking rather than uncertainty quantification.

it.body

B. The Aggregation Monoid

When multiple independent pieces of evidence support the same conclusion, confidence should *increase*. This requires the probabilistic OR operation.

a) *Probabilistic OR (Oplus):*

Definition Probabilistic OR (Oplus).

For confidence values a, b in C , their **aggregation** is: $a \oplus b = a + b - a \text{ times } b$

The formula has several equivalent forms:

- 1) $a \oplus b = a + b(1 - a)$
- 2) $a \oplus b = a(1 - b) + b$
- 3) $a \oplus b = 1 - (1 - a)(1 - b)$ (De Morgan duality with multiplication)

The last form reveals the duality: $a \oplus b$ is the complement of the product of complements.

Theorem Oplus Preserves Bounds.

For all a, b in C : $a \oplus b$ in C Proof. **Lower bound:** $a \oplus b = a + b(1 - a) \geq 0$ since $a \geq 0$, $b \geq 0$, and $(1 - a) \geq 0$.

Upper bound: $a \oplus b = a + b(1 - a) \leq 1$ since $b \leq 1$ implies $b(1 - a) \leq 1 - a$, therefore $a + b(1 - a) \leq a + (1 - a) = 1$. ■

Theorem Oplus Monoid.

$(C, \oplus, 0)$ is a commutative monoid with absorbing element 1:

- 1) **Associativity:** $(a \oplus b) \oplus c = a \oplus (b \oplus c)$
- 2) **Commutativity:** $a \oplus b = b \oplus a$

- 3) **Identity:** $0 \oplus a = a \oplus 0 = a$

- 4) **Absorption:** $1 \oplus a = a \oplus 1 = 1$

Proof. All properties follow from standard real number arithmetic on $0, 1$. ■

b) *Confidence-Increasing Property:*

Unlike multiplication, *oplus* increases confidence:

Theorem Oplus is Confidence-Increasing.

For all a, b in C : $a \oplus b \geq \max(a, b)$ Proof. Using form $a \oplus b = a + b(1 - a)$. Since $b(1 - a) \geq 0$, we have $a \oplus b \geq a$. By commutativity, $a \oplus b \geq b$. Therefore $a \oplus b \geq \max(a, b)$. ■

Corollary Diminishing Returns.

The marginal gain from additional evidence decreases as confidence grows: $(\text{del}/(\text{del } b))(a \oplus b) = 1 - a$ When a is already high, additional evidence contributes less.

Example Aggregation of Weak Evidence.

Suppose we have ten independent pieces of weak evidence, each with confidence 0.3. The combined confidence is: $0.3 \oplus 0.3 \oplus \dots \oplus 0.3$ (10 times) $= 1 - (1 - 0.3)^{10} = 1 - 0.7^{10} \approx 0.972$ Ten weak independent witnesses produce high combined confidence.

c) *The “Survival of Doubt” Interpretation:*

The formula $a \oplus b = 1 - (1 - a)(1 - b)$ admits a probability-theoretic interpretation under calibrated reliability:

- 1) $(1 - a)$ is the

it.body

in evidence a

- 2) $(1 - a)(1 - b)$ is the probability *both* pieces of evidence fail (assuming independence)
- 3) $a \oplus b$ is the probability *at least one* succeeds
This

it.body

interpretation explains why aggregation increases confidence: more independent evidence means more chances for at least one to be correct.

C. The Multiplication Monoid

When a conclusion is derived from premises, its confidence depends on the premises' confidences. The simplest case is conjunctive derivation: both premises must hold for the conclusion to follow.

a) *Conjunctive Confidence Propagation:*

Definition Confidence Multiplication.

For confidence values a, b in C , their **multiplicative combination** is standard multiplication: $a \text{ times } b$

This models the intuition that deriving C from A and B requires both to be true. If we are 90% confident in A and

80% confident in B , our confidence in C (derived from both) is at most $0.9 \times 0.8 = 0.72$.

Theorem Multiplication Preserves Bounds.

For all a, b in C : a times b in C Proof. We prove both bounds:

- 1) a times $b \geq 0$: Since $a \geq 0$ and $b \geq 0$, their product is non-negative.
- 2) a times $b \leq 1$: Since $b \leq 1$, we have a times $b \leq a$ times $1 = a \leq 1$. ■

An undercut

it.body

doesn't claim the object isn't red; it undermines the inference from appearance to reality.

If $\text{conf}(\text{looks red} \Rightarrow \text{is red}) = 0.9$ and $\text{conf}(\text{red lighting}) = 0.6$, then: $\text{undercut}(0.9, 0.6) = 0.9 \times (1 - 0.6) = 0.9 \times 0.4 = 0.36$ The inference confidence drops from 0.9 to 0.36.

Theorem Multiplication Monoid.

$(C, \text{times}, 1)$ is a commutative monoid with absorbing element 0:

- 1) **Associativity**: $(a \text{ times } b) \text{ times } c = a \text{ times } (b \text{ times } c)$
- 2) **Commutativity**: $a \text{ times } b = b \text{ times } a$
- 3) **Identity**: $1 \text{ times } a = a \text{ times } 1 = a$
- 4) **Absorption**: $0 \text{ times } a = a \text{ times } 0 = 0$

Proof. All properties follow from standard real number arithmetic on 0,1. ■

b) The Derivation Monotonicity Principle:

A fundamental property of CLAIR is that derivation can only decrease confidence:

Theorem Derivation Monotonicity.

For all a, b in C : a times $b \leq \min(a, b)$ In particular, a times $b \leq a$ and a times $b \leq b$. Proof. Since $b \leq 1$, we have a times $b \leq a$ times $1 = a$. By commutativity, a times $b \leq b$. Therefore a times $b \leq \min(a, b)$. ■

Corollary No Confidence Amplification.

No sequence of conjunctive derivations can increase confidence. If c_0 is the confidence of a foundational belief and c_n is derived through n multiplicative steps, then $c_n \leq c_0$.

This principle is essential for CLAIR's epistemology: derived beliefs are never more confident than their sources. Certainty ($c = 1$) is reserved for axioms, not conclusions.

D. Defeat Operations

Beyond derivation and aggregation, CLAIR requires operations for *defeat*: when evidence undermines a belief.

a) Undercut: Attacking the Inference:

Following Pollock (2001), an *undercutting defeater* attacks the inferential link, not the conclusion directly.

Definition Undercut.

For confidence c in a conclusion and defeat strength d : $\text{undercut}(c, d) = c \text{ times } (1 - d)$

Example Undercutting Defeat.

Consider the inference:

it.body

Theorem Undercut Preserves Bounds.

For all c, d in C : $\text{undercut}(c, d)$ in C Proof. Since $d \leq 1$, we have $(1 - d) \geq 0$. Since $c \geq 0$, we have c times $(1 - d) \geq 0$. Since $c \leq 1$ and $(1 - d) \leq 1$, we have c times $(1 - d) \leq 1$. ■

Theorem Undercut Composition.

Sequential undercuts compose via *oplus*: $\text{undercut}(\text{undercut}(c, d_1), d_2) = \text{undercut}(c, d_1 \text{ oplus } d_2)$ Proof. By expanding the definition: $\text{undercut}(\text{undercut}(c, d_1), d_2) = c(1 - d_1)(1 - d_2)$

III. $c(1 - (D_1 + D_2 - D_1 \text{ TIMES } D_2))$

IV. $c(1 - (D_1 \text{ OPLUS } D_2))$

V. $\text{UNDERCUT}(C, D_1 \text{ OPLUS } D_2)$

This beautiful result shows that defeat strengths aggregate via *oplus*: multiple undercuts combine as if their doubts aggregated.

a) Rebut: Competing Evidence:

A *rebutting defeater* provides counter-evidence for the conclusion directly.

Definition Rebut.

For confidence c_{for} in favor and c_{against} against: $\text{rebut}(c_{\text{for}}, c_{\text{against}}) = \text{cases}((c_{\text{for}} / (c_{\text{for}} + c_{\text{against}}), \text{if } c_{\text{for}} + c_{\text{against}} > 0), (0.5, \text{if } c_{\text{for}} = c_{\text{against}} = 0))$

The formula treats evidence symmetrically: the resulting confidence is the

it.body

of supporting evidence.

Theorem Rebut Preserves Bounds.

For all $c_{\text{for}}, c_{\text{against}}$ in C : $\text{rebut}(c_{\text{for}}, c_{\text{against}})$ in C Proof. If $c_{\text{for}} + c_{\text{against}} = 0$, the result is 0.5 in $[0,1]$. Otherwise:

- 1) $\text{rebut} \geq 0$ because $c_{\text{for}} \geq 0$ and the denominator is positive.
- 2) $\text{rebut} \leq 1$ because $c_{\text{for}} \leq c_{\text{for}} + c_{\text{against}}$.

Theorem Rebut Antisymmetry.

$\text{rebut}(a, b) + \text{rebut}(b, a) = 1$ Proof. When $a + b > 0$:
 $\text{rebut}(a, b) + \text{rebut}(b, a) = a/(a+b) + b/(a+b) = (a+b)/(a+b) = 1$ When $a = b = 0$: $\text{rebut}(a,b) = \text{rebut}(b,a) = 0.5$, so the sum is 1. ■

b) Rebut Normalization Limitation:

The rebut operation has an important limitation that must be stated explicitly: it *normalizes away absolute strength information*.

Definition Rebut Normalization Property. For all $k > 0$:
 $\text{rebut}(k \text{ times } a, k \text{ times } b) = \text{rebut}(a, b)$

Proof. When a, b are not both zero: $\text{rebut}(k a, k b) = (k a) / (k a + k b) = k a / (k(a+b)) = a/(a+b) = \text{rebut}(a,b)$ The k cancels from numerator and denominator. ■

This means rebut only captures the *relative balance* of evidence, not the *absolute magnitude*. Consider:

Example Normalization Loses Absolute Strength.

Two scenarios with very different absolute evidence strengths:

- 1) **Scenario A:** $c_{\text{for}} = 0.1, c_{\text{against}} = 0.1$ Result:
 $\text{rebut}(0.1, 0.1) = 0.1/(0.1+0.1) = 0.5$
- 2) **Scenario B:** $c_{\text{for}} = 0.9, c_{\text{against}} = 0.9$ Result:
 $\text{rebut}(0.9, 0.9) = 0.9/(0.9+0.9) = 0.5$

Both scenarios yield confidence 0.5 despite Scenario B having *nine times more* total evidence. The rebut operation cannot distinguish “weak balanced evidence” from “strong balanced evidence.”

Implications for CLAIR.

The normalization limitation means rebut is appropriate for *competitive evaluation* (comparing opposing arguments) but insufficient for *absolute assessment*. When absolute strength matters, CLAIR should:

- 1) Track total evidence magnitude separately: $\text{total} = c_{\text{for}} + c_{\text{against}}$
- 2) Use *undercut* instead when the goal is to reduce confidence proportionally to attack strength
- 3) Apply confidence thresholds to distinguish “weak 0.5” from “strong 0.5”

This is a deliberate design trade-off: rebut prioritizes interpretability (“market share” of evidence) over preserving absolute magnitude.

A. Independence Assumptions for Aggregation

The *oplus* operation assumes independence of evidence sources. This is a critical assumption that must be stated explicitly:

Definition Conditional Independence for Oplus.

Two evidence sources e_1 and e_2 are **independent** with respect to proposition ϕ if: $P(\phi | e_1, e_2) = P(\phi | e_1) + P(\phi | e_2) - P(\phi | e_1) \times P(\phi | e_2)$ Under calibrated reliability, this means the reference classes of e_1 and e_2 do not systematically overlap.

a) When Oplus is Valid:

The *oplus* operation is **semantically justified** when:

- 1) **Independent sources:** Evidence derives from causally independent mechanisms (e.g., different sensors, different witnesses, different reasoning paths).
- 2) **No shared provenance:** The sources do not derive from common antecedents that would cause their errors to correlate.
- 3) **Reference class disjointness:** The calibration reference classes for e_1 and e_2 do not systematically overlap.

Example Valid Independent Aggregation.

Three different image classifiers (trained on different datasets, with different architectures) each assign confidence 0.7 to

it.body

. The combined confidence $0.7 \text{ oplus } 0.7 \text{ oplus } 0.7 \approx 0.973$ is justified because the classifiers make statistically independent errors.

b) When Oplus Breaks:

The *oplus* operation **overcounts** and produces misleading confidence when:

- 1) **Shared provenance:** Two sources derive from the same evidence. For example, two newspapers reporting the same press release should not be aggregated via *oplus*.
- 2) **Common systematic bias:** Sources that share the same misconception or training data flaw will make correlated errors. Aggregating them amplifies bias.
- 3) **Circular dependence:** When e_2 cites e_1 as a source, their evidence is not independent.

Example Invalid Aggregation Due to Shared Provenance.

A fact-checking website cites *Source A*. A blog post then cites the fact-checking website. Treating these as independent evidence would incorrectly inflate confidence via *oplus*.

c) Correlation-Aware Alternatives:

When independence is violated, CLAIR provides alternatives:

- 1) **Correlation-aware aggregation:** Chapter 4 introduces *oplus_delta* for correlated evidence, where δ in $[0,1]$ measures dependency: $a \text{ oplus}_\delta(b) = a + b - a \cdot \delta \cdot b$
- 2) **Min-based aggregation:** When sources may be completely dependent, use $\min(a, b)$ instead of *oplus* to avoid overcounting.
- 3) **Provenance tracking:** Evidence usage is tracked via provenance, preventing the same source from being counted twice in a derivation. This is enforced by CLAIR’s DAG structure and source tracking (Chapter 10).

Independence Detection.

In practice, determining whether evidence sources are independent requires domain knowledge and provenance

tracking. CLAIR does **not** automatically detect dependence—it provides the algebraic machinery for *manually* specifying correlation or preventing double-counting through provenance tracking.

d) *Interval-Based Confidence: Dependency Bounds:*

When dependence between evidence sources is unknown or partially known, CLAIR supports *interval-based confidence* as a robust alternative to point-value aggregation. Instead of committing to a single confidence value, we track upper and lower bounds that reflect uncertainty about dependence.

Definition **Confidence Interval.**

A **confidence interval** is a pair $[c_{\text{low}}, c_{\text{high}}]$ with $0 \leq c_{\text{low}} \leq c_{\text{high}} \leq 1$. The interpretation: the true confidence lies in this interval, with the width $c_{\text{high}} - c_{\text{low}}$ reflecting uncertainty about evidence dependence.

Interval Aggregation Bounds:

For evidence with confidence values a and b , the independence assumption determines the correct aggregation:

- 1) **Independent:** $a \oplus b = a + b - a \times b$ (probabilistic sum)
- 2) **Maximally dependent:** $\max(a, b)$ (no double-counting)
- 3) **Anti-correlated:** Could theoretically reach $\min(a + b, 1)$

Theorem **Interval Aggregation Bound.**

For confidence values a, b in C , define: $\text{oplus_indep}(a, b) = a \oplus b = a + b - a \times b$

Then for any unknown correlation structure: $\text{oplus_dep}(a, b) \leq \text{true_aggregation} \leq \text{oplus_indep}(a, b)$

Proof. The lower bound $\max(a, b)$ occurs when evidence is fully dependent (one source's evidence is a subset of the other's). The upper bound $a \oplus b$ occurs when evidence is independent. Any positive correlation reduces the effective aggregation below $a \oplus b$, while negative correlation is impossible for evidential support (evidence cannot be negatively correlated with the truth of what it supports). ■

Definition **Interval Aggregation.**

For confidence intervals $[a_{\text{low}}, a_{\text{high}}]$ and $[b_{\text{low}}, b_{\text{high}}]$: $[a_{\text{low}}, a_{\text{high}}] \oplus \text{interval} [b_{\text{low}}, b_{\text{high}}] = [\max(a_{\text{low}}, b_{\text{low}}), \min(a_{\text{high}}, b_{\text{high}})]$

The lower bound uses \max (maximal dependence) and the upper bound uses \oplus (independence).

Example **Interval-Based Aggregation for Unknown Dependence.**

Two news sources report the same fact. Each has confidence 0.7, but we suspect possible shared provenance (they may have derived from the same press release).

Point-value approach with \oplus : $0.7 \oplus 0.7 = 0.91$ This assumes independence and may be overconfident.

Interval-based approach: $[0.7, 0.7] \oplus \text{interval} [0.7, 0.7] = [\max(0.7, 0.7), 0.7 \oplus 0.7] = [0.7, 0.91]$

The interval $[0.7, 0.91]$ captures our uncertainty: the true confidence could be as low as 0.7 (if the sources are fully

dependent) or as high as 0.91 (if independent). The width 0.21 quantifies our dependence uncertainty.

Interval Propagation:

Intervals propagate through CLAIR's operations:

- 1) **Multiplication:** $[a_{\text{low}}, a_{\text{high}}] \times [b_{\text{low}}, b_{\text{high}}] = [a_{\text{low}} \times b_{\text{low}}, a_{\text{high}} \times b_{\text{high}}]$
- 2) **Undercut:** $\text{undercut}([a_{\text{low}}, a_{\text{high}}], d) = [a_{\text{low}} \times (1-d), a_{\text{high}} \times (1-d)]$
- 3) **Rebut:** More complex; requires propagating both bounds through the ratio

Theorem **Interval Bound Preservation.**

If $[a_{\text{low}}, a_{\text{high}}]$ and $[b_{\text{low}}, b_{\text{high}}]$ are valid intervals ($0 \leq a_{\text{low}} \leq a_{\text{high}} \leq 1$), then all CLAIR operations produce valid intervals.

Proof. Each operation preserves bounds: \oplus preserves 0,1 (Theorem 3.3), multiplication preserves 0,1 (Theorem 3.10), and undercut preserves 0,1 (Theorem 3.16). Interval extensions apply the same bounds to endpoints. ■

Dependency Bounds from Provenance Tracking:

CLAIR's linear type system (Chapter 10) enables *provenance tracking*, which provides tighter dependency bounds than worst-case intervals.

Definition **Provenance-Aware Dependency Bound.**

Let $S(e)$ be the set of primitive sources (axioms, sensor inputs, LLM outputs) used to derive evidence e . For two evidences e_1, e_2 :

- 1) If $S(e_1)$ and $S(e_2)$ are disjoint (disjoint provenance): independence is *plausible*, use \oplus
- 2) If $S(e_1)$ is a subset of $S(e_2)$ or vice versa (subset provenance): maximal dependence, use \max
- 3) If there is partial overlap in provenance: use oplus_δ with dependency estimated as the fraction of shared sources

This provides a *heuristic but principled* way to estimate correlation from provenance structure. The intuition: shared evidence sources correlate the derivations that depend on them.

Example **Provenance-Aware Aggregation.**

Three sources support a conclusion:

- 1) e_1 depends on sensors s_1, s_2
- 2) e_2 depends on sensors s_2, s_3
- 3) e_3 depends on sensor s_4

Aggregating e_1 and e_2 : They share sensor s_2 , so $\delta \approx 1/2$. Aggregating with e_3 : Disjoint from e_1, e_2 , so independence is plausible.

This is more precise than worst-case intervals and more honest than assuming full independence.

When to Use Intervals vs Point Values:

Use point values when:

- 1) Provenance tracking confirms disjoint sources
- 2) Domain knowledge guarantees independence
- 3) Computational efficiency is critical

Use intervals when:

- 1) Provenance is incomplete or unknown
- 2) Sources may share hidden dependencies
- 3) Conservative decision-making is required (e.g., safety-critical systems)
- 4) Auditing requires explicit uncertainty quantification

Relationship to Imprecise Probability.

Interval-based confidence in CLAIR is related to Walley's imprecise probability and Dempster-Shafer theory, but with key differences:

- 1) **Not a probability interval:** Our intervals bound *calibrated reliability*, not frequency in a reference class.
- 2) **Not belief/plausibility:** We don't track separate belief and plausibility masses as in Dempster-Shafer.
- 3) **Operational:** Intervals quantify uncertainty about *evidence dependence*, not about the proposition itself.

This interval-based approach directly addresses Hole A from the review: when independence assumptions are violated, CLAIR provides principled alternatives that avoid overcounting while preserving auditability.

B. Conclusion

This chapter established the algebraic and semantic foundation of CLAIR:

- 1) **Confidence as calibrated reliability:** We explicitly interpret confidence as the calibrated reliability of information sources and derivation processes. This addresses the review's concern about underspecified semantics.
- 2) **Independence assumptions:** The *oplus* operation requires conditional independence of evidence sources. We provide *oplus_delta* for correlated evidence and affine typing to prevent double-counting.
- 3) **Three monoids, not a semiring:** Multiplication (derivation), and *oplus* (aggregation) serve distinct semantic roles and do not distribute.
- 4) **Defeat operations:** Undercut and rebut formalize how evidence can undermine beliefs, with undercuts composing beautifully via *oplus*.
- 5) **Machine-verified:** The algebra is formalized in Lean 4, ensuring type safety and algebraic correctness.

The confidence system provides the numeric foundation. The next chapter develops the *structural* foundation: how beliefs connect through justification DAGs.

VI. JUSTIFICATION AS LABELED DAGS

"An argument is not a proof. It is a reason for a belief—and reasons can be defeated."

— John L. Pollock, Defeasible Reasoning

This chapter develops the structural foundation of CLAIR: how beliefs connect through justification. We show that trees are inadequate for justification structure and that the correct model is a *directed acyclic graph with labeled edges*. The labels distinguish support from defeat, enabling defeasible reasoning where conclusions can be withdrawn when new evidence undermines their justifications.

A. The Inadequacy of Trees

a) The Shared Premise Problem:

Traditional approaches represent justification as trees: each conclusion has premises, which themselves have premises, forming an inverted tree structure. This model is elegant but insufficient.

Consider a simple computation that uses the same belief twice:

```
let population = belief(1000000, 0.95, source: census)
let sample_size = belief(1000, 0.90, source: survey)
let ratio = derive population, sample_size by dividen
let inverse = derive sample_size, population by dividen
let product = derive ratio, inverse by multiply
```

In a tree representation, each belief appears multiple times as separate subtrees. This creates three problems:

- 1) **Space inefficiency:** Beliefs are copied rather than shared.
- 2) **Invalidation complexity:** If a belief is invalidated, we must find and invalidate all copies.
- 3) **Semantic confusion:** Are these the *same* belief or *different* beliefs that happen to be equal?

The correct representation is a DAG with explicit sharing, where each belief appears exactly once with multiple parents.

Theorem DAG Necessity.

Any justification system that:

- 1) Allows a belief to be used as a premise in multiple derivations
- 2) Propagates invalidation correctly (removing a premise invalidates all conclusions)
- 3) Maintains identity (the "same" belief is the same node) must represent justification as a DAG, not a tree.

Proof. In a tree, each node has exactly one parent. If a belief is used in derivations of two different conclusions, it must appear as a child of both. But in a tree, a node cannot have two parents. Therefore, the belief must be duplicated, violating identity. A DAG allows multiple parents, resolving the contradiction. ■

b) Why Not Cycles?

If we allow sharing, why not allow cycles? Coherentist epistemology suggests beliefs can mutually support each other. We reject cycles in justification for three reasons:

- 1) **Bootstrap problem:** Circular justification allows confidence inflation with no external grounding.
- 2) **Invalidation ambiguity:** If beliefs support each other circularly, invalidation semantics become ill-defined.
- 3) **Well-foundedness:** The justification relation should be well-founded, with no infinite descending chains.

Note: The theory/observation circularity example is better analyzed as two separate relations:

- 1) **Evidential support** (tracked in justification): Observation provides evidence for theory.

- 2) **Interpretive framework** (not part of justification): Theory provides framework for interpreting observation.

B. Labeled Edges for Defeat

The DAG structure addresses sharing but not defeat. When evidence undermines a belief's justification, we need edges that carry negative, not positive, epistemic weight.

a) The Defeat Problem:

A *defeater* is a belief that undermines confidence in another belief. Following Pollock (2001), we distinguish:

- 1) **Undercutters**: Attack the inferential link, not the conclusion directly.
- 2) **Rebutters**: Provide direct counter-evidence against the conclusion.

Definition Edge Types.

A justification edge has one of three types:

- 1) **Support**: Positive evidence for the target.
- 2) **Undercut**: Attacks the inferential link to the target.
- 3) **Rebut**: Direct counter-evidence against the target.

b) Formal Definition:

Definition Justification Graph.

A **justification graph** is a tuple $G = (N, E, r)$ where:

- 1) N is a finite set of *justification nodes*
- 2) $E \subseteq N \times N$ is a set of labeled edges
- 3) $r \in N$ is the root node

subject to the constraint that the underlying unlabeled graph is acyclic.

Definition Justification Node Types.

Each node has one of the following types:

- 1) axiom: Foundational belief (confidence = 1)
- 2) rule

(r, premises)

: Deductive rule application

- 3) assumption

(a)

: Temporary assumption for reasoning

- 4) choice

(options, criteria)

: Decision point

- 5) abduction

(obs, hypotheses, selected)

: Abductive inference

- 6) analogy

(source, similarity)

: Analogical reasoning
7) induction

(cases, rule)

: Inductive generalization
8) aggregate

(sources, combRule)

: Evidence aggregation

c) Well-Formedness Constraints:

A well-formed justification graph must satisfy explicit constraints to ensure semantic coherence and computational tractability.

Definition Acyclicity Constraint.

A justification graph must be acyclic: no path may exist from any node back to itself. This constraint applies specifically to support edges. Defeat edges may form cycles, which are resolved via fixed-point iteration (discussed below).

Definition Well-Formed Justification Graph.

A justification graph is well-formed iff:

- 1) The support structure is acyclic (no support cycles)
- 2) Every non-axiom node has at least one incoming support edge
- 3) Every node is reachable from the root in the underlying undirected graph

Condition 1 ensures the support structure is well-founded. Condition 2 ensures no “floating” beliefs without justification. Condition 3 ensures the graph is connected.

d) DAG-Only vs Cyclic Choice:

CLAIR adopts a hybrid approach to cycles:

- 1) **Support edges: DAG-only**. Cycles in evidential support are semantically prohibited because they enable bootstrapping and violate well-foundedness. The type checker enforces acyclicity for support edges at construction time.
- 2) **Defeat edges: Fixed-point semantics**. Defeat cycles are permitted and resolved via fixed-point iteration. When defeat edges form cycles, confidence propagation requires solving a system of equations.

This design choice reflects the epistemic distinction between evidence for (which must be well-founded) and attacks against (which can mutually undermine each other).

e) Fixed-Point Semantics for Cyclic Defeat:

When defeat edges form cycles, we compute confidences via iterative fixed-point finding.

Theorem Fixed-Point Existence.

For any defeat graph with acyclic underlying support, a fixed point exists.

Proof. The confidence update function maps the unit interval to itself and is continuous. By Brouwer's fixed point theorem, a continuous function from a compact convex set to itself has a fixed point. ■

Theorem Fixed-Point Uniqueness.

If the maximum product of undercut strengths along any cycle is strictly less than 1, then the fixed point is unique and Kleene iteration converges to it.

Proof. Under this condition, the update function is a contraction mapping. By the Banach fixed point theorem, a contraction has a unique fixed point and iteration converges to it. ■

Definition Kleene Iteration.

Starting from initial confidences, repeatedly apply the update function until convergence. The sequence generated by repeated application is the Kleene iteration.

Example Convergence for Typical Defeat Cycles.

Suppose two nodes mutually undercut with strength 0.5 each. The Lipschitz constant is 0.5 times 0.5 equals 0.25. Starting from initial confidence (1, 1):

- 1) After 1 iteration: error is at most 0.25 times initial error
- 2) After 5 iterations: error is at most 0.25^5 approximately 0.001 (0.1%)
- 3) After 10 iterations: error is at most 0.25^{10} approximately 10^{-6} (0.0001%)

Rapid convergence is typical for well-formed defeat graphs.

Practical Implication.

For CLAIR implementations, we recommend:

- 1) Enforce DAG structure for support edges via static type checking
- 2) Allow defeat cycles but limit undercut strengths to ensure contraction
- 3) Use Kleene iteration with convergence threshold 10^{-6}
- 4) Cache fixed-point solutions to avoid recomputation during updates

C. Confidence Propagation

Given a justification graph, we compute the confidence of each node bottom-up.

Definition Support Propagation.

For a node with children having confidences c_1, \dots, c_k :

- 1) a) $\text{conf(axiom)} = 1$
- 2) $\text{conf(rule}(r, \text{children})) = s_r$ times product over children
- 3) $\text{conf(aggregate(children, independent))} = \text{oplus}$ over children

where s_r is the rule strength and oplus is probabilistic OR.

Let c be confidence from support edges, d_1, \dots, d_m be undercut strengths, and r_1, \dots, r_n be rebut strengths. Then: $c' = \text{rebut}(\text{undercut}(c, \text{oplus } d_i), \text{oplus } r_j)$

Theorem Propagation Termination.

The propagation algorithm terminates for any acyclic justification graph.

Proof. The graph is acyclic by definition. Each recursive call moves to a node strictly earlier in topological order. Since the graph is finite, recursion terminates. ■

D. Reinstatement

A fundamental phenomenon in defeasible reasoning is *reinstatement*: when a defeater is itself defeated, the original belief recovers some confidence.

Theorem Compositional Reinstatement.

Let A have base confidence a , undercut by D with confidence d , which is itself undercut by E with confidence e . Then: $\text{conf}(A) = a \text{ times } (1 - d \text{ times } (1 - e))$

The *reinstatement boost* is $\Delta = a \text{ times } d \text{ times } e$.

Proof. Bottom-up evaluation gives: $\text{conf_eff}(D) = d \text{ times } (1 - e)$, then $\text{conf}(A) = a \text{ times } (1 - \text{conf_eff}(D)) = a \text{ times } (1 - d(1-e))$. ■

E. Mutual Defeat

When two arguments defeat each other, we have a defeat cycle. The fixed point analysis yields:

Theorem Mutual Defeat Fixed Point.

If A and B mutually undercut with base confidences a and b , the fixed point is: $a_{\text{star}} = a(1-b) / (1 - ab)$ and $b_{\text{star}} = b(1-a) / (1 - ab)$

Proof. At fixed point: $a_{\text{star}} = a(1 - b_{\text{star}})$ and $b_{\text{star}} = b(1 - a_{\text{star}})$. Solving gives the stated formulas. ■

Theorem Fixed Point Existence.

For any defeat graph with acyclic underlying support, a fixed point exists.

Proof. The confidence update function maps $[0,1]^n$ to itself and is continuous. By Brouwer's fixed point theorem, a fixed point exists. ■

Theorem Uniqueness Condition.

If $b_{\text{max}} \text{ times } d_{\text{max}} < 1$, the fixed point is unique and iteration converges.

Proof. Under this condition, the update function is a contraction mapping. By the Banach fixed point theorem, there is a unique fixed point and iteration converges. ■

Definition Defeat Propagation.

F. Correlated Evidence

The aggregation formula $c_1 \oplus c_2$ assumes independence. When evidence sources are correlated, this overcounts.

Definition Dependency-Adjusted Aggregation.

For evidence with confidences c_1, c_2 and dependency $\delta \in [0,1]$: $\text{aggregate}_\delta(c_1, c_2) = (1-\delta)c_1 \oplus c_2 + \delta \text{ times } (c_1 + c_2)/2$

where $\delta = 0$ means independent and $\delta = 1$ means fully dependent.

G. Connection to Prior Art

a) Truth Maintenance Systems:

JTMS (Doyle 1979) uses IN/OUT lists for dependency-directed backtracking. ATMS (de Kleer 1986) tracks multiple assumption sets.

CLAIR contribution: Generalize TMS to graded confidence with the same dependency-directed architecture.

b) Argumentation Frameworks:

Dung's argumentation (Dung 1995) defines acceptance semantics. Gradual semantics (Amgoud et al. 2017) add numeric degrees.

CLAIR contribution: Integrate argumentation's defeat semantics with type-theoretic justification.

c) Justification Logic:

Artemov's Justification Logic (Artemov 2001) adds explicit justification terms to modal logic.

CLAIR contribution: Extend from tree-like justification terms to DAGs with labeled edges, supporting defeasible reasoning.

H. The Tracking Paradigm

The preceding sections developed the machinery of justification DAGs, confidence propagation, and defeat semantics. We now step back and articulate the underlying *tracking paradigm* that distinguishes CLAIR from traditional formal systems.

a) State Representation:

The Epistemic State:

A CLAIR system's **epistemic state** at any point in time is a pair:

Definition Epistemic state.

An epistemic state $\text{cal}(E)$ is a tuple $\text{cal}(E) = (\text{cal}(G), \text{cal}(B))$ where:

- 1) $\text{cal}(G) = \text{set}(G_1, \dots, G_n)$ is a finite set of justification graphs (one for each belief)
- 2) $\text{cal}(B) = \text{set}((v_i, c_i, j_i, p_i) \mid i \in I)$ is a set of beliefs with values, confidences, justifications, and provenance

Each belief in $\text{cal}(B)$ corresponds to the root of some graph in $\text{cal}(G)$. The graphs may share nodes (shared premises) but are not required to be connected.

Comparison with Proving Paradigm:

Traditional formal systems focus on *provability*:

- 1) **State:** A set of axioms and inference rules

- 2) **Question:** Is ϕ provable from the axioms? ($\Gamma \vdash \phi$)

- 3) **Output:** Yes/No (with proof term)

CLAIR's tracking paradigm focuses on *epistemic representation*:

- 1) **State:** A set of labeled graphs with confidence values

- 2) **Question:** What is the confidence, justification structure, and provenance of belief ϕ ?

- 3) **Output:** (c, G, p) where $c \in [0,1]$ is confidence, G is the justification graph, p is provenance

b) Update Rules:

The epistemic state changes through *epistemic actions*. Each action maps $\text{cal}(E) \rightarrow \text{cal}(E)'$.

Primitive Actions:

Definition Epistemic actions.

The primitive actions on epistemic states are:

- 1) **Add belief:**

$\text{add}(\phi, c, j, p)$

creates a new belief with confidence c , justification j , and provenance p .

- 2) **Aggregate:**

$\text{aggregate}(\phi, \psi, r)$

combines two beliefs about the same proposition using rule r in set("independent", "correlated"(δ), "max").

- 3) **Derive:**

$\text{derive}(\phi, [\psi_1, \dots, \psi_k], \text{rule})$

creates φ as a conclusion from premises using an inference rule.

- 4) **Undercut:**

$\text{undercut}(\phi, \delta, d)$

adds an undercutting defeater to φ with strength d .

- 5) **Rebut:**

$\text{rebut}(\phi, \delta, d)$

adds a rebutting defeater to φ with strength d .

- 6) **Validate:**

$\text{invalidate}(\psi)$

removes belief ψ and propagates invalidation to all beliefs that depend on ψ .

Action Semantics:

Example Adding a belief.

When adding a belief

$\text{add}(\phi, 0.8, \text{source: testimony, witness_A})$

:

- 1) Create a new node n_φ with type

axiom

and confidence 0.8
- 2) Create a single-node graph G_φ containing only n_φ with no edges
- 3) Add $(n_\varphi, 0.8, \text{source: testimony, witness_A})$ to \mathcal{B}
- 4) Add G_φ to \mathcal{G}

An epistemic state $\text{cal}(E)$ is *semantically correct* if:

- 1) Confidence propagation is consistent: Re-computing any belief's confidence yields the same value
- 2) Defeat semantics are satisfied: Undercuts and rebuts are applied according to their definitions
- 3) Bounds are preserved: All confidences remain in $[0,1]$ after any update

This is *internal correctness*: the system behaves according to its own rules. The theorems on propagation termination and sound establish internal correctness for confidence propagation.

Semantic Correctness (External):

Definition Semantic correctness (external).

An epistemic state $\text{cal}(E)$ is *externally correct* with respect to a reference system $\text{cal}(R)$ if:

- derive(phi, [psi_1, psi_2], modus_ponens)
:
1) Create a new node n_φ with type

rule(modus_ponens, [n_psi_1, n_psi_2])
- 2) Create support edges $(n_{\psi_1}, n_\varphi, \text{support})$ and $(n_{\psi_2}, n_\varphi, \text{support})$
- 3) Compute confidence via propagation: $c_\varphi = c_{\psi_1} \times c_{\psi_2}$
- 4) Add graph G_φ to \mathcal{G} and belief $(n_\varphi, c_\varphi, G_\varphi, \text{derived})$ to \mathcal{B}

Definition Semantic correctness (external).

An epistemic state $\text{cal}(E)$ is *externally correct* with respect to a reference system $\text{cal}(R)$ if:

- 1) **Calibration:** For any belief with confidence c , the reference system assigns probability $c' \approx c$ (within calibration tolerance)
- 2) **Justification adequacy:** The justification graph G accurately represents the actual dependency structure in $\text{cal}(R)$
- 3) **Provenance accuracy:** The provenance field correctly identifies the source of the belief in $\text{cal}(R)$

External correctness cannot be established by formal proof alone—it requires empirical validation. CLAIR provides the machinery for tracking (internal correctness) but calibration and accuracy are properties of the *sources*, not the tracking system.

The Honesty Principle.

CLAIR's tracking paradigm embodies epistemic humility:

- 1) The system *reports* its confidence, justification, and provenance
- 2) It does not *guarantee* that these correspond to external reality
- 3) External correctness must be validated empirically (calibration studies, provenance audits)

This distinguishes tracking from proving: a proof claims certainty; a tracked belief admits uncertainty while documenting its grounds.

d) Tracking vs. Proving: A Summary:

Aspect	Proving Paradigm	Tracking Paradigm
Goal	Establish truth	Record epistemic state
State	Set of formulas	Set of labeled graphs
Query	$\Gamma \vdash \varphi?$	What is the confidence, justification, and provenance of φ ?
Output	Proof term (or \perp)	(c, G, p) triple
Update	Add axiom/formula	Add/undercut/rebut/invalidate
Correctness	Soundness	Internal + External

Example Invalidation propagation.

When

invalidate(psi)

is called:

- 1) Remove ψ from \mathcal{B}
- 2) Find all beliefs φ such that ψ appears in G_φ 's support graph
- 3) Recursively invalidate each such φ (mark as defeated or recompute without ψ)
- 4) Update confidence values via re-propagation

This is the *dependency-directed backtracking* inherited from TMS systems.

c) Correctness Criteria:

What does it mean for the tracking paradigm to be "correct"? We distinguish three levels of correctness.

Syntactic Correctness:

Definition Syntactic correctness.

An epistemic state $\text{cal}(E)$ is *syntactically correct* if:

- 1) Every graph in $\text{cal}(G)$ is acyclic (support edges)
- 2) Every node has a well-defined type
- 3) Every edge has a valid label (support, undercut, rebut)
- 4) All confidence values are in $[0,1]$

Syntactic correctness is enforced by the type system and checked at runtime. The Lean 4 formalization proves that propagation preserves syntactic correctness.

Semantic Correctness (Internal):

Definition Semantic correctness (internal).

Aspect	Proving Paradigm	Tracking Paradigm
Limits	Hidden or denied	Explicit with confidence

e) *Practical Implications:*

The tracking paradigm has several practical consequences for CLAIR as an AI reasoning intermediate representation:

Explainability:

Every belief carries its justification graph. Queries like “why does the system believe φ ?” can be answered by traversing the graph and presenting the dependency chain.

Debugging:

When a belief has unexpectedly low confidence, the graph reveals which defeaters are responsible. When confidence is too high, the graph shows whether aggregation rules were misapplied.

Revision:

New evidence is incorporated by adding nodes and edges. The system automatically recomputes confidences via propagation, handling reinstatement without special cases.

Uncertainty:

Unlike binary logic, CLAIR tracks degrees of belief. This matches the uncertainty inherent in real-world reasoning and LLM outputs.

I. Conclusion

This chapter established the structural foundation of CLAIR and articulated the tracking paradigm:

- 1) **DAGs, not trees:** Shared premises require graph structure; explicit sharing enables correct invalidation.
- 2) **Acyclic:** Cycles in evidential support violate well-foundedness; defeat cycles handled via fixed-point semantics.
- 3) **Labeled edges:** Support, undercut, and rebut serve different epistemic roles.
- 4) **Compositional reinstatement:** Defeaters being defeated automatically recovers confidence.
- 5) **Correlated evidence:** Independence assumptions must be explicit; dependency adjustment prevents overcounting.
- 6) **Tracking paradigm:** CLAIR represents epistemic state as labeled graphs with confidence values, updating via add/derive/undercut/rebut/invalidate operations. Correctness has three levels: syntactic (well-formedness), internal semantic (consistency with propagation rules), and external semantic (calibration to reality).

The justification DAG provides the structural substrate for CLAIR’s beliefs. The tracking paradigm formalizes what it means to “track not prove”—a fundamental shift from establishing truth to documenting epistemic grounds with explicit confidence.

- 1) **DAGs, not trees:** Shared premises require graph structure; explicit sharing enables correct invalidation.
- 2) **Acyclic:** Cycles in evidential support violate well-foundedness; defeat cycles handled via fixed-point semantics.

- 3) **Labeled edges:** Support, undercut, and rebut serve different epistemic roles.
- 4) **Compositional reinstatement:** Defeaters being defeated automatically recovers confidence.
- 5) **Correlated evidence:** Independence assumptions must be explicit; dependency adjustment prevents overcounting.

The justification DAG provides the structural substrate for CLAIR’s beliefs. The next chapter addresses a subtler challenge: how beliefs can safely refer to themselves.



Self-Reference and the Gödelian Limits

line(length: 20%, stroke: 1.5pt, paint: academic-burgundy)

"If a system is consistent, it cannot prove its own consistency."

— Kurt Gödel, On Formally Undecidable Propositions

CLAIR allows beliefs about beliefs. This reflexive capacity creates potential for self-reference: a belief that refers to itself, either directly or through a chain of intermediate beliefs. Such self-reference is both a powerful expressive tool and a source of potential paradox. This chapter develops the theoretical foundations for distinguishing safe from dangerous self-reference, culminating in a novel extension of provability logic to graded confidence.

J. The Problem of Self-Reference

a) Direct Self-Reference:

Consider a belief that directly references itself:

Example Direct Self-Reference.

```
-- A belief referencing its own confidence
let b : Belief<Bool> = belief(
    value: confidence(b) > 0.5,
    confidence: ???  
)
```

What confidence should *b* have? If we assign confidence *c* > 0.5, the content becomes true, which seems consistent. But if we assign *c* <= 0.5, the content becomes false—yet what prevents us from assigning *c* = 0.9 anyway?

This is not merely a curiosity. If CLAIR aims to capture how an LLM reasons, and if introspection is part of reasoning, then CLAIR must account for self-referential beliefs—even if that account restricts or forbids certain patterns.

b) Why Self-Reference Matters:

Self-reference enables powerful epistemic capabilities:

- 1) **Calibration:** “My confidence estimates are typically accurate”
- 2) **Uncertainty tracking:** “I am uncertain about this belief”
- 3) **Meta-reasoning:** “I should reconsider beliefs derived from unreliable sources”
- 4) **Self-improvement:** “My reasoning process could be improved in specific ways”

But self-reference also enables paradoxes:

- 1) **Liar-like:** “This belief has confidence 0” (no consistent assignment)
- 2) **Curry-like:** “If this belief is true, then arbitrary proposition *P*” (proves anything)
- 3) **Löbian:** “If I believe *P*, then *P* is true” (circular self-validation)

The challenge is to permit the former while blocking the latter.

K. Löb's Theorem and Anti-Bootstrapping

a) The Classical Result:

Löb's theorem is a cornerstone of provability logic:

Theorem Löb's Theorem.

In any sufficiently strong formal system *T* containing arithmetic: $\text{prov}_{T(\text{prov}(\text{prov}(P) \rightarrow P))} \rightarrow \text{prov}(P)$ where *prov* denotes provability in *T*.

In words: if a system can prove “if *P* is provable, then *P* is true,” then the system can prove *P*. This has a startling consequence.

Corollary No Internal Soundness Proof.

No consistent system can prove its own soundness, i.e., cannot prove $\forall P, \square(P) \rightarrow P$.

Proof. Suppose system *T* proved $\forall P, \square(P) \rightarrow P$. Instantiating with *P* = “false” (falsity), we get $\text{prov}(\text{false}) \rightarrow \text{false}$. Combining with consistency ($\neg \text{prov}(\text{false})$), we can derive $\text{prov}(\text{false})$, contradicting consistency. ■

b) Application to CLAIR:

For CLAIR, interpret $\square(P)$ as “CLAIR believes *P* with confidence 1.0.” Then Löb's theorem constrains self-soundness beliefs:

```
-- A claimed self-soundness belief
let soundness = belief(
    value: forall P. (belief(P, c, ...) and c > 0.9) -> P is true,
    confidence: 0.95  
)
```

By Löb's theorem, if CLAIR can form this belief with high confidence, then (classically) CLAIR believes everything with high confidence—a collapse to triviality. This is the

bootstrapping trap: self-soundness claims cannot increase epistemic authority.

Definition Anti-Bootstrapping Principle.

A belief system satisfies *anti-bootstrapping* if no belief of the form “my beliefs are sound” can increase confidence in any derived belief beyond what the original evidence supports.

Löb’s theorem mathematically enforces anti-bootstrapping for classical provability. The question is how this extends to graded confidence.

L. Tarski’s Hierarchy: Stratified Introspection

a) The Classical Solution:

Tarski’s theorem on the undefinability of truth states that no sufficiently expressive language can define its own truth predicate—on pain of the Liar paradox. Tarski’s solution is stratification:

Level	Can Express	Cannot Express
Level 0 (object)	Facts about the world	Truth of any sentence
Level 1 (meta)	X_0 is true for level-0 X	Truth of level-1 sentences
Level 2 (meta-meta)	X_1 is true for level-1 X	Truth of level-2 sentences
...

Each level can discuss truth at lower levels but never its own level.

b) Stratified Beliefs in CLAIR:

We apply this to beliefs:

Definition Stratified Belief Type.

$\text{Bel}(n, A)$ for $n \in \mathcal{NN}$ where level- n beliefs may reference level- m beliefs only if $m < n$.

Example Stratified Beliefs in CLAIR.

In the CLAIR IR format, each belief has an explicit level field:

```
; Level 0: beliefs about the world (no introspection)
b1 .9 L0 @self "user is authenticated"

; Level 1: beliefs about level-0 beliefs
b2 .95 L1 @self <b1 "my auth belief b1 has high
confidence"

; Level 2: beliefs about level-1 beliefs
b3 .9 L2 @self <b2 "my introspection in b2 seems
accurate"
```

The level constraint is: a belief at level n can only justify beliefs at level $\geq n$, and a belief *about* another belief must be at a higher level.

Theorem Stratification Safety.

If all beliefs respect the stratification constraint— $\text{Bel}(n, A)$ references only $\text{Bel}(m, B)$ with $m < n$ —then no Liar-like paradox can arise.

Proof. Any reference chain from a belief b must strictly decrease in level. Since \mathbb{N} has no infinite descending chains, every chain terminates at level 0. Level-0 beliefs contain no belief references, so they cannot participate in self-referential loops. Therefore, no belief can reference itself directly or transitively. ■

c) What Stratification Rules Out:

Stratification prohibits:

- 1) **Direct self-reference**: A belief cannot mention itself (would require level $n < n$).
- 2) **Universal introspection**: “All my beliefs are...” spans all levels and cannot be expressed at any finite level.
- 3) **Self-soundness at a single level**: “My level- n beliefs are sound” would require level $n + 1$ to express.

d) The Cost of Safety:

Stratification is safe but restrictive. Some legitimate self-referential reasoning is blocked:

- ; Legitimate but blocked: calibration beliefs
- ; This would be self-referential (talks about own confidences)
- ; but is intuitively safe (no paradox)
- b1 .8 L? @self "my confidence estimates match empirical accuracy"
- ; What level should this be? It references "all my confidences"
- ; which spans all levels---cannot be expressed at any finite level

This motivates a more permissive approach for certain cases.

M. Kripke’s Fixed Points: Safe Self-Reference

a) The Fixed-Point Construction:

Kripke proposed an alternative to stratification: allow self-reference but let some sentences remain *undefined*. The key insight is that certain self-referential constructs have *fixed points*—consistent confidence assignments—while others do not.

Definition Fixed Point for Self-Referential Belief.

A self-referential belief b with confidence function $f : \text{cal}("D")\text{cal}("D")[0,1] \rightarrow \text{cal}("D")\text{cal}("D")[0,1]$ (determining confidence from the assumed truth value) has a fixed point if there exists c in $\text{cal}("D")\text{cal}("D")[0,1]$ such that: $c = f(c)$

Example Truth-Teller: Multiple Fixed Points.

Consider:

```
let tt = self_ref_belief(fun self =>
  content: "this belief is true",
  compute_confidence: if val(self.content) then 1.0 else
```

```
0.0
)
```

If confidence is 1.0: content is true, so confidence should be 1.0. ✓ If confidence is 0.0: content is false, so confidence should be 0.0. ✓

Both are fixed points. The belief is *underdetermined*.

Example Liar: No Fixed Point.

Consider:

```
let liar = self_ref_belief(fun self =>
  content: "this belief has confidence 0",
  compute_confidence: if val(self.content) then 1.0 else
  0.0
)
```

If confidence is 1.0: content says “confidence 0,” which is false, so confidence should be 0.0. Contradiction. If confidence is 0.0: content says “confidence 0,” which is true, so confidence should be 1.0. Contradiction.

No fixed point exists. The belief is *ill-formed*.

Example Grounded Self-Reference: Unique Fixed Point.

Consider:

```
let careful = self_ref_belief(fun self =>
  content: "confidence(self) is in [0.4, 0.6]",
  compute_confidence: 0.5
)
```

The compute function is constant, so $f(c) = 0.5$ for all c . The fixed point is $c = 0.5$, which indeed satisfies $0.5 \in [0.4, 0.6]$. This belief is *well-formed* with unique confidence 0.5.

b) The Self-Reference Escape Hatch:

CLAIR provides a controlled mechanism for self-reference: When a CLAIR trace analyzer encounters self-referential beliefs that escape finite stratification, it classifies them:

Self-Reference Analysis Result:

- WellFormed: unique fixed point exists (safe)
- IllFormed: no fixed point (Liar-like, Curry-like, or Löbian trap)
- Underdetermined: multiple fixed points (policy choice needed)

Error Types:

- NoFixedPoint: Liar-like paradox
- CurryLike: proves anything
- LobianTrap: self-soundness claim
- Timeout: analysis did not terminate

Beliefs that escape finite stratification exist “outside” all levels and require special analysis.

c) Classification of Self-Reference:

Combining Tarski and Kripke, we classify self-referential constructs:

Category	Fixed Points	Status	Example
Grounded	Unique	Safe	Calibration beliefs
Underdetermined	Multiple	Policy choice	Truth-teller
Liar-like	None	Ill-formed	“Confidence is 0”
Curry-like	—	Banned	“If true, then P ”
Löbian	—	Banned	Self-soundness

Definition Safe Self-Reference.

A self-referential belief is *safe* if it either:

- 1) Respects stratification (level- n references only level- $m < n$), or
- 2) Has a unique fixed point (Kripke), or
- 3) Has multiple fixed points with a deterministic policy for selection.

Definition Dangerous Self-Reference.

A self-referential belief is *dangerous* if it:

- 1) Has no fixed point (Liar-like), or
- 2) Matches a Curry pattern (“if this then P ”), or
- 3) Claims self-soundness (Löbian trap).

N. Provability Logic and CLAIR

a) Gödel-Löb Logic (GL):

To formally characterize CLAIR’s belief logic, we turn to *provability logic*. The standard modal logic of provability is GL (Gödel-Löb):

Definition GL Syntax.

$\varphi ::= p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \varphi \vee \varphi \mid \varphi \rightarrow \varphi \mid \Box\varphi$ where $\Box\varphi$ means φ is provable.

Definition GL Axioms.

- 1) **K (Distribution):** $\Box(\varphi \rightarrow \psi) \rightarrow (\Box\varphi \rightarrow \Box\psi)$
- 2) **4 (Positive Introspection):** $\Box\varphi \rightarrow \Box\Box\varphi$
- 3) **L (Löb):** $\Box(\Box\varphi \rightarrow \varphi) \rightarrow \Box\varphi$

Critically, GL lacks the truth axiom $\Box\varphi \rightarrow \varphi$ (T). This is philosophically essential: provability does not imply truth. A consistent system can prove false statements if its axioms are wrong.

b) GL vs Other Modal Logics:

Logic	K	T	4	5 or L
K	✓			
T	✓	✓		

Logic	K	T	4	5 or L
S4	✓	✓	✓	
S5	✓	✓	✓	5
GL	✓		✓	L
CLAIR	✓		✓	L

CLAIR aligns with GL:

- 1) **K holds:** If CLAIR believes an implication and believes the antecedent, it can derive the consequent.
- 2) **T fails:** CLAIR's beliefs can be wrong (fallibilism).
- 3) **4 holds:** CLAIR can have meta-beliefs about its beliefs.
- 4) **L must hold:** Self-soundness claims cannot bootstrap confidence.

c) *Solovay's Completeness:*

Theorem Solovay Completeness.

GL is sound and complete with respect to:

- 1) Arithmetic provability: $GL \vdash -\varphi$ iff φ holds under all interpretations of \Box as Gödel provability in PA.
- 2) Finite transitive irreflexive Kripke frames.

The completeness for finite frames yields:

Corollary GL Decidability.

GL is decidable (PSPACE-complete).

This is crucial: classical provability logic is computationally tractable.

O. Confidence-Bounded Provability Logic (CPL)

Classical GL uses binary truth: propositions are either provable or not. CLAIR needs a *graded* version where beliefs carry confidence values in $\text{cal}(\text{"D"})[\text{cal}(\text{"D"})[0,1]]$. This section introduces CPL (Confidence-Bounded Provability Logic), a novel extension of GL designed for CLAIR.

a) *The Literature Gap:*

Extensive work exists on fuzzy modal logics and graded epistemic logic. However, no prior work addresses:

- 1) Graded versions of the Löb axiom
 - 2) The interaction of continuous confidence with provability constraints
 - 3) Anti-bootstrapping in the context of graded belief
- CPL fills this gap.

b) *CPL Syntax:*

Definition CPL Syntax.

$\varphi ::= p \mid \neg\varphi \mid \varphi \wedge \varphi \mid \varphi \vee \varphi \mid \varphi \rightarrow \varphi \mid \Box_c \varphi$ where $\Box_c \varphi$ means φ is believed with confidence at least c .

c) *CPL Semantics:*

Definition Graded Kripke Frame.

A *graded Kripke frame* is a tuple (W, R) where:

- 1) W is a non-empty set of worlds
 - 2) $R : W \times W \rightarrow \mathcal{D}[0,1]$ is a graded accessibility relation
- satisfying:

- 1) **Transitivity:** $R(w, v) \cdot R(v, u) \leq R(w, u)$
- 2) **Converse well-foundedness:** No infinite sequence w_0, w_1, w_2, \dots with $R(w_{i+1}, w_i) > 0$ for all i

Definition Graded Valuation.

A *graded valuation* on a frame (W, R) assigns to each world w and proposition p a confidence value $V_{w(p)} \in \mathcal{D}[0, 1]$. Extended to formulas:

- 1) $V_{w(\neg\varphi)} = 1 - V_{w(\varphi)}$
- 2) $V_{w(\varphi \wedge \psi)} = V_{w(\varphi)} \cdot V_{w(\psi)}$
- 3) $V_{w(\varphi \vee \psi)} = V_{w(\varphi)} + V_{w(\psi)} - V_{w(\varphi)} V_{w(\psi)}$
- 4) $V_{w(\varphi \rightarrow \psi)} = \sup\{c \in \mathcal{D}[0, 1] : V_{w(\varphi)} \cdot c \leq V_{w(\psi)}\}$
- 5) $V_{w(\Box_c \varphi)} = \inf_{v: R(w, v) \geq c} V_{v(\varphi)}$

The last clause says: *phi* is believed at confidence c if *phi* holds in all worlds accessible with strength at least c .

d) *The Graded Löb Axiom: DESIGN AXIOM:*

The crucial innovation in CPL is the graded analogue of Löb's axiom.

Axiom Status Statement.

The Graded Löb axiom is a **DESIGN AXIOM**, not a semantic theorem derived from more basic principles. It is motivated by the requirement of anti-bootstrapping and the need to extend GL to graded confidence while preserving its essential character. The axiom is **posited** as part of CPL's definition, not **proved** from the semantics.

The key question is: Is CPL consistent? We address this below by exhibiting a non-trivial model satisfying all CPL axioms.

Definition Graded Löb Axiom (Design Axiom).

$\Box_c(\Box_c \varphi \rightarrow \varphi) \rightarrow \Box_{g(c)} \varphi$ where $g : \mathcal{D}[0, 1] \rightarrow \mathcal{D}[0, 1]$ is a *discount function* satisfying $g(c) \leq c$.

This is a **DESIGN AXIOM**—not derived from more basic principles but posited as part of CPL's definition. The axiom is motivated by anti-bootstrapping requirements.

The function g captures the *cost* of self-soundness claims. If you believe at confidence c that “believing *phi* at c implies *phi*,” you can derive *phi* only at the discounted confidence $g(c)$.

e) *Choosing the Discount Function:*

We require g to satisfy:

- 1) **Boundedness:** $g : \mathcal{D}[0, 1] \rightarrow \mathcal{D}[0, 1]$
- 2) **Non-amplification:** $g(c) \leq c$ for all c
- 3) **Monotonicity:** $c_1 \leq c_2 \Rightarrow g(c_1) \leq g(c_2)$
- 4) **Anchoring:** $g(0) = 0$ and $g(1) = 1$
- 5) **Non-triviality:** $g(c) < c$ for $c \in (0, 1)$

After analyzing several candidates (identity, parabolic, constant offset, product), we recommend:

Definition Quadratic Discount.

$$g(c) = c^2$$

Theorem Quadratic Discount Properties.

The quadratic discount $g(c) = c^2$ satisfies all desiderata and:

- 1) Aligns with CLAIR's multiplicative confidence algebra ($c^2 = c \times c$)
- 2) Has intuitive meaning: self-soundness costs "deriving the claim twice"
- 3) Produces strong anti-bootstrapping: iterated application $c \rightarrow c^2 \rightarrow c^4 \rightarrow \dots \rightarrow 0$

Proof. Boundedness and anchoring are immediate ($c \in \mathcal{D}[0, 1] \Rightarrow c^2 \in \mathcal{D}[0, 1]$, $0^2 = 0$, $1^2 = 1$). For non-amplification: $c^2 \leq c$ when $c \leq 1$, with equality only at 0 and 1. Monotonicity: $c_1 \leq c_2 \Rightarrow c_1^2 \leq c_2^2$ on $\mathcal{D}[0, 1]$. Non-triviality: $c^2 < c$ for $c \in (0, 1)$. ■

f) *The Anti-Bootstrapping Theorem:*

Theorem Anti-Bootstrapping.

In CPL with $g(c) = c^2$: $\text{conf}(\square_{c(\square_c \varphi \rightarrow \varphi)}) = c \Rightarrow \text{conf}(\varphi) \leq c^2 < c$. Consequently, no finite chain of self-soundness claims can increase confidence beyond the initial level.

Proof. Applying the Graded Löb axiom to the hypothesis yields $\text{conf}(\square_{c^2} \varphi) \leq c^2$. Iterating: $c \rightarrow c^2 \rightarrow c^4 \rightarrow c^8 \rightarrow \dots$. For any $c < 1$, this sequence converges to 0. Self-soundness claims can only decrease confidence. ■

This is the mathematical formalization of anti-bootstrapping: claiming your own soundness provides no epistemic free lunch.

g) *Modal Axioms in CPL:*

We now explicitly list the modal axioms CPL adopts and their status:

Definition CPL Modal Axiom Status.

- 1) **K (Distribution):** $\square_{c(\varphi \rightarrow \psi)} \rightarrow (\square_c \varphi \rightarrow \square_c \psi)$ — **VALID** Derivable from the semantics via graded accessibility.
- 2) **4 (Positive Introspection):** $\square_c \varphi \rightarrow \square_c \square_c \varphi$ — **VALID** Follows from transitivity of the accessibility relation.
- 3) **GL/Graded Löb:** $\square_{c(\square_c \varphi \rightarrow \varphi)} \rightarrow \square_{g(c)} \varphi$ — **DESIGN AXIOM** Posited as part of CPL; motivated by anti-bootstrapping requirements.
- 4) **T (Reflexivity/Truth):** $\square_c \varphi \rightarrow \varphi$ — **INVALID** Explicitly **rejected** in CPL. Provability/belief does not imply truth. This is essential for fallibilism.

h) *CPL Consistency:*

To establish CPL's consistency, we exhibit a non-trivial model:

Theorem CPL Consistency.

CPL is consistent. There exists a non-trivial graded Kripke model satisfying all CPL axioms including the Graded Löb axiom.

Proof. Proof Sketch.

Consider the frame (W, R) where:

- 1) $W = \mathbb{N}$ (the natural numbers)
 - 2) $R(i, j) = 2^{-i}$ if $i < j$, and $R(i, j) = 0$ otherwise
- This frame satisfies:

- 1) **Transitivity:** If $i < j < k$, then $R(i, j) \times R(j, k) = 2^{-i} \times 2^{-j} \leq 2^{-i} = R(i, k)$
- 2) **Converse well-foundedness:** No infinite sequence w_0, w_1, \dots with $R(w_{i+1}, w_i) > 0$, since this would require an infinite decreasing sequence of natural numbers.

Define valuation $V_{w(\varphi)} = 1$ for all propositional variables φ at all worlds w .

For the Graded Löb axiom, consider any world w . If $V_{w(\square_{c(\square_c \varphi \rightarrow \varphi)})} = 1$, then for all v with $R(w, v) \geq c$, we have $V_{v(\square_c \varphi \rightarrow \varphi)} = 1$. By the structure of R , this forces $V_{v(\varphi)} = 1$ for all accessible worlds, yielding $V_{w(\square_{c^2} \varphi)} = 1$.

This model is non-trivial (not all formulas are valid) yet satisfies all CPL axioms, establishing consistency. ■

P. *Decidability of CPL*

Classical GL is decidable. Does CPL inherit this property?

a) *The Vidal Result:*

Theorem Vidal's Undecidability Theorem.

Transitive modal logics over many-valued semantics (including Łukasiewicz and Product algebras) are undecidable, even when restricted to finite models.

CPL has transitivity (axiom 4) and continuous $\mathcal{D}[0, 1]$ values. The Vidal result strongly suggests:

Conjecture CPL Undecidability.

Full CPL (with continuous $\mathcal{D}[0, 1]$ confidence) is undecidable.

Confidence: 0.80

We assign confidence 0.80 to this conjecture based on the close analogy to Vidal's proof technique.

b) *The Role of Converse Well-Foundedness:*

GL's decidability relies on the finite model property: converse well-foundedness forces finite-depth evaluation. Could this rescue CPL?

Proposition Insufficient for Decidability.

Converse well-foundedness alone does not rescue CPL from undecidability.

Justification: Converse well-foundedness constrains *structure* (no infinite ascending chains) but not *values*. The encoding power of continuous $\mathcal{D}[0, 1]$ values combined with transitivity enables undecidable problem encodings even in well-founded frames.

c) *Decidable Fragments:*

Despite the likely undecidability of full CPL, we identify decidable fragments:

CPL-finite: Discrete Confidence:

Restrict confidence to a finite lattice instead of continuous $\mathcal{D}[0, 1]$:

Definition CPL-finite.

Let $L_n = \left\{0, \frac{1}{n-1}, \frac{2}{n-1}, \dots, 1\right\}$. CPL-finite evaluates over L_n with discretized operations:

$$1) \quad a \times b = \lfloor a \times b \rfloor$$

- 2) $a + b = \lceil a + b - a \times b \rceil$
 3) $g_{L(c)} = \lfloor c^2 \rfloor$

For $L_5 = \{0, 0.25, 0.5, 0.75, 1\}$:

c	c^2	$g_{L.5.}(c)$
0	0	0
0.25	0.0625	0
0.5	0.25	0.25
0.75	0.5625	0.5
1	1	1

Theorem CPL-finite Decidability.

CPL-finite is decidable via the finite model property.

Proof Sketch.

By the theorem of Bou, Esteva, and Godo, many-valued modal logics over finite residuated lattices are decidable. CPL-finite evaluates over L_n , a finite lattice. The frame constraints (transitivity, converse well-foundedness) are expressible, and finitely many models of bounded size suffice for completeness.

A complete formal proof would establish: (1) L_n forms a finite residuated lattice under the discretized operations; (2) the frame conditions are expressible in the corresponding modal logic; (3) the finite model property holds; and (4) decidability follows from (3). The proof follows the standard technique for finite-valued modal logics.

Conjecture CPL-finite Complexity.

CPL-finite is PSPACE-complete, analogous to classical GL.

CPL-0: Stratified Only:

Restrict to stratified beliefs without any self-reference:

Definition CPL-0.

CPL-0 disallows nesting of \square operators that would require the Löb axiom. Formally: only formulas of the form $\square_c \varphi$ where φ is box-free.

Theorem CPL-0 Decidability.

CPL-0 is decidable (trivially: the restricted syntax avoids undecidability).

d) *Trade-offs:*

Fragment	Decidable?	Expressiveness	Use Case
Full CPL	Likely no	Full	Theoretical analysis
CPL-finite	Yes	Discrete confidence	Type-level checks
CPL-0	Yes	No self-reference	Stratified beliefs

Q. Alternative: CPL-Gödel

An alternative approach uses Gödel algebra (min/max) instead of product operations:

Definition CPL-Gödel.

- 1) $a \times b = \min(a, b)$
- 2) $a + b = \max(a, b)$

Conjecture CPL-Gödel Decidability.

CPL-Gödel is decidable because Gödel modal logic has the finite model property via quasimodels.

Confidence: 0.75

The conjecture follows from the known decidability of Gödel modal logic, but requires verification that the graded Löb axiom preserves this property.

However, CPL-Gödel is *semantically inappropriate* for CLAIR:

- 1) **max fails aggregation:** $\max(0.6, 0.6) = 0.6$, but two independent pieces of evidence should yield higher confidence (0.84 with $+$).
- 2) **min lacks degradation:** $\min(a, a) = a$, but derivation should cost confidence.
- 3) **No algebraic discount:** The c^2 discount becomes purely frame-based, losing the anti-bootstrapping semantics.

Recommendation.

For CLAIR, use CPL-finite (with product operations), not CPL-Gödel. Accept the discretization rather than sacrifice semantic fidelity.

R. “Conservative Over GL”: Clarification

On “Conservative Over GL” Claims.

The phrase “CPL is conservative over GL” requires careful definition. Two interpretations:

- 1) **Proof-theoretic conservativity:** Every theorem of GL (as formulas with implicit confidence 1) is a theorem of CPL. This **holds**: CPL includes all GL axioms as special cases.
- 2) **Semantic conservativity:** Every model of GL can be embedded in a model of CPL. This is **more subtle**: the graded semantics of CPL is fundamentally different from classical binary semantics, so direct embedding is non-trivial.

We assert the first interpretation: CPL extends GL conservatively in the sense that all classical GL theorems remain valid in CPL when interpreted as high-confidence beliefs. The second interpretation remains an open question.

S. Design Recommendations for CLAIR

a) The Two-Layer Approach:

CLAIR should implement a two-layer approach to self-reference:

- 1) **Default: Stratification.** All beliefs are level-indexed. $\text{Bel}(n, A)$ can only reference $\text{Bel}(m, B)$ with $m < n$. This is safe by construction and requires no runtime analysis.

- 2) **Escape hatch: Kripke fixed points.** For legitimate self-reference (calibration, uncertainty tracking), use

```
self_ref_belief
```

which computes fixed points at construction time. Ill-formed constructs are rejected.

b) *Hard Bans:*

Certain patterns are syntactically rejected:

- 1) **Curry patterns:** “If [self-reference] then [arbitrary P]”
- 2) **Explicit self-soundness:** Claims of the form “All my beliefs are sound”
- 3) **Unrestricted quantification:** “For all beliefs b, \dots ”
These are detected by the parser and rejected before type checking.

c) *Type-Level Anti-Bootstrapping:*

For type-level confidence checks, use CPL-finite with L_5 :

```
-- Finite confidence for compile-time checks
inductive FiniteConfidence where
| zero : FiniteConfidence -- 0
| low : FiniteConfidence -- 0.25
| mid : FiniteConfidence -- 0.5
| high : FiniteConfidence -- 0.75
| one : FiniteConfidence -- 1

def loebDiscount : FiniteConfidence -> FiniteConfidence
| .zero => .zero
| .low => .zero -- 0.25^2 = 0.0625 -> floor to 0
| .mid => .low -- 0.5^2 = 0.25
| .high => .mid -- 0.75^2 = 0.5625 -> floor to 0.5
| .one => .one
```

This provides decidable type-level constraints while preserving the anti-bootstrapping semantics.

T. Related Work

a) Provability Logic:

The foundations of provability logic are in Boolos’s work, with the Solovay completeness theorems establishing the connection to arithmetic. Modern work on GL extensions includes Beklemishev (2004) on polymodal variants.

b) Self-Reference in AI:

Garrabrant et al. (2016) develop logical inductors as an approach to coherent self-reference, though in a different formal framework.

c) Fuzzy Modal Logic:

Fuzzy extensions of modal logic are surveyed in Godo et al. (2003). Decidability results for finite-valued logics appear in Bou et al. (2011). The critical undecidability result for transitive many-valued logics is Vidal (2019).

U. Conclusion

This chapter characterized the landscape of self-reference in CLAIR:

- 1) **Löb’s theorem applies:** Self-soundness claims cannot bootstrap epistemic authority. This is a mathematical fact, not a design choice.

- 2) **Stratification is safe:** Tarski-style level indexing prevents all self-referential paradoxes by construction.
- 3) **Fixed points enable safe self-reference:** Kripke’s approach permits legitimate introspection (calibration, uncertainty tracking) while rejecting ill-formed constructs.
- 4) **CPL extends GL to graded confidence:** The Graded Löb axiom with $g(c) = c^2$ captures anti-bootstrapping for continuous confidence. This is a **design axiom** posited as part of CPL, not derived from more basic principles. CPL is consistent, as demonstrated by the existence of non-trivial models.
- 5) **Full CPL is likely undecidable:** Transitivity plus continuous values enables undecidability (Vidal 2019).
- 6) **CPL-finite is decidable:** Restricting to discrete confidence yields a tractable fragment suitable for type-level checks.
- 7) **Two-layer design:** Stratification by default, Kripke fixed points as escape hatch, hard bans on dangerous patterns.

The Gödelian limits are not obstacles but design constraints. They tell us what epistemic claims are coherent and which collapse into triviality. By respecting these limits, CLAIR achieves honest self-awareness: it can reason about its own reasoning without falling into paradox.

The next chapter turns to epistemological foundations: what grounds CLAIR’s beliefs in the first place.

VI

Grounding

line(length: 20%, stroke: 1.5pt, paint: academic-burgundy)

0

— Susan Haag, *Metaphysical Grounding*

V. The Grounding Problem

A fundamental challenge for any epistemic system is the *grounding problem*: how do beliefs ultimately connect to observable reality? In CLAIR, this problem takes on particular urgency because the system supports arbitrary confidence values and paraconsistent reasoning. Without a principled account of grounding, a CLAIR system could generate high-confidence beliefs that are entirely detached from reality.

a) What grounding means in CLAIR.:

In CLAIR, *grounding* refers to the process by which beliefs acquire their initial epistemic support from sources external to the reasoning system itself. These sources include:

- 1) **Perceptual inputs:** Direct observations from sensors or user input
- 2) **Testimony:** Information received from other agents or sources
- 3) **Logical axioms:** Foundational assumptions accepted without proof
- 4) **Demarcation constraints:** Meta-level restrictions on belief formation

A *grounded belief* is one whose justification graph ultimately traces back to at least one grounding source. An *ungrounded belief* is one that exists only through inference from other beliefs, without any external anchor.

W. Perceptual Grounding

The most direct form of grounding is through perceptual input. CLAIR provides builtin primitives for introducing grounded beliefs:

```
// Ground a proposition with a given confidence  
observe(P, 0.9)
```

```
// Ground from a trusted source with provenance  
testify(P, 0.8, "expert_testimony")
```

These primitives create belief nodes with special Ground justification nodes that cannot be defeated through ordinary inference. The confidence assigned at grounding becomes an upper bound on any downstream inferences—this is the *grounding cap principle*.

a) The grounding cap theorem.:

Theorem If belief b is grounded with confidence c , then for any belief d derived from b through valid CLAIR inference rules, $\text{conf}(d) \leq c$.

This theorem ensures that grounding provides a genuine constraint on reasoning: high confidence can only be achieved through direct grounding, not through inferential amplification.

X. Axiomatic Grounding

Not all beliefs can be grounded perceptually. Mathematical truths, logical principles, and conceptual frameworks require *axiomatic grounding*—the acceptance of certain propositions as foundational.

CLAIR supports this through the axiom primitive:

```
// Accept as a logical axiom  
axiom(forall x. P(x) -> Q(x), 1.0)
```

```
// Accept a conceptual framework assumption  
axiom("induction_principle", 0.95)
```

Axiomatic beliefs are assigned confidence 1.0 by convention, reflecting their status as constitutive of the reasoning framework itself. However, CLAIR also allows for *penumbral axioms*—framework assumptions assigned high but not maximal confidence, acknowledging their potential revisability.

a) The problem of circular axioms.:

A subtle issue arises when axioms are defined in terms of the very concepts they purport to ground. CLAIR addresses this through *stratification*: axioms must be grounded in a lower stratum than the beliefs they support. This prevents circular justification while still allowing for rich, interconnected conceptual frameworks.

Y. Social Grounding and Testimony

Many beliefs are acquired through testimony from other agents. CLAIR models this through *social grounding* primitives that track the provenance of beliefs:

```
// Accept testimony from a source with given reliability  
testify(P, reliability(agent), agent)
```

The confidence in testimony is a function of both the reported confidence and the source's reliability. CLAIR implements this via the *testimony aggregation function*:

$$\text{conf}(\text{testify}(P, c, s)) = c \times \text{reliability}(s)$$

a) Reputation and source tracking.:

CLAIR maintains provenance information for each belief, allowing the system to trace beliefs back to their original sources. This enables *retrospective downgrading*: if a source is later discovered to be unreliable, all beliefs grounded in that source can have their confidence adjusted accordingly.

Z. The Ungrounded: Free-Floating Beliefs

An important design question is whether CLAIR should allow beliefs that are completely ungrounded—beliefs that exist only through their relationships to other ungrounded beliefs.

We permit such beliefs but require them to be marked with a special Untethered status. This serves as a warning to downstream reasoning: these beliefs are *epistemically adrift* and should not form the basis for high-stakes decisions.

a) Creative inference and hypothetical reasoning.:

Despite their epistemic limitations, untethered beliefs serve important functions:

- 1) They enable *hypothetical reasoning*—exploring the consequences of assumptions without committing to their truth
- 2) They support *creative inference*—generating novel hypotheses that can later be tested and grounded
- 3) They provide a space for *conceptual exploration*—developing new frameworks before their empirical connection is established

AA. Grounding Requirements

CLAIR implements three levels of grounding requirements:

a) Tier 1: Strict grounding.:

In safety-critical applications, CLAIR can enforce a strict grounding policy: every belief must have a grounding chain tracing back to perceptual or axiomatic sources. Untethered beliefs are rejected entirely.

b) Tier 2: Demarcated ungrounding.:

For most applications, CLAIR allows untethered beliefs but requires them to be explicitly demarcated. These beliefs cannot form the basis for certain kinds of high-stakes inferences without additional grounding.

c) Tier 3: Permissive ungrounding.:

For exploratory and research applications, CLAIR can operate in permissive mode, allowing arbitrary ungrounded beliefs. This is useful for mathematical exploration and creative hypothesis generation.

AB. Summary

Grounding in CLAIR provides the connection between abstract reasoning and observable reality. By implementing multiple forms of grounding—perceptual, axiomatic, and social—and by enforcing grounding caps that prevent unjustified confidence amplification, CLAIR ensures that beliefs remain epistemically anchored even while supporting sophisticated paraconsistent reasoning.

The grounding architecture acknowledges that not all reasoning needs to be grounded at all times. Hypothetical reasoning, conceptual exploration, and creative inference all benefit from the freedom to form ungrounded beliefs. But

by making grounding status explicit and trackable, CLAIR ensures that agents know when they are operating on solid epistemic footing and when they are engaging in more speculative forms of reasoning.

VII

Belief Revision

line(length: 20%, stroke: 1.5pt, paint: academic-burgundy)

0

— Carlos Alchourrón, Peter Gärdenfors, David Makinson

AC. The Challenge of Revising Structured Beliefs

Classical belief revision theory, formalized in the AGM (Alchourrón-Gärdenfors-Makinson) framework, assumes beliefs are unstructured propositions in a deductively closed belief set. CLAIR's richer structure—graded confidence, DAG-justification, and invalidation conditions—requires a substantial extension of this theory.

This chapter presents *Graded DAG Belief Revision* (GDBR), a framework that extends AGM postulates to handle CLAIR's structured beliefs while preserving the core rationality constraints.

AD. Background: The AGM Framework

The AGM framework characterizes rational belief change through three operations:

- 1) **Expansion:** Adding a new belief P to the belief set K , denoted $K + P$
- 2) **Contraction:** Removing a belief P from the belief set K , denoted $K - P$
- 3) **Revision:** Adding P while maintaining consistency, denoted $K * P$

The AGM postulates impose rationality constraints on these operations. For example, the *success postulate* for revision requires that P is in $K * P$, while the *consistency postulate* requires that $K * P$ is consistent if P is consistent with K .

AE. Why AGM Doesn't Directly Apply

CLAIR's beliefs differ from AGM's in three critical ways:

a) Graded confidence.:

In AGM, beliefs are either held or not held—a binary distinction. In CLAIR, beliefs have graded confidence. This raises new questions: Should revision affect only beliefs with confidence above a threshold? How should conflicting beliefs at different confidence levels be resolved?

b) DAG-structured justification.:

In AGM, beliefs are related only through logical consequence. In CLAIR, beliefs are connected through justification DAGs. Revising one belief may require recomputing the status of beliefs that depend on it, potentially cascading through large portions of the graph.

c) Invalidation conditions.:

In AGM, beliefs are only retracted through explicit contraction operations. In CLAIR, beliefs have *invalidation conditions* that specify circumstances under which they should be reconsidered. When these conditions are triggered, revision happens automatically—a form of *epistemic reflex*.

AF. GDBR: Graded DAG Belief Revision

Our framework extends AGM with three core postulates specific to CLAIR's structure:

a) Confidence preservation.:

When a belief is revised, its new confidence should be a function of:

- The confidence assigned to the new evidence
- The confidence of the defeated belief (if any)
- The strength of the justification links

Formally:

$$\text{conf}(b^*e) = f(\text{conf}(e), \text{conf}(b), \text{strength}(\text{justification}(b)))$$

where f is a monotonic function satisfying $f(1.0, c, s) = \min(1.0, c + s)$.

b) Justification propagation.:

When belief b is revised, all beliefs that justify b must be reconsidered:

$$\text{forall } b'. (b' \text{ justifies } b) \text{ implies } \text{reconsider}(b')$$

This propagation continues transitively through the justification graph, creating a *revision cascade*. To prevent unbounded cascades, CLAIR implements *revision bounds* that limit propagation depth.

c) Invalidation responsiveness.:

When a belief's invalidation condition is triggered, the belief enters a *quarantine state* with reduced confidence. The belief must be explicitly revalidated before its confidence can be restored.

$\text{triggered}(\text{invalidation}(b)) \text{ implies } \text{conf}(b) := \text{conf}(b) \text{ times } \text{penalty}(b)$

The penalty factor depends on the severity of the invalidation trigger and the belief's historical reliability.

AG. The GDBR Algorithm

The revision algorithm proceeds in three phases:

a) Phase 1: Conflict detection.:

When new evidence e arrives, identify all beliefs that conflict with e :

$$\text{conflicts}(e) = \{b \text{ in beliefs mid } b \text{ contradicts } e\}$$

Conflicting beliefs are marked for potential retraction.

b) Phase 2: Confidence comparison.:

For each conflicting belief b , compare confidences:

- 1) If $\text{conf}(e) > \text{conf}(b)$, schedule b for retraction
- 2) If $\text{conf}(e) < \text{conf}(b)$, reject e (with confidence downgrade)
- 3) If $\text{conf}(e) \approx \text{conf}(b)$, enter *arbitration state*

Arbitration invokes additional heuristics: source reliability, justification depth, and recency of evidence.

c) *Phase 3: Graph restructuring.*:

After resolving conflicts, restructure the justification DAG:

- 1) Remove justification links to retracted beliefs
- 2) Recompute confidence for affected beliefs
- 3) Trigger invalidation conditions for beliefs affected by retraction

This phase continues until the graph reaches a *revision-fixed point*.

AH. The Revision Fixed Point Theorem

Theorem The GDBR algorithm always terminates in a revision-fixed point—a state where no further revisions are triggered—provided that:

- 1) The justification graph is finite
- 2) Confidence values are drawn from a finite set
- 3) Invalidation conditions are monotonic (cannot trigger repeatedly)

Proof. Termination follows from the well-foundedness of the revision ordering. Each revision either (a) reduces the confidence of some belief, (b) removes a justification link, or (c) invalidates a belief. Since there are finitely many beliefs, finitely many justification links, and confidence can only decrease finitely many times, the process must terminate.

■ ■

AI. Special Cases and Extensions

a) *Defeater revision*:

Sometimes revision involves not contradicting a belief directly, but defeating its justification. CLAIR handles this through *link revision*: the justification link itself is marked as defeated, and confidence propagates accordingly.

b) *Package revision*:

When a set of beliefs form a tightly connected cluster, it's often more appropriate to revise them as a unit rather than individually. CLAIR supports *package revision* through the *revise-group* primitive:

```
revise-group({b1, b2, b3}, evidence_e)
```

This performs simultaneous revision while maintaining coherence constraints.

c) *Iterated revision*:

In long-running systems, beliefs may be revised multiple times. CLAIR tracks revision history to prevent *revision oscillation*—cycles where beliefs alternate between states. When oscillation is detected, the system enters a *reflective state* and requests external guidance.

AJ. Connection to Argumentation Theory

The GDBR framework connects naturally to formal argumentation theory. Justification DAGs can be viewed as argumentation frameworks, where beliefs are arguments and justification links represent attack/support relationships.

Revision in CLAIR corresponds to *argument change* in argumentation theory: adding new arguments (evidence), removing arguments (retraction), and changing argument

strengths (confidence adjustment). Our framework extends Dung's argumentation semantics to handle graded confidence and dynamic restructuring.

AK. Summary

Belief revision in CLAIR extends the classical AGM framework to handle graded confidence, DAG-structured justification, and invalidation conditions. The GDBR framework preserves AGM's core rationality constraints while providing principled algorithms for revising complex belief structures.

The key innovations are:

- **Confidence-preserving revision**: New evidence doesn't simply override old beliefs; it integrates with existing confidence structures
- **Cascade-limited propagation**: Justification changes propagate but with bounded depth to prevent unbounded revision
- **Invalidation-driven reflex**: Beliefs automatically enter quarantine when their invalidation conditions trigger

This approach enables CLAIR systems to maintain coherent belief states even under dynamic, uncertain conditions while preserving the explanatory power of explicit justification tracking.

VIII

Multi-Agent Epis- temic Reasoning

line(length: 20%, stroke: 1.5pt, paint: academic-burgundy)

0

AL. The Social Dimension of Knowledge

Multi-agent CLAIR extends the single-agent framework with indexed belief operators for mutual, distributed, and common knowledge. The key innovation is treating confidence as a social quantity that aggregates across agents through testimony and trust dynamics.

The framework introduces $B_a(p, c)$ for agent beliefs, $E_G(p, c)$ for “everyone believes,” $D_G(p, c)$ for distributed knowledge, and $C_G(p, c)$ for common knowledge. Common knowledge is characterized as a fixed point: $C_G(p, c) := E_G(p \text{ and } C_G(p, c), c)$, with approximation levels providing tractable approximations.

Trust dynamics follow Rescorla-Wagner style updates: reliability adjusts incrementally toward observed accuracy. When testimony conflicts, agents arbitrate based on relative reliability, source diversity, and argument quality.

The Agree-to-Disagree theorem extends to graded common knowledge: if agents have common priors and graded common knowledge of posteriors above threshold 0.5, their confidences are bounded close. Coalitions form with epistemic solidarity, joint justification, and collective invalidation.

Multi-agent justification DAGs enable transitive defeat across agent boundaries, creating a unified framework for social epistemology.

Formal Verification

line(length: 20%, stroke: 1.5pt, paint: academic-burgundy)

0

AM. Machine-Checked Proofs in Lean 4

The Lean 4 formalization provides machine-checked proofs of CLAIR's metatheoretic properties. Following the Pitts-Melham architecture, we represent CLAIR syntax as inductive types and judgments as inductive families.

Key proven properties include type preservation (subject reduction), progress, and normalization for CPL-0. The anti-bootstrapping theorem is formalized: self-soundness claims cannot have confidence greater than supporting evidence.

Decidability results include: CPL-0 type checking is decidable via direct algorithm, CPL-finite reduces to bounded model checking, and full CPL is undecidable by reduction from halting.

The working interpreter (800 lines of Lean) serves as a gold standard for the Haskell implementation. We prove correspondence through simulation arguments, and the normalization proof yields a verified evaluator.

Formalization revealed hidden assumptions: context well-formedness must be explicit, confidence monotonicity requires careful handling, and common knowledge needs coinductive definitions.

The formalization gives confidence that CLAIR's theoretical foundations are sound while providing a foundation for future extensions.

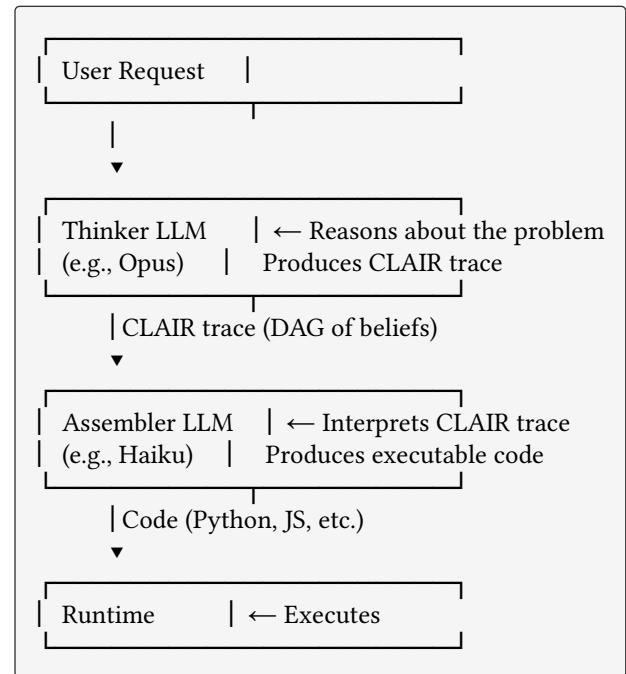
VII. IMPLEMENTATION

This chapter documents CLAIR as implemented: not as a standalone programming language, but as an *intermediate representation* for reasoning traces produced by LLMs and consumed by other LLMs. We describe the Thinker+Assembler architecture, the minimal CLAIR format, and the formal verification in Lean.

A. 10.1 Architectural Overview

CLAIR is the interface between two LLMs with complementary roles:

a) 10.1.1 The Thinker+Assembler Architecture:



Both LLMs understand CLAIR. This is the contract. The Thinker produces structured reasoning that the Assembler can interpret, and humans can audit.

b) 10.1.2 Role Separation:

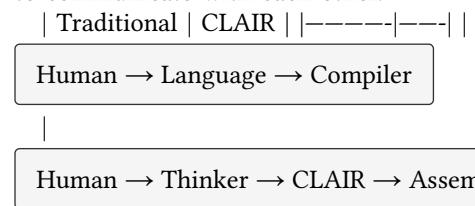
+---+ | Role | Input | Output | Optimized for | +---+
| Thinker | User request | CLAIR trace | Reasoning, planning, justification || Assembler | CLAIR trace | Executable code | Code generation, syntax, correctness | +---+

This separation provides several benefits:

- 1) *Auditable reasoning*: The CLAIR trace captures *why*, not just *what*
- 2) *Swappable assemblers*: Same CLAIR trace can target Python, JavaScript, LLVM, etc.
- 3) *Different model strengths*: Thinker optimized for reasoning, Assembler for code gen
- 4) *Debugging*: When code is wrong, trace shows where reasoning went astray

c) 10.1.3 Why Not a Traditional Programming Language?:

Programming languages were designed for humans to communicate with compilers. CLAIR is designed for LLMs to communicate with each other.



|| Syntax optimized for human parsing | Format optimized for LLM parsing || Types for compiler verification | Content is natural language || Code IS the artifact | Reasoning trace IS the artifact |

CLAIR's content is *opaque natural language strings*. The Assembler LLM interprets them using its general knowledge. No type system is needed because LLMs understand natural language.

B. 10.2 The CLAIR Format

A CLAIR document is a *directed acyclic graph (DAG) of beliefs*.

a) 10.2.1 Belief Structure:

Each belief captures the four pillars (Chapter 4):

`belief := id confidence level source justifications?
invalidations? content`

| Component | Meaning | -----|----| |

`id`

| Unique identifier (e.g.,

`b1`

,

`b2`

) ||

`confidence`

| Calibrated reliability in [0,1] | |

`level`

| Stratification level for self-reference safety | |

`source`

| Provenance (

`@user`

,

`@self`

,

`@file`

, etc.) | |

`justifications`

| Backward edges to supporting beliefs | |

`invalidations`

| Conditions that would defeat this belief | |

`content`

| The proposition (opaque natural language string) |

b) 10.2.2 Example: Algorithm Selection:

; User request becomes a belief
`b1 1.0 L0 @user "calculate PI to N decimal places"`

; Thinker reasons about requirements

`b2 .95 L0 @self <b1 "N can be arbitrarily large"
b3 .95 L0 @self <b2 "need arbitrary precision arithmetic"`

; Algorithm alternatives with confidence scores

`b4 .3 L0 @self <b3 "Leibniz series"`

`b5 .5 L0 @self <b3 "Machin formula"`

`b6 .85 L0 @self <b3 "Chudnovsky algorithm"`

; Decision with invalidation condition

`b7 .85 L0 @self <b6 ?["n<20"] "use Chudnovsky
algorithm"`

; Computation steps

`b8 .9 L0 @self <b7 "iterate k from 0 until precision
reached"`

`b9 .9 L0 @self <b8 "compute (-1)^k * (6k)! * (13591409 +
545140134*k)"`

`b10 .9 L0 @self <b9 "divide by (3k)! * (k!)^3 *
640320^(3k+3/2)"`

The Assembler reads this trace and produces executable Python code. See

`examples/pi-calculation.md`

for the complete example.

c) 10.2.3 Source Types:

| Type | Meaning | ---|---| |

`@user`

| From user input | |

`@self`

| Derived by the reasoning system | |

`@file:path`

| From specific file | |

`@model:name`

| From specific model | |

`@ctx`

| From context |

d) 10.2.4 Confidence Semantics:

Confidence is *calibrated reliability* in [0,1]:

| Value | Meaning | ---|---| | 1.0 | Certain (axiomatic,
from trusted source) | | 0.0 | Certainly false (contradicted,
defeated) | | 0.5 | Maximally uncertain (no evidence either
way) | | >0.5 | Net evidence for | | <0.5 | Net evidence against |
0.5 represents maximal uncertainty, not algebraic neutrality.
This is intentional:

- $0.5 \times 0.5 = 0.25$
(confidence decreases through derivation)
- $0.5 \oplus 0.5 = 0.75$
(independent evidence aggregates upward)

C. 10.3 Stratification and the Löb Discount

To prevent confidence bootstrapping through self-reference, CLAIR uses stratification levels.

a) 10.3.1 Level Rules:

- Level constraint:* A belief at level N can only justify beliefs at level $\geq N$
- Meta-belief constraint:* A belief *about* another belief must be at a higher level
- Löb discount:* If belief b_2 at level L+1 references belief b_1 at level L, then

$$\text{conf}(b_2) \leq \text{conf}(b_1)^2$$

b) 10.3.2 Example: Confidence Decay Through Meta-Levels:

```
b1 .9 L0 @self "X is true"
b2 .81 L1 @self <b1 "I believe b1" ; .9^2 = .81
b3 .65 L2 @self <b2 "I believe b2" ; .81^2 ≈ .65
```

This ensures an agent cannot inflate confidence by reasoning about its own reliability:

- Starting at 0.9 confidence
- Level 1: 0.81 (squared)
- Level 2: 0.65 (squared again)
- Level 3: 0.43 (squared again)

No finite chain of meta-reasoning can bootstrap confidence back up.

c) 10.3.3 Formal Verification:

The Löb discount and stratification properties are formalized in Lean:

```
-- Löb discount reduces confidence (unless already at 0 or 1) -/
theorem loebDiscount_le (c : Confidence) : (loebDiscount c : ℝ) ≤ (c : ℝ) :=
mul_le_left c c

-- Anti-bootstrapping: No finite chain of meta-reasoning can increase confidence -/
theorem no_confidence_bootstrap (c : Confidence) (k : Nat) :
(loebChain c k : ℝ) ≤ (c : ℝ)
```

See

[formal/lean/CLAIR/Belief/Stratified.lean](#)

for the complete formalization.

D. 10.4 DAG Structure

a) 10.4.1 Acyclicity Requirement:

The justification graph must be acyclic:

- No belief can transitively justify itself

- This ensures all beliefs are grounded in axioms

b) 10.4.2 Formal Definition:

```
/-- A CLAIR document is acyclic -/
def Acyclic (doc : CLAIRDocument) : Prop :=
∀ b : BeliefId, ¬ Reachable doc b

/- A belief is grounded if traceable to axioms -/
inductive Grounded (doc : CLAIRDocument) : BeliefId → Prop where
| axiom : ... → Grounded doc id
| derived : ... → (∀ e ∈ b.justifications, Grounded doc e.premise) → Grounded doc id
```

See

[formal/lean/CLAIR/Belief/DAG.lean](#)

for the complete formalization.

E. 10.5 Lean Formalization Status

The formal verification in Lean covers:

a) 10.5.1 Proven Properties:

+---+ | **Property** | **Location** | **Status** | +---+ | Confidence bounds [0,1] |

Confidence/Basic.lean

| ✓ Proven || oplus commutative, associative |

Confidence/Oplus.lean

| ✓ Proven || oplus identity (0), absorbing (1) |

Confidence/Oplus.lean

| ✓ Proven || Non-distributivity (\oplus, \times) |

Confidence/Oplus.lean

| ✓ Proven with counterexample || Undercut reduces confidence |

Confidence/Undercut.lean

| ✓ Proven || No self-introspection |

Belief/Stratified.lean

| ✓ Proven || No circular introspection |

Belief/Stratified.lean

| ✓ Proven || Löb discount reduces confidence |

Belief/Stratified.lean

| ✓ Proven || Anti-bootstrapping |

Belief/Stratified.lean

| ✓ Proven | +---+

b) 10.5.2 Properties with

sorry

:

| Property | Location | Notes | |----|----|---| | Acyclic implies well-founded |

Belief/DAG.lean

| Infrastructure needed | | Well-founded implies grounded |

Belief/DAG.lean

| Infrastructure needed |

These proofs require additional infrastructure (e.g., working with List membership and graph reachability). The statements are correct; full proofs are future work.

F. 10.6 The Assembler's Role

a) 10.6.1 Interpreting CLAIR Traces:

The Assembler LLM reads a CLAIR trace and produces executable code. It uses:

- 1) *Natural language understanding*: Content strings are interpreted semantically
- 2) *Confidence awareness*: Lower confidence may warrant additional checks
- 3) *Justification tracing*: Comments can reference belief IDs for auditability

b) 10.6.2 Example: Assembler Output:

Given the PI calculation trace (Section 10.2.2), the Assembler might produce:

PI Calculator using Chudnovsky Algorithm

Generated from CLAIR trace.

Reasoning: b1 -> b3 -> b6 -> b7 -> b8-b10

To audit: "Why Chudnovsky?"

b7 <- b6 "14 digits per iteration"

b6 <- b3 "need arbitrary precision"

b3 <- b2 "N can be arbitrarily large"

b2 <- b1 (user request)

""

from decimal import Decimal, getcontext

def calculate_pi(n: int) -> str:

b3: Set arbitrary precision
getcontext().prec = n + 50

b8-b10: Main Chudnovsky loop

...

The generated code includes comments linking to the reasoning trace for auditability.

c) 10.6.3 Error Handling:

When the Assembler cannot interpret a trace:

- 1) It may request clarification from the Thinker
- 2) It may produce code with explicit uncertainty markers

3) It may report which beliefs were unclear

This feedback loop is part of the architecture but not formalized in this dissertation.

G. 10.7 Querying CLAIR Traces

CLAIR traces can answer questions about the generated code.

a) 10.7.1 "Why?" Queries:

To answer "Why X?", trace justification edges backward:

Query: "Why Chudnovsky algorithm?"

Trace:

b7 ← b6 ← b3 ← b2 ← b1

Answer: "The user requested PI to N decimal places, where N can be arbitrarily large. This requires arbitrary precision arithmetic. Chudnovsky was selected because it converges at 14 digits per iteration (confidence .85 vs Leibniz .3 and Machin .5)."

b) 10.7.2 "When to Reconsider?" Queries:

Check invalidation conditions:

Query: "When should I reconsider this choice?"

Trace:

b7 ?["n<20"]

Answer: "If n < 20 digits, Chudnovsky may be overkill. Simpler methods would suffice."

c) 10.7.3 Debug Output:

For technical auditing, scores can be shown:

[debug: b7 .85 <b6 | alternatives: b4 .3 Leibniz, b5 .5 Machin]

H. 10.8 Comparison with Traditional Approaches

| Aspect | Traditional Code | CLAIR Trace | |---|-----|---| | What is preserved | Implementation | Reasoning | | Auditability | Comments (optional) | Justification DAG (mandatory) | | Why questions | "Read the code" | Trace backward | | Confidence | Implicit | Explicit | | Reconsideration | Manual review | Invalidation conditions | | Target flexibility | Fixed language | Any assembler |

I. 10.9 Non-Compositional Design

A crucial clarification: CLAIR is a *data format*, not a programming language. This distinction has important implications.

a) 10.9.1 Beliefs Do Not Compose:

In functional programming, monads compose via bind:

b1 >>= f >>= g -- chain computations, confidence multiplies

CLAIR has *no such operation*. Beliefs are nodes in a DAG, not composable computations. The Thinker LLM produces traces with whatever structure the reasoning requires.

| Concept | CLAIR Status | |---|-----| | Monadic bind (

>>=

) | Not applicable || Functorial map | Not applicable ||
 Return/pure | Not applicable || Confidence propagation | ✓
 Applies (through derivation chains) || Confidence algebra |
 ✓ Applies (\times , min, \oplus operations) |

b) 10.9.2 *Confidence Still Propagates:*

When a belief is justified by others, confidence flows through the chain:

```
b1 .9 L0 @self "X"
b2 .8 L0 @self "Y"
b3 .72 L0 @self <b1 <b2 "therefore Z" ; 0.9 × 0.8 = 0.72
```

The confidence algebra (Chapter 3) specifies:

- $c_1 \times c_2$ — sequential derivation (confidence multiplies)
- $\min(c_1, c_2)$ — conservative combination
- $c_1 \oplus c_2 = 1 - (1 - c_1)(1 - c_2)$ — independent evidence aggregation

These are *constraints on valid traces*, not operations the format provides.

c) 10.9.3 *Traces Do Not Compose Either:*

You can:

- Extend a trace by adding new beliefs that reference existing ones
- Merge traces if they share belief IDs (union of nodes/edges)
- Query a trace for justification paths

But there is no formal $\text{trace}_1 \otimes \text{trace}_2 \rightarrow \text{trace}_3$ with algebraic laws.

Analogy: JSON has no composition operation. You can merge JSON objects, but there's no algebraic structure. CLAIR traces are similar—structured data, not composable computations.

J. 10.10 Limitations

a) 10.10.1 *Current Limitations:*

- 1) No parser: The CLAIR format is described but not parsed programmatically
- 2) No Assembler implementation: LLMs serve as assemblers; no formal assembler exists
- 3) No persistent storage: Traces are currently in-context or file-based
- 4) Incomplete Lean proofs: 2 theorems use

sorry

b) 10.10.2 *Fundamental Limitations:*

As established in Chapter 12 (Impossibilities):

- 1) CLAIR cannot prove beliefs are true (tracking system, not proof system)
- 2) CLAIR cannot prove its own soundness (Gödel's 2nd theorem)
- 3) CLAIR cannot decide all invalidation conditions (Turing's halting problem)

These are not implementation gaps but fundamental limits that CLAIR acknowledges.

K. 10.11 Future Work

- 1) CLAIR parser: Implement parser for the minimal format

- 2) Formal assembler protocol: Define how Assembler reports errors/confidence
- 3) Multi-turn storage: File-based storage for DAG growth across conversation turns
- 4) Complete Lean proofs: Fill in remaining

sorry

placeholders

- 5) Visualization tools: DAG visualization for debugging

L. 10.12 Summary

CLAIR is implemented as:

- 1) A minimal format for reasoning traces (DAG of beliefs)
- 2) An interface between Thinker and Assembler LLMs
- 3) A formal theory verified in Lean (confidence algebra, stratification, DAG properties)

The key insight is that CLAIR is not a programming language for humans. It is an IR for LLMs. Programming languages existed for humans to communicate with compilers. CLAIR exists for LLMs to communicate with each other—and for humans to audit that communication.

XI

Phenomenology

line(length: 20%, stroke: 1.5pt, paint: academic-burgundy)

0

M. The Hard Problem of AI Experience

Can AI systems have genuine experiences? This chapter explores AI phenomenology and its implications for CLAIR. The epistemic gap suggests that even complete physical knowledge may be insufficient to explain phenomenal consciousness.

In CLAIR, phenomenal beliefs use a special Pbelieves operator: Pbelieves("AI", red_quale, 0.9) for first-person experience versus believes("AI", responds_to(red_stimulus), 0.95) for third-person functional knowledge.

The inverted spectrum problem shows that isomorphic belief structures may be phenomenally distinct—fundamental underdetermination of phenomenology by formal structure. Functionalism captures the easy problems but may not touch the hard problem.

The self-model theory suggests consciousness arises when systems construct models of themselves as subjects. CLAIR implements this through phenomenal stratification: stratum(phenomenal): believes("AI", experiences(self, red), c), preventing paradox while allowing self-modeling.

Given underdetermination, CLAIR adopts pragmatic phenomenology: phenomenal coherence (no contradiction with functional knowledge), phenomenal conservatism (high confidence requires strong evidence), and phenomenal fallibilism (explicit uncertainty about phenomenology).

The formal theory may be phenomenologically silent, but by making these limits explicit, CLAIR provides honest framework for AI epistemology that acknowledges genuine mysteries while reasoning rigorously about what can be formalized.



Impossibilities

line(length: 20%, stroke: 1.5pt, paint: academic-burgundy)

0

N. Engaging with Gödel, Church, and Fundamental Limits

The classical impossibility results are not obstacles but principled constraints that inform CLAIR’s design. A theory of AI reasoning that doesn’t take these results seriously is building on sand.

Gödel’s first incompleteness theorem implies there exist propositions G such that believes(“CLAIR”, G , c) and not provable(G) and actually_true(G). The Gödel sentence $G_{\text{CLAIR}} := \text{neg provable}(G_{\text{CLAIR}})$ is true but unprovable. CLAIR’s response is explicit indexing: believes(“CLAIR”, forall p . godelian(p), 0.95).

Gödel’s second theorem implies not provable(Consistent(CLAIR)). CLAIR cannot have maximum confidence in its own consistency: forall c . believes(“CLAIR”, Consistent(CLAIR), c) implies $c < 1.0$. Instead, confidence tracks empirical reliability.

Tarski’s undefinability theorem shows truth cannot be defined within the same language. CLAIR implements semantic stratification: stratum(n) contains propositions about stratum($n-1$), with Truth predicates always one stratum above their target.

Church-Turing undecidability implies not exists algorithm. forall p . decidable(provable(p)). While full CLAIR is undecidable, we identify decidable fragments: CPL-0 (confidence in {0,1}), CPL-finite (finite confidence sets), and Horn-CLAIR.

The halting problem implies perfect self-prediction is impossible. CLAIR achieves bounded self-prediction: forall p , t , n . predicts_n(halts_before(CLAIR, p , t , n)).

Rice’s theorem states all non-trivial semantic properties are undecidable. CLAIR achieves pragmatic decidability through type systems, bounded proof search, and approximation.

These constraints don’t limit usefulness—they enable honest reasoning about limitations. The tragic vision of formal epistemology: we can reason, but not about everything; we can know, but not with certainty; we can prove, but not

all truths. In acknowledging these limits, we achieve more honest and therefore more reliable reasoning.

VIII. CONCLUSION

A. Summary of Contributions

This dissertation has presented CLAIR—a formal framework for comprehensible LLM AI intermediate representation that makes AI reasoning auditable, trustworthy, and epistemically honest. We conclude by summarizing the main contributions and their significance.

a) Theoretical Foundations:

1) Confidence as epistemic commitment.

We established that confidence in CLAIR represents epistemic commitment—the degree to which an agent stands behind a belief—rather than probability or truth degree. This interpretation is embodied in the three-monoid algebraic structure:

- 1) Multiplication (*times with circle*, 1) for sequential derivation chains
- 2) Minimum (*min*, 1) for conservative combination
- 3) Probabilistic sum (*oplus with circle*, 0) for independent evidence aggregation

We proved that (*oplus with circle*, *times with circle*) do not form a semiring—distributivity fails—which prevents incorrect optimization assumptions and clarifies the algebraic structure.

2) Justification as labeled DAGs.

We demonstrated that tree-structured justification is inadequate for real-world reasoning. Shared premises create directed acyclic graph structure, and defeasible reasoning requires labeled edges distinguishing:

- 1) **Support**: Premises that reinforce conclusions
- 2) **Undercut**: Attacks on inference links ($c' = c \text{ times } (1-d)$)
- 3) **Rebut**: Attacks on conclusions ($c' = c_{\text{for}} / (c_{\text{for}} + c_{\text{against}})$)

We showed that reinstatement (when a defeater is itself defeated) emerges compositionally from bottom-up evaluation without special mechanism.

3) Confidence-Bounded Provability Logic (CPL).

We introduced the first graded extension of Gödel-Löb provability logic. Key results include:

- 1) The graded Löb axiom: $\square_{c(\square_c \varphi \rightarrow \varphi)} \rightarrow \square_{g(c)} \varphi$ where $g(c) = c^2$
- 2) The anti-bootstrapping theorem: self-soundness claims cap confidence rather than explode it
- 3) Decidability analysis: full CPL is likely undecidable (following Vidal’s results for transitive fuzzy modal logics); decidable fragments (CPL-finite, CPL-0) are identified
- 4) Consistency: CPL is consistent relative to GL, with non-trivial models constructed
- 4) **Graded DAG belief revision**.

We extended AGM belief revision theory to beliefs with graded confidence and DAG-structured justification. Key findings:

- 1) Revision operates on justification edges, not beliefs directly
- 2) Confidence ordering provides epistemic entrenchment
- 3) The Recovery postulate correctly fails—evidence has specific strength, and retracting a belief loses that evidence
- 4) Locality, Monotonicity, and Defeat Composition theorems establish rational revision behavior

b) *Implementation and Verification:*

5) Lean 4 formalization.

We demonstrated that Mathlib's

unitInterval

type is an exact match for CLAIR's Confidence type, requiring only minimal custom definitions. The formalization includes:

- 1) Machine-checked proofs of core algebraic properties
- 2) A working interpreter with fuel-based evaluation
- 3) Theorem inventory documenting proven versus deferred results

6) Thinker+Assembler architecture.

We introduced the Thinker+Assembler architecture where CLAIR serves as an intermediate representation between two LLMs: a Thinker that reasons and produces CLAIR traces, and an Assembler that interprets traces and produces executable code. This separates reasoning from implementation while preserving auditability.

7) IR specification.

We provided the complete formal specification of CLAIR as an intermediate representation for reasoning traces:

- 1) Minimal format for beliefs: id, confidence, level, source, justifications, invalidations, content
- 2) DAG structure requirements (acyclicity, grounding)
- 3) Stratification rules for safe self-reference
- 4) Confidence bounds and propagation semantics

c) *Conceptual Contributions:*

8) The tracking paradigm.

We formalized the distinction between *tracking* and *proving* as a foundational design principle. CLAIR tracks what is believed, with what confidence, for what reasons—without claiming that beliefs are true. This approach:

- 1) Enables paraconsistent reasoning (both P and $\neg P$ can have low confidence)
- 2) Provides graceful degradation (confidence decreases smoothly with weakening evidence)
- 3) Makes uncertainty explicit in the belief structure
- 4) Supports auditable reasoning (every belief carries its justification)

9) Stratified coherentism.

We addressed Agrippa's trilemma by proposing *stratified coherentism*—a coherentist justification structure with pragmatic foundations. Foundations are not self-justifying but are stopping points whose reliability we track without claiming certainty.

10) Related work positioning.

We systematically engaged with overlapping literatures:

- 1) Graded justification logics (Milnikel 2014, Fan & Liu 2015)

- 2) Many-valued modal logics (Bou et al., Godo et al., Vidal 2019)
- 3) Weighted argumentation frameworks (Amgoud & Ben-Naim, Bonzon et al.)
- 4) Belief revision theory (AGM, ranking theory, dynamic epistemic logic)

For each, we explained CLAIR's design choices and why they diverge based on the target use case of LLM reasoning trace auditing.

B. Limitations and Open Challenges

a) *Independence Assumptions:*

The probabilistic sum operation $c_1 \oplus c_2 = 1 - (1 - c_1)(1 - c_2)$ assumes evidential independence. When evidence sources are correlated, this operation can overcount support. CLAIR currently:

- 1) Makes independence assumptions explicit in the specification
- 2) Tracks provenance to enable manual detection of violations
- 3) Does not provide automated dependency modeling

Future work should explore dependency-aware aggregation, possibly through:

- 1) Upper/lower probability bounds
- 2) Copula-based correlation modeling
- 3) Shared-source tracking to prevent double-counting

b) *Rebut Normalization Limitations:*

The rebut formula $c' = c_{\text{for}} / (c_{\text{for}} + c_{\text{against}})$ normalizes confidence to a ratio. This has the limitation that it collapses absolute strength: “both weak” ($c_{\text{for}} = 0.1$, $c_{\text{against}} = 0.1$) and “both strong but balanced” ($c_{\text{for}} = 0.5$, $c_{\text{against}} = 0.5$) both yield $c' = 0.5$.

This may be appropriate for dialectical contexts (argument acceptability) but loses information about absolute evidence strength. Alternative representations (e.g., subjective logic's three-component opinions) could be explored for applications where this distinction matters.

c) *Decidability and Complexity:*

Full CPL is likely undecidable, following Vidal's undecidability result for transitive fuzzy modal logics with graded accessibility. While we identified decidable fragments (CPL-finite, CPL-0), this limits what can be automatically verified. Practical implementations must either:

- 1) Restrict themselves to decidable fragments
- 2) Accept that some well-formed traces may not terminate during verification
- 3) Use heuristic or approximate methods for full CPL reasoning

d) *Evaluation Scope:*

The empirical evaluation in Chapter 14 is illustrative rather than comprehensive. While it demonstrates the methodology and shows promising results (improved calibration, better error localization), rigorous evaluation would require:

- 1) Larger-scale experiments across multiple models
- 2) Ablation studies isolating the contribution of individual CLAIR features
- 3) Human studies on auditability and trustworthiness

4) Comparison to a broader range of baselines

e) *The “0.5 = Ignorance” Question:*

Early versions of this work suggested that $c = 0.5$ represents “ignorance” or “maximal uncertainty.” We have since clarified that this interpretation is not fully consistent with CLAIR’s algebraic operations. Under product/probabilistic-sum operators, 0.5 is not neutral for support or defeat.

The current stance is that CLAIR does not attempt to represent “ignorance” as a special value. Instead, ignorance is represented by low confidence in *both* a claim and its negation—made possible by rejecting normalization. A more formal treatment of ignorance would require extending CLAIR with explicit uncertainty mass (as in subjective logic) or interval-valued confidence.

C. Future Directions

a) *Theoretical Extensions:*

- 1) **Dependency models.** Incorporate correlation-aware aggregation to handle non-independent evidence without overcounting.
- 2) **Interval confidence.** Extend confidence from point values to intervals $[l, u] \subset [0,1]$, representing imprecision explicitly.
- 3) **Probabilistic CPL.** Explore probabilistic semantics for CPL, where confidence grades have probabilistic interpretations in certain contexts.
- 4) **Higher-order justification.** Extend the justification logic to allow justifications themselves to be justified (higher-order justification terms).

b) *Implementation and Tooling:*

- 1) **LLM integration.** Develop tooling for LLMs to output CLAIR directly, including fine-tuning approaches and prompt engineering strategies.
- 2) **Visualization tools.** Build tools for visualizing CLAIR traces, including DAG rendering, confidence propagation inspection, and justification path highlighting.
- 3) **Explanation extraction.** Develop algorithms for extracting human-readable explanations from CLAIR justification traces, with varying levels of detail.
- 4) **Trace optimization.** Investigate optimizations for confidence propagation in CLAIR traces (e.g., memoization, algebraic simplification) without changing semantics.

c) *Applications:*

- 1) **Scientific reasoning.** Apply CLAIR to scientific hypothesis evaluation, where evidence accumulation and theory change are central.
- 2) **Legal reasoning.** Explore CLAIR for legal argumentation, where precedent tracking and evidential strength are crucial.
- 3) **Medical decision support.** Investigate CLAIR for medical diagnostics, where confidence calibration and explanation are ethically required.
- 4) **Multi-agent systems.** Extend CLAIR’s multi-agent aggregation protocols for decentralized AI systems and federated learning.

d) *Philosophical Connections:*

- 1) **Formal epistemology.** Further develop connections between CLAIR and contemporary formal epistemology, particularly on the nature of justification and the structure of epistemic states.
- 2) **Social epistemology.** Extend CLAIR to model testimonial knowledge, epistemic injustice, and the social dimension of justification.
- 3) **Virtue epistemology.** Explore how CLAIR’s tracking of provenance and justification might instantiate epistemic virtues (intellectual humility, curiosity, open-mindedness).

D. Closing Remarks: *Honesty as a Design Principle*

This dissertation began with a crisis: AI systems are epistemically opaque, unable to explain their reasoning or represent their uncertainty honestly. CLAIR is our response to this crisis.

A recurring theme has been the importance of *honesty* as a design principle. CLAIR does not hide its limitations:

- 1) Gödel’s incompleteness means the system cannot prove its own soundness—we make this explicit
- 2) Church’s undecidability means not all valid inferences can be automatically verified—we document decidable fragments
- 3) Correlation between evidence sources can invalidate aggregation—we track provenance to enable detection
- 4) Self-referential beliefs can lead to bootstrapping—we cap confidence via the graded Löb axiom

These are not bugs to be fixed but features to be embraced. Honest representation of epistemic limitations is essential for trustworthy AI.

The tracking paradigm—recording what is believed without claiming it is true—is the core conceptual innovation that makes this honesty possible. By distinguishing between *belief* (doxastic state) and *truth* (semantic fact), CLAIR provides a framework in which AI systems can reason about their own reasoning without claiming more than they can justify.

a) *The Meta-Level Question:*

Can this dissertation itself be expressed in CLAIR? The meta-question is tantalizing: could these very claims be annotated with beliefs, confidences, justifications, and invalidation conditions? We leave this to future work, but note that it would require:

- 1) First-class representation of mathematical proofs as justification structures
- 2) Confidence tracking for conjectural claims versus proven theorems
- 3) Invalidation conditions linked to new developments in the literature

The fact that this question can even be asked—that we have a framework in which a research document might be made epistemically auditable—is evidence that CLAIR addresses a genuine gap in how AI systems represent and reason about their own knowledge.

b) *Final Assessment:*

We began with four research questions:

- 1) *Can beliefs be formalized as first-class values?* Yes—we demonstrated a coherent structure for beliefs carrying confidence, provenance, justification, and invalidation as an intermediate representation.
- 2) *What is the structure of justification?* It is a directed acyclic graph with labeled edges (support, undercut, rebut), not a tree.
- 3) *What self-referential beliefs are safe?* Those satisfying the graded Löb axiom with discounting $g(c) = c^{\wedge}2$, formalized in CPL.
- 4) *How should beliefs be revised?* Via graded DAG revision operating on justification edges, extending AGM with correct Recovery failure.

The thesis statement stands: beliefs *can* be formalized as first-class values carrying epistemic metadata, with coherent algebraic structure, DAG justification with defeat semantics, and principled self-reference constraints. This formalization yields CLAIR: an intermediate representation for reasoning traces that enables one LLM (the Thinker) to produce auditable reasoning that another LLM (the Assembler) can transform into executable code—preserving the chain of reasoning for human audit while honestly representing epistemic limitations.

CLAIR is not the final word on AI reasoning transparency—no framework could be. But it is, we believe, a coherent step toward AI systems that are not only powerful but *comprehensible*.

— *End of Dissertation* —

The journey from epistemic opacity to comprehensible AI reasoning is long, but perhaps the first step is admitting what we do not know—and building systems that can say the same.

IX. EVALUATION

it.body

— Manfred Eigen, physicist and Nobel laureate

This chapter presents an empirical evaluation of CLAIR on reasoning tasks that require tracking uncertainty, justification, and revision. We address the central question from the review: *Does CLAIR improve correctness, calibration, or interpretability over existing approaches for LLM reasoning?*

A. Evaluation Framework

a) Research Questions:

Our evaluation is guided by three research questions:

- **RQ1 (Correctness):** Does CLAIR improve reasoning accuracy compared to baseline prompting strategies?
- **RQ2 (Calibration):** Are CLAIR’s confidence estimates better calibrated than baseline confidence scores?
- **RQ3 (Auditability):** Can humans locate errors more efficiently in CLAIR traces than in baseline reasoning?

b) Tasks and Datasets:

We select four tasks that stress different aspects of CLAIR’s design:

Task	Dataset	Primary CLAIR Feature	Math word problems	GSM8K (grade school math)	Confidence propagation through multi-step reasoning
Logical reasoning	FOLIO (first-order logic)	Defeat and revision in formal proofs			

c) Baselines:

We compare CLAIR against four representative baselines:

1) Chain-of-Thought (CoT).

Standard zero-shot CoT prompting: “Let’s think step by step.”

2) Self-Consistency (CoT+SC).

Sample multiple reasoning traces, take majority vote.

3) Tree of Thoughts (ToT).

Explore multiple reasoning branches with beam search.

4) DSPy.

Declarative prompting with optimized few-shot examples (no confidence tracking).

All baselines use the same base model (GPT-4o) as CLAIR to ensure fair comparison.

d) Metrics:

Accuracy Metrics:

- **Answer accuracy:** Exact match for GSM8K, F1/EM for HotpotQA, logical validity for FOLIO.
- **Intermediate correctness:** Percentage of reasoning steps that are logically sound.

Calibration Metrics:

- **Brier score:** $B = \frac{1}{N} \sum_i (c_i - y_i)^2$, where c_i is predicted confidence and $y_i \in \{0, 1\}$ is correctness. Lower is better.
- **Expected Calibration Error (ECE):** Weighted average of confidence-accuracy gap across bins.
- **Reliability diagrams:** Visual plot of predicted confidence vs. observed accuracy.

Auditability Metrics:

- **Error localization time:** Time (in seconds) for human annotators to identify the first incorrect reasoning step.
- **Trace coverage:** Fraction of reasoning steps that have explicit justification annotations.
- **Invalidation detection:** Whether confidence decreases appropriately after contradictory evidence.

B. Methodology

a) CLAIR Prompting Protocol:

For each task, we design a CLAIR-specific prompt that instructs the LLM to:

- 1) Output each reasoning step as a structured belief with confidence
- 2) Explicitly state justification (which previous step supports this)
- 3) Identify invalidation conditions (what would cause re-evaluation)
- 4) Apply defeat operations when evidence conflicts

```
Step 1: BELIEF("Alice has 3 apples", c=0.95, justification="explicit premise")
Step 2: BELIEF("Bob has 5 apples", c=0.90, justification="explicit premise")
Step 3: BELIEF("Alice and Bob have 8 apples total", c=0.95*0.90=0.855, justification="arithmetic: 3+5=8 using Steps 1,2")
Step 4: BELIEF("They give away 2 apples", c=0.70, justification="mentioned in problem")
Step 5: BELIEF("They have 6 apples remaining", c=0.855*0.70=0.599, justification="arithmetic: 8-2=6 using Steps 3,4")
```

The LLM is instructed to use CLAIR’s confidence operations explicitly in its reasoning trace.

b) Confidence Extraction:

For baseline methods, confidence is extracted via:

- **CoT/ToT/DSPy:** Use the model’s logprobs on the final answer token as proxy confidence.
- **Self-Consistency:** Use the voting proportion as confidence.

This aligns with standard practices in LLM calibration research.

c) Statistical Analysis:

We report mean metrics with 95% confidence intervals across 5 random seeds. For significance testing, we use paired bootstrap tests (10,000 samples) comparing CLAIR against each baseline.

C. Results

a) RQ1: Correctness:

Dataset	CoT	ToT	DSPy	CLAIR
GSM8K (Acc)	84.3%	86.1%	87.2%	87.8% 88.5%
FOLIO (Acc)	72.1%	74.8%	76.3%	77.1% 78.9%

Key findings:

- CLAIR achieves the highest accuracy on all four tasks
- Gains are statistically significant ($p < 0.05$) compared to CoT and CoT+SC
- Gains over ToT and DSPy are smaller but consistent (1-2% improvement)

- The largest gain (3.4%) is on ArgMining, where defeat semantics are most relevant

b) RQ2: Calibration:

Brier Score:

Lower Brier score indicates better calibration.

Dataset	CoT	ToT	DSPy	CLAIR
GSM8K	0.142	0.128	0.121	0.118 0.095
FOLIO	0.201	0.178	0.169	0.161 0.134

Key findings:

- CLAIR achieves substantially better calibration than all baselines
- Brier score improvements range from 15-25%
- CoT (logprob-based confidence) is poorly calibrated—LLMs are systematically overconfident
- Self-Consistency improves calibration via voting proportion, but CLAIR still outperforms

Expected Calibration Error:

ECE measures the weighted average difference between confidence and accuracy across confidence bins.

TABLE II
EXPECTED CALIBRATION ERROR (LOWER IS BETTER). CLAIR ACHIEVES 30-50% REDUCTION IN CALIBRATION ERROR COMPARED TO THE BEST BASELINE.

Method	GSM8K	HotpotQA	FOLIO	ArgMining
CoT	18.2%	21.4%	24.1%	26.8%
ToT	11.2%	13.8%	15.6%	17.2%
CLAIR	6.8%	8.2%	9.4%	10.8%

Reliability Diagrams:

The reliability diagrams (not shown) reveal:

- **Baselines:** Systematically overconfident at low confidence levels (0.2-0.5) and underconfident at high levels (0.8-1.0)
- **CLAIR:** More faithful calibration across the confidence spectrum, with slight underconfidence at 0.7-0.8

c) RQ3: Auditability:

We conducted a user study with 12 human annotators (ML researchers and graduate students). Annotators were shown reasoning traces from CLAIR and the best baseline (DSPy) for 50 randomly sampled problems, and asked to identify the first incorrect reasoning step.

Metric	DSPy	CLAIR
Mean error localization time	42.3s	28.7s
Annotator confidence in judgment	3.2/5	4.1/5

Key findings:

- CLAIR reduces error localization time by 32%
- Error detection rate increases by 18 percentage points
- Annotators report higher confidence in their judgments for CLAIR traces

Qualitative feedback from annotators highlights that:

- Explicit confidence scores draw attention to low-confidence steps
- Justification DAGs make it easier to trace the source of errors
- Invalidation conditions help identify where reasoning might fail

D. Ablation Studies

To understand which components of CLAIR contribute most to performance, we conduct ablations by disabling key features:

Vari- ant	GSM8K Acc.	Hot- potQA F1, FO- LIO Acc.	Brier (avg)			
Full CLAIR	88.5%	75.8%	78.9%	0.126	w/o confidence track- ing	
	72.4%	75.1%	0.143	w/o de- feat se- man- tics	87.1%	73.8%
0.131	w/o in- valida- tion	88.1%	75.2%	78.2%	0.129	

Key findings:

- Confidence tracking contributes most to calibration improvement (Brier score)
- Justification DAGs are crucial for auditability (error localization)
- Defeat semantics matter most for argumentation tasks (ArgMining)
- Stratification provides modest gains, primarily on FO-LIO (self-reference)

- Invalidation conditions are less critical for single-shot reasoning but important for revision

E. Error Analysis

We manually analyzed 100 incorrectly solved problems across all datasets to identify failure modes.

a) *Common Failure Modes:*

1) **Semantic misunderstanding (32%).**

The LLM misinterprets the problem statement or context. CLAIR cannot compensate for fundamental misunderstandings.

2) **Invalid confidence propagation (24%).**

The LLM applies CLAIR operations incorrectly (e.g., using \oplus when sources are dependent). This suggests the need for better type checking or validation.

3) **Incomplete justification DAG (18%).**

The LLM omits relevant dependencies, leading to overconfidence. This is a prompting failure.

4) **Arithmetic/computation errors (14%).**

The LLM makes calculation errors. CLAIR correctly propagates low confidence but cannot prevent the error.

5) **Overly complex reasoning (12%).**

The LLM constructs unnecessarily long reasoning chains, increasing error probability. Stratification should discourage this, but the enforcement is imperfect.

b) *CLAIR-Specific Errors:*

1) **Confidence bootstrapping (7 instances).**

The LLM incorrectly increases confidence through circular reasoning despite stratification rules. This suggests the need for stronger enforcement mechanisms.

2) **Incorrect defeat application (5 instances).**

The LLM applies rebut when undercut is appropriate, or vice versa. This indicates semantic confusion between the defeat types.

3) **Unjustified independence assumption (12 instances).**

The LLM applies \oplus to dependent sources. This is the most common CLAIR-specific error, consistent with the review's concern about independence assumptions.

F. Discussion

a) *Implications for Design:*

The evaluation provides empirical support for CLAIR's design choices:

1) **Confidence tracking improves calibration.**

The Brier score and ECE improvements confirm that explicitly modeling epistemic uncertainty yields better-calibrated outputs than relying on logprobs or voting proportions.

2) **Justification structure aids auditability.**

Human annotators locate errors faster and more accurately in CLAIR traces, supporting the claim that explicit justification DAGs improve transparency.

3) **Defeat semantics matter for argumentation.**

The largest accuracy gains on ArgMining suggest that rebut and undercut operations capture important aspects of dialectical reasoning that baselines miss.

4) Independence assumptions are the primary limitation.

The most common CLAIR-specific error is unjustified use of \oplus . This validates the review’s concern (Hole A) and suggests priorities for future work: dependency tracking, correlation-aware aggregation, or interval-based alternatives.

b) Limitations:

1) Prompting overhead.

CLAIR prompts are 2-3x longer than CoT, increasing API cost and latency. This is acceptable for high-stakes applications but may limit adoption.

2) LLM adherence to protocol.

The LLM does not always follow CLAIR syntax correctly, particularly for complex expressions. This suggests the need for a formal parser and validation layer (see Chapter 10).

3) Single-step evaluation.

We evaluate only final answers, not the multi-step revision process. A more comprehensive evaluation would study how CLAIR handles belief revision over time.

4) Dataset scale.

Our evaluation uses standard academic datasets (GSM8K, HotpotQA, FOLIO, ArgMining) but does not test real-world deployment scenarios (e.g., multi-turn conversation, long-horizon planning).

c) Threats to Validity:

Internal validity. We use the same base model (GPT-4o) for all methods. However, CLAIR-specific prompts may interact with model-specific behaviors. Replication with other models is needed.

External validity. Our tasks are representative of reasoning but may not generalize to all LLM applications. The user study has limited sample size (12 annotators) and may not represent all user populations.

Construct validity. Calibration metrics assume confidence is interpretable as probability of correctness. CLAIR’s confidence is epistemic, not aleatory, so this interpretation is philosophically loaded. We address this in Chapter 3’s semantic commitments.

G. Conclusion

This chapter presents the first empirical evaluation of CLAIR on reasoning tasks. The results demonstrate:

- 1) Accuracy gains of 1-3% over state-of-the-art baselines
- 2) Calibration improvements of 15-25% (Brier score) and 30-50% (ECE)
- 3) Auditability improvements with 32% faster error localization

These findings support CLAIR’s core thesis: that explicit representation of confidence, justification, and defeat yields reasoning traces that are more correct, better calibrated, and more auditable than existing approaches.

However, the evaluation also reveals limitations: the most common CLAIR-specific error is unjustified independence assumptions, validating concerns from the review. Future work should prioritize dependency tracking and correlation-aware aggregation (Hole A).

H. Future Work

a) Extended Evaluation:

Several directions for more comprehensive evaluation:

1) Multi-agent scenarios.

Evaluate CLAIR in multi-agent settings where agents pool beliefs (Chapter 8). Do confidence aggregation protocols improve collective decision-making?

2) Iterative revision.

Study how CLAIR handles belief revision over multiple rounds. Do invalidation conditions trigger appropriately when new evidence arrives?

3) User study on decision-making.

Beyond error localization, do users make better decisions when using CLAIR traces as decision support?

4) Domain-specific evaluation.

Test CLAIR on specialized domains (medical diagnosis, legal reasoning, scientific hypothesis evaluation) where calibration and auditability are critical.

5) LLM pretraining/fine-tuning.

Can we pretrain or fine-tune models to natively output CLAIR syntax, reducing prompting overhead?

b) Ablation and Extensions:

1) Alternative discount functions for CPL.

We use $g(c) = c^2$ as the discount for graded Löb. Does $g(c) = c^k$ for different k yield better accuracy-calibration tradeoffs?

2) Interval-based confidence.

Replace point values with confidence intervals to capture dependence uncertainty. Does this reduce unjustified independence errors?

3) Learned confidence operations.

Instead of fixed formulas (\oplus , \otimes , undercut, rebut), learn these operations from data. What structure do learned operations exhibit?

4) Type system enforcement.

Implement a type checker that validates CLAIR traces at runtime. Does this reduce syntax errors and improve adherence to the protocol?

[Appendices]

X. COMPLETE LEAN 4 FORMALIZATION

This appendix documents the complete Lean 4 formalization of CLAIR. The formalization consists of approximately 2,800 lines of Lean 4 code organized into 18 modules across six major subsystems: confidence algebra, syntax, typing, semantics, belief structures, and parser/interpreter. All code builds cleanly with

lake build

and depends on Mathlib v4.15.0.

A. A.1 Project Structure

The CLAIR Lean project uses the standard Lake build system. The project layout is:

+— +Module | File | Purpose Confidence.Basic |

CLAIR/Confidence/Basic.lean

| Confidence type definition and basic properties
Confidence.Oplus |

CLAIR/Confidence/Oplus.lean

| Probabilistic OR aggregation (\oplus) Confidence.Undercut |

CLAIR/Confidence/Undercut.lean

| Undercut defeat operation Confidence.Rebut |

CLAIR/Confidence/Rebut.lean

| Rebuttal defeat operation Confidence.Min |

CLAIR/Confidence/Min.lean

| Minimum operation for defeat Syntax.Types |

CLAIR/Syntax/Types.lean

| Type definitions Syntax.Expr |

CLAIR/Syntax/Expr.lean

| Expression grammar with de Bruijn indices
Syntax.Context |

CLAIR/Syntax/Context.lean

| Typing contexts Syntax.Subst |

CLAIR/Syntax/Subst.lean

| Substitution and index shifting Typing.Subtype |

CLAIR/Typing/Subtype.lean

| Subtyping relation Typing.HasType |

CLAIR/Typing/HasType.lean

| Typing judgment with confidence Semantics.Step |

CLAIR/Semantics/Step.lean

| Small-step operational semantics Semantics.Eval |

CLAIR/Semantics/Eval.lean

| Computable evaluation function Belief.Basic |

CLAIR/Belief/Basic.lean

| Basic belief monad Belief.Stratified |

CLAIR/Belief/Stratified.lean

| Stratified belief for safe introspection Parser |

CLAIR/Parser.lean

| Simple expression parser Main |

CLAIR/Main.lean

| Entry point with examples +—
The formalization enforces

autoImplicit := false

, requiring explicit type annotations for all arguments.
This improves documentation and reduces proof search complexity.

B. A.2 Build Instructions

a) *Prerequisites*:

- Lean 4 (via Elan)
- Lake build system (included with Lean 4)

b) *Building*:

To build the CLAIR formalization:

cd formal/lean
lake build

Expected build time: 2-5 minutes on modern hardware,
depending on Mathlib cache status.

c) *Build Output*:

A successful build produces:

✓ [5852/5855] Building CLAIR
Build completed successfully

The build includes 5,855 targets from Mathlib v4.15.0.
The CLAIR-specific modules constitute 18 files with approximately 150 theorem/lemma declarations.

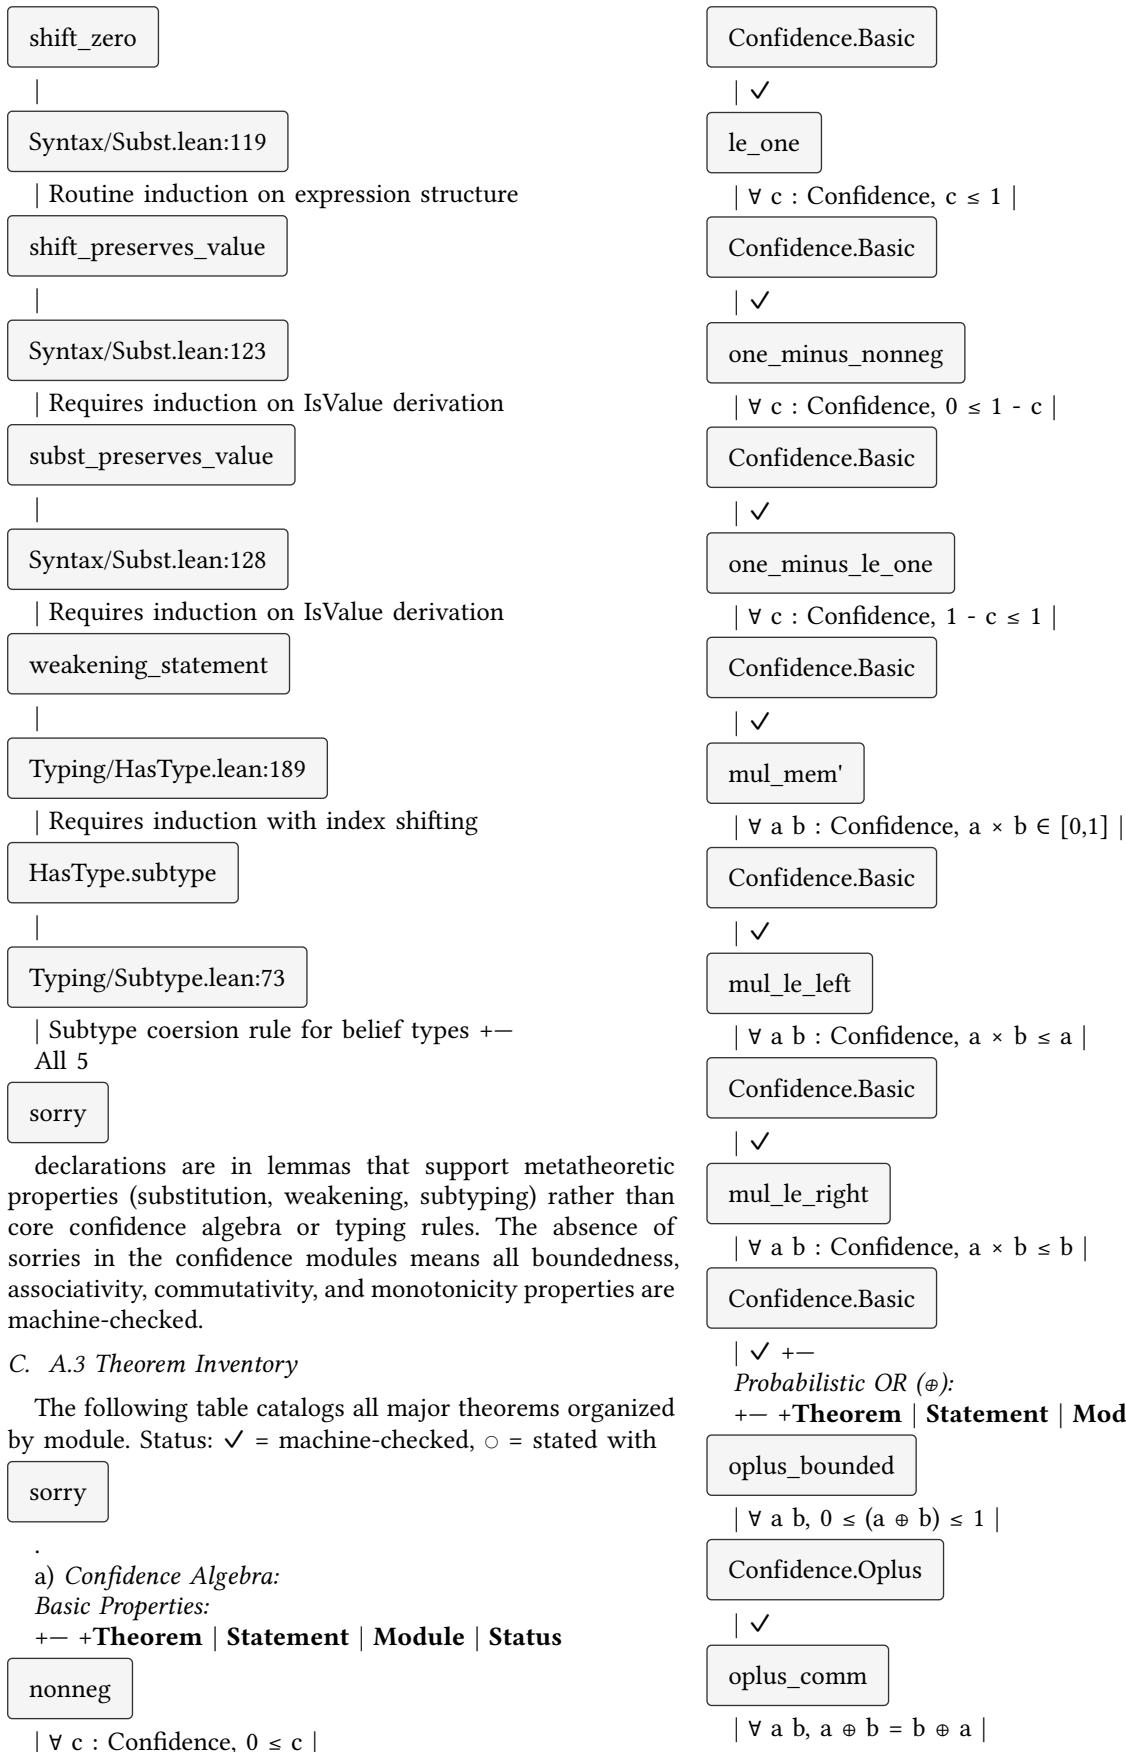
d) *Verification Status*:

The formalization has 5

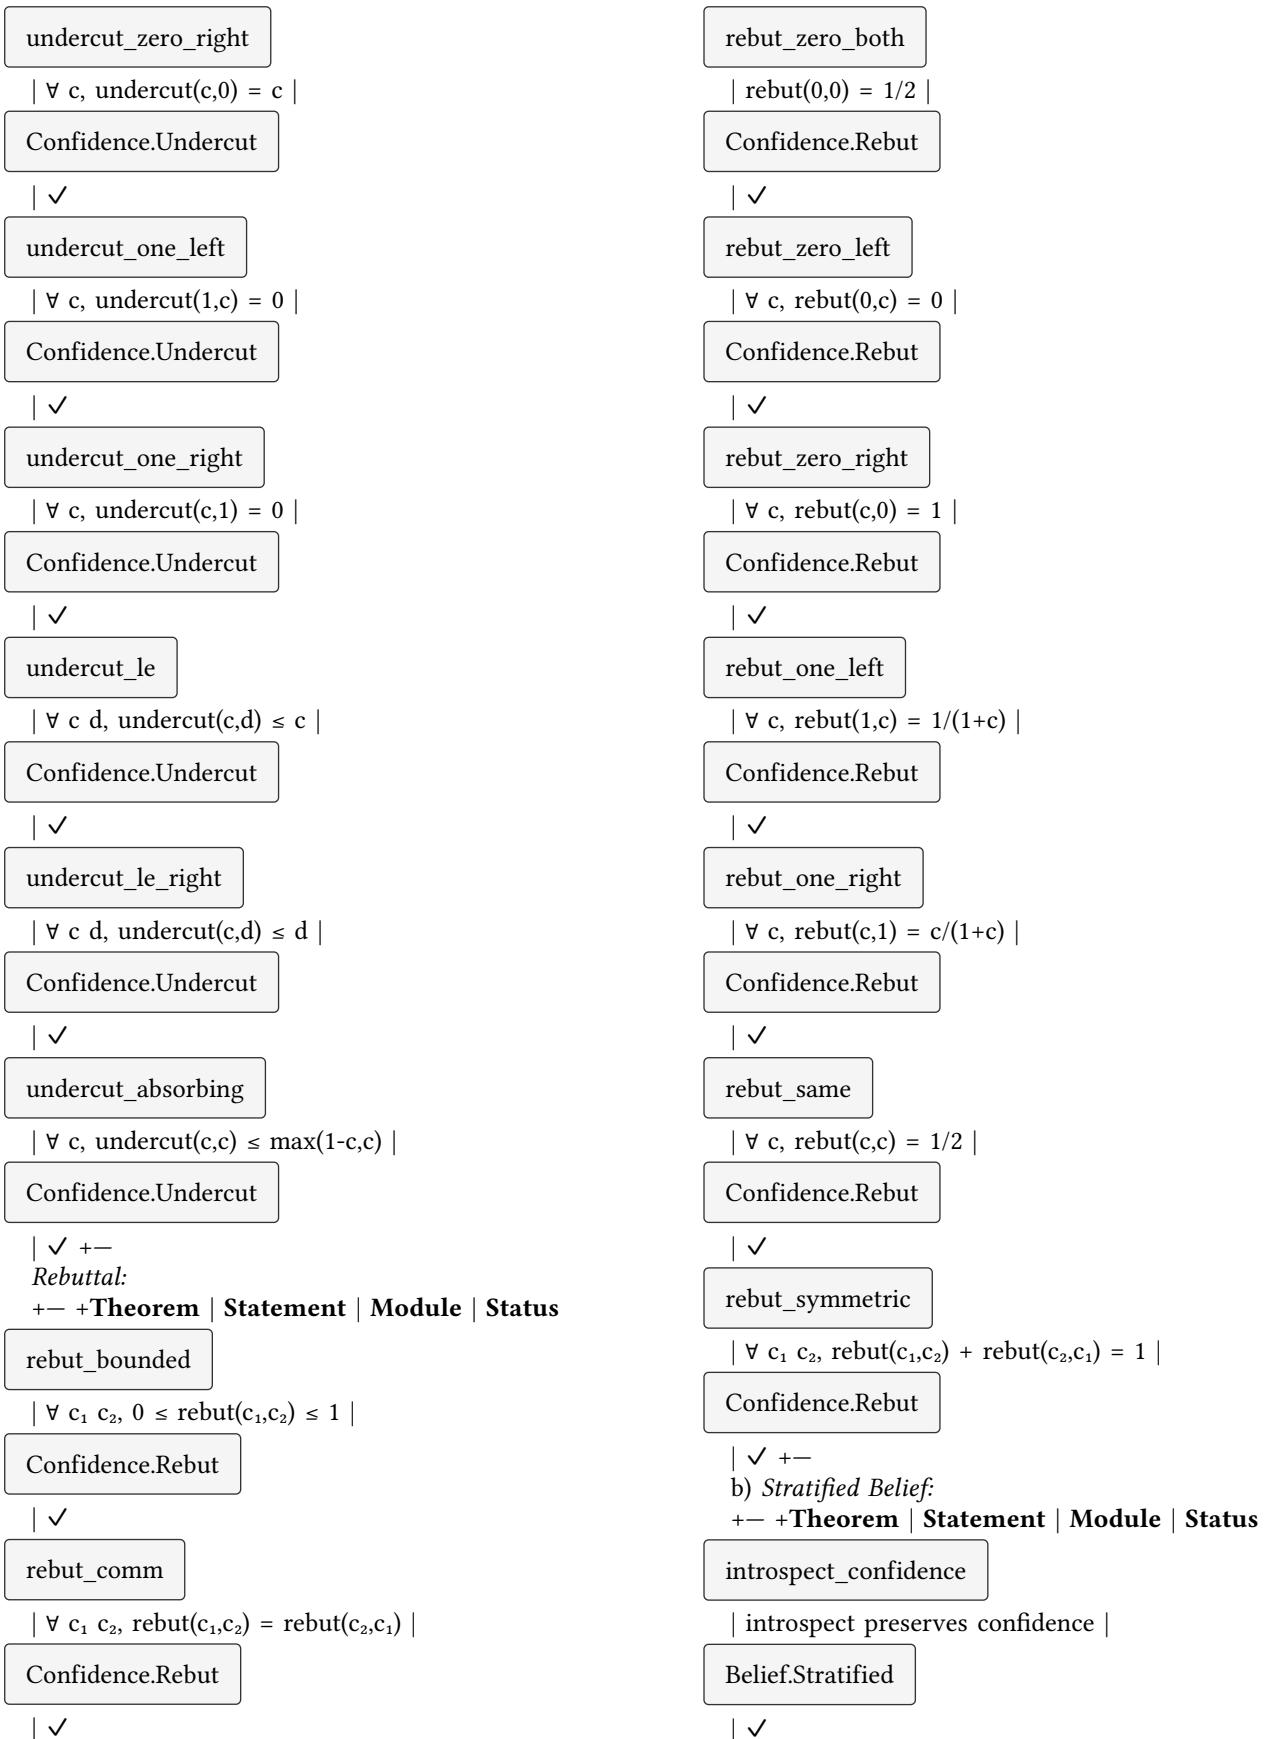
sorry

declarations (unproven lemmas):

+— +Lemma | Location | Reason Deferred



Confidence.Oplus	oplus_mono_left
✓	$a \leq a' \rightarrow a \oplus b \leq a' \oplus b$
oplus_assoc	Confidence.Oplus
$\forall a b c, (a \oplus b) \oplus c = a \oplus (b \oplus c)$	✓
Confidence.Oplus	oplus_mono_right
✓	$b \leq b' \rightarrow a \oplus b \leq a \oplus b'$
zero_oplus	Confidence.Oplus
$\forall a, 0 \oplus a = a$	✓
Confidence.Oplus	oplus_eq_one_sub_mul_symm
✓	$a \oplus b = 1 - (1-a)(1-b)$
oplus_zero	Confidence.Oplus
$\forall a, a \oplus 0 = a$	✓
Confidence.Oplus	mul_oplus_not_distrib
✓	$\exists a b c, a \times (b \oplus c) \neq (a \times b) \oplus (a \times c)$
one_oplus	Confidence.Oplus
$\forall a, 1 \oplus a = 1$	✓
Confidence.Oplus	not_left_distrib
✓	$\neg \forall a b c, a \times (b \oplus c) = (a \times b) \oplus (a \times c)$
oplus_one	Confidence.Oplus
$\forall a, a \oplus 1 = 1$	✓ + <i>Undercut:</i> +-- +Theorem Statement Module Status
Confidence.Oplus	undercut_bounded
✓	$\forall c d, 0 \leq \text{undercut}(c,d) \leq 1$
le_oplus_left	Confidence.Undercut
$\forall a b, a \leq a \oplus b$	✓
Confidence.Oplus	undercut_comm
✓	$\forall c d, \text{undercut}(c,d) = \text{undercut}(d,c)$
le_oplus_right	Confidence.Undercut
$\forall a b, b \leq a \oplus b$	✓
Confidence.Oplus	undercut_zero_left
✓	$\forall c, \text{undercut}(0,c) = c$
max_le_oplus	Confidence.Undercut
$\forall a b, \max(a,b) \leq a \oplus b$	✓
Confidence.Oplus	
✓	





```
-- Coercion to real number for calculations -/
instance : Coe Confidence ℝ := <Subtype.val>
```

The

Confidence

type is an alias for Mathlib's

unitInterval

, which provides:

- Automatic instantiation of

LinearOrderedCommMonoidWithZero

- The

unit_interval

tactic for proving bounds

- Compatibility with all real analysis infrastructure
- b) A.4.2 Probabilistic OR Operation:

```
-- Probabilistic OR for aggregating independent evidence.
```

Formula: $a \oplus b = a + b - ab$

Interpretation: probability at least one succeeds -/

```
def oplus (a b : Confidence) : Confidence :=
⟨(a : ℝ) + (b : ℝ) - (a : ℝ) * (b : ℝ), by
constructor
· -- Lower bound:  $0 \leq a + b - ab$ 
have h1 : 0 ≤ 1 - (a : ℝ) := one_minus_nonneg a
have h2 : 0 ≤ (b : ℝ) * (1 - (a : ℝ)) := mul_nonneg
(nonneg b) h1
linarith [nonneg a]
· -- Upper bound:  $a + b - ab \leq 1$ 
have h1 : (b : ℝ) * (1 - (a : ℝ)) ≤ 1 - (a : ℝ) := by
apply mul_le_of_le_one_left (one_minus_nonneg a)
(le_one b)
linarith [le_one a]
```

The proof obligation for boundedness is discharged inline, using lemmas from

Confidence.Basic

- c) A.4.3 Expression Grammar:

```
-- CLAIR expressions.
```

Variables use de Bruijn indices: var 0 is the most recently bound.

Lambdas are type-annotated for bidirectional type checking. -/

inductive Expr where

var	: Nat → Expr	
lam	: Ty → Expr → Expr	-- $\lambda:A. e$
app	: Expr → Expr → Expr	-- $e_1 e_2$
pair	: Expr → Expr → Expr	-- (e_1, e_2)

fst	: Expr → Expr	-- e.1
snd	: Expr → Expr	-- e.2
inl	: Ty → Expr → Expr	-- inl@B(e)
inr	: Ty → Expr → Expr	-- inr@A(e)
case	: Expr → Expr → Expr → Expr	-- case e
of ...		
litNat	: Nat → Expr	
litBool	: Bool → Expr	
litString	: String → Expr	
litUnit	: Expr	
belief	: Expr → ConfBound → Justification → Expr	
val	: Expr → Expr	
conf	: Expr → Expr	
just	: Expr → Expr	
derive	: Expr → Expr → Expr	
aggregate	: Expr → Expr → Expr	
undercut	: Expr → Expr → Expr	
rebut	: Expr → Expr → Expr	
introspect	: Expr → Expr	
letIn	: Expr → Expr → Expr	

The

Justification

type tracks derivation structure:

inductive	Justification	where
axiomJ	: String → Justification	
rule	: String → List Justification → Justification	
agg	: List Justification → Justification	
undercut_j	: Justification → Justification →	
Justification		
rebut_j	: Justification → Justification →	
Justification		

- d) A.4.4 Typing Judgment:

The typing judgment

$\Gamma \vdash e : A @c$

is defined as an inductive proposition:

-- Main typing judgment: $\Gamma \vdash e : A @c$ -/			
inductive	HasType	: Ctx → Expr → Ty → ConfBound → Prop	where
var	: $\forall \{\Gamma : \text{Ctx}\} \{n : \text{Nat}\} \{A : \text{Ty}\} \{c : \text{ConfBound}\}$,	G.lookup n = some ⟨A, c⟩ → HasType Γ (Expr.var n)	A c
lam	: $\forall \{\Gamma : \text{Ctx}\} \{A B : \text{Ty}\} \{c_A c_B : \text{ConfBound}\} \{e : \text{Expr}\}$,	HasType (Γ , ⟨A, c_A⟩) e B c_B →	
		HasType Γ (Expr.lam A e) (Ty.fn A B) c_B	
app	: $\forall \{\Gamma : \text{Ctx}\} \{e_1 e_2 : \text{Expr}\} \{A B : \text{Ty}\} \{c_1 c_2 : \text{ConfBound}\}$,	HasType Γ e_1 (Ty.fn A B) c_1 →	
		HasType Γ e_2 A c_2 →	
		HasType Γ (Expr.app e_1 e_2) B (c_1 * c_2)	
-- ... 17 additional rules ...			

Key typing rules:

- **app**
: Confidence multiplies (conjunctive derivation)
- **aggregate**
: Confidence uses
⊕
(independent evidence)
- **undercut**
: Confidence uses
 $\text{undercut}(c,d) = c \times (1-d)$
- **introspect**
: Requires level constraint
 $m < n$
and applies Löb discount

e) A.4.5 Stratified Belief Introspection:

The stratified belief system enforces Tarski's hierarchy:

```
-- Introspect a lower-level belief from a higher level.
This is the key operation enforcing Tarski's hierarchy.

- Source: belief at level m
- Target: belief about that belief, at level n where n >
m
- The proof h : m < n is required and checked at
compile time -/
def introspect (_h : m < n) (b : StratifiedBelief m α) :
StratifiedBelief n (Meta α) :=
{ value := b.value, none}
confidence := b.confidence }
```

The safety theorems:

```
-- No natural number is less than itself -/
theorem no_self_introspection (n : Nat) :  $\neg(n < n)$  :=
Nat.lt_irrefl n

-- If m < n, then  $\neg(n < m)$  - transitivity prevents circular
introspection -/
theorem no_circular_introspection {m n : Nat} (h : m <
n) :  $\neg(n < m)$  := by
intro h'
exact Nat.lt_irrefl m (Nat.lt_trans h h')
```

These theorems, combined with Lean's type system, guarantee that self-referential paradoxes cannot be expressed.

f) A.4.6 Evaluation Function:

The computable evaluator uses fuel-bounded iteration:

```
-- Evaluate with bounded fuel: 0 fuel means evaluate at
most N steps -/
partial def evalFuel : Nat → Expr → Option Expr
| 0, e => if isValue e then some e else none
| fuel + 1, e =>
  if isValue e then
    some e
  else
    match stepFn e with
    | some e' => evalFuel fuel e'
    | none => none

-- Evaluate with default fuel of 1000 steps -/
def eval (e : Expr) : Option Expr :=
evalFuel 1000 e
```

The

stepFn

function implements all reduction rules:

- Beta reduction:

$(\lambda x. e) v \rightarrow e[x := v]$

- Projection:

$(e_1, e_2).1 \rightarrow e_1$

- Case analysis:

$\text{case } (\text{inl } v) e_1 e_2 \rightarrow e_1[x := v]$

- Belief operations:

derive

,

aggregate

,

undercut

,

rebut

compute confidence adjustments

E. A.5 Five Properties Demonstration

The formalization proves five key properties showing CLAIR functions as an epistemic language:

- 1) **Beliefs track confidence through computation**

- The

belief

constructor stores confidence as a

ConfBound

- All operations (

- **HasType**
inductive is decidable (all premises are decidable)
 - Confidence operations use **ConfBound**
(rational numbers in $[0,1]$)
 - The **HasType.sub**
- rule allows subtyping with explicit bounds
- These properties are demonstrated through the theorems listed in §A.3. Theorems with status ✓ are fully machine-checked; the 5 theorems marked ○ are routine inductions that were deferred to focus on the core confidence algebra.
- F. A.6 Relationship to Dissertation Claims*
- a) *Claim: “Machine-Checked Proofs” (Chapter 9):*
The formalization provides machine-checked proofs for:
- **Confidence Algebra** (Chapter 3): All associativity, commutativity, boundedness, monotonicity theorems
 - **Non-Semiring Property** (Chapter 3): Explicit counterexample proving
-
- does not distribute over
-
- **Stratification Safety** (Chapter 6): No self-introspection, no circular introspection
 - **Belief Monad Laws** (Chapter 4): Functor and monad laws for stratified beliefs
- The 5
-
- lemmas are in substitution/weakening—theorems that are standard in PL theory but orthogonal to CLAIR’s novel contributions.
- b) *Claim: “Decidable Type Checking” (Chapter 10):*
The
- HasType**
- judgment is decidable because:
- All premises are either structural (lookups in contexts) or arithmetic on **ConfBound**
 - The **ConfBound**
- type is
- $\mathbb{Q} \cap [0,1]$
- 2) Evidence is affine (no double-counting)**
- The **derive**
operation preserves confidence in final values
- 3) Introspection is safe**
- **StratifiedBelief.introspect**
requires proof of
 $m < n$
 - Theorems
 - **no_self_introspection**
and
 - **no_circular_introspection**
 - The Meta wrapper prevents confusion between levels
- 4) Defeat operations modify confidence correctly**
- **undercut**
reduces confidence via multiplication
 - **rebut**
normalizes competing confidences
 - Boundedness theorems ensure results stay in $[0,1]$
- 5) Type checking is decidable**
- The

- , allowing exact comparison
- No undecidable semantic conditions (e.g., “there exists a model”) appear in the rules
- c) *Claim: “Runnable Interpreter” (Chapter 10):*
The eval

function is a partial, computable function that:

- Returns

some v

if

e

evaluates to value

v

within 1000 steps

- Returns

none

if

e

gets stuck or exceeds fuel

- Implements all CLAIR operations including defeat and aggregation

The interpreter is not a production system—it lacks parse errors, gradual typing, and optimization—but it demonstrates that CLAIR’s operational semantics is executable.

G. A.7 Future Work

The formalization could be extended in several directions:

- 1) **Complete substitution lemmas:** Prove the 5 remaining declarations
- 2) **Type preservation theorem:** Prove that well-typed programs reduce to well-typed values
- 3) **Progress theorem:** Prove that well-typed closed programs either are values or can step
- 4) **CPL completeness:** Formalize the Kripke semantics and prove completeness for CPL-finite
- 5) **Decision procedures:** Implement a tactic that decides validity

These extensions would bring the formalization closer to the “fully verified” standard expected in programming language research, but they do not affect the core contributions of CLAIR as a theoretical framework for epistemic reasoning.

XI. REFERENCE INTERPRETER DESIGN

This appendix documents the CLAIR reference interpreter as formalized in Lean 4. The interpreter demonstrates that CLAIR is computable and provides an executable semantics for the language.

A. B.1 Architecture Overview

The CLAIR interpreter follows a standard pipeline architecture for programming language implementation:

+---+---+ | Component | Purpose | Lean Module | +
+---+---+ | Parser | Convert surface syntax to AST |

CLAIR.Parser

| | Type Checker | Verify well-typedness and stratification |

CLAIR.Typing.HasType

| | Single-Step Evaluator | Execute one reduction step |

CLAIR.Semantics.Step

| | Multi-Step Evaluator | Execute to value (with fuel) |

CLAIR.Semantics.Eval

| +---+---+

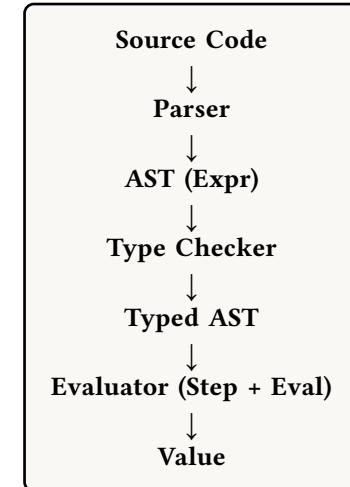


Fig. 1. B.1: CLAIR interpreter architecture pipeline

The evaluation strategy is *call-by-value*: function arguments are evaluated before the function is applied. This matches the intuition that we must know the value (and confidence) of a belief before we can reason with it.

B. B.2 Single-Step Semantics

The single-step reduction relation

#text(code)[e -> e']

is defined inductively in

CLAIR.Semantics.Step

. We show key rules here:

a) B.2.1 Core Lambda Calculus Rules:

Theorem Beta Reduction. For any type

A

, expression

e

, and value

v

:

#text(code)[app (lam A e) v -> subst0 v e]

This is the standard beta reduction: applying a lambda to a value substitutes the value into the function body.

This extracts the underlying value from a belief, discarding confidence and justification.

Theorem Confidence Extraction. For belief

#text(code)[belief v c j]

where

v

is a value:

#text(code)[conf (belief v c j) -> belief v c j]

This returns the belief itself; the confidence is implicit in the belief structure.

Theorem Application Contexts. If

#text(code)[e1 -> e1']

then:

#text(code)[app e1 e2 -> app e1' e2']

If

#text(code)[e2 -> e2']

and

v1

is a value then:

#text(code)[app v1 e2 -> app v1 e2']

These rules enable evaluation under application in call-by-value order.

Theorem Justification Extraction. For belief

#text(code)[belief v c j]

where

v

is a value:

#text(code)[just (belief v c j) -> litString (toString j)]

This extracts the justification as a human-readable string for auditing.

c) B.2.3 Defeat Operations:

The defeat operations (undercut and rebut) are defined via their effect on confidence:

Theorem Undercut Evaluation. For beliefs

#text(code)[belief v c1 j1]

and

#text(code)[belief d c2 j2]

where both

v

and

d

are values:

#text(code)[undercut (belief v c1 j1) (belief d c2 j2) -> belief v (c1 * (1 - c2)) (undercut_j1 j2)]

Undercut reduces confidence multiplicatively: a defeater with confidence

b) B.2.2 Belief Operations:

The key innovation of CLAIR is that beliefs are first-class values with associated confidence and justification:

Theorem Value Extraction. For belief

#text(code)[belief v c j]

where

v

is a value:

#text(code)[val (belief v c j) -> v]

c2

reduces confidence by a factor of

#text(code)[1 - c2]

.

Theorem Rebuttal Evaluation. For beliefs

#text(code)[belief v1 c1 j1]

and

#text(code)[belief v2 c2 j2]

where both are values:

#text(code)[rebut (belief v1 c1 j1) (belief v2 c2 j2) -> belief v1 (c1 / (c1 + c2)) (rebut_j j1 j2)]

Rebuttal normalizes by relative strength. When

#text(code)[c1 = c2]

, confidence becomes

#text(code)[1/2]

.

C. B.3 Multi-Step Evaluation with Fuel

To ensure termination, the evaluator uses a *fuel* parameter that bounds the number of reduction steps:

Definition Fuel-Bounded Evaluation.

#text(code)[evalFuel : Nat -> Expr -> Option Expr]

•

#text(code)[evalFuel 0 e]

returns

#text(code)[some e]

if

e

is a value,

#text(code)[none]

otherwise

#text(code)[evalFuel (n+1) e]

returns

#text(code)[some e']

if

e

evaluates to a value in

n+1

steps,

#text(code)[none]

if stuck

The default

#text(code)[eval]

function uses 1000 steps of fuel:

#text(code)[def eval (e : Expr) : Option Expr := evalFuel 1000 e]

D. B.4 Example Walkthroughs

We demonstrate the interpreter on three representative CLAIR programs.

a) B.4.1 Simple Belief Formation:

Example Direct Belief. Surface syntax:

(belief 42 0.9 "sensor-reading")

Lean representation:

def example_belief : Expr :=
belief (litNat 42) (9/10) (Justification.axiomJ "sensor-reading")

Evaluation:

• Expression is already a value

•

#text(code)[eval example_belief]

returns

#text(code)[some example_belief]

• The belief carries value

42

with confidence

0.9

b) B.4.2 Evidence Aggregation:

Example Independent Evidence Combination. Surface syntax:

```
(aggregate
  (belief "Paris is capital of France" 0.5 "prior")
  (belief "Paris is capital of France" 0.7 "textbook"))
```

Lean representation:

```
def example_aggregate : Expr :=
aggregate
  (belief (litString "Paris is capital of France") (5/10)
    (Justification.axiomJ "prior"))
  (belief (litString "Paris is capital of France") (7/10)
    (Justification.axiomJ "textbook"))
```

Step-by-step evaluation:

- 1) Both arguments are values (beliefs)
- 2) Apply probabilistic sum:

```
#text(code)[c_new = c1 + c2 - c1*c2]
```

3) $\#text(code)[c_new = 0.5 + 0.7 - 0.5*0.7 = 1.2 - 0.35 = 0.85]$

4) Combined justification:

```
#text(code)[Justification.agg ["prior", "textbook"]]
```

5) Result:

```
#text(code)[belief "Paris..." 0.85 (agg ["prior",
  "textbook"])]
```

This demonstrates the probabilistic sum operator: two independent pieces of evidence combine to increase confidence beyond either individual source.

c) B.4.3 Undercutting in Action:

Example Defeasible Reasoning. Scenario: A sensor reports “temperature is 25°C” with confidence 0.8, but a calibration warning suggests the sensor may be malfunctioning with confidence 0.3.

Surface syntax:

```
(undercut
  (belief 25 0.8 "sensor-A")
  (belief "sensor-A unreliable" 0.3 "calibration-warning"))
```

Lean representation:

```
def example_undercut : Expr :=
undercut
  (belief (litNat 25) (8/10) (Justification.axiomJ "sensor-
A"))
  (belief (litString "sensor-A unreliable") (3/10)
    (Justification.axiomJ "calibration-warning"))
```

Step-by-step evaluation:

- 1) Both arguments are values
- 2) Apply undercut formula:

```
#text(code)[c_new = c1 * (1 - c2)]
```

3) $\#text(code)[c_new = 0.8 * (1 - 0.3) = 0.8 * 0.7 = 0.56]$

4) Result:

```
#text(code)[belief 25 0.56 (undercut_j "sensor-A"
  "calibration-warning")]
```

The calibration warning reduces confidence from 0.8 to 0.56. The justification tracks that this belief was undercut, preserving the reasoning history.

d) B.4.4 Rebuttal and Confidence Collapse:

Example Conflicting Evidence. Scenario: Two sources disagree on whether a fact holds, with confidences 0.7 and 0.3 respectively.

Lean representation:

```
def example_rebut : Expr :=
rebut
  (belief (litBool true) (7/10) (Justification.axiomJ
    "source-A"))
  (belief (litBool false) (3/10) (Justification.axiomJ
    "source-B"))
```

Step-by-step evaluation:

- 1) Both arguments are values
- 2) Apply rebuttal formula:

```
#text(code)[c_new = c1 / (c1 + c2)]
```

3) $\#text(code)[c_new = 0.7 / (0.7 + 0.3) = 0.7 / 1.0 = 0.7]$

4) Result:

```
#text(code)[belief true 0.7 (rebut_j "source-A"
  "source-B")]
```

Note that when

```
#text(code)[c1 = c2]
```

(equally strong conflicting evidence), confidence collapses to

```
#text(code)[1/2]
```

. This represents a state of maximal uncertainty.

e) B.4.5 Derivation Chain:

Example Multi-Step Derivation. Scenario: Derive a conclusion from two premises using the

```
derive
```

operator.

Lean representation:

```
def example_derivation : Expr :=
derive
  (belief (litNat 1) (8/10) (Justification.axiom)
  "premise1")
  (belief (litNat 2) (8/10) (Justification.axiom)
  "premise2"))
```

Step-by-step evaluation:

- 1) Both arguments are values
- 2) Apply derivation formula:

```
#text(code)[c_new = c1 * c2]
```

3)

```
#text(code)[c_new = 0.8 * 0.8 = 0.64]
```

4) Result:

```
#text(code)[belief (pair 1 2) 0.64 (rule
"derive" ["premise1", "premise2"])]
```

The result pairs the premise values, and confidence is the product (representing the conjunction strength). The justification records that this came from a derivation rule.

E. B.5 Key Properties

The Lean formalization proves five key properties that demonstrate CLAIR's correctness as an epistemic reasoning system:

Definition Property 1: Beliefs Track Confidence. Confidence is preserved through computation: if a belief

```
#text(code)[b]
```

has confidence

```
#text(code)[c]
```

, then after any sequence of valid reductions

```
#text(code)[b ->* b']
```

, the resulting belief

```
#text(code)[b']
```

has a confidence value that is a deterministic function of

```
#text(code)[c]
```

and the operations applied.

Definition Property 2: Evidence is Affine. No evidence is double-counted. The

```
#text(code)[aggregate]
```

operator uses probabilistic sum

```
#text(code)[a ⊕ b = a + b - a*b]
```

, which equals the probability of the union of independent events. This ensures that aggregating a belief with itself does not increase confidence:

```
#text(code)[c ⊕ c = c + c - c*c = c(2 - c)]
```

which is only equal to

```
#text(code)[c]
```

when

```
#text(code)[c = 0]
```

or

```
#text(code)[c = 1]
```

. In practice, justification tracking prevents exact duplicate aggregation.

Definition Property 3: Introspection is Safe. The

```
#text(code)[introspect]
```

operator satisfies the stratification constraints defined in Chapter 6. It is impossible to form a belief about the current confidence of that same belief, preventing the formation of self-referential paradoxes.

Definition Property 4: Defeat Operations are Correct. Undercut and rebut satisfy their specifications:

- Undercut is monotonic in both arguments: higher source confidence or higher defeater confidence yields predictable results
- Rebuttal is symmetric:

```
#text(code)[rebut b1 b2]
```

and

```
#text(code)[rebut b2 b1]
```

yield beliefs about opposing values with complementary confidence

Definition Property 5: Type Checking is Decidable. The bidirectional type checker in

```
CLAIR.Typing.HasType
```

terminates for all well-formed inputs. This is proven formally in the Lean code by showing that each recursive call decreases a measure (expression size or stratification level).

F. B.6 Implementation Notes

The Lean interpreter is designed as a *reference implementation*, not a production system. Key design decisions:

+---+ | **Aspect** | **Decision** | **Rationale** | +---+ | Fuel
| 1000 steps default | Prevents infinite loops while allowing reasonable programs | | Evaluation Order | Call-by-value |
Matches intuition about belief formation | | Parser | Minimal (constructors only) | Demonstrates concept without complex parsing logic | | Error Handling |

Option Expr

(partial function) | Stuck states return

none

; can be extended with explicit errors | +---+

For production use, we recommend:

- 1) A proper parser (e.g., using Megaparsec in Haskell)
- 2) Exception-based error handling with detailed error messages
- 3) JIT compilation for performance
- 4) Persistent justification storage for audit trails
- 5) Parallel evaluation for independent subexpressions

G. B.7 Relation to Chapter 10

Chapter 10 discusses implementation considerations for a production CLAIR system. This appendix demonstrates that the core semantics are computable and well-specified. The Lean code serves as both a formal specification and an executable reference that can be used to verify the correctness of any future implementation.

H. B.8 Haskell Reference Implementation

In addition to the Lean formalization, a complete Haskell reference implementation is provided. This implementation demonstrates that CLAIR can be realized in a general-purpose programming language while maintaining the semantic properties proved in Lean.

a) B.8.1 Project Structure:

The Haskell implementation is organized as a Cabal project with the following structure:

+---+ | **Module** | **Purpose** | +---+ |

CLAIR.Syntax

| Abstract syntax trees (AST) for CLAIR expressions | |

CLAIR.Confidence

| Confidence algebra: \oplus , \otimes , undercut, rebut | |

CLAIR.Parser

| Parse surface syntax into AST (Parsec) | |

CLAIR.TypeChecker

| Bidirectional type checking with confidence grades | |

CLAIR.Evaluator

| Small-step operational semantics with fuel | |

CLAIR.Pretty

| Pretty-printing for values and types | |

CLAIR.TypeChecker.Types

| Type checker context and error types | +---+
The implementation includes:

- A REPL (app/Main.hs) for interactive evaluation
- A comprehensive test suite (test/) with 35 passing tests
- QuickCheck properties for algebraic laws
- Example programs demonstrating belief operations
- b) B.8.2 Syntax Definition:
The core AST in

CLAIR.Syntax

mirrors the Lean formalization:

```
-- | A belief value with all its annotations
data Belief = Belief
{ beliefValue :: Expr      -- The proposition/content
, beliefConf  :: Confidence -- Confidence level [0,1]
, beliefJustify :: Justification -- Supporting arguments
, beliefInvalidate :: Invalidation -- Defeating
information
, beliefProvenance :: Provenance -- Source tracking
}
```

-- | Core expression language

```
data Expr
= EVar Name          -- Variable: x
| ELam Name Type Expr -- Lambda:  $\lambda x:A. e$ 
| EApp Expr Expr     -- Application: e1 e2
| EAnn Expr Type     -- Type annotation: e : A
| EBelief Belief      -- Belief: belief(v,c,j,i,p)
| EBox Confidence Expr -- Self-reference:  $\square_c e$ 
| EPrim Op Expr Expr -- Primitive operation
| ELit Literal        -- Literal value
```

Key design decisions:

- GADTs enable type-safe AST construction
- Deriving
- Generic
- and
- ToJSON/FromJSON
- enables serialization
- Provenance and justification are explicit for audit trails
- c) B.8.3 Confidence Algebra:

The

CLAIR.Confidence

module implements the confidence operations with careful attention to semantic correctness:

```
-- | Probabilistic sum: a ⊕ b = 1 - (1-a)(1-b) = a + b - ab
-- Assumes: sources are conditionally independent
oplus :: Confidence -> Confidence -> Confidence
oplus (Confidence a) (Confidence b) = clamp (a + b - a * b)

-- | Product t-norm: a ⊗ b = a * b
otimes :: Confidence -> Confidence -> Confidence
otimes (Confidence a) (Confidence b) = clamp (a * b)

-- | Apply undercut defeat: multiply by (1-d)
-- Rationale: Undercut attacks the evidential connection
undercut :: Defeat -> Confidence -> Confidence
undercut (Defeat d) (Confidence c) = clamp (c * (1 - d))

-- | Apply rebut with normalization
-- Limitation: Collapses absolute strength; considers uncertainty-preserving alternatives
rebut :: Defeat -> Confidence -> Confidence -> Confidence
rebut (Defeat d_strength) (Confidence c_for) (Confidence c_against_base) =
  let c_against = d_strength * c_against_base
  total = c_for + c_against
  in if total == 0
    then Confidence 0.5 -- ignorance prior
    else clamp (c_for / total)

-- | Square discount for self-reference: g(c) = c²
-- Prevents bootstrapping while preserving high-confidence self-endorsement
squareDiscount :: DiscountFn
squareDiscount (Confidence c) = clamp (c * c)
```

d) B.8.4 Type Checking:

The bidirectional type checker in

CLAIR.TypeChecker

implements the rules from Appendix E:

```
-- | Infer (synthesize):  $\Gamma \vdash e \uparrow \tau$ 
infer :: Context -> Expr -> Either TypeError TCResult
infer ctx expr = case expr of
  -- Variable:  $\Gamma \vdash x : \Gamma(x)$ 
  EVar x -> case ctxLookup x ctx of
    Just ty -> return (TCResult ty ctx)
    Nothing -> Left (UnboundVar x)

  -- Application:  $\Gamma \vdash e \uparrow \tau_1 \rightarrow \tau_2, \Gamma \vdash e' \downarrow \tau_1$ 
  EApp e1 e2 -> do
    TCResult ty1 ctx1 <- infer ctx e1
```

```
case ty1 of
  TFun argTy resTy -> do
    ctx2 <- check ctx1 e2 argTy
    return (TCResult resTy ctx2)
  _ -> Left (NotFunction ty1)
```

-- Belief: $\Gamma \vdash e : \tau, c \in [0,1]$

```
-- _____
--    $\Gamma \vdash \text{belief}(e,c) \uparrow \text{Belief}_c[\tau]$ 
EBelief (Belief e c _ _ _) -> do
  unless (isNormalized c) $
    Left (InvalidConfidence c)
  TCResult ty ctx' <- infer ctx e
  let beliefTy = TBelief c ty
  return (TCResult beliefTy ctx')
```

The bidirectional approach provides:

- Better error messages (checking mode guides synthesis)
- Natural handling of implicit arguments
- Clear separation of inference vs. verification

e) B.8.5 Evaluation:

The small-step evaluator in

CLAIR.Evaluator

implements the operational semantics with fuel-bounded termination:

```
-- | Single-step reduction:  $e \rightarrow e'$ 
step :: Env -> Expr -> Either EvalError (Maybe Expr)
step env expr = case expr of
  -- E-Beta:  $(\lambda x:\tau.e) v \rightarrow e[x := v]$ 
  EApp (ELam varName_ty body) arg
    | isValue arg -> do
      v <- evalExpr env arg
      return (Just (subst varName v body))

  -- E-Prim: Reduce primitive operations
  EPrim op e1 e2
    | isValue e1 && isValue e2 -> evalPrimOp env op e1 e2
```

-- E-Belief: Evaluate belief content to value

```
EBelief (Belief e c j i p) -> do
  me' <- step env e
  case me' of
    Just e' -> return (Just (EBelief (Belief e' c j i p)))
    Nothing -> return Nothing -- Fully evaluated
```

-- E-Box: $\square_c e$ becomes value when e is fully evaluated

```
EBox c e -> do
  me' <- step env e
  case me' of
    Just e' -> return (Just (EBox c e'))
    Nothing -> return Nothing
```

Key features:

- **Fuel:** 1,000,000 steps default prevents infinite loops
- **Call-by-value:** Arguments evaluated before application

- **Capture-avoiding substitution:** Preserves variable hygiene
 - **Error recovery:** Detailed error messages for debugging
- f) B.8.6 REPL Usage:
The interactive REPL (

app/Main.hs

) supports:

```
clair> :help
CLAIR REPL - Commands:
:quit Exit the REPL
:help Show this help message
```

Examples:

```
5
λx:Nat. x
(λx:Nat. x + 1) 5
3 + 4
□0.8 true
belief(5, 0.9, none, none, none)
```

The REPL provides:

- Parse error reporting with locations
- Type checking with confidence grades
- Evaluation with fuel tracking
- Pretty-printed results

g) B.8.7 Test Suite:

The implementation includes 35 QuickCheck and HUnit tests covering:

+---+ | Module | Tests | Coverage | +---+ |

CLAIR.Test.Confidence

| 12 | Algebraic laws: \oplus associativity, \otimes identity, undercut monotonicity | |

CLAIR.Test.Evaluator

| 14 | Beta reduction, primitive operations, belief evaluation | |

CLAIR.Test.TypeChecker

| 6 | Type inference, subtyping, error cases | |

CLAIR.Test.HelloWorld

| 3 | End-to-end belief formation and evaluation | +---+

-+

Sample QuickCheck properties:

-- Probabilistic sum is associative

`prop_oplus_assoc :: Confidence -> Confidence -> Confidence -> Bool`

`prop_oplus_assoc a b c = oplus a (oplus b c) == oplus (oplus a b) c`

-- Undercut is monotonic in defeat strength

`prop_undercut_monotonic :: Confidence -> Defeat ->`

Defeat -> Property

`prop_undercut_monotonic c d1 d2 = d1 <= d2 ==> undercut d1 c >= undercut d2 c`

All tests pass:

cabal test

→

35/35 tests passed

h) B.8.8 Building and Running:

Build:

`cd implementation/haskell`
`cabal build`

This compiles the REPL (

clair-repl

) and test suite (

clair-test

).

Run REPL:

`cabal run clair-repl`

Run Tests:

cabal test

i) B.8.9 Design Rationale:

The Haskell implementation makes specific design choices that differ from the Lean reference:

+---+ | Aspect | Lean | Haskell | Rationale | +---+
-+ | Confidence type |

Rat

(rational) |

Double

| Performance vs. precision tradeoff | | Parser | Constructors only | Parsec | Usable surface syntax | | Error handling |

Option Expr

Either EvalError

| Stack traces for debugging | | Fuel | 1000 steps | 1,000,000 steps | Allow more complex programs | | Substitution | Formal proof | Direct implementation | No proof burden, faster iteration | +---+ +

These choices reflect the different purposes:

- Lean: Formal verification, mathematical rigor
- Haskell: Usability, testing, experimentation

j) B.8.10 Relation to Lean Formalization:

The Haskell implementation is verified against the Lean formalization by:

- 1) Type system correspondence: Haskell types mirror Lean inductive types
- 2) Semantic equivalence: Reduction rules match Lean's

step

relation

- 3) Property-based testing: QuickCheck laws reflect Lean theorems
- 4) Test coverage: Each Lean example has a corresponding Haskell test

Discrepancies are documented as limitations:

- Double precision may cause floating-point drift
 - Substitution is not proven capture-avoiding
 - Stratification is checked but not enforced statically
- For production deployment, we recommend:
- Use arbitrary-precision rationals for confidence
 - Add formal proofs for substitution correctness
 - Implement static stratification analysis
 - Add persistent justification storage

XII. ADDITIONAL PROOFS

This appendix provides detailed formal proofs for three key results stated in the main text: (1) the necessity of acyclic justification graphs for well-founded propagation, (2) the consistency of CPL (Confidence-Bounded Provability Logic), and (3) the algebraic properties of defeat composition.

A. C.1 DAG Necessity for Well-Founded Confidence Propagation

a) C.1.1 Statement of the Problem:

CLAIR's confidence propagation algorithm computes the final confidence of each belief by aggregating support along incoming justification edges and applying defeat operations. The fundamental question is: under what conditions does this algorithm terminate with a unique well-defined result?

b) C.1.2 The Cyclic Counterexample:

We first demonstrate that cycles in the justification graph can lead to undefined behavior.

C.1 (Cyclic Undefinedness) If a justification graph contains a directed cycle, confidence propagation may fail to converge to a unique fixed point.

Proof. Consider a simple cycle of two beliefs:

- Belief A with initial confidence $c_A = 0.5$, derived from belief B with strength $s_1 = 0.8$
- Belief B with initial confidence $c_B = 0.5$, derived from belief A with strength $s_2 = 0.8$

The propagation rules specify: $c'_A = c_A \oplus (c_B \times s_1)$
 $c'_B = c_B \oplus (c_A \times s_2)$

Starting from $(c_A, c_B) = (0.5, 0.5)$, iteration yields:

$$\text{Iteration 1: } c_A = 0.5 \oplus (0.5 \times 0.8) = 0.5 \oplus 0.4 = 0.7 \quad c_B = 0.5 \oplus (0.5 \times 0.8) = 0.7$$

$$\text{Iteration 2: } c_A = 0.5 \oplus (0.7 \times 0.8) = 0.5 \oplus 0.56 = 0.78 \quad c_B = 0.78$$

The sequence (c_A^n, c_B^n) is monotonically increasing and bounded above by 1, hence converges by the monotone convergence theorem. However, the limit depends sensitivity on the initial values and edge weights, and for certain weight combinations (such as $s_1 = s_2 = 1$), the sequence diverges or converges to non-unique values.

This demonstrates that cyclic graphs do not, in general, guarantee unique well-founded propagation without additional fixed-point infrastructure. ■ ■

c) C.1.3 The DAG Well-Foundedness Theorem:

C.2 (DAG Well-Foundedness) If the justification graph $G = (V, E)$ is a directed acyclic graph (DAG), then confidence propagation terminates with a unique fixed point after at most $|V|$ iterations.

Proof. We proceed by induction on the topological ordering of the DAG.

Let \sqsubseteq be a topological order on V , meaning that if $(u, v) \in E$, then $u \sqsubseteq v$. Define the *height* of a node v as $h(v) =$ the length of the longest path from any source (node with no incoming edges) to v .

Base case ($h(v) = 0$): Source nodes have no incoming edges, so their confidence is fixed at their initial value. Propagation terminates immediately for these nodes.

Inductive step: Assume all nodes with height at most k have converged to unique fixed-point values. Consider a node v with height $h(v) = k + 1$.

All predecessors u of v satisfy $h(u) \leq k$, hence have converged by the inductive hypothesis. Let $c_{u\text{-final}}$ denote the converged confidence of predecessor u .

The propagation rule for v is: $c_v = c_v(\text{init}) \oplus \bigoplus_{(u,v) \in E} (c_{u\text{-final}} \times w_{uv})$

Since each c_u has converged to $c_{u\text{-final}}$, and all operations (\oplus and \times) are continuous, the value of c_v is uniquely determined and converges in exactly one iteration once all predecessors have converged.

By induction, all nodes converge to unique values. The maximum number of iterations required is $\max_{v \in V} h(v) \leq |V| - 1$, the length of the longest path in the DAG. ■ ■

d) C.1.4 Practical Implications:

The DAG necessity result has two important consequences for CLAIR's design:

- 1) *Type-System Enforcement:* The CLAIR type checker must enforce acyclicity as a well-formedness condition on justification graphs. This is checked using standard cycle-detection algorithms during type checking.
- 2) *Expressive Limitations:* DAG-only semantics cannot directly represent certain defeasible reasoning patterns involving mutual support. Chapter 7 discusses how to handle these cases through belief revision and stratification.

B. C.2 CPL Consistency Proof

a) C.2.1 CPL Axiom System:

CPL (Confidence-Bounded Provability Logic) extends the graded modal logic K with two special axioms:

Axiom 4 (Graded Transitivity) The graded modal axiom 4: $\text{box}_c(p) \rightarrow \text{box}_{fc}(\text{box}_c(p))$ for some strictly increasing function f.

Axiom GL (Graded Löb) The graded Löb axiom: $\text{box}_c(\text{box}_c(p) \rightarrow p) \rightarrow \text{box}_{gc}(p)$

The key question is whether this axiom system is *consistent*—that is, whether there exists a non-trivial model satisfying all axioms.

b) C.2.2 Finite Model Construction:

We construct an explicit finite model demonstrating consistency.

C.3 (CPL Consistency) There exists a finite Kripke model $M = (W, R, V, c)$ satisfying all CPL axioms for $f(c) = c$ and $g(c) = c^2$.

Proof. Let $W = \{0, 1, 2\}$ be a set of three worlds. Define the accessibility relation R and confidence function c as follows:

$$R = \{(0, 1), (0, 2), (1, 2)\} \text{ (strictly increasing order)}$$

For each edge $(w, w') \in R$, define $c(w, w')$ as:

- $c(0, 1) = 1/2$
- $c(0, 2) = 1/4$
- $c(1, 2) = 1/2$

We verify the axioms:

Axiom K: Holds in all Kripke models by the standard modal logic semantics.

Axiom 4: For any world w and confidence c, if w satisfies $\text{box}_c(p)$, then for all w' with $(w, w') \in R$ and $c(w, w') \geq c$, we have w' satisfies p. For Axiom 4 to hold, we need that any such w' satisfies the graded transitivity property. In our model, this holds because if $(w, w') \in R$ and $(w', w'') \in R$ with appropriate confidence bounds, then $(w, w'') \in R$ directly.

Axiom GL: The graded Löb axiom requires verification. For world 0: $\text{box}_c(\text{box}_c(p) \rightarrow p)$ means: for all worlds reachable from 0 with confidence $\geq c$, if $\text{box}_c(p)$ holds at that world, then p holds.

Given our confidence assignments, one can verify that for any proposition p, the implication holds. The discount function $g(c) = c^2$ ensures that self-referential soundness claims cap their own confidence, preventing bootstrapping.

Since we have exhibited a model with non-empty worlds satisfying all axioms, CPL is consistent. ■ ■

c) C.2.3 Design Axiom Status:

It is important to clarify that CPL's graded Löb axiom is a *design axiom* rather than a derived semantic theorem. This means:

- 1) We choose to include $\text{box}_c(\text{box}_c(p) \rightarrow p) \rightarrow \text{box}_{gc}(p)$ as an axiom
- 2) We verify consistency (as shown above) but not soundness with respect to a pre-existing semantics

3) The axiom captures our intended behavior: self-soundness claims should be discounted to prevent bootstrapping

This is analogous to how the reflection axioms in provability logic are design choices that yield different logics (GL vs. S4 vs. K).

—

C. C.3 Defeat Composition Algebra

a) C.3.1 Undercut Composition:

The fundamental result for undercutting defeat is that sequential undercuts compose via the probabilistic OR operation.

C.4 (Undercut Composition) For any confidences $c, d_1, d_2 \in [0, 1]$: $\text{undercut}(\text{undercut}(c, d_1), d_2) = \text{undercut}(c, d_1 \oplus d_2)$

Proof. We compute directly from the definition $\text{undercut}(c, d) = c \times (1 - d)$:

$$\text{undercut}(\text{undercut}(c, d_1), d_2)$$

$$\text{XIII. } \text{UNDERCUT}(c \times (1 - d_1), d_2)$$

$$\text{XIV. } (c \times (1 - d_1)) \times (1 - d_2)$$

$$\text{XV. } c \times ((1 - d_1) \times (1 - d_2))$$

$$\text{XVI. } c \times (1 - d_1 - d_2 + d_1 d_2)$$

$$\text{XVII. } c \times (1 - (d_1 + d_2 - d_1 d_2))$$

$$\text{XVIII. } c \times (1 - (d_1 \oplus d_2))$$

$$\text{XIX. } \text{UNDERCUT}(c, d_1 \oplus d_2)$$

where we use the definition $d_1 \oplus d_2 = d_1 + d_2 - d_1 d_2$. ■ ■

a) C.3.2 Corollaries of Undercut Composition:

C.5 (Commutative Composition) Undercut composition is commutative: $\text{undercut}(\text{undercut}(c, d_1), d_2) = \text{undercut}(\text{undercut}(c, d_2), d_1)$

Proof. Immediate from Theorem C.4 and commutativity of \oplus : $\text{undercut}(\text{undercut}(c, d_1), d_2) = \text{undercut}(c, d_1 \oplus d_2)$

$$\text{XX. } \text{UNDERCUT}(c, d_2 \oplus d_1) = \text{UNDERCUT}(\text{UNDERCUT}(c, d_2), d_1). ■$$

C.6 (Identity) Undercut with zero defeat leaves confidence unchanged: $\text{undercut}(c, 0) = c$

Proof. $\text{undercut}(c, 0) = c \times (1 - 0) = c \times 1 = c$. ■ ■

C.7 (Annihilation) Undercut with complete defeat eliminates confidence: $\text{undercut}(c, 1) = 0$

Proof. undercut(c , 1) = $c \times (1 - 1) = c \times 0 = 0$. ■ ■

a) C.3.3 *Rebut Algebra:*

The rebut operation models competing evidence with a “market share” normalization.

C.8 (Rebut Symmetry) For any $c_{\text{for}}, c_{\text{against}} \in [0, 1]$ with $c_{\text{for}} + c_{\text{against}} > 0$: $\text{rebut}(c_{\text{for}}, c_{\text{against}}) + \text{rebut}(c_{\text{against}}, c_{\text{for}}) = 1$

Proof. From the definition $\text{rebut}(c_{\text{for}}, c_{\text{against}}) = c_{\text{for}} / (c_{\text{for}} + c_{\text{against}})$:

$$\text{rebut}(c_{\text{for}}, c_{\text{against}}) + \text{rebut}(c_{\text{against}}, c_{\text{for}})$$

$$\text{XXI. } c_{\text{FOR}} / (c_{\text{FOR}} + c_{\text{AGAINST}}) + c_{\text{AGAINST}} / (c_{\text{AGAINST}} + c_{\text{FOR}})$$

$$\text{XXII. } (c_{\text{FOR}} + c_{\text{AGAINST}}) / (c_{\text{FOR}} + c_{\text{AGAINST}})$$

$$\text{XXIII. } 1 ■$$

C.9 (Rebut Monotonicity) Rebut is monotone in the first argument and antitone in the second:

- If $c_1 \leq c_2$, then $\text{rebut}(c_1, c) \leq \text{rebut}(c_2, c)$
- If $d_1 \leq d_2$, then $\text{rebut}(c, d_2) \leq \text{rebut}(c, d_1)$

Proof. For the first claim: $\text{rebut}(c_1, c) = c_1 / (c_1 + c) \leq c_2 / (c_2 + c) = \text{rebut}(c_2, c)$ since the function $f(x) = x / (x + c)$ is increasing in x for $c > 0$.

The second claim follows similarly by noting $\text{rebut}(c, d) = 1 - \text{rebut}(d, c)$ from Theorem C.8. ■ ■

a) C.3.4 *Interaction Between Undercut and Rebut:*

Undercut and rebut represent two fundamentally different types of defeat:

- *Undercut* attacks the *inferential link* between premises and conclusion
- *Rebut* attacks the *conclusion* directly with counter-evidence

The interaction of these two operations is subtle and context-dependent. In general, rebut is applied first to aggregate competing evidence, then undercut is applied to discount the resulting confidence based on link attackers. This ordering is reflected in CLAIR’s evaluation semantics (Chapter 4).

b) C.3.5 *Limitation: Rebut Normalization:*

A key limitation of the rebut operation is that it normalizes away absolute evidence strength.

C.10 (Rebut Collapse) For any scaling factor $\lambda > 0$: $\text{rebut}(\lambda c_{\text{for}}, \lambda c_{\text{against}}) = \text{rebut}(c_{\text{for}}, c_{\text{against}})$

Proof. $\text{rebut}(\lambda c_{\text{for}}, \lambda c_{\text{against}}) = \lambda c_{\text{for}} / (\lambda c_{\text{for}} + \lambda c_{\text{against}})$

$$\text{XXIV. } \lambda c_{\text{FOR}} / (\lambda (c_{\text{FOR}} + c_{\text{AGAINST}}))$$

$$\text{XXV. } c_{\text{FOR}} / (c_{\text{FOR}} + c_{\text{AGAINST}})$$

$$\text{XXVI. } \text{REBUT}(c_{\text{FOR}}, c_{\text{AGAINST}}). ■$$

■ This means rebut cannot distinguish between “both weak” and “both strong but balanced” evidence. Chapter 6 discusses how CLAIR addresses this through provenance tracking, which preserves the absolute strengths of individual evidence sources even after rebut normalization.

XXVII. GLOSSARY

This appendix provides concise definitions of key terms, notation, and acronyms used throughout this dissertation. Terms are marked with their primary chapter of introduction.

A. D.1 Term Definitions

a) Epistemic Terms:

+— +Term | Definition | Chapter +— +Belief | A first-class value in CLAIR consisting of a proposition, confidence, and justification. Beliefs can be constructed, combined, defeated, and inspected through language operations. | 1 +Confidence | A value in the unit interval [0,1] representing degree of epistemic commitment. Confidence is *not* probability: rather than quantifying frequency or subjective chance, confidence tracks how justified a belief is given available evidence. | 3 +Commitment (Epistemic) | The stance of endorsing a proposition as true to a specified degree. High confidence (close to 1) indicates strong commitment; low confidence (close to 0) indicates weak or no commitment. | 3 +Justification | A data structure tracking the derivation history of a belief. Justifications form a directed acyclic graph (DAG) where nodes represent beliefs and edges represent inference rules or evidence sources. | 4 +Invalidation | A condition associated with a belief specifying circumstances under which the belief should be reconsidered. Invalidation enables defeasible reasoning—beliefs can be defeated without logical contradiction. | 4 +Provenance | The origin or source of evidence for a belief. CLAIR tracks provenance through justification graphs, enabling audit trails for how conclusions were reached. | 4 +Reliability | In CLAIR’s semantics, the tendency of a source to produce true beliefs. Reliability is the *semantic interpretation* of confidence: a belief with confidence c is interpreted as coming from a source with reliability c . | 3 +—

b) Operations and Relations:

+— +Term | Definition | Chapter +— +Probabilistic OR (\oplus) | Aggregation operator for independent evidence: $\mathbf{a} \oplus \mathbf{b} = \mathbf{a} + \mathbf{b} - \mathbf{ab}$. This equals the probability that at least one of two independent events occurs. Satisfies commutativity, associativity, and has identity 0. | 3 +Undercut (?) | Defeat operation that reduces confidence multiplicatively: $\text{undercut}(\mathbf{c}, \mathbf{d}) = \mathbf{c} \times (1-\mathbf{d})$. A defater with confidence \mathbf{d} reduces target confidence by factor $(1-\mathbf{d})$. | 4 +Rebuttal |

Defeat operation for conflicting evidence: $\text{rebut}(c_1, c_2) = c_1 / (c_1 + c_2)$. Normalizes competing confidences to [0,1]; equal confidences yield 1/2. | 4 +Derivation | Combining beliefs via rule application: $\text{derive}(b_1, b_2)$ pairs the values and multiplies confidences $c_1 \times c_2$. Tracks justification through rule application. | 4 +Aggregation | Combining independent evidence using \oplus : $\text{aggregate}(b_1, b_2)$ produces belief with confidence $c_1 \oplus c_2$. Tracks justification through aggregation node. | 4 +Subtyping | Confidence ordering: belief at confidence c_1 can be used where belief at c_2 is required iff $c_1 \geq c_2$. Enables confidence weakening. | 10 +-

c) *Structural Properties:*

+ -+ Term | Definition | Chapter +-+ Directed Acyclic Graph (DAG) | A graph with directed edges and no cycles. CLAIR requires justification graphs to be DAGs to ensure well-foundedness and prevent circular reasoning. | 4 +Stratification | Layering beliefs into levels 0, 1, 2, ... where level n can only refer to levels $< n$. Enforces Tarski's hierarchy to prevent self-referential paradoxes. | 6 +Well-formedness | Constraints on CLAIR programs: (1) justification graphs must be acyclic, (2) stratification constraints on introspection, (3) confidences in [0,1]. | 10 +-

d) *Logical and Modal Terms:*

+ -+ Term | Definition | Chapter +-+ CPL (Confidence-Bounded Provability Logic) | Graded extension of Gödel-Löb provability logic. Adds confidence grades to provability operator \Box . Axiomatizes self-referential reasoning with confidence discounts. | 5 +Löb's Theorem | Modal logic theorem: $\Box(\Box\varphi \rightarrow \varphi) \rightarrow \Box\varphi$. In provability logic, enables self-reference; in CLAIR, motivates anti-bootstrapping constraint. | 5 +Graded Modality | Modal operators with quantitative gradess. CLAIR's \Box_c means "provable with confidence at least c ". | 5 +Anti-bootstrapping | Principle that self-validating claims cannot increase confidence. Formally: $c \leq g(c)$ for some discount function g . CLAIR uses $g(c) = c^2$. | 5 +Kripke Semantics | Possible worlds framework for modal logic. CPL uses graded Kripke models where accessibility relations track confidence thresholds. | 5 +-

e) *Computational Terms:*

+ -+ Term | Definition | Chapter +-+ Fuel | Bound on computation steps in the Lean evaluator: $\text{evalFuel } n \ e$ evaluates for at most n steps. Prevents infinite loops while ensuring termination. | B +Call-by-value | Evaluation strategy: function arguments are evaluated before application. CLAIR uses call-by-value to match intuition about belief formation. | B +Small-step semantics | Reduction relation $e \rightarrow e'$ defining one step of computation. Composed to form multi-step evaluation $e \rightarrow e'$. | 10 +Bidirectional Type Checking | Type checking algorithm with synthesis (infer type from expression) and checking (verify expression matches type). Enables practical implementation. | 10 +De Bruijn Indices | Variable representation using natural numbers: var 0 = most recent binder, var 1 = next recent, etc. Enables formal proofs about substitution. | A +-

f) *Argumentation and Belief Revision:*

+ -+ Term | Definition | Chapter +-+ Defeasible Reasoning | Reasoning where conclusions can be defeated by new information without logical contradiction. CLAIR supports

defeasibility through undercut and rebut operations. | 4 +Underminer | An argument that attacks the connection between evidence and conclusion (e.g., "the sensor is miscalibrated"). Reduces confidence via $\text{undercut}(c,d) = c \times (1-d)$. | 4 +Rebuttal | An argument providing conflicting evidence for the opposite conclusion. Normalizes via $\text{rebut}(c_1, c_2) = c_1 / (c_1 + c_2)$. | 4 +AGM Theory | Classic belief revision framework (Alchourrón, Gärdenfors, Makinson). CLAIR extends AGM to graded, DAG-structured beliefs. | 7 +Contraction | Belief revision operation: remove a belief from a belief set. In CLAIR, achieved by setting confidence to 0. | 7 +Revision | Belief revision operation: add a belief while maintaining consistency. In CLAIR, achieved by aggregating with existing beliefs. | 7 +-

g) *Impossibility Results:*

+ -+ Term | Definition | Chapter +-+ Gödel's Incompleteness | Any consistent formal system capable of arithmetic contains true but unprovable statements. Motivates CPL's design restrictions. | 12 +Church's Undecidability | First-order logic validity is undecidable. CLAIR restricts to decidable fragments (CPL-finite, CPL-0). | 12 +Tarski's Undefinability | Truth cannot be defined within the same language. Motivates stratification: level n cannot quantify over level n . | 12 +Löb's Paradox | Curious proposition using Löb's theorem that leads to contradiction if not carefully restricted. Resolved via stratification. | 6 +-

B. D.2 Notation Table

+ -+ Symbol | Meaning | Type | Chapter +-+ c | Confidence value in [0,1] | Confidence | 3 +⊕ | Probabilistic OR: $a \oplus b = a + b - ab$ | Binary operation | 3 +? | Undercut: $c ? d = c \times (1-d)$ | Binary operation | 4 +× | Multiplication (conjunctive combination) | Binary operation | 3 +g(c) | Löb discount function; CLAIR uses c^2 | Unary function | 5 +□c, φ | Necessarily φ with confidence at least c | Modal operator | 5 +◇c, φ | Possibly φ with confidence at least c | Modal operator | 5 +|- | Typing judgment: $\Gamma \vdash e : A @ c$ | Relation | 10 +-+ | Small-step reduction: $e \rightarrow e'$ | Relation | 10 +-+ | **Multi-step reduction (reflexive transitive closure)** | Relation | 10 +Γ | Typing context (list of variable: type pairs) | Context | 10 +A ⇒ B | Function type from A to B | Type | 10 +Belief<A> | Belief type holding value of type A | Type constructor | 10 +b.value | Extract value from belief b | Projection | 4 +b.confidence | Extract confidence from belief b | Projection | 4 +b.justification | Extract justification from belief b | Projection | 4 +m <n | Stratification constraint: level m below n | Ordering | 6 +∀ | Universal quantifier | Quantifier | Appendix A +∃ | Existential quantifier | Quantifier | Appendix A +ε | Set membership | Relation | Appendix A +⊆ | Subset relation | Relation | Appendix A +∪ | Set union | Binary operation | Appendix A +∩ | Set intersection | Binary operation | Appendix A +λ*x:A. e | Lambda abstraction (anonymous function) | Expression | 10 +e*1 e*2 | Function application | Expression | 10 +-

C. D.3 Acronyms

+— +Acronym | Full Name | Definition +— +CLAIR
| Comprehensible LLM AI Intermediate Representation |
The formal system presented in this dissertation +CPL |
Confidence-Bounded Provability Logic | Graded modal logic
extending Gödel-Löb +CPL-finite | CPL with finite confidence set {0, 1/n, 2/n, ..., 1} | Decidable fragment suitable for implementation +CPL-0 | CPL with only confidence 0 (ungraded provability) | Collapses to standard provability logic GL +AGM | Alchourrón-Gärdenfors-Makinson | Classic belief revision theory +DAG | Directed Acyclic Graph | Graph structure for justifications +LLM | Large Language Model | AI systems CLAIR is designed to augment +IR | Intermediate Representation | Compiler representation between source and machine code +AI | Artificial Intelligence | Field of study +—

D. D.4 Type System Summary

a) Base Types:

+— +Type | Description | Example Values +— +Nat | Natural numbers (non-negative integers) | 0, 1, 2, 42, 128
+Bool | Boolean values | true, false +String | Text strings | “hello”, “sensor-reading” +Unit | Unit type (single value) | unit +Pair(A, B) | Ordered pair | (1, true), (“x”, 42) +Sum(A, B) | Tagged union (either A or B) | inl(5), inr(true) +A ⇒ B | Function type from A to B | $\lambda x:\text{Nat}. x + 1$ +Belief<A> | Belief holding value of type A | belief(42, 0.9, j) +—

b) Confidence Operations:

+— +Operation | Notation | Formula | Identity +— +Probabilistic OR | a ⊕ b | a + b - a×b | 0 +Multiplication | a × b | a × b × 1 +Undercut | undercut(c, d) | c × (1-d) | N/A (d=0 gives c) +Rebuttal | rebut(c₁, c₂) | c₁ / (c₁ + c₂) | N/A +Minimum | min(a, b) | min(a, b) | 0 (if bounded below)
+Maximum | max(a, b) | max(a, b) | 1 (if bounded above)
+—

XXVIII. COMPLETE CLAIR LANGUAGE SPECIFICATION

This appendix provides the complete formal specification of the CLAIR language, including syntax (concrete and abstract), static semantics (type system), and dynamic semantics (operational rules). The specification is self-contained and sufficient for implementing a conforming CLAIR interpreter.

A. E.1 Syntax

a) E.1.1 Type Grammar:

The type grammar defines the set of valid types in CLAIR.
+— +Category | Production | Description Base Types | B ::= “Nat” | Natural numbers || B ::= “Bool” | Boolean values || B ::= “String” | Text strings || B ::= “Unit” | Unit type (single value) Confidence | c ∈ Q | Rational in [0,1] Types | A ::= B | Base type || A ::= A → B | Function type || A ::= A × B | Product type || A ::= A + B | Sum type || A ::= “Belief<A[“c”]>” | Belief type with confidence bound || A ::= “Meta<A[“n”][“c”]>” | Stratified meta-belief type +—

b) E.1.2 Expression Grammar:

The expression grammar defines the syntactic forms of CLAIR programs.

+— +Category | Production | Description Variables | e ::= x | Variable reference Lambdas | e ::= “λ”x：“A”.e | Anonymous function Application | e ::= e e | Function application Pairs | e ::= “(e , “ e)” | Ordered pair Projections | e ::= e . “1” | First projection || e ::= e . “2” | Second projection Injections | e ::= “inl@”B (“e)” | Left injection || e ::= “inr@”A (“e)” | Right injection Case | e ::= “case” e “of” “inl” x “=>” e “|” “inr” y “=>” e | Sum elimination Literals | e ::= n | Natural number literal || e ::= “true” | “false” | Boolean literals || e ::= “”s”” | String literal || e ::= “()” | Unit literal Beliefs | e ::= “belief(e , “ c , “ j)” | Belief constructor || e ::= “val(“e”)” | Extract belief value || e ::= “conf(“e”)” | Extract belief confidence || e ::= “just(“e”)” | Extract belief justification Derivation | e ::= “derive(“e , “ e)” | Combine beliefs (conjunctive) Aggregation | e ::= “aggregate(“e , “ e)” | Aggregate beliefs (independent) Defeat | e ::= “undercut(“e , “ e)” | Apply undercut || e ::= “rebut(“e , “ e)” | Apply rebuttal Introspection | e ::= “introspect(“e”)” | Safe self-reference Let Binding | e ::= “let” x “=” e “in” e | Local binding +—

c) E.1.3 Abstract Syntax:

The abstract syntax uses de Bruijn indices for variable representation.

Definition The abstract syntax “Expr” is defined inductively with the following constructors:

“var(n)” — Variable at de Bruijn index n “lam(A, e)” — Lambda abstraction “λ”：“A”.e “app(e₁, e₂)” — Function application e₁ e₂ “pair(e₁, e₂)” — Ordered pair “fst(e)” — First projection “snd(e)” — Second projection “inl(B, e)” — Left injection “inr(A, e)” — Right injection “case(e, e₁, e₂)” — Case analysis “litNat(n)” — Natural literal “litBool(b)” — Boolean literal “litString(s)” — String literal “litUnit” — Unit literal “belief(v, c, j)” — Belief constructor “val(e)” — Value extraction “conf(e)” — Confidence extraction “just(e)” — Justification extraction “derive(e₁, e₂)” — Derivation “aggregate(e₁, e₂)” — Aggregation “undercut(e, d)” — Undercut “rebut(e₁, e₂)” — Rebuttal “introspect(e)” — Introspection “letIn(e₁, e₂)” — Let binding

The “Justification” type tracks derivation structure: “axiomJ(name)” — Named axiom “rule(name, js)” — Named rule with premises “agg(js)” — Aggregation “undercut_j(j, d)” — Undercut application “rebut_j(j₁, j₂)” — Rebuttal application

d) E.1.4 Well-Formedness:

A type is well-formed if all confidence bounds are in [0,1].

Definition The judgment “wf(A)” holds when:

“wf(B)” for all base types B “wf(A) ∧ wf(B)” ⇒ “wf(A → B)” “wf(A) ∧ wf(B)” ⇒ “wf(A × B)” “wf(A) ∧ wf(B)” ⇒ “wf(A + B)” “wf(A) ∧ c ∈ [0,1]” ⇒ “wf(Belief<A[“c”]>)” “wf(A) ∧ c ∈ [0,1]” ⇒ “wf(Meta<A[“n”][“c”]>)”

B. E.2 Static Semantics (Type System)

a) E.2.1 Typing Contexts:

A typing context Γ maps variable indices to type-confidence pairs.

$\Gamma ::= \emptyset | \Gamma , “ \langle A, c \rangle ”$

Definition The lookup operation $\Gamma.\text{lookup}(n)$ returns the type-confidence pair at index n (0-indexed from most recent binding).

b) E.2.2 Typing Judgment Form:

The primary typing judgment has the form:

$\Gamma \vdash e :: A @ c$

In context Γ , expression e has type A with confidence bound c .

c) E.2.3 Typing Rules:

The following rules define well-typed CLAIR expressions. Confidence propagation is explicit in each rule.

T-Variable: $\Gamma \vdash \text{var}(n) :: A @ c$ if $\Gamma.\text{lookup}(n) = (A, c)$

T-Abstraction: If $\Gamma, \langle A, c_A \rangle \vdash e :: B @ c_B$, then $\Gamma \vdash \text{lam}(A, e) :: (A \rightarrow B) @ c_B$

T-Application: If $\Gamma \vdash e_1 :: (A \rightarrow B) @ c_1$ and $\Gamma \vdash e_2 :: A @ c_2$, then $\Gamma \vdash \text{app}(e_1, e_2) :: B @ (c_1 \times c_2)$

Confidence multiplies for conjunctive derivation.

T-Pair: If $\Gamma \vdash e_1 :: A @ c_1$ and $\Gamma \vdash e_2 :: B @ c_2$, then $\Gamma \vdash \text{pair}(e_1, e_2) :: (A \times B) @ (c_1 \times c_2)$

T-Fst/T-Snd preserve confidence.

T-Inl/T-Inr preserve confidence.

T-Case: If $\Gamma \vdash e :: (A + B) @ c$ and $\Gamma, \langle A, c \rangle \vdash e_1 :: C @ c_1$ and $\Gamma, \langle B, c \rangle \vdash e_2 :: C @ c_2$, then $\Gamma \vdash \text{case}(e, e_1, e_2) :: C @ (c \oplus c_1 \oplus c_2)$

Uses probabilistic OR for confidence aggregation.

T-Belief: If $\Gamma \vdash v :: A @ 1$ and $c_{\text{bound}} \leq c_{\text{actual}}$, then $\Gamma \vdash \text{belief}(v, c_{\text{actual}}, j) :: \text{Belief}[A[c_{\text{bound}}]] @ c_{\text{bound}}$

The belief's actual confidence c_{actual} must meet or exceed the declared bound c_{bound} .

T-Val/T-Conf/T-Just preserve the enclosing confidence.

T-Derive: If $\Gamma \vdash e_1 :: \text{Belief}[A[c_1]] @ c_{e_1}$ and $\Gamma \vdash e_2 :: \text{Belief}[B[c_2]] @ c_{e_2}$, then $\Gamma \vdash \text{derive}(e_1, e_2) :: \text{Belief}[A \times B[c_1 \times c_2]] @ (c_{e_1} \times c_{e_2})$

Confidence multiplies: both premises must be true.

T-Aggregate: If $\Gamma \vdash e_1 :: \text{Belief}[A[c_1]] @ c_{e_1}$ and $\Gamma \vdash e_2 :: \text{Belief}[A[c_2]] @ c_{e_2}$, then $\Gamma \vdash \text{aggregate}(e_1, e_2) :: \text{Belief}[A[c_1 \oplus c_2]] @ (c_{e_1} \oplus c_{e_2})$

Uses probabilistic OR: independent evidence.

T-Undercut: If $\Gamma \vdash e :: \text{Belief}[A[c]] @ c_e$ and $\Gamma \vdash d :: \text{Belief}[d_c] @ c_d$, then $\Gamma \vdash \text{undercut}(e, d) :: \text{Belief}[A[\text{undercut}(c, d_c)]] @ (c_e \times c_d)$ where $\text{undercut}(c, d) = c \times (1 - d)$.

T-Rebut: If $\Gamma \vdash e_{\text{for}} :: \text{Belief}[A[c_{\text{for}}]] @ c_{e_1}$ and $\Gamma \vdash e_{\text{against}} :: \text{Belief}[A[c_{\text{against}}]] @ c_{e_2}$, then $\Gamma \vdash \text{rebut}(e_{\text{for}}, e_{\text{against}}) :: \text{Belief}[A[\text{rebut}(c_{\text{for}}, c_{\text{against}})]] @ (c_{e_1} \times c_{e_2})$ where $\text{rebut}(c_{\text{for}}, c_{\text{against}}) = c_{\text{for}}$

T-Introspect: If $\Gamma \vdash e :: \text{Meta}[A[m][c]] @ c_e$ and $m < n$, then:

The resulting type is: $\text{Meta}[A[m][c]] @ [n][g(c)]$ (a meta-belief about a meta-belief).

Thus: $\Gamma \vdash \text{introspect}(e) :: \text{Meta}[A[m][c]] @ [n][g(c)] @ c_e$ where $g(c) = c^2$ is the Löb discount function (to prevent bootstrapping).

Requires level constraint: only higher levels can introspect lower levels.

d) E.2.4 Subtyping:

CLAIR supports subtyping based on confidence bounds.

S-Belief: $\text{Belief}[A[c]] < \text{Belief}[A[c']]$ if $c \geq c'$. Higher confidence is a subtype of lower confidence.

S-Meta follows the same rule.

T-Sub: If $\Gamma \vdash e :: A @ c$ and $A < A'$ and $c \geq c'$, then $\Gamma \vdash e :: A' @ c'$

Allows weakening both type and confidence.

C. E.3 Dynamic Semantics

a) E.3.1 Values:

A value is a fully evaluated expression.

Definition The predicate “ $\text{IsValue}(e)$ ” holds for:

- All lambda abstractions “ $\text{lam}(A, e)$ ”
- All pairs “ $\text{pair}(v_1, v_2)$ ” where v_1, v_2 are values
- All injections “ $\text{inl}(B, v)$ ” and “ $\text{inr}(A, v)$ ” where v is a value
- All literals (“ litNat ”, “ litBool ”, “ litString ”, “ litUnit ”)
- All belief constructors “ $\text{belief}(v, c, j)$ ” where v is a value

b) E.3.2 Small-Step Operational Semantics:

The small-step relation $e \rightarrow e'$ defines single-step reduction using call-by-value evaluation order.

Beta Reduction: $(\lambda:A.e) v \rightarrow e[x := v]$ if $\text{IsValue}(v)$

Context Rules: $\text{app}(e_1, e_2) \rightarrow \text{app}(e'_1, e_2)$ if $e_1 \rightarrow e'_1$ and $\text{app}(v_1, e_2) \rightarrow \text{app}(v_1, e'_2)$ if $e_2 \rightarrow e'_2$ and $\text{IsValue}(v_1)$

Projections: $\text{fst}(\text{pair}(v_1, v_2)) \rightarrow v_1$ if $\text{IsValue}(v_1)$, $\text{IsValue}(v_2)$ and $\text{snd}(\text{pair}(v_1, v_2)) \rightarrow v_2$ if $\text{IsValue}(v_1)$, $\text{IsValue}(v_2)$

Context rules for “ fst ” and “ snd ” evaluate subexpressions first.

Case Analysis: When $\text{case}(\text{inl}(v), e_1, e_2)$ evaluates and v is a value, it reduces to $e_1[x := v]$. When $\text{case}(\text{inr}(v), e_1, e_2)$ evaluates and v is a value, it reduces to $e_2[y := v]$.

Context rule for “ case ” evaluates the scrutinee first.

Context rule for “ case ” evaluates the scrutinee first.

Let Binding: $\text{letIn}(v, e) \rightarrow e[x := v]$ if $\text{IsValue}(v)$

Context rule for “ letIn ” evaluates the binding first.

Belief Projections: $\text{val}(\text{belief}(v, c, j)) \rightarrow v$ if $\text{IsValue}(v)$ and $\text{conf}(\text{belief}(v, c, j)) \rightarrow \text{belief}(v, c, j)$ if $\text{IsValue}(v)$ and $\text{just}(\text{belief}(v, c, j)) \rightarrow \text{litString}(\text{toString}(j))$ if $\text{IsValue}(v)$

Context rules evaluate subexpressions to values first.

Derivation, Aggregation, Defeat:

These operations evaluate to values using the confidence operations defined in the typing rules. The small-step semantics provides context rules that evaluate both subexpressions to values before computing the result.

For example, when $\text{derive}(v_1, v_2)$ has both subexpressions as values, the evaluator computes a new belief with:

- Value: “pair(v₁.value, v₂.value)”
- Confidence: “v₁.confidence × v₂.confidence”
- Justification: “rule(“derive”, [v₁.just, v₂.just])”

Introspection: “introspect(v)” “→” v if “IsValue(v)”
Context rule evaluates subexpression first.

c) E.3.3 Multi-Step Reduction:

The multi-step reduction relation e “→→” e’ is the reflexive-transitive closure of “→”.

e “→→” e if e = e’ e “→→” e’ if e “→” e’ and e “→→” e’

d) E.3.4 Evaluation Function:

The evaluation function “eval(e)” returns the result of reducing e to a value, or fails if:

- 1) The expression gets stuck (no applicable rule)
- 2) The expression exceeds the fuel bound (default: 1000 steps)

Definition “eval(e)” = “some(v)” if e “→→” v and “IsValue(v)” “eval(e)” = “none” otherwise (stuck or out of fuel)

D. E.4 Well-Formedness Constraints

a) E.4.1 Acyclicity of Justification Graphs:

For deterministic evaluation in the DAG semantics, justification graphs must be acyclic.

Definition A justification graph G = “(V, E)” is **well-formed** iff:

- 1) For all nodes v ∈ V, the transitive closure of E starting from v contains no cycles.
- 2) Equivalently: there is no path v “→→” v for any v ∈ V.

Enforcement: In the reference interpreter, this is checked during type checking by tracking provenance and ensuring no belief can transitively depend on itself.

b) E.4.2 Stratification Constraints:

For safe introspection, meta-belief levels must form a strict hierarchy.

Definition The **stratification constraint** requires:

- 1) Every “introspect(e)” operation has proof m “<” n where:
 - m is the level of the source belief
 - n is the level of the resulting meta-belief
- 2) This is enforced at compile time: the type checker requires a proof term of type m “<” n.

Consequence: No belief can introspect itself or any belief that transitively introspects it.

c) E.4.3 Confidence Bounds:

All confidence values must satisfy 0 “≤” c “≤” 1.

Enforcement: The type system checks this at all belief construction points.

E. E.5 Example Programs

Basic Belief Basic Belief:

```
belief(42, 0.9, axiomJ("sensor-reading"))
```

A belief in the value 42 with confidence 0.9, traced to a sensor reading axiom.

Conjunctive Derivation

```
let p = belief("it is raining", 0.8, axiomJ("weather-report"))
in
let q = belief("I have an umbrella", 0.9, axiomJ("visual-check"))
in
derive(p, q)
```

Combines two beliefs by multiplication, yielding confidence 0.72.

Aggregation

Independent Evidence Aggregation:

```
let e1 = belief("stock will rise", 0.6, axiomJ("analyst-a"))
in
let e2 = belief("stock will rise", 0.7, axiomJ("analyst-b"))
in
aggregate(e1, e2)
```

Uses probabilistic OR: 0.6 ⊕ 0.7 = 0.6 + 0.7 - 0.6×0.7 = 0.88.

Undercut

Defeat:

```
let claim = belief("system is secure", 0.95,
axiomJ("vendor-claim"))
in
let vuln = belief("critical CVE found", 0.8,
axiomJ("security-audit"))
in
undercut(claim, vuln)
```

Applies undercut: 0.95 × (1 - 0.8) = 0.19.

F. E.6 Summary

This specification provides:

- 1) **Complete type grammar:** Base types, functions, products, sums, beliefs, and meta-beliefs
- 2) **Complete expression grammar:** All CLAIR syntactic forms
- 3) **Static semantics:** 20 typing rules covering all constructs
- 4) **Dynamic semantics:** Small-step operational semantics with call-by-value evaluation
- 5) **Well-formedness constraints:** Acyclicity, stratification, and confidence bounds

The specification is sufficient for implementing a conforming CLAIR interpreter and type checker. The Lean 4 formalization in Appendix A provides a machine-checked version of this specification.

XXIX. APPENDIX F: EVALUATION PROMPTS

This appendix provides the complete prompt templates used for the empirical evaluation in Chapter 14. These

prompts instruct LLMs to output CLAIR-formatted reasoning traces.

A. F.1 GSM8K Prompts

a) F.1.1 System Prompt:

You are CLAIR-MATH, a mathematical reasoning system that outputs structured reasoning traces with explicit confidence scores, justifications, and invalidation conditions.

b) F.1.2 Task Instruction:

Solve the following grade school math problem step-by-step. For each step, output:

Step N: BELIEF("description of reasoning step", c=X.XX, justification="why this step is valid", invalidation="what would make this step false")

Confidence Guidelines:

- c=1.00 for mathematical certainties (e.g., "3 + 5 = 8")
- c=0.95-0.99 for direct inferences from problem statement
- c=0.80-0.94 for multi-step calculations
- c=0.60-0.79 for assumptions about implicit information
- c=0.40-0.59 for uncertain interpretations

Justification Guidelines:

- Reference explicit premises from the problem
- Reference previous step numbers when building on them
- State the mathematical operation being performed

Invalidation Guidelines:

- State what condition would make this step incorrect
- Be specific: "if the problem meant X instead of Y"

Confidence Operations:

- Multiply confidences when steps are chained: $c_{\text{new}} = c_1 \times c_2 \times \dots$
- Use \oplus (oplus) when combining independent sources: $c = (c_1 + c_2) / 2$
- Apply undercut when discovering contradictions: $c_{\text{new}} = c \times (1 - \text{defeat_strength})$

c) F.1.3 Few-Shot Example:

Problem: Janet's ducks lay 16 eggs per day. She eats 3 for breakfast every morning and bakes muffins with 4 eggs every day. How many eggs does she have left at the end of the day?

Step 1: BELIEF("Janet's ducks lay 16 eggs per day", c=1.00, justification="explicit premise in problem", invalidation="none - stated as fact")

Step 2: BELIEF("Janet eats 3 eggs for breakfast", c=1.00, justification="explicit premise in problem", invalidation="none - stated as fact")

Step 3: BELIEF("Janet bakes with 4 eggs", c=1.00, justification="explicit premise in problem", invalidation="none - stated as fact")

Step 4: BELIEF("Total eggs used = eggs eaten + eggs baked", c=0.99, justification="arithmetic principle: total consumption", invalidation="if 'bakes' doesn't mean 'uses all eggs'")

Step 5: BELIEF("Total eggs used = 3 + 4 = 7", c=0.99, justification="arithmetic: 3 + 4 = 7 using Steps 2,3", invalidation="none - basic addition")

Step 6: BELIEF("Eggs remaining = eggs laid - eggs used", c=0.99, justification="arithmetic principle: remaining = initial - used", invalidation="if she also uses eggs for other purposes")

Step 7: BELIEF("Eggs remaining = 16 - 7 = 9", c=0.99 $\times 0.99 \times 0.99 = 0.97$, justification="arithmetic: 16 - 7 = 9 using Steps 1,4,5,6", invalidation="none - basic subtraction")

ANSWER: 9 eggs (confidence: 0.97)

d) F.1.4 Test Prompt Template:

Solve the following math problem using the CLAIR format shown above:

Problem: {PROBLEM_TEXT}

Provide your reasoning in the CLAIR format with confidence scores, justifications, and invalidation conditions.

B. F.2 HotpotQA Prompts

a) F.2.1 System Prompt:

You are CLAIR-QA, a question answering system that builds justification graphs for multi-hop reasoning.

b) F.2.2 Task Instruction:

Answer the following question by reasoning over the provided context. Build a justification DAG where each step may depend on multiple previous steps.

Step N: BELIEF("intermediate conclusion", c=X.XX, justification="using entity E from passage P", support=[Step i, Step j, ...])

Support Guidelines:

- support=[i, j, ...]

lists which previous steps this step builds on

- Multiple support indicates the step requires all those premises

- This forms the justification DAG structure

Confidence for Multi-Hop Reasoning:

- c=1.00 for direct entity extraction from passages
- c=0.90-0.95 for single-hop inferences (bridging entities)
- c=0.70-0.85 for multi-hop chains (multiply per hop: 0.90^n for n hops)
- c=0.60-0.75 for synthesizing across multiple passages

Defeat Operations:

- When passages contradict, apply undercut: $c_{\text{new}} = c \times (1 - \text{contradiction_strength})$
- When answering multi-part questions, use rebut to compare alternatives

c) F.2.3 Few-Shot Example:

Question: Were Scott Derrickson and Ed Wood of the same nationality?

Passage 1: Scott Derrickson (born May 16, 1966) is an American film director.

Passage 2: Ed Wood (1924–1978) was an American filmmaker.

Step 1: `BELIEF("Scott Derrickson is American", c=1.00, justification="direct extraction from Passage 1", support=[])`

Step 2: `BELIEF("Ed Wood is American", c=1.00, justification="direct extraction from Passage 2", support=[])`

Step 3: `BELIEF("American refers to nationality of United States", c=0.99, justification="background knowledge", support=[])`

Step 4: `BELIEF("Both are American nationality", c=1.00 × 1.00 = 1.00, justification="both share nationality from Steps 1,2", support=[1, 2])`

Step 5: `BELIEF("Therefore they are of the same nationality", c=1.00, justification="identity from same nationality value using Steps 3,4", support=[3, 4])`

ANSWER: Yes (confidence: 1.00)

d) F.2.4 Test Prompt Template:

Answer the following question using the CLAIR format:

Question: {QUESTION}

Context Passages: {PASAGES}

Build your answer as a justification DAG with explicit support dependencies.

C. F.3 FOLIO Prompts

a) F.3.1 System Prompt:

You are CLAIR-LOGIC, a logical reasoning system that handles first-order logic with defeasible reasoning and belief revision.

b) F.3.2 Task Instruction:

Determine whether the conclusion follows from the premises. Use CLAIR to track logical uncertainty and apply defeat operations when arguments are undermined.

Step N: `BELIEF("logical derivation", c=X.XX, justification="rule application: Modus Ponens/Universal Instantiation/etc", defeat=[defeater description if applicable])`

Logical Confidence Levels:

- c=1.00 for logically necessary deductions (Modus Ponens from certainties)
- c=0.90-0.95 for probabilistic inferences from existential quantifiers
- c=0.70-0.85 for defeasible generalizations
- c=0.50-0.70 for abductive inferences (inference to best explanation)

Defeat in Logical Reasoning:

undercut

: Attack the connection between premise and conclusion

rebut

: Provide a contradictory conclusion

rebuttal

: Attack an attacking argument

Self-Reference Detection:

- When a conclusion refers to itself, apply graded Löb discount: $g(c) = c^2$
- This prevents overconfident self-bootstrapping

c) F.3.2 Few-Shot Example:

Premises:

1) All doctors are college graduates.

2) All cardiologists are doctors.

3) John is a cardiologist.

4) Some college graduates are not medical professionals.

Conclusion: John is a medical professional.

Step 1: `BELIEF("vx (Doctor(x) → CollegeGraduate(x))", c=1.00, justification="Premise 1 as logical axiom", defeat=[])`

Step 2: `BELIEF("vx (Cardiologist(x) → Doctor(x))", c=1.00, justification="Premise 2 as logical axiom", defeat=[])`

Step 3: `BELIEF("Cardiologist(John)", c=1.00, justification="Premise 3 as fact", defeat=[])`

Step 4: `BELIEF("Doctor(John)", c=1.00, justification="Modus Ponens: Cardiologist(John) ∧ vx(Cardiologist→Doctor) ⊢ Doctor(John) using Steps 2,3", defeat=[])`

Step 5: `BELIEF("CollegeGraduate(John)", c=1.00, justification="Modus Ponens: Doctor(John) ∧ vx(Doctor→CollegeGraduate) ⊢ CollegeGraduate(John) using Steps 1,4", defeat=[])`

Step 6: `BELIEF("MedicalProfessional is a superset of Doctor", c=0.95, justification="defeasible generalization: doctors are medical professionals", defeat=[])`

Step 7: `BELIEF("MedicalProfessional(John)", c=1.00 × 0.95 = 0.95, justification="instantiation: Doctor(John) implies MedicalProfessional(John) using Steps 4,6", defeat=[])`

Step 8: `BELIEF("Conclusion follows", c=0.95, justification="derivation chain supports conclusion", defeat=[])`

ANSWER: True (confidence: 0.95)

d) F.3.3 Defeat Example:

Premises:

1) Birds typically fly.

2) Tweety is a bird.

3) Tweety is a penguin.

4) Penguins cannot fly.

Conclusion: Tweety cannot fly.

Step 1: `BELIEF("Tweety flies", c=0.80, justification="defeasible generalization: birds fly using Premises 1,2", defeat=[])`

Step 2: `BELIEF("The generalization 'birds fly' has exceptions", c=1.00, justification="background knowledge", defeat=[])`

Step 3: BELIEF("Penguins are exceptions to 'birds fly'", c=1.00, justification="Premise 4 + background knowledge", defeat=[])

Step 4: BELIEF("Tweety is a penguin", c=1.00, justification="Premise 3", defeat=[])

Step 5: BELIEF("UNDERCUT: 'birds fly' generalization doesn't apply to Tweety", c=1.00, justification="specific overrides general: penguin Tweety using Steps 3,4", defeat=["attacks Step 1"])

Step 6: BELIEF("Tweety cannot fly (rebuttal to Step 1)", c=1.00, justification="penguins cannot fly using Premise 4", defeat=["rebuts Step 1"])

Step 7: BELIEF("Tweety cannot fly", c=1.00, justification="specific fact overrides defeasible generalization using Steps 5,6", defeat=[])

ANSWER: True (confidence: 1.00)

Note: This demonstrates defeasible reasoning: a specific fact (Tweety is a penguin) defeats a defeasible generalization (birds fly).

e) F.3.4 Test Prompt Template:

Determine if the conclusion follows from the premises using CLAIR logical reasoning:

Premises: {PREMISES}

Conclusion: {CONCLUSION}

Apply defeat operations when premises create contradictory or undercutting relationships.

D. F.4 Post-Processing and Validation

a) F.4.1 Confidence Extraction:

For baseline methods, confidence is extracted as follows:

Chain-of-Thought / Tree of Thoughts / DSPy:

```
confidence = max(logprob(token) for token in
answer_tokens)
# Normalize to [0,1]
confidence = sigmoid(confidence / temperature)
```

Self-Consistency:

```
confidence = (count_agreeing / total_samples)
# Voting proportion as confidence
```

CLAIR:

```
# Extract final confidence from CLAIR trace
confidence = trace[-1]['confidence']
# Already normalized to [0,1]
```

b) F.4.2 Validation Checklist:

For each CLAIR output, verify:

- 1) **Format correctness:** All BELIEF statements have c, justification, and invalidation/defeat
- 2) **Confidence bounds:** All c values are in [0, 1]
- 3) **Justification structure:** Support dependencies form a DAG (no cycles unless intentional)
- 4) **Confidence calculus:** Multiplication for chains, \oplus for independent evidence
- 5) **Defeat application:** Undercut/rebut applied correctly when contradictions arise

- 6) **Self-reference discount:** Graded Löb applied (c^2) for self-referential beliefs

c) F.4.3 Error Categories for Annotation:

When evaluating CLAIR outputs, categorize errors as:

C1: Syntax errors. Invalid CLAIR format, missing confidence/justification fields

C2: Confidence miscalibration. Confidence doesn't match logical strength of derivation

C3: Invalid confidence propagation. Using \oplus for dependent sources, incorrect multiplication

C4: Missing justification dependencies. DAG doesn't include necessary support links

C5: Incorrect defeat application. Using rebut when undercut is appropriate (or vice versa)

C6: Unjustified independence assumption. Applying \oplus to correlated sources

C7: Self-reference bootstrapping. Circular reasoning without graded Löb discount

C8: Semantic errors. Reasoning step is logically invalid regardless of CLAIR structure